# Explaining away ambiguity: Learning verb selectional preference with Bayesian networks\*

### Massimiliano Ciaramita and Mark Johnson

Cognitive and Linguistic Sciences Box 1978, Brown University Providence, RI 02912, USA

massimiliano\_ciaramita@brown.edu mj@cs.brown.edu

#### Abstract

This paper presents a Bayesian model for unsupervised learning of verb selectional preferences. For each verb the model creates a Bayesian network whose architecture is determined by the lexical hierarchy of Wordnet and whose parameters are estimated from a list of verbobject pairs found from a corpus. "Explaining away", a well-known property of Bayesian networks, helps the model deal in a natural fashion with word sense ambiguity in the training data. On a word sense disambiguation test our model performed better than other state of the art systems for unsupervised learning of selectional preferences. Computational complexity problems, ways of improving this approach and methods for implementing "explaining away" in other graphical frameworks are discussed.

# 1 Selectional preference and sense ambiguity

Regularities of a verb with respect to the semantic class of its arguments (subject, object and indirect object) are called **selectional preferences** (SP) (Katz and Fodor, 1964; Chomsky, 1965; Johnson-Laird, 1983). The verb *pilot* carries the information that its object will likely be some kind of *vehicle*; subjects of the verb *think* tend to be *human*; and subjects of the verb *bark* tend to be *dogs*. For the sake of simplicity we will focus on the verb-object relation although the techniques we will describe can be applied to other verb-argument pairs.

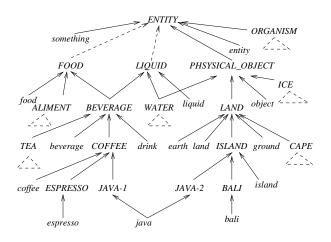


Figure 1: A portion of Wordnet.

Models of the acquisition of SP are important in their own right and have applications in Natural Language Processing (NLP). The selectional preferences of a verb can be used to infer the possible meanings of an unknown argument of a known verb; e.g., it might be possible to infer that xxxx is a kind of dog from the following sentence: "The xxxx barked all night". In parsing a sentence selectional preferences can be used to rank competing parses, providing a partial measure of semantic well-formedness. Investigating SP might help us to understand the structure of the mental lexicon.

Systems for unsupervised learning of SP usually combine statistical and knowledge-based approaches. The knowledge-base component is typically a database that groups words into classes. In the models we will see, the knowledge base is Wordnet (Miller, 1990). Wordnet groups nouns into classes of synonyms representing concepts, called **synsets**, e.g., {car, auto, automobile, ...}. A noun that belongs to several synsets is **ambiguous**. A tran-

<sup>\*</sup> We would like to thank the Brown Laboratory for Linguistic Information Processing; Thomas Hofmann; Elie Bienenstock; Philip Resnik, who provided us with training and test data; and Daniel Garcia for his help with the SMILE library of classes for Bayesian networks that we used for our experiments. This research was supported by NSF awards 9720368, 9870676 and 9812169.

sitive and asymmetrical relation, **hyponymy**, is defined between synsets. A synset is a hyponym of another synset if the former has the latter as a broader concept; for example, *BEV-ERAGE* is a hyponym of *LIQUID*. Figure 1 depicts a portion of the hierarchy.

The statistical component consists predicate-argument pairs extracted from a corpus in which the semantic class of the words is not indicated. A trivial algorithm might get a list of words that occurred as objects of the verb and output the semantic classes the words belong to according to Wordnet. For example, if the verb drink occurred with water and water  $\in LIQUID$ , the model would learn that drink selects for LIQUID. As Resnik (1997) and Abney and Light (1999) have found, the main problem these systems face is the presence of ambiguous words in the training data. If the word java also occurred as an object of drink, since  $java \in BEVERAGE$ and  $java \in ISLAND$ , this model would learn that drink selects for both BEVERAGE and ISLAND.

More complex models have been proposed. These models, though, deal with word sense ambiguity by applying an *unselective* strategy similar to the one above; i.e., they assume that ambiguous words provide equal evidence for all their senses. These models choose as the concepts the verb selects for those that are in common among several words (e.g., *BEVERAGE* above). This strategy works to the extent that these overlapping senses are also the concepts the verb selects for.

# 2 Previous approaches to learning selectional preference

### 2.1 Resnik's model

Ours system is closely related to those proposed in (Resnik, 1997) and (Abney and Light, 1999). The fact that a predicate p selects for a class c, given a syntactic relation r, can be represented as a relation, selects(p, r, c); e.g., that eat selects for FOOD in object position can be represented as selects(eat, object, FOOD). In (Resnik, 1997) selectional preference is quantified by comparing the prior distribution of a given class c appearing as an argument, P(c), and the conditional probability of the same class given a predicate and a syntac-

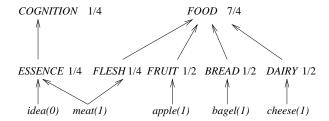


Figure 2: Simplified Wordnet. The numbers next to the synsets represent the values of freq(p, r, c) estimated using (3), the numbers in parentheses represent the values of freq(p, r, w).

tic relation P(c|p,r), e.g., P(FOOD) and P(FOOD|eat, object). The relative entropy between P(c) and P(c|p,r) measures how much the predicate constrains its arguments:

$$S(p,r) = D(P(c|p,r) || P(c))$$
 (1)

Resnik defines the **selectional association** of a predicate for a particular class c to be the portion of the selectional preference strength due to that class:

$$A(p, r, c) = \frac{1}{S(p, r)} P(c|p, r) \log \frac{P(c|p, r)}{P(c)}$$
 (2)

Here the main problem is the estimation of P(c|p,r). Resnik suggests as a plausible estimator  $\hat{P}(c|p,r) \stackrel{\text{def}}{=} freq(p,r,c)/freq(p,r)$ . But since the model is trained on data that are not sense-tagged, there is no obvious way to estimate freq(p,r,c). Resnik suggests considering each observation of a word as evidence for each of the classes the word belongs to,

$$freq(p, r, c) \approx \sum_{w \in c} \frac{count(p, r, w)}{classes(w)}$$
 (3)

where count(p, r, w) is the number of times the word w occurred as an argument of p in relation r, and classes(w) is the number of classes w belongs to. For example, suppose the system is trained on (eat, object) pairs and the verb occurred once each with meat, apple, bagel, and cheese, and Wordnet is simplified as in Figure 2. An ambiguous word like meat provides evidence also for classes that appear unrelated to those selected by the verb. Resnik's assumption is that only the classes selected by the verb will be associated with each

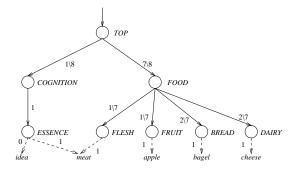


Figure 3: The HMM version of the simple example.

of the observed words, and hence will receive the highest values for P(c|p,r). Using (3) we find that the highest frequency is in fact associated with FOOD:  $freq(eat, object, food) \approx \frac{1}{4} + \frac{1}{2} + \frac{1}{2} + \frac{1}{2} = \frac{7}{4}$  and P(FOOD|eat) = 0.44. However, some evidence is found also for COGNITION:  $freq(eat, object, cognition) \approx \frac{1}{4}$  and P(COGNITION|eat) = 0.06.

### 2.2 Abney and Light's approach

Abney and Light (1999) pointed out that the distribution of senses of an ambiguous word is not uniform. They noticed also that it is not clear how the probability P(c|p,r) is to be interpreted since there is no explicit stochastic generation model involved.

They proposed a system that associates a Hidden Markov Model (HMM) with each predicate-relation pair (p,r). Transitions between synset states represent the hyponymy relation, and  $\varepsilon$ , the empty word, is emitted with probability 1; transitions to a final state emit a word w with probability  $0 \le P(w) \le 1$ . Transition and emission probabilities are estimated using the EM algorithm on training data that consist of the nouns that occurred with the verb. Abney and Light's model estimates P(c|p,r) from the model trained for (p,r); the distribution P(c) can be calculated from a model trained for all nouns in the corpus.

This model did not perform as well as expected. An ambiguous word in the model can be generated by more than one state sequence. Abney and Light discovered that the EM algorithm finds parameter values that associate some probability mass with all the transitions in the multiple paths that lead to an ambigu-

ous word. In other words, when there are several state sequences for the same word, EM does not select one of them over the others.<sup>1</sup> Figure 3 shows the parameters estimated by EM for the same example as above. The transition to the COGNITION state has been assigned a probability of 1/8 because it is part of a possible path to meat. The HMM model does not solve the problem of the unselective distribution of the frequency of occurrence of an ambiguous word to all its senses. Abney and Light claimed that this is a serious problem, particularly when the ambiguous word is a frequent one, and caused the model to learn the wrong selectional preferences. To correct this undesirable outcome they introduced some smoothing and balancing techniques. However, even with these modifications their system's performance was below that achieved by Resnik's.

### 3 Bayesian networks

Bayesian network (Pearl, 1988), or Bayesian belief network (BBN), consists of a set of variables and a set of directed edges connecting the variables. The variables and the edges define a directed acyclic graph (DAG) where each variable is represented by a node. Each variable is associated with a finite number of (mutually exclusive) states. To each variable A with parents  $B_1, ..., B_n$  is attached a conditional probability table (CPT)  $P(A|B_1,...,B_n)$ . Given a BBN, Bayesian inference can be used to estimate marginal and posterior probabilities given the evidence at hand and the information stored in the CPTs, the **prior prob**abilities, by means of Bayes' rule, P(H|E) = $\frac{P(H)P(E|H)}{D(E)}$ , where H stands for hypothesis and P(E)E for evidence.

Bayesian networks display an extremely interesting property called **explaining away**. Word sense ambiguity in the process of learning SP defines a problem that might be solved by a model that implements an explaining away strategy. Suppose we are learning the selectional preference of *drink*, and the network in Figure 4 is the

<sup>&</sup>lt;sup>1</sup>As a matter of fact, for this HMM there are (infinitely) many parameter values that maximize the likelihood of the training data; i.e., the parameters are not identifiable. The intuitively correct solution is one of them, but so are infinitely many other, intuitively incorrect ones. Thus it is no surprise that the EM algorithm cannot find the intuitively correct solution.

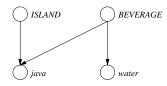


Figure 4: A Bayesian network for word ambiguity.

knowledge base. The verb occurred with java and water. This situation can be represented as a Bayesian network. The variables ISLAND and BEVERAGE represent concepts in a semantic hierarchy. The variables java and water stand for possible instantiations of the concepts. All the variables are Boolean; i.e., they are associated with two states, true or false. Suppose the following CPTs define the priors associated with each node. The unconditional probabilities are P(I = true) = P(B = true) = 0.01 and P(I = false) = P(B = false) = 0.99, and the CPTs for the child nodes are

	$P(X = x   Y_1 = y_1, Y_2 = y_2)$			
	I, B	$I, \neg B$	$\neg I, B$	$\neg I, \neg B$
j = true	0.99	0.99	0.99	0.01
j = false	0.01	0.01	0.01	0.99
w = true	0.99	0.99	0.01	0.01
w = false	0.01	0.01	0.99	0.99

These values mean that the occurrence of either concept is a priori unlikely. If either concept is true the word java is likely to occur. Similarly, if BEVERAGE occurs it is likely to observe also the word water. As the posterior probabilities show, if java occurs, the beliefs in both concepts increase: P(I|j) = P(B|j) = 0.3355. However, water provides evidence for BEVERAGE only. Overall there is more evidence for the hypothesis that the concept being expressed is BEV-ERAGE and not ISLAND. Bayesian networks implement this inference scheme; if we compute the conditional probabilities given that both words occurred, we obtain P(B|j, w) = 0.98 and P(I|j,w) = 0.02. The new evidence caused the "island" hypothesis to be explained away!

#### 3.1 The relevance of priors

Explaining away seems to depend on the specification of the prior probabilities. The priors

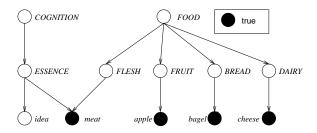


Figure 5: A Bayesian network for the simple example.

define the background knowledge available to the model relative to the conditional probabilities of the events represented by the variables, but also about the joint distributions of several events. In the simple network above, we defined the probability that either concept is selected (i.e., that the corresponding variable is true) to be extremely small. Intuitively, there are many concepts and the probability of observing any particular one is small. This means that the joint probability of the two events is much higher in the case in which only one of them is true (0.0099) than in the case in which they are both true (0.0001). Therefore, via the priors, we introduced a bias according to which the hypothesis that one concept is selected will be favored over two co-occurring ones. This is a general pattern of Bayesian networks; the prior causes simpler explanations to be preferred over more complex ones, and thereby the explaining away effect.

# 4 A Bayesian network approach to learning selectional preference

### 4.1 Structure and parameters of the model

The hierarchy of nouns in Wordnet defines a DAG. Its mapping into a BBN is straightforward. Each word or synset in Wordnet is a **node** in the network. If A is a hyponym of B there is an **arc** in the network from B to A. All the variables are Boolean. A synset node is *true* if the verb selects for that class. A word node is *true* if the word can appear as an argument of the verb. The priors are defined following two intuitive principles. First, it is *unlikely* that a verb a *priori* selects for any particular synset. Second, if a verb does select for a synset, say *FOOD*, then it is *likely* that it also selects for

<sup>&</sup>lt;sup>2</sup>I, B, j and w abbreviate ISLAND, BEVERAGE, java and water, respectively.

its hyponyms, say FRUIT. The same principles apply to words: it is *likely* that a word appears as an argument of the verb if the verb selects for any of its possible senses. On the other hand, if the verb does not select for a synset, it is *unlikely* that the words instantiating the synset occur as its arguments. "Likely" and "unlikely" are given numerical values that sum up to 1. The following table defines the scheme for the CPTs associated with each node in the network;  $p_i(X)$  denotes the *ith* parent of the node X.

	$P(X = x   p_1(X) \lor, \dots, \lor p_n(X) = true)$		
x = true	likely		
x = false	unlikely		
	$P(X = x   p_1(X) \land, \dots, \land p_n(X) = false)$		
x = true	unlikely		
x = false	likely		

For the root nodes, the table reduces to the unconditional probability of the node. we can test the model on the simple example seen earlier.  $W^+$  is the set of words that occurred with the verb. The nodes corresponding to the words in  $W^+$  are set to true and the others left unset. For the previous example  $W^+ = \{meat, apple, bagel, cheese\}$ , and the corresponding nodes are set to true, as depicted in Figure 5. With likely and unlikely respectively equal to 0.99 and 0.01, the posterior probabilities are P(F|m,a,b,c) = 0.9899and P(C|m, a, b, c) = 0.0101. Explaining away works. The posterior probability of COGNI-TION gets as low as its prior, whereas the probability of *FOOD* goes up to almost 1. A Bayesian network approach seems to actually implement the *conservative* strategy we thought to be the correct one for unsupervised learning of selectional restrictions.

### 4.2 Computational issues in building BBNs based on Wordnet

The implementation of a BBN for the whole of Wordnet faces computational complexity problems typical of graphical models. A densely connected BBN presents two kinds of problems. The first is the storage of the CPTs. The size of a CPT grows exponentially with the number of parents of the node.<sup>4</sup> This problem can be

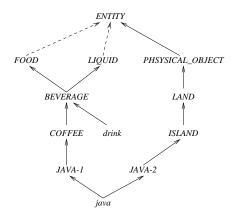


Figure 6: The subnetwork for *drink*.

solved by optimizing the representation of these tables. In our case most of the entries have the same values, and a compact representation for them can be found (much like the one used in the **noisy-OR** model (Pearl, 1988)).

A harder problem is performing inference. The graphical structure of a BBN represents the dependency relations among the random variables of the network. The algorithms used with BBNs usually perform inference by dynamic programming on the triangulated moral graph. A lower bound on the number of computations that are necessary to model the joint distribution over the variables using such algorithms is  $2^{|n|+1}$ , where n is the size of the maximal boundary set according to the visitation schedule.

#### 4.3 Subnetworks and balancing

Because of these problems we could not build a single BBN for Wordnet. Instead we simplified the structure of the model by building a smaller subnetwork for each predicate-argument pair. A subnetwork consists of the union of the sets of ancestors of the words in  $W^+$ . Figure 6 provides an example of the union of these "ancestral subgraphs" of Wordnet for the words java and drink (compare it with Figure 1).

This simplification does not affect the computation of the distributions we are interested in; that is, the *marginals* of the synset nodes. A BBN provides a compact representation for the joint distribution over the set of variables

senses. The size of its CPT is therefore  $2^{26}$ . Storing a table of float numbers for this node alone requires around  $(2^{26})8 = 537$  MBytes of memory.

<sup>&</sup>lt;sup>3</sup>F, C, m, a, b and c respectively stand for FOOD, COGNITION, meat, apple, bagel and cheese

<sup>&</sup>lt;sup>4</sup>Some words in Wordnet have more than 20 senses. For example, *line* in Wordnet is associated with 25

in the network. If  $N = X_1, ..., X_n$  is a Bayesian network with variables  $X_1, ..., X_n$ , its joint distribution P(N) is the product of all the conditional probabilities specified in the network,

$$P(N) = \prod_{j} P(X_j | pa(X_j))$$
 (4)

where pa(X) is the set of parents of X. A BBN generates a factorization of the joint distribution over its variables. Consider a network of three nodes A, B, C with arcs from A to B and C. Its joint distribution can be characterized as P(A, B, C) = P(A)P(B|A)P(C|A). If there is no evidence for C the joint distribution is

$$P(A, B, C) = P(A)P(B|A) \sum_{C} P(C|A)$$
$$= P(A)P(B|A)$$
$$= P(A, B)$$
(5)

The node C gets marginalized out. Marginalizing over a childless node is equivalent to removing it with its connections from the network. Therefore the subnetworks are equivalent to the whole network; i.e., they have the same joint distribution.

Our model computes the value of P(c|p,r), but we did not compute the prior P(c) for all nouns in the corpus. We assumed this to be a constant, equal to the *unlikely* value, for all classes. In a BBN the values of the marginals increase with their distance from the root nodes. To avoid undesired bias (see table of results) we defined a balancing formula that adjusted the conditional probabilities of the CPTs in such a way that we got all the marginals to have approximately the same value.<sup>5</sup>

### 5 Experiments and results<sup>6</sup>

#### 5.1 Learning of selectional preferences

When trained on predicate-argument pairs extracted from a large corpus, the San Jose Mercury Corpus, the model gave very good results. The corpus contains about 1.3 million verbobject tokens. The obtained rankings of classes according to their posterior marginal probabilities were good. Table 1 shows the top and the

Ranking	Synset	P(c p,r)
1	VEHICLE	0.9995
2	VESSEL	0.9893
3	AIRCRAFT	0.9937
4	AIRPLANE	0.9500
5	SHIP	0.9114
255	CONCEPT	0.1002
256	LAW	0.1001
257	PHILOSOPHY	0.1000
258	JURISPRUDENCE	0.1000

Table 1: Results for (maneuver, object).

bottom of the list of synsets for the verb maneuver. The model learned that maneuver "selects" for members of the class VEHICLE and of other plausible classes, hyponyms of VEHICLE. It also learned that the verb does not select for direct objects that are members of classes, like CONCEPT or PHILOSOPHY.

### 5.2 Word sense disambiguation test

A direct evaluation measure for unsupervised learning of SP models does not exist. These models are instead evaluated on a word-sense disambiguation test (WSD). The idea is that systems that learn SP produce word sense disambiguation as a side-effect. Java might be interpreted as the island or the beverage, but in a context like "the tourists flew to Java" the former seems more correct, because fly could select for geographic locations but not for beverages. A system trained on a predicate p should be able to disambiguate arguments of p if it has learned its selectional restrictions.

We tested our model using the test and training data developed by Resnik (see Resnik, 1997). The same test was used in (Abney and Light, 1999). The training data consists of predicate-object counts extracted from 4/5 of the Brown corpus (about 1M words). The test set consists of predicate-object pairs from the remaining 1/5 of the corpus, which has been manually sense-annotated by Wordnet researchers. The results are shown in Table 2. The baseline algorithm chooses at random one of the multiple senses of an ambiguous word. The "first sense" method always chooses the most frequent sense (such a system should be trained on sense-tagged data). Our model per-

 $<sup>^5</sup>$ More details can be found in an extended version of the paper: www.cog.brown.edu/ $\sim$ massi/.

<sup>&</sup>lt;sup>6</sup>For these experiments we used values for the *likely* and *unlikely* parameters of 0.9 and 0.1, respectively.

Method	Result
Baseline	28.5%
Abney and Light (HMM smoothed)	35.6%
Abney and Light (HMM balanced)	42.3%
Resnik	44.3%
BBN (without balancing)	45.6%
BBN (with balancing)	51.4%
First Sense	82.5%

Table 2: Results

formed better than the state of the art models for unsupervised learning of SP. It seems to define a better estimator for P(c|p, r).

It is remarkable that the model achieved this result making only a limited use of distributional information. A noun is in  $W^+$  if it occurred at least once in the training set, but the system does not know if it occurred once or several times; either it occurred or it didn't. The model did not suffer too much from this limitation during this task. This is probably due to the sparseness of the training data for the test. For each verb the average number of object types is 3.3, for each of them the average number of tokens is 1.3; i.e., most of the words in the training data only occurred once. For this training set we also tested a version of the model that built a word node for each observed object token and therefore integrated the distributional information. On the WSD test it performed exactly the same as the simpler version. When trained on the San Jose Mercury Corpus the model performed worse on the WSD test (35.8%). This is not too surprising considering the differences between the SJM and the Brown corpora: the former, a recent newswire corpus; the latter, an older, balanced corpus. Another important factor is the different relevance of distributional information. The training data from the SJM Corpus are much richer and noisier than the Brown data. Here the frequency information is probably crucial; however, in this case we could not implement the simple scheme above.

#### 5.3 Conclusion

Explaining away implements a cognitively attractive and successful strategy. A straightforward improvement would be for the model to

make full use of the distributional information present in the training data; we only partially achieved this. Bayesian networks are usually confronted with a single presentation of evidence. Their extension to multiple evidence is not trivial. We believe the model can be extended in this direction. Possibly there are several ways to do so (multinomial sampling, dedicated implementations, etc.). However, we believe that the most relevant finding of this research might be that "explaining away" is not only a property of Bayesian networks but of Bayesian inference in general and that it might be implementable in other kinds of graphical models. We observed that the property seems to depend on the specification of the prior probabilities. We found that the HMM model of (Abney and Light, 1999) was unidentifiable; that is, there are several solutions for the parameters of the model, including the desired one. Our intuition is that it should be possible to implement "explaining away" in a HMM with priors, so that it would prefer only one or a few solutions over many. This model would have also the advantage of being computationally simpler.

### References

- S. Abney and M. Light. 1999. Hiding a semantic hierarchy in a Markov model. In *Proceedings of the Workshop on Unsupervised Learning in Natural Language Processing*, ACL.
- N. Chomsky. 1965. Aspects of the Theory of Syntax. MIT Press, Cambridge, MA.
- P. N. Johnson-Laird. 1983. Mental Models: Towards a Cognitive Science of Language, Inference, and Consciousness. Harvard University Press.
- J. J. Katz and J. A. Fodor. 1964. The structure of a semantic theory. In J. J. Katz and J. A. Fodor, editors, *The Structure of Language*. Prentice-Hall.
- G. Miller. 1990. Wordnet: An on-line lexical database. *International Journal of Lexicography*, 3(4).
- J. Pearl. 1988. Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference. Morgan Kaufman.
- P. Resnik. 1997. Selectional preference and sense disambiguation. In *Proceedings of the ANLP-97 Workshop: Tagging Text with Lexical Semantics: Why, What, and How?*