

## Chapter 1

# AN OVERVIEW OF SOCIAL TAGGING AND APPLICATIONS

Manish Gupta

*University of Illinois at Urbana Champaign*

gupta58@illinois.edu

Rui Li

*University of Illinois at Urbana Champaign*

ruili1@illinois.edu

Zhijun Yin

*University of Illinois at Urbana Champaign*

zyin3@illinois.edu

Jiawei Han

*University of Illinois at Urbana Champaign*

hanj@cs.uiuc.edu

**Abstract** Social tagging on online portals has become a trend now. It has emerged as one of the best ways of associating metadata with web objects. With the increase in the kinds of web objects becoming available, collaborative tagging of such objects is also developing along new dimensions. This popularity has led to a vast literature on social tagging. In this survey paper, we would like to summarize different techniques employed to study various aspects of tagging. Broadly, we would discuss about properties of tag streams, tagging models, tag semantics, generating recommendations using tags, visualizations of tags, applications of tags, integration of different tagging systems and problems associated with tagging usage. We would discuss topics like why people tag, what influences the choice of tags, how to model the tagging process, kinds of tags, different power laws observed in tagging domain, how tags are cre-

ated and how to choose the right tags for recommendation. Metadata generated in the form of tags can be efficiently used to improve web search, for web object classification, for generating ontologies, for enhanced browsing etc. We would discuss these applications and conclude with thoughts on future work in the area.

**Keywords:** Social tagging, bookmarking, tagging, social indexing, social classification, collaborative tagging, folksonomy, folk classification, ethnoclassification, distributed classification, folk taxonomy

## 1. Introduction

Social tagging became popular with the launch of sites like Delicious and Flickr. Since then, different social systems have been built that support tagging of a variety of resources. Given a particular web object or resource, tagging is a process where a user assigns a tag to an object. On Delicious, a user can assign tags to a particular bookmarked URL. On Flickr, users can tag photos uploaded by them or by others. Whereas Delicious allows each user to have her personal set of tags per URL, Flickr has a single set of tags for any photo. On blogging sites like Blogger, Wordpress, Livejournal, blog authors can add tags to their posts. On micro-blogging sites like Twitter, hash tags are used within the tweet text itself. On social networking sites like Facebook, Orkut, etc., users often annotate parts of the photos. Users can also provide tagging information in other forms like marking something as “Like” on Facebook. Upcoming event sites can allow users to comment on and tag events. Recently, tripletags (tags of the format namespace:key=value (e.g., geo:lat=53.1234) are becoming popular. Such a syntax can improve the usability of tags to a large extent. Using rel-tags<sup>1</sup>, a page can indicate that the destination of that hyperlink is an author-designated tag for the current page. Rel-tags have been used by various implementation sites to tag blogs, music, links, news articles, events, listings, etc. Citation websites have tags attached to publication entries. Cataloging sites like LibraryThing and Shelfari allow users to tag books. Social news sites like Digg, SlashDot allow users to attach tags to news stories. Yelp, CitySearch and other such business/product reviews sites allow users to attach their reviews and other users to select tags to rate reviews too. Multimedia objects like podcasts, live casts, videos and music can also be tagged on sites like Youtube, imeem, Metacafe, etc. On Yahoo! Answers, you can tag an answer as positive or negative depending on how helpful it was. Tags are often used to collect such binary or multi-valued ratings or categorical decisions from users. Tags are omni-present on the web. But what led to the emergence of tagging based systems? As

we shall see in this section, tags are a better way of generating metadata and prevent problems associated with fixed taxonomies in social systems.

## **Problems with Metadata Generation and Fixed Taxonomies**

Different web portals focus on sharing of different types of objects like images, news articles, bookmarks, etc. Often to enrich the context related to these objects and thereby support more applications like search, metadata needs to be associated with these objects. However, manual metadata creation is costly in terms of time and effort [35]. Also, vocabulary of this metadata may be completely different from that of system designer or content producers or taxonomy creators or eventual users. Besides associating metadata to the objects, building a taxonomy for these social sharing systems may be useful in classifying and organizing the objects. But fixed static taxonomies are rigid, conservative, and centralized [41]. Items do not always fit exactly inside one and only one category. Hierarchical classifications are influenced by the cataloguer's view of the world and, as a consequence, are affected by subjectivity and cultural bias. Rigid hierarchical classification schemes cannot easily keep up with an increasing and evolving corpus of items. Social systems need to hire expert cataloguers who can use same thinking and vocabulary as users and can build taxonomies which can be stable over time. Once such a hierarchy is created, the object creators can be asked to assign a fixed category to the object, in the hierarchy. This can induce "post activation analysis paralysis" [21] into the user. By their very nature, hierarchies tend to establish only one consistent authoritative structured vision. This implies a loss of precision, erases difference of expression, and does not take into account the variety of user needs and views.

## **Folksonomies as a Solution**

Folksonomies and social tagging help in preventing these problems and hence provide a simpler, cheaper and a more natural way of organizing web objects. A folksonomy (folk (people) + taxis (classification) + nomos (management)) is a user-generated classification, emerging through bottom-up consensus. The term was coined by Thomas Vander Wal in the AIFIA mailing list to mean the wide-spreading practice of collaborative categorization using freely chosen keywords by a group of people cooperating spontaneously. A folksonomy can be defined as a collection of a set of users, set of tags, set of resources or objects, and a ternary relation between users, tags and resources with

a time dimension [12]. Unlike formal taxonomies, folksonomies have no explicitly defined relationship between terms. All terms belong to a flat namespace, i.e., there is no hierarchy. Since users themselves tag the objects, folksonomies directly reflect the vocabulary of users [41]. Hence, a folksonomy is a simple, emergent and iterative system. It helps create the most popular way of organizing objects referred to as desire lines<sup>2</sup>. Apart from this, tagging provides no barriers to entry or cooperation and hence involves low cognitive cost. Tagging helps users get immediate feedback. This feedback loop leads to a form of asymmetric communication between users through metadata. The users of a system negotiate the meaning of the terms in the folksonomy, whether purposefully or not, through their individual choices of tags to describe objects for themselves. Further, folksonomies are inclusive, i.e., they include terms related to popular topics and also terms related to long tail topics. With appropriate browsing support, interlinking related tag sets is wonderful for finding things unexpectedly in a general area.

In summary, folksonomies are a trade-off between traditional structured centralized classification and no classification or metadata at all. Their advantage over traditional top-down classification is their capability of matching users' real needs and language, not their precision. Building, maintaining, and enforcing a sound controlled vocabulary is often too expensive in terms of development time and presents a steep learning curve to the user to learn the classification scheme. In other words, folksonomies are better than nothing, when traditional classification is not viable.

## Outline

In this survey paper, we present a systematic detailed study of tagging literature. We first list different user motivations and different ways of tagging web objects in Section 2. There have been a lot of generative models proposed to understand the process of tagging. We present a summary of such models in Section 3. Section 4 describes various parameters for tagging system design. In Section 5, we present a summarization of work done on analysis of tagging distributions, identification of tag semantics, expressive power of tags versus keywords. Appropriate rendering of tags can provide useful information to users. Different visualization schemes like tag clouds have been explored to support browsing on web portals. We present some works related to such visualization studies in Section 6. When a user wishes to attach tags to an object, the system can recommend some tags to the user. A user can select one of those tags or come up with a new one. In Sec-

tion 7, we discuss different ways of generating tag recommendations. In Section 8, we describe different applications for which tags can be used. An integration of such folksonomies can help in solving the problem of sparsity of tags associated with Web objects. We describe some works about integration of different folksonomies in Section 9. Usage of tags involves a lot of problems like sparsity, ambiguities and canonicalization. We list these problems in Section 10. Finally, we conclude with thoughts on future work in Section 11.

## 2. Tags: Why and What?

Since 2005, there have been works describing why people tag and what the tags mean. We briefly summarize such works [1, 34, 16, 61, 7, 21, 48, 35, 27, 17] below. We provide a detailed classification of user tagging motivations and also list different kinds of tags in this section.

### Different User Tagging Motivations

- **Future Retrieval:** Users can tag objects aiming at ease of future retrieval of the objects by themselves or by others. Tags may also be used to incite an activity or act as reminders to oneself or others (e.g., the “to read” tag). These descriptive tags are exceptionally helpful in providing metadata about objects that have no other tags associated.
- **Contribution and Sharing:** Tags can be used to describe the resource and also to add the resource to conceptual clusters or refined categories for the value of either known or unknown audience.
- **Attract Attention:** Popular tags can be exploited to get people to look at one’s own resources.
- **Play and Competition:** Tags can be based on an internal or external set of rules. In some cases, the system devises the rules such as the ESP Game’s incentive to tag what others might also tag. In others, groups develop their own rules to engage in the system such as when groups seek out all items with a particular feature and tag their existence.
- **Self Presentation (Self Referential Tags):** Tags can be used to write a user’s own identity into the system as a way of leaving their mark on a particular resource. E.g., the “seen live” tag in Last.FM marks an individual’s identity or personal relation to the resource. Another example are tags beginning with “my” like “mystuff”.

- **Opinion Expression:** Tags can convey value judgments that users wish to share with others (e.g., the “elitist” tag in Yahoo!’s Podcast system is utilized by some users to convey an opinion). Sometimes people tag simply to gain reputation in the community.
- **Task Organization:** Tags can also be used for task organization e.g., “toread”, “jobsearch”, “gtd” (got to do), “todo”.
- **Social Signalling:** Tags can be used to communicate contextual information about the object to others.
- **Money:** Some sites like Squidoo and Amazon Mechanical Turk pay users for creating tags.
- **Technological Ease:** Some people tag because the current technology makes it easy to upload resources with tags to the web. E.g. drag-and-drop approach for attaching labels to identify people in photos. The latest photo browser commercial packages, such as Adobe Photoshop Album, adopted similar methods to support easy labeling of photos. With ‘Phonetags’<sup>3</sup>, a listener hears a song on the radio, uses her cell phone to text back to a website with tags and star ratings. Later, returning to the website, the user can type in her phone number and see the songs she had bookmarked.

## Kinds of Tags

- **Content-Based Tags:** They can be used to identify the actual content of the resource. E.g., Autos, Honda Odyssey, batman, open source, Lucene.
- **Context-Based Tags:** Context-based tags provide the context of an object in which the object was created or saved, e.g., tags describing locations and time such as San Francisco, Golden Gate Bridge, and 2005-10-19.
- **Attribute Tags:** Tags that are inherent attributes of an object but may not be able to be derived from the content directly, e.g., author of a piece of content such as Jeremy’s Blog and Clay Shirky. Such tags can be used to identify what or who the resource is about. Tags can also be used to identify qualities or characteristics of the resource (e.g., scary, funny, stupid, inspirational).
- **Ownership Tags:** Such tags identify who owns the resource.

- **Subjective Tags:** Tags that express user's opinion and emotion, e.g., funny or cool. They can be used to help evaluate an object recommendation (item qualities). They are basically put with a motivation of self-expression.
- **Organizational Tags:** Tags that identify personal stuff, e.g., my-paper or mywork, and tags that serve as a reminder of certain tasks such as to-read or to-review. This type of tags is usually not useful for global tag aggregation with other users' tags. These tags are intrinsically time-sensitive. They suggest an active engagement with the text, in which the user is linking the perceived subject matter with a specific task or a specific set of interests.
- **Purpose Tags:** These tags denote non-content specific functions that relate to an information seeking task of users (e.g., learn about LaTeX, get recommendations for music, translate text).
- **Factual Tags:** They identify facts about an object such as people, places, or concepts. These are the tags that most people would agree to apply to a given object. Factual tags help to describe objects and also help to find related objects. Content-based, context-based and objective, attribute tags can be considered as factual tags. Factual tags are generally useful for learning and finding tasks.
- **Personal Tags:** Such tags have an intended audience of the tag applier themselves. They are most often used to organize a user's objects (item ownership, self-reference, task organization).
- **Self-referential tags:** They are tags to resources that refer to themselves. e.g., Flickr's "sometait hurts"<sup>4</sup> - for "so meta it hurts" is a collection of images regarding Flickr, and people using Flickr. The earliest image is of someone discussing social software, and then subsequent users have posted screenshots of that picture within Flickr, and other similarly self-referential images.
- **Tag Bundles:** This is the tagging of tags that results in the creation of hierarchical folksonomies. Many taggers on Delicious have chosen to tag URLs with other URLs, such as the base web address for the server (e.g., a C# programming tutorial might be tagged with <http://www.microsoft.com>).

## Categorizers Versus Describers

Taggers can be divided into two main types [28]: categorizers and describers. Categorizer users are the ones who apply tags such that

the objects are easier to find later for personal use. They have their own vocabulary. Sets in Delicious is a perfect example of metadata by categorizers. On the other hand, describer users tag objects such that they are easier to be searched by others. Often tags to a single object would contain many synonyms. Vocabulary of a describer is much larger compared to an average categorizer. But a categorizer has her own limited personal vocabulary and subjective tags. ESP game is a perfect example of metadata creation by describers. Categorizers and describers can be identified using these intuitions:

- The more the number of tags that were only used once by a user, the higher the probability that the user is a describer.
- The faster the tagging vocabulary increases, the more likely it is that the person is a describer.
- A categorizer tends to achieve tag entropy that is as low as possible because he tries to “encode” her resources in a good and balanced way.

These intuitions can be formalized as metrics like tag ratio (ratio between tags and resources), orphaned tags (proportion of tags which annotate only a small amount of resources) and tag entropy (reflects the effectiveness of the encoding process of tagging).

## Linguistic Classification of Tags

Based on linguistics, tags can be classified as follows [56].

- Functional: Tags that describe the function of an object. (e.g., weapon)
- Functional collocation: These are defined by function but in addition, they have to be collected in a place (and/or time). (e.g., furniture, tableware)
- Origin collocation: Tags that describe why things are together? (e.g., garbage, contents, dishes (as in “dirty dishes” after a meal)).
- Function and origin: Tags that describe why an object is present, what is the purpose, or where did it come from. (e.g., “Michelangelo” and “medieval” on an image of a painting by Michelangelo)
- Taxonomic: They are words that can help in classifying the object into an appropriate category. (e.g., “Animalia” or “Chordata” tag to an image of a heron)



- Adjective: They describe the object that denotes the resource. (e.g., “red”, “great”, “beautiful”, “funny”)
- Verb: These are action words. (e.g., “explore”, “todo”, “jumping”)
- Proper name: Most of the tags are of this category. (e.g., “New Zealand”, “Manhattan bridge”)

## Game Based Tagging

In the ESP game, the players cannot see each other’s guesses. The aim is to enter the same word as your partner in the shortest possible time. Peekaboom takes the ESP Game to the next level. Unlike the ESP Game, it is asymmetrical. To start, one user is shown an image and the other sees an empty blank space. The first user is given a word related to the image, and the aim is to communicate that word to the other player by revealing portions of the image. So if the word is “eye” and the image is a face, you reveal the eye to your partner. But the real aim here is to build a better image search engine: one that could identify individual items within an image. PhotoPlay[13] is a computer game designed to be played by three to four players around a horizontal display. The goal for each player is to build words related to any of the four photos on the display by selecting from a 7x7 grid of letter tiles. All these games, help in tagging the resources.

### Problems with Game-Based Tagging

Game-based tagging mechanisms may not provide high quality tags [30]. Maximizing your scores in the game means sacrificing a lot of valuable semantics. People tend to write very general properties of an image rather than telling about the specifics or details of the image. E.g., colors are great for matching, but often are not the most critical or valuable aspects of the image. The labels chosen by people trying to maximize their matches with an anonymous partner are not necessarily the most “robust and descriptive” labels. They are the easiest labels, the most superficial labels, the labels that maximize the speed of a match rather than the quality of the descriptor. In addition, they are words that are devoid of context or depth of knowledge. Tagging for your own retrieval is different than tagging for retrieval by people you know and even more different than tagging for retrieval in a completely uncontextualized environment.

## 3. Tag Generation Models

In order to describe, understand and analyze tags and tagging systems, various tag generation models have been proposed. These models

study various factors that influence the generation of a tag, such as the previous tags suggested by others, users' background knowledge, content of the resources and the community influences. In this section, we present different models which have been proposed in the literature, and discuss advantages and disadvantages of these models.

## Polya Urn Generation Model

Intuitively, the first factor that influences the choice of tags is the previous tag assignments. The amount of effort required to tag items may affect an individual's decision to use tags. Using suggested tags rather than one's own requires less effort. Pirolli and Card's theory of information foraging [40] suggests greater adoption of suggested tags because people adapt their behavior to optimize the information/effort ratio. Users cost-tune their archives by spending the least amount of effort needed to build up enough structure to support fast retrieval of their most useful resources. Based on this intuition, various models based on the stochastic Polya urn process have been proposed.

**Basic Polya Urn Model.** Golder and Huberman [16] propose a model based on a variation of the stochastic Polya urn model where the urn initially contains two balls, one red and one black. In each step of the simulation, a ball is selected from the urn and then it is put back together with a second ball of the same color. After a large number of draws the fraction of the balls with the same color stabilizes but the fractions converge to random limits in each run of the simulation.

This model successfully captures that previously assigned tags are more likely to be selected again. However, this basic model fails to capture that new tags will also be added into the system. So several extensions of this model have been proposed later.

**Yule-Simon Model.** Yule-Simon model [50] assumes that at each simulation step a new tag is invented and added to the tag stream with a low probability of  $p$ . This leads to a linear growth of the distinct tags with respect to time and not to the typical continuous, but declining growth. Yule-Simon model can be described as follows. At each discrete time step one appends a word to the text: with probability  $p$  the appended word is a new word, that has never occurred before, while with probability  $1 - p$  the word is copied from the existing text, choosing it with a probability proportional to its current frequency of occurrence. This simple process produces frequency-rank distributions with a power law tail whose exponent is given by  $a = 1 - p$ .

Cattuto et al. [10] study the temporal evolution of the global vocabulary size, i.e., the number of distinct tags in the entire system, as well as the evolution of local vocabularies, that is, the growth of the number of distinct tags used in the context of a given resource or user. They find that the number  $N$  of distinct tags present in the system is  $N(T) \propto T^\gamma$ , with  $\gamma < 1$ . The rate at which new tags appear at time  $T$  scales as  $T^{\gamma-1}$ , i.e., new tags appear less and less frequently, with the invention rate of new tags monotonically decreasing very slowly towards zero. This sub-linear growth is generally referred to as Heaps' law.

**Yule-Simon Model with Long Term Memory.** Cattuto et al. [11] propose a further variation of the Simon model. It takes the order of the tags in the stream into account. Like the previous models, it simulates the imitation of previous tag assignments but instead of imitating all previous tag assignments with the same probability it introduces a kind of long-term memory. Their model can be stated as follows: the process by which users of a collaborative tagging system associate tags to resources can be regarded as the construction of a “text”, built one step at a time by adding “words” (i.e., tags) to a text initially comprised of  $n_0$  words. This process is meant to model the behavior of an effective average user in the context identified by a specific tag. At a time step  $t$ , a new word may be invented with probability  $p$  and appended to the text, while with probability  $1 - p$  one word is copied from the existing text, going back in time by  $x$  steps with a probability that decays as a power law,  $Q_t(x) = a(t)/(x + \tau)$ .  $a(t)$  is a normalization factor and  $\tau$  is the characteristic time-scale over which the recently added words have comparable probabilities. Note that  $Q_t(x)$  returns a power law distribution of the probabilities. This Yule-Simon model with long term memory successfully reproduces the characteristic slope of the frequency-rank distribution of co-occurrence tag streams but it fails to explain the distribution in resource tag streams as well as the decaying growth of the set of distinct tags because it leads to a linear growth.

**Information Value Based Model.** Halpin et al. [18] present a model which does not only simulate the imitation of previous tag assignments but it also selects tags based on their information value. The information value of a tag is 1 if it can be used for only selecting appropriate resources. A tag has an information value of 0 if it either leads to the selection of no or all resources in a tagging system. They empirically estimate the information value of a tag by retrieving the number of webpages that are returned by a search in Delicious with the tag. Besides of the selection based on the information value, the

model also simulates the imitation of previous tag assignments using the Polya urn model. They model tag selection as a linear combination of information value and preferential attachment models. Probability of a tag  $x$  being reinforced or added can be expressed as  $P(x) = \lambda \times P(I(x)) + (1 - \lambda) \times P(a) \times P(o) \times P(\frac{R(x)}{\sum R(i)})$  where  $\lambda$  is used to weigh the factors.  $P(a)$  is the probability of a user committing a tagging action at any time  $t$ .  $P(n)$  determines the number  $n$  of tags a user is likely to add at once based on the distribution of the number of tags a given user employs in a single tagging action. An old tag is reinforced with constant probability  $P(o)$ . If the old tag is added, it is added with a probability  $\frac{R(x)}{\sum R(i)}$  where  $R(x)$  is the number of times that particular previous tag  $x$  has been chosen in the past and  $\sum R(i)$  is the sum of all previous tags. Overall, the proposed model leads to a plain power law distribution of the tag frequencies and to a linear growth of the set of distinct tags. It thus only partially reproduces the frequency-rank distributions in co-occurrence and resource tag streams and it is not successful in reproducing the decaying tag growth.


**Fine-tuning by Adding More Parameters.** Klaas et al. [12] present the following model: The simulation of a tag stream always starts with an empty stream. Then, in each step of the simulation, with probability  $I$  (0.6-0.9) one of the previous tag assignments is imitated. With probability  $BK$ , the user selects an appropriate tag from her background knowledge about the resource. It corresponds to selecting an appropriate natural language word from the active vocabulary of the user. Each word  $t$  has been assigned a certain probability with which it gets selected, which corresponds to the probability with which  $t$  occurs in the Web corpus. The parameter  $n$  represents the number of popular tags a user has access to. In case of simulating resource streams,  $n$  will correspond to the number of popular tags shown. (e.g.,  $n = 7$  for Delicious). In case of co-occurrence streams  $n$  will be larger because the union of the popular tags of all resources that are aggregated in the co-occurrence stream will be depicted over time. Furthermore, the parameter  $h$  can be used for restricting the number of previous tag assignments which are used for determining the  $n$  most popular tags. The probability of selecting the concrete tag  $t$  from the  $n$  tags is then proportional to how often  $t$  was used during the last  $h$  tag assignments. Using these parameters, the authors describe a model that reproduces frequency rank for both tag co-occurrence and resource streams and also simulates the tag vocabulary growth well.

## Language Model

The content of resource would affect generation of tags. Hence, tagging process can also be simulated using a language model like the latent Dirichlet allocation model [6]. Tagging is a real-world experiment in the evolution of a simple language [7]. Zhou et al. [63] propose a probabilistic generative model for generation of document content as well as associated tags. This helps in simultaneous topical analysis of terms, documents and users. Their user content annotation model can be explained as follows. For document content, each observed term  $\omega$  in document  $d$  is generated from the source  $x$  (each document  $d$  maps one-to-one to a source  $x$ ). Then from the conditional probability distribution on  $x$ , a topic  $z$  is drawn. Given the topic  $z$ ,  $\omega$  is finally generated from the conditional probability distribution on the topic  $z$ . For document tags, similarly, each observed tag word  $\omega$  for document  $d$  is generated by user  $x$ . Specific to this user, there is a conditional probability distribution of topics, from which a topic  $z$  is then chosen. This hidden variable of topic again finally generates  $\omega$  in the tag.

Figure 1.1 shows the user content annotation model using the plate notation.

Figure 1.1. User content annotation model




Unknown distributions  $\theta$  and  $\phi$  can be learnt using EM. But EM can lead to local maxima or may be expensive. Hence, they use Gibbs sampling. Rather than the parameters, the posteriors are evaluated. Posteriors include:  $P(d, z | w)$  and  $P(x, z | w)$  alongwith  $P(d | z)$ ,  $P(x | z)$ ,  $P(z | w)$ . Number of topics are determined using the perplexity measure.

## Other Influence Factors

Besides above models, researchers [48, 34] have also observed that there are other factors which are likely to influence how people apply the tags.

Sen et al. [48] mention three factors that influence people’s personal tendency (their preferences and beliefs) to apply tags: (1) their past tagging behaviors, (2) community influence of the tagging behavior of other members, and (3) tag selection algorithm that chooses which tags to display. New users have an initial personal tendency based on their experiences with other tagging systems, their comfort with technology, their interests and knowledge. Personal tendency evolves as people interact with the tagging system. Figure 1.2 shows how these factors affect the tagging behaviour.

Figure 1.2. Factors affecting tagging behavior



Experiments with Movielens dataset reveal the following. Once a user has applied three or more tags, the average cosine similarity for the  $n$ th tag application is more than 0.83. Moreover, similarity of a tag application to the user’s past tags continues to rise as users add more tags. Besides reusing tag classes, users also reuse individual tags from their vocabulary. Community influence on a user’s first tag is stronger for users who have seen more tags. The tag selection algorithm influences the distribution of tag classes (subjective, factual, and personal).

The community influence on tag selection in Flickr has been studied by Marlow et al. [34]. One feature of the contact network is a user’s ability to easily follow the photos being uploaded by their friends. This provides a continuous awareness of the photographic activity of their Flickr contacts, and by transitivity, a constant exposure to tagging practices. Do these relationships affect the formation of tag vocabularies, or are individuals guided by other stimuli? They find that the random users are much more likely to have a smaller overlap in common tags, while contacts are more distributed, and have a higher overall mean. This

result shows a relationship between social affiliation and tag vocabulary formation and use even though the photos may be of completely different subject matter. This commonality could arise from similar descriptive tags (e.g., bright, contrast, black and white, or other photo features), similar content (photos taken on the same vacation), or similar subjects (co-occurring friends and family), each suggesting different modes of diffusion.

Apart from the different aspects mentioned above, user tagging behaviors can be largely dictated by the forms of contribution allowed and the personal and social motivations for adding input to the system [34].

#### 4. Tagging System Design

What are the different parameters that should be considered when designing a social tagging system? In this section, we present the some design parameters [48, 34].

- **Tag Sharing:** What are the different privacy levels supported by the system for sharing? (public, private, groups).
- **Tag Selection/Tagging Support:** This includes tag recommendation algorithm used, number of recommendations shown, meta information shown. Different categories are a. blind tagging, where a tagging user cannot view tags assigned to the same resource by other users while tagging (e.g., Delicious) b. viewable tagging, where the user can see the tags already associated with a resource (e.g., Yahoo! Podcasts) c. suggestive tagging, where the system suggests possible tags to the user (e.g., Yahoo! MyWeb2.0). The implication of suggested tagging may be a quicker convergence to a folksonomy. In other words, a suggestive system may help consolidate the tag usage for a resource, or in the system, much faster than a blind tagging system would. As for viewable tagging, implications may be overweighting certain tags that were associated with the resource first, even if they would not have arisen otherwise.
- **Item Ownership/Tagging Rights:** Tagging system could allow only owner to tag (Technorati) or anyone (Amazon) or may support different levels of permissions for people to tag (Flickr). Figure 1.3 from [19] shows some examples. The system can specify who may remove a tag, whether no one (e.g., Yahoo! Podcasts), anyone (e.g., Odeo), the tag creator (e.g., Last.fm) or the resource owner (e.g., Flickr). Tags that are assigned to a photo may be radically divergent depending on whether the tagging is performed by

the photographers, their friends, or strangers looking at their photos.

Figure 1.3. Tagging Rights

Tag User	Others	Technorati HTML Meta tags (Wikipedia)
	Self	Flickr CiteULike Connotea Delicious Frassle, Furl Simpy, Spurl unalog
		Self      Others Content Creator

- **Tag Scope/Tag Aggregation:** Tag scope could be broad or narrow. System with broad tag scope follows a bag model and may allow for a multiplicity of tags for the same resource (<user, item, tag> maintained e.g. Delicious). System with narrow tag scope follows a set model approach and asks all users to collectively tag an individual resource, thus denying any repetition (<item, tag> maintained e.g., Technorati, Flickr). In the case that a bag model is being used, the system has the ability to use aggregate statistics for a given resource to present users with the collective opinions of the taggers; thus accurately finding relationships between users, tags, and resources.
- **Tag Format:** Some tagging systems may support multi-word tags, tag normalization and other metadata like notes in Delicious.
- **Type of Object:** An object to be tagged can be web page, bibliographic material, video, image, user etc. Tags given to textual resources may differ from tags for resources/objects with no such textual representation, like images or audio.
- **Source of Material:** Resources to be tagged can be supplied by the participants (e.g., YouTube, Flickr, Technorati, Upcoming), by the system (e.g., ESP game, Last.fm, Yahoo! Podcasts), or, alternatively, a system can be open for tagging of any web resource (e.g., Delicious, Yahoo! MyWeb2.0).
- **Resource Connectivity:** Resources in the system can be linked to each other independent of the user tags. Connectivity can be roughly categorized as linked, grouped or none. E.g., web pages are connected by directed links; Flickr photos can be assigned to



groups; and events in Upcoming have connections based on the time, city and venue associated with the event. Implications for resultant tags and usefulness may include convergence on similar tags for connected resources.

- **Social Connectivity:** Some systems allow users within the system to be linked together. Like resource connectivity, social connectivity can also be defined as linked, grouped, or none. Many other dimensions are present in social networks, e.g., whether links are typed (like in Flickr's contacts/friends model) and whether links are directed, where a connection between users is not necessarily symmetric (in Flickr, for example, none of the link types is symmetric). Implications of social connectivity include the adoption of localized folksonomies based on social structure in the system.
- **User Incentives:** Users may tag just for socialization, money, for fun while playing etc. as mentioned in section 2.

## 5. Tag analysis

To better understand social tagging data, a lot of research has been done in analyzing a variety of properties of social tagging data, such as how tags are distributed and their hidden semantics. In the following subsections we present some major analysis and results.

### Tagging Distributions

Researchers began their study with analyzing tags distribution in tagging systems. They found that most of them are power law distributions, which is one of prominent features of a social tagging system.

**Tagging System Vocabulary.** As has been noted by different studies on a variety of datasets, total number of distinct tags in the system with respect to time follows a power law. However, recent studies have shown that this vocabulary growth is somewhat sublinear.

**Resource's Tag Growth.** For a single resource over time, vocabulary growth for tags also follows power law with exponent  $2/3$  [10]. Frequency-rank distribution of tag streams also follows a power law [12]. For some webpages tagged on Delicious, tag frequency (sorted) versus tag rank for a web page is a decreasing graph with a sudden drop between rank 7 and 10 [12]. This may be due to an artifact of the user interface of Delicious. The graph of probability distribution of number

of tags contained in a posting versus the number of tags displays an initial exponential decay with typical number of tags as 3-4 and then becomes a power law tail with exponent as high as -3.5 [10].

Researchers have also observed convergence of the tag distributions. In [18], Halpin et al. observe that majority of sites reach their peak popularity, the highest frequency of tagging in a given time period, within 10 days of being saved on Delicious (67% in the data set of Golder and Huberman [16]) though some sites are rediscovered by users (about 17% in their data set), suggesting stability in most sites but some degree of burstiness in the dynamics that could lead to a cyclical relationship to stability characteristic of chaotic systems. They also plot KL divergence between the tag frequency distributions for a resource versus the time. The curve drops very steeply. For almost all resources the curve reaches zero at about the same time. In the beginning few weeks, curve is quite steep and slowly becomes gentle as time progresses. Golder and Huberman also find that the proportion of frequencies of tags within a given site stabilizes over time.

Cattuto et al. [10] have shown the variation of the probability distribution of the vocabulary growth exponent  $\gamma$  for resources, as a function of their rank. The curve for the 1000 top-ranked (most bookmarked) resources closely fits a Gaussian curve at  $\gamma \approx 0.71$ . This indicates that highly bookmarked resources share a characteristic law of growth. On computing the distribution  $P(\gamma)$  for less and less popular resources, the peak shifts towards higher values of  $\gamma$  and the growth behavior becomes more and more linear.

Wetzker et al. [57] also show that most popular URLs disappear after peaking. They also point out that some of the tags can peak periodically, e.g., Christmas.

**User Tag Vocabulary Growth.** There are also studies that focus on tags applied by a specific user. Golder and Huberman [16] show that certain users' sets of distinct tags grow linearly as new resources are added. But Marlow et al. [34] find that for many users, such as those with few distinct tags in the graph, distinct tag growth declines over time, indicating either agreement on the tag vocabulary, or diminishing returns on their usage. In some cases, new tags are added consistently as photos are uploaded, suggesting a supply of fresh vocabulary and constant incentive for using tags. Sometimes only a few tags are used initially with a sudden growth spurt later on, suggesting that the user either discovered tags or found new incentives for using them.

## Identifying Tag Semantics

Intuitively, tags as user generated classification labels are semantically meaningful. So, research has been done for exploring the semantics of tags. These research works include three aspects: (1) Identifying similar tags, (2) mapping tags to taxonomies, and (3) extracting certain types of tags.

**Analysis of Pairwise Relationships between Tags.** In order to measure similarity of tags beyond words, researchers proposed various models to explore tags' similarity. Most of them are based on a simple assumption that tags that are similar may be used to tag the same resources, and similar resource would be tagged by similar tags. Therefore, inter tag correlation graph (tag as nodes, edges between two tags if they co-occur, weight on edge = cosine distance measure using number of times a tag was used) can be built for a tagging system. An analysis of the structural properties of such tag graphs may provide important insights into how people tag and how semantic structures emerge in distributed folksonomies. A simple approach would be measuring tags similarity based on the number of common web pages tagged by them. In section 7, we show how analysis of co-occurrence of tags can be used to generate tag recommendations.

**Extracting Ontology from Tags.** Another line of research for identifying semantics of tags is mapping tags to an existing ontology. Being able to automatically classify tags into semantic categories allows us to understand better the way users annotate media objects and to build tools for viewing and browsing the media objects. The simplest approach is based on string matching. Sigurbjörnsson et al. [49] map Flickr tags onto WordNet semantic categories using straight forward string matching between Flickr tags and WordNet lemmas. They found that 51.8% of the tags in Flickr can be assigned a semantic category using this mapping. To better assign tags to a category, content of resources associated with a given tag could be used. Overell et al. [38] designed a system to auto-classify tags using Wikipedia and Open Directory. They used structural patterns like categories and templates that can be extracted from resource metadata to classify Flickr tags. They built a classifier to classify Wikipedia articles into eleven semantic categories (act, animal, artifact, food, group, location, object, person, plant, substance and time). They map Flickr tags to Wikipedia articles using anchor texts in Wikipedia. Since they have classified Wikipedia articles, Flickr tags can be categorized using the same classification. They

classify things as what, where and when. They show that by deploying ClassTag they can improve the classified portion of the Flickr vocabulary by 115%. Considering the full volume of Flickr tags, i.e., taking tag frequency into account, they show that with ClassTag nearly 70% of Flickr tags can be classified. Figure 1.4 shows an overview of their system. Figure 1.5 shows the classification of Flickr tags.

Figure 1.4. Overview of the ClassTag system

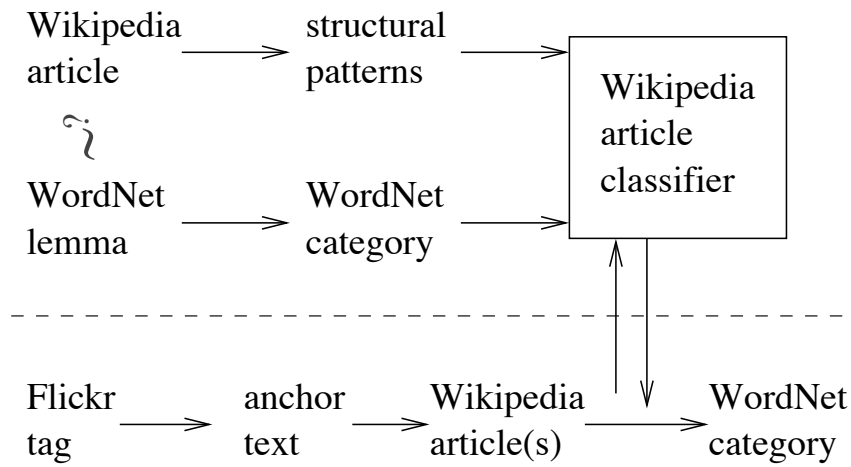



Figure 1.5. Classification of Flickr tags using ClassTag system



**Extracting Place and Event Semantics.** Tags also contain specific information, such as locations or events. Rattenbury et al. [43] study the problem of extracting place and event semantics for Flickr tags. They analyze two methods inspired by burst-analysis techniques


(popular in signal processing) and one novel method: Scale-structure Identification. The location,  $l_p$ , (consisting of latitude-longitude coordinates) associated with photo  $p$  generally marks where the photo was taken; but sometimes marks the location of the photographed object. The time,  $t_p$ , associated with photo  $p$  generally marks the photo capture time; but occasionally refers to the time the photo was uploaded to Flickr. They aim to determine, for each tag in the dataset, whether the tag represents an event (or place). The intuition behind the various methods they present is that an event (or place) refers to a specific segment of time (or region in space). The number of usage occurrences for an event tag should be much higher in a small segment of time than the number of usage occurrences of that tag outside the segment. The scale of the segment is one factor that these methods must address; the other factor is calculating whether the number of usage occurrences within the segment is significantly different from the number outside the segment. The Scale-structure Identification method performs a significance test that depends on multiple scales simultaneously and does not rely on apriori defined time segments. The key intuition is: if tag  $x$  is an event then the points in  $T_x$ , the time usage distribution, should appear as a single cluster at many scales. Interesting clusters are the ones with low entropy. For place identification,  $L_x$  is used rather than  $T_x$ . Periodic events have strong clusters, at multiple scales, that are evenly spaced apart in time. Practically, because tags occur in bursts, a periodic tag should exhibit at least three strong clusters (to rule out tags that just happened to occur in two strong temporal clusters but are not truly periodic). Overall, their approach has a high precision however a large proportion of tags remain unclassified.

## Tags Versus Keywords

To identify the potential of tags in being helpful for search, there have been works that compare tags with keywords. As shown in Figure 1.6, given a web document, the “most important” words (both wrt  $tf$  as well as  $tf \times idf$ ) of the document are generally covered by the vocabulary of user-generated tags [32]. This means that the set of user-generated tags has the comparable expression capability as the plain English words for web documents. Li et al. [32] found that most of the missed keywords are misspelled words or words invented by users, and usually cannot be found in dictionary.


Further, they define tag match ratio  $e(T, U)$  for tag set  $T$  associated with a URL  $U$  as ratio of weights of the tags of a particular URL that can be matched by the document.  $e(T, U) = \frac{\sum_{k|t_k \in U} w(t_k)}{\sum_i w(t_i)}$ . Here,  $w(t)$  is

Figure 1.6. Tag coverage for important keywords



the weight of tag  $t$ , i.e., the frequency of tag  $t$  in the data set. The tag match ratio represents the ratio of important tags of a URL matched by the document. Figure 1.7 shows the distribution of tag match ratio for URLs in their Delicious dataset.

Figure 1.7. Distribution of the tag match ratio



Besides this, tags often have much more expressive power. E.g., consider the Google home page. It does not mention the phrase “search engine” anywhere. But “searchengine” can be found as a tag against the bookmarked URL

<http://www.google.com/ig> on Delicious.

## 6. Visualization of tags

Social tagging is one of the most important forms of user generated content. Appropriate rendering of tags can provide useful information to users. Tag clouds have been explored to support browsing on web portals, and various tag selection methods for tag clouds have been developed. Much work has been done to identify hierarchy of tags in the tag cloud construction. Visualization schemes stress on display formats

and evolutionary aspect of tag clouds, etc. We review these in this section.

## **Tag Clouds for Browsing/Search**

Tag cloud, a visual depiction of user-generated tags, is used to facilitate browsing and search process of the tags. Sinclair and Cardew-Hall [51] discuss situations when using tag clouds is better than doing search. They conducted an experiment, giving participants the option of using a tag cloud or a traditional search interface to answer various questions. They found that participants preferred the search interface if the information need is specific while they preferred the tag cloud if the information-seeking task was more general. It is partly because tags are good at broad categorization, as opposed to specific characterization of content. In total, the number of questions answered using the tag cloud was greater than those answered using the search box. The tag cloud provides a visual summary of the contents of the database. A number of participants commented that tag clouds gave them an idea of where to begin their information seeking. The tag cloud helps the information seeker to get familiar with the domain and allows more focused queries based on the information gained from browsing. It appears that scanning the tag cloud requires less cognitive load than formulating specific query terms.

Tag clouds have their disadvantages too. First, tag clouds obscure useful information by being skewed towards what is popular. Second, answering a question using the tag cloud required more queries per question than the search box. Third, many of the participants commented that the tag cloud did not allow them to narrow their search enough to answer the given questions. On average, roughly half of the articles in the dataset remain inaccessible from the tag cloud. Most tagging systems mitigate this by including tag links when viewing individual articles, thus exposing some of the less popular tags. However, this is not necessarily useful when someone is seeking specific information.

Millen et al. [36] did experiments on Dogear system where clicking a tag leads to a view of bookmarks that are associated with that tag. They found that the most frequent way to browse bookmarks is by clicking on another person's name, followed by browsing bookmarks by selecting a specific tag from the system-wide tag cloud. It is considerably less common for a user to select a tag from another user's tag cloud and very less chances of using more advanced browsing of tag intersections.

## Tag Selection for Tag Clouds

Since there is only limited display space for tags in tag clouds, how to select the appropriate tags is a challenging task. Hassan-Montero and Herrero-Solana [20] describe a system for bi-dimensional visualization of tag clouds. Tag selection is based on usefulness as determined by: (1) its capacity to represent each resource as compared to other tags assigned to the same resource, (2) the volume of covered resources as compared to other tags, (3) its capacity to cover these resources less covered by other tags. Semantic relationships among tags are defined in terms of their similarity, quantified by means of the Jaccard coefficient. K-means clustering is then applied on tag similarity matrix, with an apriori chosen number of clusters and a fixed number of selected relevant tags. They apply Multidimensional Scaling, using Pearson's correlation as the similarity function, on a tag-to-tag correlation matrix. MDS creates a bi-dimensional space, which is then visualized through a fish-eye system. Alphabetical-based schemes are useful for know-item searching, i.e., when user knows previously what tag she is looking for, such as when user browses her personal tag cloud. They propose a tag cloud layout based on the assumption that clustering techniques can improve tag clouds' browsing experience. The display method is similar to traditional tag cloud layout, with the difference that tags are grouped with semantically similar tags, and likewise clusters of tags are displayed near semantically similar clusters. Similar tags are horizontally neighbors, whereas similar clusters are vertically neighbors. Clustering offers more coherent visual distribution of tags than traditional alphabetical arrangements, allowing to differentiate among main topics in tag cloud, as well as to infer semantic knowledge from the neighbors' relationships.

Begelman et al. [4] propose a clustering algorithm to find strongly related tags. The algorithm is based on counting the number of co-occurrences of any pair of tags and a cut-off point is determined when the co-occurrence count is significant enough to be used. To determine this cutoff point, they start from the tail on the right end and seek the point where the first derivative of the count has its first high peak (that is when the second derivative goes from positive to negative) and check if the peak was high enough. This results in a sparse matrix that represents tags, so that the value of each element is the similarity of the two tags. Using this definition of similarity, they design an inter-tag correlation network graph. They then cluster this graph using spectral bisection and modularity function.



## **Tag Hierarchy Generation**

Beyond the flat structure of tags, hierarchical structure also exists in the tagging space. Caro et al. [9] present the tagFlake system, which supports semantically informed navigation within a tag cloud. The system organizes tags extracted from textual content in hierarchical organizations, suitable for navigation, visualization, classification and tracking. It extracts the most significant tag/terms from text documents and maps them onto a hierarchy in such a way that descendant terms are contextually dependent on their ancestors within the given corpus of documents. This provides tagFlake with a mechanism for enabling navigation within the tag space and for classification of the text documents based on the contextual structure captured by the generated hierarchy.

Li et al. [31] present Effective Large Scale Annotation Browser (ELSABer), to browse large-scale social annotation data. ELSABer helps the users browse a huge number of annotations in a semantic, hierarchical and efficient way. ELSABer has the following features: (1) the semantic relations between annotations are explored for browsing of similar resources; (2) the hierarchical relations between annotations are constructed for browsing in a top-down fashion; (3) the distribution of social annotations is studied for efficient browsing. The hierarchy structure is determined by a decision tree with the features including tag coverage, URL intersection rate, inverse-coverage rate, etc.

## **Tag Clouds Display Format**

Tags clouds can be displayed in different formats. Bielenberg and Zacher [5] have proposed circular clouds, as opposed to the typical rectangular layout, where the most heavily weighted tags appear closer to the center. Font size and distance to the center represent the importance of a tag, but distance between tags does not represent their similarity.

Owen and Lemire [26] present models and algorithms to improve the display of tag clouds that consist of in-line HTML, as well as algorithms that use nested tables to achieve a more general two-dimensional layout in which tag relationships are considered. Since the font size of a displayed tag is usually used to show the relative importance or frequency of the tag, a typical tag cloud contains large and small text interspersed. A consequence is wasteful white space. To handle the space waste problem, the authors propose the classic electronic design automation (EDA) algorithm, min-cut placement, for area minimization and clustering in tag clouds. For the large clumps of white space, the solution is a hybrid of the classic Knuth-Plass algorithm for text justification, and a book-

placement exercise considered by Skiena. The resulting tag clouds are visually improved and tighter.

## Tag Evolution Visualization

Other than the text information, tags usually have the time dimension. To visualize the tag evolution process is an interesting topic. Dubinko et al. [15] consider the problem of visualizing the evolution of tags within Flickr. An animation provided via Flash in a web browser allows the user to observe and interact with the interesting tags as they evolve over time. The visualization is made up of two interchangeable metaphors - the ‘river’ and the ‘waterfall’. The visualization provides a view of temporal evolution, with a large amount of surface data easily visible at each timestep. It allows the user to interact with the presentation in order to drill down into any particular result. It remains “situated” in the sense that the user is always aware of the current point of time being presented, and it provides random access into the time stream so that the user can reposition the current time as necessary. There are two novel contributions in their algorithm. The first is a solution to an interval covering problem that allows any timescale to be expressed efficiently as a combination of a small number of pre-defined timescales that have been pre-computed and saved in the “index” structure. The second contribution is an extension of work on score aggregation allowing data from the small number of pre-computed timescales to be efficiently merged to produce the optimal solution without needing to consume all the available data. The resulting visualization is available at Taglines<sup>5</sup>. In some cases, the user may seek data points that are particularly anomalous, while in other cases it may be data points that are highly persistent or that manifest a particular pattern. The authors focus on one particular notion of “interesting” data: the tags during a particular period of time that are most representative for that time period. That is, the tags that show a significantly increased likelihood of occurring inside the time period, compared to outside.

Russel [44] has proposed Cloudalicious<sup>6</sup>, a tool to study the evolution of the tag cloud over time. Cloudalicious takes a request for a URL, downloads the tagging data from Delicious, and then graphs the collective users tagging activity over time. The y-axis shows the relative weights of the most popular tags for that URL. As the lines on the graph move from left to right, they show signs of stabilization. This pattern can be interpreted as the collective opinion of the users. Diagonal lines are the most interesting elements of these graphs as they suggest that

the users doing the tagging have changed the words used to describe the site.

## Popular Tag Cloud Demos

Some demos for visualizing tags are also available on the Web. Grafolicious<sup>7</sup> produces graphs illustrating when and how many times a URL has been bookmarked in Delicious. HubLog<sup>8</sup> gives a graph of related tags connected with the given tags. Although these demos gave a vivid picture of social annotations in different aspects, their goals are not to help users to browse annotations effectively. PhaseTwo<sup>9</sup> aims at creating visually pleasant tag clouds, by presenting tags in the form of seemingly random collections of circles with varying sizes: the size of the circle denotes its frequency. Delicious also provides its own tag cloud view<sup>10</sup>. Tag.alicio.us<sup>11</sup> operates as a tag filter, retrieving links from Delicious according to tag and time constraints (e.g., tags from this hour, today, or this week). Extisp.icio.us<sup>12</sup> displays a random scattering of a given user's tags, sized according to the number of times that the user has reused each tag, and Facetious<sup>13</sup> was a reworking of the Delicious database, which made use of faceted classification, grouping tags under headings such as “by place” (Iraq, USA, Australia), “by technology” (blog, wiki, website) and “by attribute” (red, cool, retro). Tag clouds have also been integrated inside maps for displaying tags having geographical information, such as pictures taken at a given location.

## 7. Tag recommendations

The tagging system can recommend some tags to a user, and the user can select one of those tags or come up with a new one. Tag recommendation is not only useful to improve user experience, but also makes rich annotation available. There have been many studies on tag recommendation. Tags can be recommended based on their quality, co-occurrence, mutual information and object features.

### Using Tag Quality

Tag quality can guide the tag recommendation process. The tag quality can be evaluated by facet coverage and popularity, and those tags of high quality are used for recommendation. Xu et al. [61] propose a set of criteria for tag quality and then propose a collaborative tag suggestion algorithm using these criteria to discover the high-quality tags. A good tag combination should include multiple facets of the tagged objects. The number of tags for identifying an object should be minimized, and the number of objects identified by the tag combination should be

small. Note that personally used organizational tags are less likely to be shared by different users. Thus, they should be excluded from tag recommendations. The proposed algorithm employs a goodness measure for tags derived from collective user authorities to combat spam. The goodness measure is iteratively adjusted by a reward-penalty algorithm, which also incorporates other sources of tags, e.g., content-based auto-generated tags. The algorithm favors tags that are used by a large number of people, and minimizes the overlap of concepts among the suggested tags to allow for high coverage of multiple facets and honors the high correlation among tags.

### Using Tag Co-occurrences

One important criterion used for tag recommendation is tag co-occurrence. Those tags co-occurring with the existing tags of the object are used for recommendation. Sigurbjörnsson and Zwol [49] present four strategies to recommend tags. These include two co-occurrence based strategies: Jaccard similarity and an asymmetric measure  $P(t_j | t_i) = \frac{|t_i \cap t_j|}{|t_i|}$ . Tag Aggregation and promotion strategies are based on voting or weighted voting based on co-occurrence count. Given a set of user-defined tags  $U$  on an object  $O$ , they want to rank candidate tags  $C$  that can be recommended for object  $O$  based on co-occurrence counts of  $u$  and  $c$  such that  $u \in U$  and  $c \in C$ . From the tag frequency distribution, they learned that both the head and the tail of the power law would probably not contain good tags for recommendation. Considered that user-defined tags with very low collection frequency are less reliable than tags with higher collection frequency, those tags for which the statistics are more stable were promoted. They compute promotion score for a  $(u, c)$  pair by using stability of tag  $u$ , descriptiveness of tag  $c$  and rank of tag  $c$  wrt tag  $u$ . These are in turn defined using system parameters  $k_s$ ,  $k_d$  and  $k_r$ .

$$stability(u) = \frac{k_s}{k_s + abs(k_s - \log(|u|))}$$

where  $|u|$  is the collection frequency of user defined tag  $u$ .

Tags with very high frequency are likely to be too general for individual photos.

$$descriptive(c) = \frac{k_d}{k_d + abs(k_d - \log(|c|))}$$

The rank  $rank(u, c)$  of a candidate tag  $c$  for tag  $u$  is  $\frac{k_r}{k_r + r - 1}$  where  $r$  is position of tag  $c$  for a tag  $u$ . The promotion score can be defined as  $promotion(u, c) = rank(u, c) \times stability(u) \times descriptive(c)$ . Tag score is finally computed as  $score(c) = \sum_{u \in U} vote(u, c) \times promotion(u, c)$ . Here  $vote(u, c)$  is 1 if tags  $u$  and  $c$  co-occur, else 0. Tag frequency distribution follows a perfect power law, and the mid section of this power law contained the most interesting candidates for tag recommendation. They

found that locations, artifacts and objects have a relatively high acceptance ratio (user acceptance of the recommended tag). However, people, groups and unclassified tags (tags that do not appear in WordNet) have relatively low acceptance ratio.

## Using Mutual Information between Words, Documents and Tags

Mutual information is another criterion for tag recommendation. Song et al. [52] treat the tagged training documents as triplets of (words, documents, tags), and represent them as two bipartite graphs, which are partitioned into clusters by Spectral Recursive Embedding (SRE) and using Lanczos algorithm for symmetric low rank approximation for the weighted adjacency matrix for the bipartite graphs. Tags in each topical cluster are ranked by a novel ranking algorithm. A two-way Poisson Mixture Model (PMM) is proposed to model the document distribution into mixture components within each cluster and aggregate words into word clusters simultaneously. During the online recommendation stage, given a document vector, its posterior probabilities of classes are first calculated. Then based on the joint probabilities of the tags and the document, tags are recommended for this document based on their within-cluster ranking. The efficiency of the Poisson mixture model helps to make recommendations in linear-time in practice. Within a cluster, node ranking is defined by  $Rank_i = \exp(-\frac{1}{r(i)^2})$  for  $r(i) \neq 0$  where  $r(i) = np_i \times \log(nr_i)$ . N-Precision ( $np_i$ ) of a node  $i$  is the weighted sum of its edges that connect to the nodes within the same cluster, divided by the total sum of edge weights in that cluster. N-recall ( $Nr_i$ )=edges associated with node  $i$ /edges associated with node  $i$  within the same cluster.

## Using Object Features

Tag recommendation can also be performed using object features. E.g., the extracted content features from the images can be helpful in tag recommendation. In [33], Liu et al. propose a tag ranking scheme to automatically rank the tags associated with a given image according to their relevance to the image content. To estimate the tag relevance to the images, the authors first get the initial tag relevance scores based on probability density estimation, and then apply a random walk on a tag similarity graph to refine the scores. Since all the tags have been ranked according to their relevance to the image, for each uploaded image, they find the  $K$  nearest neighbors based on low-level visual features, and then the top ranked tags of the  $K$  neighboring images are collected

and recommended to the user. In [58], Wu et al. model the tag recommendation as a learning task that considers multi-modality including tag co-occurrence and visual correlation. The visual correlation scores are derived from Visual language model (VLM), which is adopted to model the content of the tags in visual domain. The optimal combination of these ranking features is learned by the Rankboost algorithm.

## 8. Applications of Tags

In this section, we would describe different applications for which tags have been used. Social tagging can be useful in the areas including indexing, search, taxonomy generation, clustering, classification, social interest discovery, etc.

### Indexing

Tags can be useful for indexing sites faster. Users bookmark sites launched by their friends or colleagues before a search engine bot can find them. Tags are also useful in deeper indexing. Many pages bookmarked are deep into sites and sometimes not easily linked to by others, found via bad or nonexistent site navigation or linked to from external pages. Carmel et al. [8] claim that by appropriately weighting the tags according to their estimated quality, search effectiveness can be significantly improved. They propose a novel framework for bookmark (a triple of document  $d$ , user  $u$ , tag  $t$ ) weighting that estimates the effectiveness of bookmarks for IR tasks as fundamental entities in social bookmarking systems.

### Search

Tags have been found useful for web search, personalized search and enterprise search. Tags offer multiple descriptions of a given resource, which potentially increases the likelihood that searcher and tagger find a common language and thus using tags, retrieval effectiveness may be enhanced. Social bookmarking can provide search data not currently provided by other sources.

Heymann et al. [23] analyze posts to Delicious: how many bookmarks exist (about 115M), how fast is it growing, and how active are the URLs being posted (quite active). They observe the following. (1) Pages posted to Delicious are often recently modified. (2) Approximately 12.5% of URLs posted by users are new, unindexed pages. (3) Roughly 9% of results for search queries are URLs present in Delicious. (4) While some users are more prolific than others, the top 10% of users only ac-

count for 56% of posts. (5) 30-40% of URLs and approximately one in eight domains posted were not previously in Delicious. (6) Popular query terms and tags overlap significantly (though tags and query terms are not correlated). (7) Most tags were deemed relevant and objective by users. (8) Approximately 120000 URLs are posted to Delicious each day. (9) There are roughly 115M public posts, coinciding with about 30-50M unique URLs. (10) Tags are present in the pagetext of 50% of the pages they annotate and in the titles of 16% of the pages they annotate. (11) Domains are often highly correlated with particular tags and vice versa.

Similarly, Heckner et al. [21] found out using a survey that Flickr and Youtube users perceive tags as helpful for IR, and show a certain tendency towards searching other collections rather than their own collections. Users reveal that Flickr search often leads to better precision and recall for pictures compared to Google search. They also point out that Flickr search is quite specific whereas with Google search you get more wild cards you don't expect based on the information on a page rather than just the tags on a photo.


**Semantic Query Expansion.** Schenkel et al. [45] develop an incremental top- $k$  algorithm for answering queries of the form  $Q(u, q_1, q_2, \dots, q_n)$  where  $u$  is the user and  $q_1, q_2, \dots, q_n$  are the query keywords. The algorithm performs two-dimensional expansions: social expansion considers the strength of relations among users, and semantic expansion considers the relatedness of different tags. It is based on principles of threshold algorithms, and folds friends and related tags into the search space in an incremental on-demand manner.

Figure 1.8 shows their social network model that incorporates both social and semantic relationships. In contrast with standard IR model, the content-based score of a document is additionally user-specific, i.e., it depends on the social context of the query initiator.

They present an algorithm, ContextMerge, to efficiently evaluate the top- $k$  matches for a query, using the social context score.

ContextMerge makes use of information that is available in social tagging systems, namely lists of documents tagged by a user and number of documents tagged with tags. It incrementally builds social frequencies by considering users that are related to the querying user in descending order of friendship similarity, computes upper and lower bounds for the social score from these frequencies, and stops the execution as soon as it can be guaranteed that the best  $k$  documents have been identified.

Figure 1.8. Social network model



**Enhanced Similarity Measure.** Estimating similarity between a query and a web page is important to the web search problem. Tags provided by web users provide different perspectives and so are usually good summaries of the corresponding web pages and provide a new metadata for the similarity calculation between a query and a web page. Furthermore, similar tags are used to tag similar web pages. Semantics of tags can be measured and enhanced similarity measure between queries and tags can be formulated.

Wu et al. [59] show how emergent semantics can be statistically derived from the social annotations. They propose to use a probabilistic generative model to model the user's annotation behavior and to automatically derive the emergent semantics of the tags. Synonymous tags are grouped together and highly ambiguous tags are identified and separated. The three entities (user, resource, tag) are represented in the same multi-dimensional vector space called the conceptual space. They extend the bigram separable mixture model to a tripartite probabilistic model to obtain the emergent semantics contained in the social annotations data. Furthermore, they apply the derived emergent semantics to discover and search shared web bookmarks. They provide a basic search model that deals with queries that are a single tag and rank semantic related resources without considering personalized information of the



user. They also extend this basic model for personalized search and to support complicated queries (where a complicated query is a boolean combination of tags and other words appearing in the resources).

Bao et al. [2] also observe that the social annotations can benefit web search in this aspect. They proposed SocialSimRank (SSR) which calculates the similarity between social annotations and web queries. Preliminary experimental results show that SSR can find the latent semantic association between queries and annotations.

**Enhanced Static Ranking.** Estimating the quality of a web page is also important to the web search problem. The amount of annotations assigned to a page indicates its popularity and indicates its quality in some sense. In order to explore social tags for measuring the quality of web pages, researchers have exploited the tagging graph.

Hotho et al. [24] present a formal model and a new search algorithm for folksonomies, called FolkRank, that exploits the structure of the folksonomy. The FolkRank ranking scheme is then used to generate personalized rankings of the items in a folksonomy, and to recommend users, tags and resources. The social network graph is defined as a tripartite graph of resources, users and tags as nodes where the edges are the relationships between resources and tags, tags and users, and users and resources. Adapted PageRank is then defined as  $\vec{w} = \alpha\vec{w} + \beta A\vec{w} + \gamma\vec{p}$  where  $\alpha + \beta + \gamma = 1$ . The FolkRank algorithm computes a topic-specific ranking in a folksonomy as follows:

- The preference vector  $\vec{p}$  is used to determine the topic. It may have any distribution of weights, as long as  $\|\vec{w}\|_1 = \|\vec{p}\|_1$  holds. Typically a single entry or a small set of entries is set to a high value, and the remaining weight is equally distributed over the other entries. Since the structure of the folksonomies is symmetric, we can define a topic by assigning a high value to either one or more tags and/or one or more users and/or one or more resources.
- Let  $\vec{w}_0$  be the fixed point with  $\beta = 1$
- Let  $\vec{w}_1$  be the fixed point with  $\beta < 1$
- $\vec{w} = \vec{w}_1 - \vec{w}_0$  is the final weight vector.

Thus, they compute the winners and losers of the mutual reinforcement of resources when a user preference is given, compared to the baseline without a preference vector. The resulting weight  $w[x]$  of an element  $x$  of the folksonomy is called the FolkRank of  $x$ . They observed that Adapted PageRank ranking contains many globally frequent tags, while the FolkRank ranking provides more personal tags.

Bao et al. [2] also propose a novel algorithm, SocialPageRank (SPR) to measure the popularity of web pages using social annotations. The algorithm also utilize the social tagging graph. The intuition behind the algorithm is the mutual enhancement relation among popular web pages, up-to-date web users and hot social annotations. SocialPageRank (SPR) captures the popularity of web pages and successfully measures the quality (popularity) of a web page from the web users' perspective.

**Personalized Search.** Furthermore, personal tags are naturally good resources for describing a person's interests. So personal search could be enhanced via exploring personal tags. Xu et al. [60] present a framework in which the rank of a web page is decided not only by the term matching between the query and the web page's content but also by the topic matching between the user's interests and the web page's topics.

The personalized search is conducted by ranking the web pages using two guidelines, term matching and topic matching. When a user  $u$  issues a query  $q$ , a web page  $p$  is ranked not only by the term similarity between  $q$  and  $p$  but also by the topic similarity between  $u$  and  $p$ . Three properties of folksonomy are studied for the topic space estimation:

(1) The categorization property. Many of the social annotations are subject descriptor keywords at various levels of specificity. The selection of proper annotations for a web page is somewhat a classification of the web page to the categories represented by the annotations.

(2) The keyword property. Annotations can be seen as good keywords for describing the respective web pages from various aspects.

(3) The structure property. In folksonomy systems, users' bookmarking actions form a cross link structure between the users and the web pages. They model the structure using a user-web page bipartite graph.

When a user  $u$  issues a query  $q$ , two search processes begin, a term matching process and a topic matching process. The term matching process calculates the similarity between  $q$  and each web page to generate a user unrelated ranked document list. The topic matching process calculates the topic similarity between  $u$  and each web page to generate a user related ranked document list. Then a merge operation is conducted to generate a final ranked document list based on the two sub ranked document lists. They adopt ranking aggregation to implement the merge operation by  $r(u, q, p) = \gamma \cdot r_{term}(q, p) + (1 - \gamma) \cdot r_{topic}(u, p)$

Topic space selection is done using (1) social annotations as topics (2) ODP categories as topics (3) interest and topic adjusting via bipartite collaborative link structure.

**Intranet (Enterprise) Search.** Dmitriev et al. [14] show how user annotations can be used to improve the quality of intranet (enterprise) search. They propose two ways to obtain user annotations, using explicit and implicit feedback, and show how they can be integrated into a search engine.

One way to obtain annotations is to let users explicitly enter annotations for the pages they browse. The implicit method of obtaining annotations is to use the queries users submit to the search engine as annotations for pages users click on. They propose several strategies to determine which pages are relevant to the query, i.e., which pages to attach an annotation to, based on clickthrough data associated with the query. There are different ways of getting feedback using this method. For every click record, the first strategy produces a (URL, Annotation) pair, where Annotation is the QueryString. This strategy is simple to implement, and gives a large number of annotations. A second strategy only produces a (URL, Annotation) pair for a click record which is the last record in a session. Annotation is still the QueryString. A query chain is a time-ordered sequence of queries, executed over a short period of time. The assumption behind using query chains is that all subsequent queries in the chain are actually refinements of the original query. The third strategy, similar to the first one, produces a (URL, Annotation) pair for every click record, but Annotation now is the concatenation of all QueryStrings from the corresponding query chain. Finally, the fourth strategy produces a (URL, Annotation) pair for a click record which is the last record in the last session in a query chain, and Annotation is, again, the concatenation of QueryStrings from the corresponding query chain.

Compared to the baseline system (search engine without annotations), they obtained around 14% improvement in accuracy when using explicit annotations and around 9.3% performance improvement when using implicit annotations.

## Taxonomy Generation

Tags can be organized in hierarchical taxonomy structure based on their semantic meanings. In this part, we introduce two methods to generate the taxonomy structure.

**Using Centrality in Tag Similarity Graph.** Heymann and Garcia-Molina [22] provide an algorithm to convert a large corpus of tags into a navigable hierarchical taxonomy. Their greedy algorithm extracts a hierarchical taxonomy using graph centrality in a tag similarity

graph. It starts with a single node tree and then it adds each tag to the tree in the decreasing order of tag centrality in the similarity graph. It decides the location of each candidate tag by its similarity to every node currently present in the tree. The candidate tag is then either added as a child of the most similar node if its similarity to that node is greater than some threshold, or it is added to the root node if there does not currently exist a good parent for that node. Furthermore, they describe some features that can help to establish hierarchical relationships: (1) density ( $\frac{\text{\#annotated objects}}{\text{\#objects}}$ ) (2) overlap ( $\frac{\text{\#shared annotated objects}}{\text{\#annotated objects}}$ ) (3) distribution of specificity in the graph that describes the detail level of tags in the system (4) agreement between users on which tags are appropriate for a given subject.

**Using Association Rule Mining.** Schmitz et al. [46] discuss how association rule mining can be adopted to analyze and structure folksonomies, and how the results can be used for ontology learning and supporting emergent semantics. Since folksonomies provide a three-dimensional dataset (users, tags, and resources) instead of a usual two-dimensional one (items and transactions), they present first a systematic overview of projecting a folksonomy onto a two-dimensional structure. They determine the set of objects, i.e., the set on which the support will be counted by a permutation  $P$  on the set 1, 2, 3. To induce an ontology from Flickr tags, Schmitz et al. [47] follow this approach. Term  $x$  potentially subsumes term  $y$  if:  $P(x \mid y \geq t)$  and  $P(y \mid x < t), D_x \geq D_{min}, D_y \geq D_{min}, U_x \geq U_{min}, U_y \geq U_{min}$  where:  $t$  is the co-occurrence threshold,  $D_x$  is the # of documents in which term  $x$  occurs, and must be greater than a minimum value  $D_{min}$ , and  $U_x$  is the # of users that use  $x$  in at least one image annotation, and must be greater than a minimum value  $U_{min}$ .

Once the co-occurrence statistics are calculated, candidate term pairs are selected using the specified constraints. A graph of possible parent-child relationships is then built, and the co-occurrence of nodes with ascendants that are logically above their parent is filtered out. For a given term  $x$  and two potential parent terms  $p_{x_i}$  and  $p_{x_j}$ , if  $p_{x_i}$  is also a potential parent term of  $p_{x_j}$  then  $p_{x_i}$  is removed from the list of potential parent terms for term  $x$ . At the same time, the co-occurrence of terms  $x, p_{x_i}$  and  $p_{x_j}$  in the given relationships indicates both that the  $x \Rightarrow p_{x_j}$  relationship is more likely than simple co-occurrence might indicate, and similarly that the  $p_{x_i} \Rightarrow p_{x_j}$  relationship should be reinforced. Weights of each of these relationships are incremented accordingly. Finally, they consider each leaf in the tree and choose the best path up to a root, given the (reinforced) co-occurrence weights for potential parents of each node,

and coalesce the paths into trees. They use this induced taxonomy to improve search by inferring parent terms for images with child terms.

## Public Library Cataloging

Folksonomies have the potential to help public library catalogues by enabling clients to store, maintain and organize items of interest in the catalogue [54]. Spiteri et al. [54] acquired tags over a thirty-day period from the daily tag logs of three folksonomy sites: Delicious, Furl and Technorati. The tags were evaluated against Section 6<sup>14</sup> (choice and form of terms) of the National Information Standards Organization (NISO) guidelines for the construction of controlled vocabularies. This evaluation revealed that the folksonomy tags correspond closely to the NISO guidelines that pertain to the types of concepts expressed by the tags, the predominance of single tags, the predominance of nouns and the use of recognized spelling. Potential problem areas in the structure of the tags pertain to the inconsistent use of the singular and plural form of count nouns, and the incidence of ambiguous tags in the form of homographs and unqualified abbreviations or acronyms. If library catalogues decide to incorporate folksonomies, they could provide clear guidelines to address these noted weaknesses, as well as links to external dictionaries and references sources such as Wikipedia to help clients disambiguate homographs and to determine if the full or abbreviated forms of tags would be preferable.

## Clustering and Classification

Tags can be used as the additional features for both clustering and classification.

### K-means Clustering in an Extended Vector Space Model.

Ramage et al. [42] explore the use of tags in K-means clustering in an extended vector space model that includes tags as well as page text. They also provide a novel generative clustering algorithm based on latent Dirichlet allocation that jointly models text and tags. They examine five ways to model a document with a bag of words  $B_w$  and a bag of tags  $B_t$  as a vector  $V$ : Words Only, Tags Only, Words + Tags, Tags as Words Times  $n$  (vector  $V$  is represented as  $B_w \cup (B_t * n)$  and vocabulary is  $W \cup T$ ), Tags as New Words (word#computer is different from tag#computer). They evaluate the models by comparing their output to an established web directory. They find that the naive inclusion of tagging data improves cluster quality versus page text alone, but a more principled inclusion can substantially improve the quality of all models

with a statistically significant absolute F-score increase of 4%. The generative model outperforms K-means with another 8% F-score increase.

**Classification of Blog Entries.** Brooks and Montanez [7] analyze the effectiveness of tags for classifying blog entries by gathering the top 350 tags from Technorati and measuring the similarity of all articles that share a tag. They find that tags are useful for grouping articles into broad categories, but less effective in indicating the particular content of an article. They show that automatically extracting highly relevant words can produce a more focused categorization of articles. They also show that clustering algorithms can be used to reconstruct a topical hierarchy among tags. They conclude that tagging does manage to group articles into categories, but that there is room for improvement.

**Web Object Classification.** Yin et al. [62] cast the web object classification problem as an optimization problem on a graph of objects and tags. They then propose an efficient algorithm which not only utilizes social tags as enriched semantic features for the objects, but also infers the categories of unlabeled objects from both homogeneous and heterogeneous labeled objects, through the implicit connection of social tags.

Let  $C$  be a category set,  $c_1, c_2, \dots, c_k$ . Every object  $u$  and every tag  $v$  is a vertex in the graph  $G(u, v \in V)$ . If an object  $u$  is associated with a tag  $v$ , there will be an edge between  $u$  and  $v$ , denoted as  $(u, v) \in E$ .  $V$  consists of four types of vertices:  $V_S$  is a set of objects of type  $S$ .  $V_T^l$  is a set of labeled objects of type  $T$ .  $V_T^u$  is a set of unlabeled objects of type  $T$ .  $V_{tag}$  is a set of tags. The problem of web object classification can then be defined as the problem of assigning category weights to each vertex in the graph.

They propose that: (1) Category assignment of a vertex in  $V_S$  should not deviate much from its original label. (2) Category assignment of the vertex in  $V_T^l$  should remain the same with its original label if it is fully trustable. Even if it is not, it should not deviate too much. (3) Category of the vertex in  $V_T^u$  should take the prior knowledge into consideration if there is any. (4) Category assignment of any vertex in graph  $G$  should be as consistent as possible to the categories of its neighbors. Using these properties, they write an objective function. By minimizing the above objective function, the optimal class distribution can be found.

## Social Interest Discovery

Tags represent common wisdom, so it can be useful for social interest discovery. Li et al. [32] propose that human users tend to use descriptive

tags to annotate the contents that they are interested in. User-generated tags are consistent with the web content they are attached to, while more concise and closer to the understanding and judgements of human users about the content.

They have developed an Internet Social Interest Discovery system, ISID, to discover the common user interests and cluster users and their saved URLs by different interest topics.

The aggregated user tags of a URL embrace different human judgments on the same subjects of the URL. This property is not possessed by the keywords of their referring web pages. Tags carrying the variation of human judgments reflects the different aspects of the same subjects. More importantly, it helps to identify the social interests in more finer granularity.

Their evaluation results show that: (1) the URLs' contents within a ISID cluster have noticeably higher similarity than that of the contents of URLs across different clusters, and (2) nearly 90% of all users have their social interests discovered by the ISID system.

ISID architecture provides the following functions. (1) Find topics of interests: For a given set of bookmark posts, find all topics of interests. Each topic of interests is a set of tags with the number of their co-occurrences exceeding a given threshold. ISID uses association rules algorithms to identify the frequent tag patterns for the posts. (2) Clustering: For each topic of interests, find all the URLs and the users such that those users have labeled each of the URLs with all the tags in the topic. For each topic, a user cluster and a URL cluster are generated. (3) Indexing: Import the topics of interests and their user and URL clusters into an indexing system for application queries.

They find that the tag-based cosine similarity is quite close to keyword based cosine similarity, indicating that tags really capture the main concepts of documents. Also their results are significantly different from those of [3]. With the tags of blog data, Bateman et al. [3], found that the average pairwise cosine similarity of the articles in tag-based clusters is only a little higher than that of randomly clustered articles, while much lower than that of articles clustered with high  $tf \times idf$  key words. However, evaluation by Li et al. [32] shows that tag-based clustering is highly accurate. The reason of this difference is that the clustering of articles in [3] is based on single tags, while topic clustering in [32] is based on multiple co-occurring tags.

## Enhanced Browsing

Social tagging results into a list of weighted tags associated with every resource. Zubiaga et al. [64] suggest alternative navigation ways using social tags: pivot browsing (moving through an information space by choosing a reference point to browse, e.g., pivoting on a tag allows to look for related tags; pivoting on a particular user), popularity driven navigation (sometimes a user would like to get documents that are popular for a known tag, e.g., retrieving only the documents where a tag has been top weighted), and filtering (social tagging allows to separate the stuff you do not want from the stuff you do want, e.g., gathering documents containing a tag but excluding another one). Currently, Wikipedia supports these navigation methods: search engine, category-driven navigation, link-driven navigation.

The tag cloud (or tag index) supports easy social navigation in that each of the tags is clickable; clicking a tag leads to a view of bookmarks that are associated with that tag. Tag clouds are either system-wide, or specific to one user, depending on the current view. Millen and Feinberg [36] study tag browsing on Dogear system. They point out that there is considerable browsing of the bookmark space by other people, other tags (everyone) and other people's tags. These results suggest widespread curiosity about what others are bookmarking. The most frequent way to browse bookmarks is by clicking on another person's name, followed by browsing bookmarks by selecting a specific tag from the system-wide tag cloud. During the trial period, 89% of individuals (2291 of 2579) clicked on URLs that had been bookmarked by another person. 74% of the total pages visited (32596 of 44144) were bookmarked by someone else. This provides considerable evidence that the Dogear service is supporting a high degree of social navigation.

## 9. Integration

There exist a large number of folksonomies dealing with similar type of objects. As a result, different folksonomies have different tags for the same object. An integration of such folksonomies can help in solving the problem of sparsity of tags associated with Web objects. Integration of folksonomies can help in creating richer user profiles. Some work has been done to integrate these taxonomies by tag co-occurrence analysis and clustering. We discuss such efforts in this section.

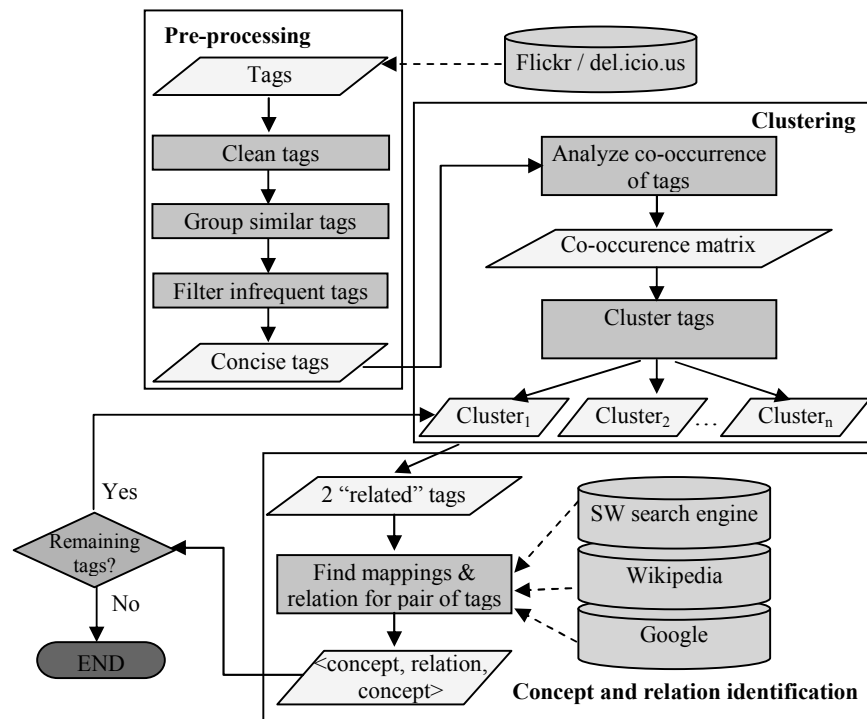


## Integration using Tag Co-occurrence Analysis and Clustering

Specia and Motta [53] tackle the problem of integrating folksonomies. They present an approach to minimize the problems of ambiguity, lack of synonymy and discrepancies in granularity by making explicit the semantics behind the tag space in social annotation systems. Using data collected from Delicious and Flickr, they use co-occurrence analysis and clustering techniques to construct meaningful groups of tags that correspond to concepts in an ontology. By exploiting external resources, such as Wikipedia, WordNet, and semantic web ontologies, meaningful relationships can be established between such tag groups.

Figure 1.9 shows the system pipeline.

Figure 1.9. Integration system architecture



To establish relationships between tags within each cluster and to refine clusters, they use the following procedure.

(1) Post each possible pair of tags within the cluster to the semantic web search engine in order to retrieve ontologies that contain both tags.

All combinations of pairs are tried, since it is not possible to know within which pairs a relation holds (look for matches with labels and identifiers).

(2) If any of the tags is not found by the search engine, consider that they can be acronyms, misspellings or variations of known terms, and look for them in additional resources like wikipedia or Google spell correction.

(3) If the two tags (or the corresponding terms selected from Wikipedia or Google) are not found together by the semantic web search engine, consider them not to be related and eliminate the pair from that cluster if they are not (possibly) related to any other tags, that is, all the combinations of pairs of tags must be searched.

(4) Conversely, if ontologies are found containing the two tags: (a) Check whether the tags were correctly mapped into elements of the ontologies. Tags can refer to the following elements: concepts, instances, or properties. (b) Retrieve information about the tags in each of the ontologies: the type of tag (concept, instance, property), its parents (up to three levels) if it is a concept or an instance, and its domain and range or value if it is a property.

(5) For each pair of tags for which the semantic web search engine retrieved information, investigate possible relationships between them: (a) A tag is an ancestor of the other. (b) A tag is the range or the value of one of the properties of the other tag. (c) Both tags have the same direct parent. (d) Both tags have the same ancestors, at the same level. (e) Both tags have the same ancestors, at different levels.

## **TAGMAS: Federated Tagging System**

Iturrioz et al. [25] propose the TAGMAS (TAG Management System) architecture: a federation system that supplies a uniform view of tagged resources distributed across a range of Web2.0 platforms. The TAGMAS system addresses the problem that users do not have consistent view of their resources or a single query end-point with which to search them. The TAGMAS architecture is based on a tagging ontology that provides a homogeneous representation of tags and tagging events. By aggregating user tagging events that span multiple sites, such as Flickr and Delicious, it is possible to query TAGMAS using SPARQL<sup>15</sup>, enabling users to find resources distributed across many sites by their tags, the date when tagged, which site they were tagged in. This declarative query is gradually transformed into a set of distinct invocations where the specificities of each folkserver (data model, location or envelop protocol) is considered. The results are then transported back where details about the envelop protocol or location are gradually removed till raw re-

sources matching the query are rendered to the user which ignore where the resource is located. This heterogeneity stems from four main sources, mainly, the data model, the API model, the enveloped model (REST, XML-RPC or SOAP) and the site place.

Applications of their system include automatic tag creation (which permit to create desktop-specific tags), folksonomy loading (which permit to import a folksonomy from a folkserver), resource annotation (where a resource can be annotated along loaded folksonomies) and resource searching (where tag-based filtering is used to locate resources regardless of where the resource is held). This facility is parameterized for the folkservers whose folksonomy has been downloaded into the desktop.

## Correlating User Profiles from Different Folksonomies

Szomszor et al. [55] compare user tag-clouds from multiple folksonomies to: (1) show how they tend to overlap, regardless of the focus of the folksonomy (2) demonstrate how this comparison helps finding and aligning the user's separate identities, and (3) show that cross-linking distributed user tag-clouds enriches users profiles. During this process, they find that significant user interests are often reflected in multiple Web2.0 profiles, even though they may operate over different domains. However, due to the free-form nature of tagging, some correlations are lost, a problem they address through the implementation and evaluation of a user tag filtering architecture.

Out of the 84851 distinct Delicious tags, and 149529 distinct Flickr tags, 28550 are used in both systems. To measure the alignment between two user tag clouds, they measure the frequency of tags common to Delicious and Flickr. As users tag more resources in Flickr and Delicious, their intersection frequency will increase. Therefore, this increases the confidence that two correlated profiles in Delicious and Flickr refer to the same individual as their total intersection frequency increases.

## 10. Tagging problems

Though tags are useful, exploiting them for different applications is not easy. Tags suffer from problems like spamming, canonicalization and ambiguity issues. Other problems such as sparsity, no consensus, etc. are also critical. In this section, we discuss these problems and suggest solutions described in the literature.

## Spamming

Spammers can mis-tag resources to promote their own interests. Wetzer et al. [57] have observed such phenomena where a single user labeled a large number of bookmarks with the same tags all referring to the same blog site. They have also observed a phenomenon where users upload thousands of bookmarks within minutes and rarely actively contribute again. They characterize the spammers as possessing these properties: very high activity, tagging objects belonging to a few domains, high tagging rate per resource, and bulk posts. To detect such spamming, they propose a new concept called diffusion of attention which helps to reduce the influence of spam on the distribution of tags without the actual need of filtering. They define the attention a tag achieves in a certain period of time as the number of users using the tag in this period. The diffusion for a tag is then given as the number of users that assign this tag for the first time. This measures the importance of an item by its capability to attract new users while putting all users on an equal footage. Every user's influence is therefore limited and a trend can only be created by user groups.

Koutrika et al. [29] study the problem of spamming extensively. How many malicious users can a tagging system tolerate before results significantly degrade? What types of tagging systems are more vulnerable to malicious attacks? What would be the effort and the impact of employing a trusted moderator to find bad postings? Can a system automatically protect itself from spam, for instance, by exploiting user tag patterns? In a quest for answers to these questions, they introduce a framework for modeling tagging systems and user tagging behavior. The framework combines legitimate and malicious tags. This model can study a range of user tagging behaviors, including the level of moderation and the extent of spam tags, and compare different query answering and spam protection schemes. They describe a variety of query schemes and moderator strategies to counter tag spam. Particularly, they introduce a social relevance ranking method for tag search results that takes into account how often a user's postings coincide with others' postings in order to determine their "reliability". They define a metric for quantifying the "spam impact" on results. They compare the various schemes under different models for malicious user behavior. They try to understand the weaknesses of existing systems and the magnitude of the tag spam problem. They also make predictions about which schemes will be more useful and which malicious behaviors will be more disruptive in practice.

## Canonicalization and Ambiguities

Ambiguity arises in folksonomies because different users apply terms to documents in different ways. Acronyms can also lead to ambiguities. Users often combine multiple words as a single tag, without spaces, e.g., ‘vertigovideostillsbbc’ on Flickr. Currently, tags are generally defined as single words or compound words, which means that information can be lost during the tagging process. Single-word tags lose the information that would generally be encoded in the word order of a phrase. There is no synonym or homonym control in folksonomies. Different word forms, plural and singular, are also often both present. Folksonomies provide no formal guidelines for the choice and form of tags, such as the use of compound headings, punctuation, word order. In addition, the different expertise and purposes of tagging participants may result in tags that use various levels of abstraction to describe a resource.

Guy and Tonkin [17] point out the existence of useless tags due to misspellings, bad encoding like an unlikely compound word grouping (e.g., TimBernersLee); tags that do not follow convention in issues such as case and number; personal tags that are without meaning to the wider community (e.g., mydog); single use tags that appear only once in the database (e.g., billybobsdog), symbols used in tags. Conventions have become popular, such as dates represented according to the ISO standard (e.g., 20051201 for “1st December, 2005”) and the use of the year as a tag. One wildly popular convention is geotagging, a simple method of encoding latitude and longitude within a single tag; this represented over 2% of the total tags sampled on Flickr. A common source of “misspelt” tags was in the transcoding of other alphabets or characters.

Zubiaga [64] suggests a solution to the canonicalization problem. To merge all forms of the same tag, the system can rely on a method like that by Librarything. This site allows users to define relations between tags, indicating that some of them have the same meaning. In his blog entry, Lars Pind [39] has suggested various ways to solve canonicalization problem, including the following: (1) suggest tags for user, (2) find synonyms automatically, (3) help user use the same tags that others use, (4) infer hierarchy from the tags, and (5) make it easy to adjust tags on old content. Quintarelli [41] mentions that the system can have a correlation feature that, given a tag, shows related tags, i.e., tags that people have used in conjunction with the given tag to describe the same item. Guy and Tonkin [17] suggest educating the users, simple errorchecking in systems when tags are entered by users, making tag suggestions (synonyms, expansion of acronyms etc.) when users submit resources (e.g., using Scrumptious, a recent Firefox extension, offers popular tags

from Delicious for every URL). They also suggest creation of discussion tools through which users can share reasons for tagging things in a certain way. More understanding of who is submitting certain tags could possibly alter personal rating of posts by other users.

## Other Problems

There are many other problems related to social tagging, including sparsity, no consensus and search inefficiency. Sparsity is related to the annotation coverage of the data set. Bao et al. [2] point out that certain pages may not be tagged at all. Users do not generally associate tags to newly emerging web pages or web pages that can be accessed easily from hub pages, or uninteresting web pages. Noll et al. [37] observe that the amount of new information provided by metadata (tags, anchor text, search keywords) is comparatively low. All three types of data stay below 6% novelty for about 90% of documents. Search keywords dominate tags which in turn dominate anchor text words. Tags are generally more diverse than anchor texts. On one hand, this result suggests that tags are noisier than anchor texts and therefore potentially less useful. On the other hand, the diversity of tags could be an advantage since it might capture information and meanings that anchor texts miss.

Halpin et al. [18] point that users may not reach a consensus over the appropriate set of tags for a resource leading to an unstable system. As Golder and Huberman [16] suggest, changes in the stability of such patterns might indicate that groups of users are migrating away from a particular consensus on how to characterize a site and its content or negotiating the changing meaning of that site. Quintarelli [41] points out that tags have no hierarchy. Folksonomies are a flat space of keywords. Folksonomies have a very low findability quotient. They are great for serendipity and browsing but not aimed at a targeted approach or search. Tags do not scale well if you are looking for specific targeted items.

## 11. Conclusion and Future Directions

In this work, we surveyed social tagging with respect to different aspects. We discussed different user motivations and different ways of tagging web objects. We presented a summary of the various tag generation models. We analyzed different tagging system parameters and tagging distributions. We then summarized ways of identifying tag semantics, ways of visualizing tags using tag clouds and ways of recommending tags to users. We presented a variety of applications for which tags have been used. Further, we discussed ways of integrating different folksonomies and problems related to the usage of tags.

Tags are taking on a new meaning as other forms of media like microblogs are gaining popularity. Below we mention a few aspects which can be a part of future research.

## **Analysis**

Most current research on tagging analysis focus on one single tag stream itself. However, as the type of user generated content evolves, tags may be different and related to different kinds of user generated data, such as microblogs and query logs. For example, How does tag growth differ in microblogs versus that for bookmarks and images. Tagging models for microblogs can be quite different from other tagging models. E.g., certain tags reach a peak on twitter quite unexpectedly. These tags don't relate to any specific events. Such varying degree of social influence when a pseudo event happens hasn't been captured by any of the tagging models, yet.

## **Improving System Design**

Current tagging systems only support a type of tags and researchers have developed mechanisms to extract hierarchical structures (ontology) from this flat tagging space. Systems can provide more functionality like hierarchical tags, say (programming/java), multi-word tags. A tagging system can also support a tag discussion forum where users can debate about the appropriateness of a tag for a resource. Structured tags can also be supported, i.e., allow people to tag different portions of a web page with different tags and assign key=value pairs rather than just "values". E.g., person="Mahatma Gandhi", location="Porbandar", year="1869", event="birth". By adding more such functionality into the system, we can expect that a more meaningful semantic structure could be extracted.

## **Personalized Tag Recommendations**

Is the user a describer or organizer? What is the context? Is she tagging on sets in Flickr or just photos in the photostream (i.e. context within the tagging site itself)? Based on her history, what is the probability that she would choose a new tag? What are the words used in her previous tags, words used in her social friends' tags? Given some tags tied to a resource, we can identify whether users prefer to repeat tags for this resource or do they like to put on new tags. Using this we can vary the tag history window size shown with the resource. Apart from tag recommendation, a recommendation system can also recommend related resources once a user selects a particular tag.

## More Applications

There are also interesting applications which are worth exploring. E.g., (1) Tagging support for desktop systems using online tags. (2) Geographical/demographical analysis of users' sentiments based on the tags they apply to products launched at a particular location. (3) Mashups by integrating resources with same/similar tags. (4) Establishing website trustworthiness based on what percent of the keywords mentioned in the <meta> tag are actual tags for web page bookmarks. (5) Summary generation using tags with NLP. (6) Intent detection and behavioral targeting based on user history of tags.

## Dealing With Problems

Sparsity, canonicalization, ambiguities in tags still remain as open problems. More work needs to be done to come up with solutions to effectively deal with them. Also, certain tags get outdated. E.g., a camera model may be marked as 'best camera'. But after two years, it no longer remains the 'best camera'. How can we clean up such kind of tag rot? Similarly, tags that haven't been repeated by another user within a time window, can be considered as personal tags and can be removed from public view.



## References

- [1] Morgan Ames and Mor Naaman. Why we tag: Motivations for annotation in mobile and online media. In *Conference on Human Factors in Computing Systems, CHI 2007*, Sam Jose, CA, April 2007.
- [2] Shenghua Bao, Guirong Xue, Xiaoyuan Wu, Yong Yu, Ben Fei, and Zhong Su. Optimizing web search using social annotations. In *WWW '07: Proceedings of the 16th international conference on World Wide Web*, pages 501–510, New York, NY, USA, 2007. ACM.
- [3] Scott Bateman, Christopher Brooks, Gordon Mccalla, and Peter Brusilovsky. Applying collaborative tagging to e-learning. In *Proceedings of the 16th International World Wide Web Conference (WWW2007)*, 2007.
- [4] Grigory Begelman, Philipp Keller, and Frank Smadja. Automated tag clustering: Improving search and exploration in the tag space, 2006.
- [5] K. Bielenberg. Groups in Social Software: Utilizing Tagging to Integrate Individual Contexts for Social Navigation. Master’s thesis, 2005.
- [6] David M. Blei, Andrew Y. Ng, and Michael I. Jordan. Latent dirichlet allocation. *J. Mach. Learn. Res.*, 3:993–1022, 2003.
- [7] Christopher H. Brooks and Nancy Montanez. Improved annotation of the blogosphere via autotagging and hierarchical clustering. In *WWW '06: Proceedings of the 15th international conference on World Wide Web*, pages 625–632, New York, NY, USA, 2006. ACM Press.
- [8] David Carmel, Haggai Roitman, and Elad Yom-Tov. Who tags the tags?: a framework for bookmark weighting. In *CIKM '09: Proceeding of the 18th ACM conference on Information and knowledge management*, pages 1577–1580, New York, NY, USA, 2009. ACM.

- [9] Luigi Di Caro, K. Seluk Candan, and Maria Luisa Sapino. Using tagflake for condensing navigable tag hierarchies from tag clouds. In Ying Li, Bing Liu, and Sunita Sarawagi, editors, *KDD*, pages 1069–1072. ACM, 2008.
- [10] Ciro Cattuto, Andrea Baldassarri, Vito D. P. Servedio, and Vittorio Loreto. Vocabulary growth in collaborative tagging systems, Apr 2007.
- [11] Ciro Cattuto, Vittorio Loreto, and Luciano Pietronero. Semiotic dynamics and collaborative tagging. *Proceedings of the National Academy of Sciences (PNAS)*, 104(5):1461–1464, January 2007.
- [12] Klaas Dellschaft and Steffen Staab. An epistemic dynamic model for tagging systems. In *HT '08: Proceedings of the nineteenth ACM conference on Hypertext and hypermedia*, pages 71–80, New York, NY, USA, 2008. ACM.
- [13] Nicholas Diakopoulos and Patrick Chiu. Photoplay: A collocated collaborative photo tagging game on a horizontal display. preprint (2007) available at <http://www.fxpall.com/publications/FXPAL-PR-07-414.pdf>.
- [14] Pavel A. Dmitriev, Nadav Eiron, Marcus Fontoura, and Eugene Shekita. Using annotations in enterprise search. In *WWW '06: Proceedings of the 15th international conference on World Wide Web*, pages 811–817, New York, NY, USA, 2006. ACM.
- [15] Micah Dubinko, Ravi Kumar, Joseph Magnani, Jasmine Novak, Prabhakar Raghavan, and Andrew Tomkins. Visualizing tags over time. In *WWW '06: Proceedings of the 15th international conference on World Wide Web*, pages 193–202, New York, NY, USA, 2006. ACM Press.
- [16] Scott Golder and Bernardo A. Huberman. The structure of collaborative tagging systems, Aug 2005.
- [17] Marieke Guy and Emma Tonkin. Folksonomies: Tidying up tags? *D-Lib Magazine*, 12, Jan 2006.
- [18] Harry Halpin, Valentin Robu, and Hana Shepherd. The complex dynamics of collaborative tagging. In *WWW '07: Proceedings of the 16th international conference on World Wide Web*, pages 211–220, New York, NY, USA, 2007. ACM.

- [19] Tony Hammond, Timo Hannay, Ben Lund, and Joanna Scott. Social bookmarking tools (i): A general review. *D-Lib Magazine*, 11(4), April 2005.
- [20] Y. Hassan-Montero and V. Herrero-Solana. Improving tag-clouds as visual information retrieval interfaces. In *InScit2006: International Conference on Multidisciplinary Information Sciences and Technologies*, 2006.
- [21] M. Heckner, M. Heilemann, and C. Wolff. Personal information management vs. resource sharing: Towards a model of information behaviour in social tagging systems. In *Int'l AAAI Conference on Weblogs and Social Media (ICWSM)*, San Jose, CA, USA, May 2009.
- [22] Paul Heymann and Hector Garcia-Molina. Collaborative creation of communal hierarchical taxonomies in social tagging systems. Technical Report 2006-10, Stanford University, April 2006.
- [23] Paul Heymann, Georgia Koutrika, and Hector Garcia-Molina. Can social bookmarking improve web search? In *WSDM '08: Proceedings of the international conference on Web search and web data mining*, pages 195–206, New York, NY, USA, 2008. ACM.
- [24] Andreas Hotho, Robert Jschke, Christoph Schmitz, and Gerd Stumme. Information retrieval in folksonomies: Search and ranking. *The Semantic Web: Research and Applications*, pages 411–426, 2006.
- [25] Jon Iturrioz, Oscar Diaz, and Cristobal Arellano. Towards federated web2.0 sites: The tagmas approach. In *Tagging and Metadata for Social Information Organization Workshop, WWW07*, 2007.
- [26] Owen Kaser and Daniel Lemire. Tag-cloud drawing: Algorithms for cloud visualization, May 2007.
- [27] Margaret E. I. Kipp and Grant D. Campbell. Patterns and inconsistencies in collaborative tagging systems : An examination of tagging practices. In *Annual General Meeting of the American Society for Information Science and Technology*. American Society for Information Science and Technology, November 2006.
- [28] Christian Körner. Understanding the motivation behind tagging. ACM Student Research Competition - Hypertext 2009, July 2009.
- [29] Georgia Koutrika, Frans A. Effendi, Zolt'n Gyöngyi, Paul Heymann, and Hector G. Molina. Combating spam in tagging systems: An evaluation. *ACM Trans. Web*, 2(4):1–34, 2008.

- [30] Liz Lawley. social consequences of social tagging. Web article, 2005.
- [31] Rui Li, Shenghua Bao, Yong Yu, Ben Fei, and Zhong Su. Towards effective browsing of large scale social annotations. In *WWW '07: Proceedings of the 16th international conference on World Wide Web*, pages 943–952, New York, NY, USA, 2007. ACM.
- [32] Xin Li, Lei Guo, and Yihong E. Zhao. Tag-based social interest discovery. In *WWW '08: Proceeding of the 17th international conference on World Wide Web*, pages 675–684, New York, NY, USA, 2008. ACM.
- [33] Dong Liu, Xian S. Hua, Linjun Yang, Meng Wang, and Hong J. Zhang. Tag ranking. In *WWW '09: Proceedings of the 18th international conference on World wide web*, pages 351–360, New York, NY, USA, April 2009. ACM.
- [34] Cameron Marlow, Mor Naaman, Danah Boyd, and Marc Davis. Ht06, tagging paper, taxonomy, flickr, academic article, toread. In *HYPERTEXT '06: Proceedings of the seventeenth conference on Hypertext and hypermedia*, pages 31–40, New York, NY, USA, 2006. ACM.
- [35] A. Mathes. Folksonomies - cooperative classification and communication through shared metadata. *Computer Mediated Communication*, December 2004.
- [36] David R. Millen and Jonathan Feinberg. Using social tagging to improve social navigation. In *Workshop on the Social Navigation and Community based Adaptation Technologies*, 2006.
- [37] Michael G. Noll and Christoph Meinel. The metadata triumvirate: Social annotations, anchor texts and search queries. *WI/IAT '08: Web Intelligence and Intelligent Agent Technology, IEEE/WIC/ACM International Conference on*, 1:640–647, 2008.
- [38] Simon Overell, Börkur Sigurbjörnsson, and Roelof van Zwol. Classifying tags using open content resources. In *WSDM '09: Proceedings of the Second ACM International Conference on Web Search and Data Mining*, pages 64–73, New York, NY, USA, 2009. ACM.
- [39] Lars Pind. Folksonomies: How we can improve the tags. Web article, 2005.
- [40] Peter Pirolli. Rational analyses of information foraging on the web. *Cognitive Science*, 29(3):343–373, 2005.

- [41] Emanuele Quintarelli. Folksonomies: power to the people.
- [42] Daniel Ramage, Paul Heymann, Christopher D. Manning, and Hector G. Molina. Clustering the tagged web. In *WSDM '09: Proceedings of the Second ACM International Conference on Web Search and Data Mining*, pages 54–63, New York, NY, USA, 2009. ACM.
- [43] Tye Rattenbury, Nathaniel Good, and Mor Naaman. Towards automatic extraction of event and place semantics from flickr tags. In *SIRIR '07: Proceedings of the 30th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 103–110, New York, NY, USA, 2007. ACM Press.
- [44] Terrell Russell. cloudalicious: folksonomy over time. In *JCDL '06: Proceedings of the 6th ACM/IEEE-CS joint conference on Digital libraries*, pages 364–364, New York, NY, USA, 2006. ACM.
- [45] Ralf Schenkel, Tom Crecelius, Mouna Kacimi, Sebastian Michel, Thomas Neumann, Josiane X. Parreira, and Gerhard Weikum. Efficient top-k querying over social-tagging networks. In *SIGIR '08: Proceedings of the 31st annual international ACM SIGIR conference on Research and development in information retrieval*, pages 523–530, New York, NY, USA, 2008. ACM.
- [46] Christoph Schmitz, Andreas Hotho, Robert Jschke, and Gerd Stumme. Mining association rules in folksonomies. In *Data Science and Classification*, pages 261–270. Springer, 2006.
- [47] Patrick Schmitz. Inducing ontology from flickr tags. In *WWW 2006 WWW 2006, May 2226, 2006, Edinburgh, UK. IW3C2*, 2006.
- [48] Shilad Sen, Shyong K. Lam, Al Mamunur Rashid, Dan Cosley, Dan Frankowski, Jeremy Osterhouse, F. Maxwell Harper, and John Riedl. tagging, communities, vocabulary, evolution. In *CSCW '06: Proceedings of the 2006 20th anniversary conference on Computer supported cooperative work*, pages 181–190, New York, NY, USA, November 2006. ACM.
- [49] Börkur Sigurbjörnsson and Roelof van Zwol. Flickr tag recommendation based on collective knowledge. In *WWW '08: Proceeding of the 17th International Conference on World Wide Web*, pages 327–336, New York, NY, USA, 2008. ACM.
- [50] Herbert A. Simon. On a class of skew distribution functions. *Biometrika*, 42(3/4):425–440, 1955.

- [51] James Sinclair and Michael Cardew-Hall. The folksonomy tag cloud: when is it useful? *J. Inf. Sci.*, 34(1):15–29, 2008.
- [52] Yang Song, Ziming Zhuang, Huajing Li, Qiankun Zhao, Jia Li, Wang C. Lee, and C. Lee Giles. Real-time automatic tag recommendation. In *SIGIR '08: Proceedings of the 31st annual international ACM SIGIR conference on Research and development in information retrieval*, pages 515–522, New York, NY, USA, 2008. ACM.
- [53] Lucia Specia and Enrico Motta. Integrating folksonomies with the semantic web. In *ESWC '07: Proceedings of the 4th European conference on The Semantic Web*, pages 624–639, Berlin, Heidelberg, 2007. Springer-Verlag.
- [54] Louise F. Spiteri. Structure and form of folksonomy tags: The road to the public library catalogue. *Webology*, 4(2, Artikel 41), 2007.
- [55] Martin Szomszor, Ivan Cantador, and Harith Alani. Correlating user profiles from multiple folksonomies. In *ACM Conference on Hypertext and Hypermedia*, June 2008.
- [56] Csaba Veres. The language of folksonomies: What tags reveal about user classification. In *Natural Language Processing and Information Systems*, volume 3999/2006 of *Lecture Notes in Computer Science*, pages 58–69, Berlin / Heidelberg, July 2006. Springer.
- [57] Robert Wetzker, Carsten Zimmermann, and Christian Bauckhage. Analyzing social bookmarking systems: A del.icio.us cookbook. In *Proceedings of the ECAI 2008 Mining Social Data Workshop*, pages 26–30. IOS Press, 2008.
- [58] Lei Wu, Linjun Yang, Nenghai Yu, and Xian S. Hua. Learning to tag. In *WWW '09: Proceedings of the 18th international conference on World wide web*, pages 361–370, New York, NY, USA, 2009. ACM.
- [59] X. Wu, L. Zhang, and Y. Yu. Exploring social annotations for the semantic web. In *Proc. of the 15th International Conference on World Wide Web (WWW'06)*, 2006.
- [60] Shengliang Xu, Shenghua Bao, Ben Fei, Zhong Su, and Yong Yu. Exploring folksonomy for personalized search. In *SIGIR '08: Proceedings of the 31st annual international ACM SIGIR conference on Research and development in information retrieval*, pages 155–162, New York, NY, USA, 2008. ACM.

- [61] Zhichen Xu, Yun Fu, Jianchang Mao, and Difu Su. Towards the semantic web: Collaborative tag suggestions. In *WWW2006: Proceedings of the Collaborative Web Tagging Workshop*, Edinburgh, Scotland, 2006.
- [62] Zhijun Yin, Rui Li, Qiaozhu Mei, and Jiawei Han. Exploring social tagging graph for web object classification. In *KDD '09: Proceedings of the 15th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 957–966, New York, NY, USA, 2009. ACM.
- [63] Ding Zhou, Jiang Bian, Shuyi Zheng, Hongyuan Zha, and C. Lee Giles. Exploring social annotations for information retrieval. In *WWW '08: Proceeding of the 17th international conference on World Wide Web*, pages 715–724, New York, NY, USA, 2008. ACM.
- [64] Arkaitz Zubiaga. Enhancing navigation on wikipedia with social tags. In *Wikimania '09*, 2009.