Word Sense Disambiguation by Web Mining for Word Co-occurrence Probabilities

Peter D. TURNEY
Institute for Information Technology
National Research Council of Canada
Ottawa, Ontario, Canada, K1A 0R6
peter.turney@nrc-cnrc.gc.ca

Abstract

This paper describes the National Research Council (NRC) Word Sense Disambiguation (WSD) system, as applied to the English Lexical Sample (ELS) task in Senseval-3. The NRC system approaches WSD as a classical supervised machine learning problem, using familiar tools such as the Weka machine learning software and Brill's rule-based part-of-speech tagger. Head words are represented as feature vectors with several hundred features. Approximately half of the features are syntactic and the other half are semantic. The main novelty in the system is the method for generating the semantic features, based on word co-occurrence probabilities. The probabilities are estimated using the Waterloo MultiText System with a corpus of about one terabyte of unlabeled text, collected by a web crawler.

1 Introduction

The Senseval-3 English Lexical Sample (ELS) task requires disambiguating 57 words, with an average of roughly 140 training examples and 70 testing examples of each word. Each example is about a paragraph of text, in which the word that is to be disambiguated is marked as the *head* word. The average head word has around six senses. The training examples are manually classified according to the intended sense of the head word, inferred from the surrounding context. The task is to use the training data and any other relevant information to automatically assign classes to the testing examples.

This paper presents the National Research Council (NRC) Word Sense Disambiguation (WSD) system, which generated our four entries for the Senseval-3 ELS task (NRC-Fine, NRC-Fine2, NRC-Coarse, and NRC-Coarse2). Our approach to the ELS task is to treat it as a classical supervised machine learning problem. Each example is represented as a feature vector with several hundred features. Each of the 57 ambiguous words is represented with a different set of features. Typically, around half of the features are syntactic and the other half are semantic. After the raw examples are converted to feature vectors, the Weka machine learning software is used to induce a model of the training data and predict the classes of the testing examples (Witten and Frank, 1999).

The syntactic features are based on part-of-speech tags, assigned by a rule-based tagger (Brill, 1994). The main innovation of the NRC WSD system is the method for generating the semantic features, which are derived from word co-occurrence probabilities. We estimated these probabilities using the Waterloo MultiText System with a corpus of about one terabyte of unlabeled text, collected by a web crawler (Clarke et al., 1995; Clarke and Cormack, 2000; Terra and Clarke, 2003).

In Section 2, we describe the NRC WSD system. Our experimental results are presented in Section 3 and we conclude in Section 4.

2 System Description

This section presents various aspects of the system in roughly the order in which they are executed. The following definitions will simplify the description.

Head Word: One of the 57 words that are to be disambiguated.

Example: One or more contiguous sentences, illustrating the usage of a head word.

Context: The non-head words in an example.

Feature: A property of a head word in a context. For instance, the feature `tag_hp1_NNP` is the property of having (or not having) a proper noun (`NNP` is the part-of-speech tag for a proper noun) immediately following the head word (`hp1` represents the location *head plus one*).

Feature Value: Features have values, which depend on the specific example. For instance, `tag_hp1_NNP` is a binary feature that has the value 1 (*true*: the following word *is* a proper noun) or 0 (*false*: the following word is *not* a proper noun).

Feature Vector: Each example is represented by a vector. Features are the dimensions of the vector

space and a vector of feature values specifies a point in the feature space.

2.1 Preprocessing

The NRC WSD system first assigns part-of-speech tags to the words in a given example (Brill, 1994), and then extracts a nine-word window of tagged text, centered on the head word (i.e., four words before and after the head word). Any remaining words in the example are ignored (usually most of the example is ignored). The window is not allowed to cross sentence boundaries. If the head word appears near the beginning or end of the sentence, where the window may overlap with adjacent sentences, special *null* characters fill the positions of any missing words in the window.

In rare cases, a head word appears more than once in an example. In such cases, the system selects a single window, giving preference to the earliest occurring window with the least nulls. Thus each example is converted into one nine-word window of tagged text. Windows from the training examples for a given head word are then used to build the feature set for that head word.

2.2 Syntactic Features

Each head word has a unique set of feature names, describing how the feature values are calculated.

Feature Names: Every syntactic feature has a name of the form *matchtype_position_model*. There are three *matchtypes*, ptag, tag, and word, in order of increasingly strict matching. A ptag match is a *partial tag match*, which counts similar part-of-speech tags, such as NN (singular noun), NNS (plural noun), NNP (singular proper noun), and NNPS (plural proper noun), as equivalent. A tag match requires exact matching in the part-of-speech tags for the word and the model. A word match requires that the word and the model are exactly the same, letter-for-letter, including upper and lower case.

There are five *positions*, hm2 (head minus two), hm1 (head minus one), hd0 (head), hp1 (head plus one), and hp2 (head plus two). Thus syntactic features use only a five-word sub-window of the nine-word window.

The syntactic feature names for a head word are generated by all of the possible legal combinations of *matchtype*, *position*, and *model*. For ptag names, the *model* can be any partial tag. For tag names, the *model* can be any tag. For word names, the *model* names are not predetermined; they are extracted from the training windows for the given head word. For instance, if a training window contains the head word followed by "of", then one of the features will be word_hp1_of.

For word names, the *model* names are not allowed to be words that are tagged as nouns, verbs, or adjectives. These words are reserved for use in building the semantic features.

Feature Values: The syntactic features are all binary-valued. Given a feature with a name of the form *matchtype_position_model*, the feature value for a given window depends on whether there is a match of *matchtype* between the word in the position *position* and the model *model*. For instance, the value of tag_hp1_NNP depends on whether the given window has a word in the position hp1 (head plus one) with a tag (part-of-speech tag) that matches NNP (proper noun). Similarly, the feature word_hp1_of has the value 1 (*true*) if the given window contains the head word followed by "of"; otherwise, it has the value 0 (*false*).

2.3 Semantic Features

Each head word has a unique set of feature names, describing how the feature values are calculated.

Feature Names: Most of the semantic features have names of the form *position_model*. The *position* names can be pre (preceding) or fol (following). They refer to the nearest noun, verb, or adjective that precedes or follows the head word in the nine-word window.

The *model* names are extracted from the training windows for the head word. For instance, if a training window contains the word "compelling", and this word is the nearest noun, verb, or adjective that precedes the head word, then one of the features will be pre_compelling.

A few of the semantic features have a different form of name, avg_*position_sense*. In names of this form, *position* can be pre (preceding) or fol (following), and *sense* can be any of the possible senses (i.e., classes, labels) of the head word.

Feature Values: The semantic features are all real-valued. For feature names of the form *position_model*, the feature value depends on the semantic similarity between the word in position *position* and the model word *model*.

The semantic similarity between two words is estimated by their Pointwise Mutual Information, $\mathrm{PMI}(w_1, w_2)$, using Information Retrieval (Turney, 2001; Terra and Clarke, 2003):

$$\mathrm{PMI}(w_1, w_2) = \log_2 \left(\frac{p(w_1 \wedge w_2)}{p(w_1)p(w_2)} \right).$$

We estimate the probabilities in this equation by issuing queries to the Waterloo MultiText System (Clarke et al., 1995; Clarke and Cormack, 2000; Terra and Clarke, 2003). Laplace smoothing is

```
weka.classifiers.meta.Bagging
  -W weka.classifiers.meta.MultiClassClassifier
    -W weka.classifiers.meta.Vote
      -B weka.classifiers.functions.supportVector.SMO
      -B weka.classifiers.meta.LogitBoost -W weka.classifiers.trees.DecisionStump
      -B weka.classifiers.meta.LogitBoost -W weka.classifiers.functions.SimpleLinearRegression
      -B weka.classifiers.trees.adtree.ADTree
      -B weka.classifiers.rules.JRip
```

Table 1: Weka (version 3.4) commands for processing the feature vectors.

applied to the PMI estimates, to avoid division by zero.

$PMI(w_1, w_2)$ has a value of zero when the two words are statistically independent. A high positive value indicates that the two words tend to co-occur, and hence are likely to be semantically related. A negative value indicates that the presence of one of the words suggests the absence of the other. Past work demonstrates that PMI is a good estimator of semantic similarity (Turney, 2001; Terra and Clarke, 2003) and that features based on PMI can be useful for supervised learning (Turney, 2003).

The Waterloo MultiText System allows us to set the neighbourhood size for co-occurrence (i.e., the meaning of $w_1 \wedge w_2$). In preliminary experiments with the ELS data from Senseval-2, we got good results with a neighbourhood size of 20 words.

For instance, if w is the noun, verb, or adjective that precedes the head word and is nearest to the head word in a given window, then the value of `pre_compelling` is $PMI(w, \text{compelling})$. If there is no preceding noun, verb, or adjective within the window, the value is set to zero.

In names of the form avg_*position_sense*, the feature value is the average of the feature values of the corresponding features. For instance, the value of `avg_pre_argument_1_10_02` is the average of the values of all of the pre_*model* features, such that *model* was extracted from a training window in which the head word was labeled with the sense `argument_1_10_02`.

The idea here is that, if a testing example should be labeled, say, `argument_1_10_02`, and w_1 is a noun, verb, or adjective that is close to the head word in the testing example, then $PMI(w_1, w_2)$ should be relatively high when w_2 is extracted from a training window with the same sense, `argument_1_10_02`, but relatively low when w_2 is extracted from a training window with a different sense. Thus avg_*position*_argument_1_10_02 is likely to be relatively high, compared to other avg_*position_sense* features.

All semantic features with names of the form *position_model* are normalized by converting them to percentiles. The percentiles are calculated separately for each feature vector; that is, each feature vector is normalized internally, with respect to its own values, not externally, with respect to the other feature vectors. The pre features are normalized independently from the fol features. The semantic features with names of the form avg_*position_sense* are calculated after the other features are normalized, so they do not need any further normalization. Preliminary experiments with the ELS data from Senseval-2 supported the merit of percentile normalization, which was also found useful in another application where features based on PMI were used for supervised learning (Turney, 2003).

2.4 Weka Configuration

Table 1 shows the commands that were used to execute Weka (Witten and Frank, 1999). The default parameters were used for all of the classifiers. Five base classifiers (-B) were combined by voting. Multiple classes were handled by treating them as multiple two-class problems, using a 1-against-all strategy. Finally, the variance of the system was reduced with bagging.

We designed the Weka configuration by evaluating many different Weka base classifiers on the Senseval-2 ELS data, until we had identified five good base classifiers. We then experimented with combining the base classifiers, using a variety of meta-learning algorithms. The resulting system is somewhat similar to the JHU system, which had the best ELS scores in Senseval-2 (Yarowsky et al., 2001). The JHU system combined four base classifiers using a form of voting, called Thresholded Model Voting (Yarowsky et al., 2001).

2.5 Postprocessing

The output of Weka includes an estimate of the probability for each prediction. When the head word is frequently labeled U (unassignable) in the training examples, we ignore U examples during training, and then, after running Weka, relabel the lowest probability testing examples as U.

System	Fine-Grained Recall	Coarse-Grained Recall
Best Senseval-3 System	72.9%	79.5%
NRC-Fine	69.4%	75.9%
NRC-Fine2	69.1%	75.6%
NRC-Coarse	NA	75.8%
NRC-Coarse2	NA	75.7%
Median Senseval-3 System	65.1%	73.7%
Most Frequent Sense	55.2%	64.5%

Table 2: Comparison of NRC-Fine with other Senseval-3 ELS systems.

3 Results

A total of 26 teams entered 47 systems (both supervised and unsupervised) in the Senseval-3 ELS task. Table 2 compares the fine-grained and coarse-grained scores of our four entries with other Senseval-3 systems.

With NRC-Fine and NRC-Coarse, each semantic feature was scored by calculating its PMI with the head word, and then low scoring semantic features were dropped. With NRC-Fine2 and NRC-Coarse2, the threshold for dropping features was changed, so that many more features were retained. The Senseval-3 results suggest that it is better to drop more features.

NRC-Coarse and NRC-Coarse2 were designed to maximize the coarse score, by training them with data in which the senses were relabeled by their coarse sense equivalence classes. The fine scores for these two systems are meaningless and should be ignored. The Senseval-3 results indicate that there is no advantage to relabeling.

The NRC systems scored roughly midway between the best and median systems. This performance supports the hypothesis that corpus-based semantic features can be useful for WSD. In future work, we plan to design a system that combines corpus-based semantic features with the most effective elements of the other Senseval-3 systems.

For reasons of computational efficiency, we chose a relatively narrow window of nine-words around the head word. We intend to investigate whether a larger window would bring the system performance up to the level of the best Senseval-3 system.

4 Conclusion

This paper has sketched the NRC WSD system for the ELS task in Senseval-3. Due to space limitations, many details were omitted, but it is likely that their impact on the performance is relatively small.

The system design is relatively straightforward and classical. The most innovative aspect of the system is the set of semantic features, which are purely corpus-based; no lexicon was used.

References

Eric Brill. 1994. Some advances in transformation-based part of speech tagging. In *Proceedings of the 12th National Conference on Artificial Intelligence (AAAI-94)*, pages 722–727.

Charles L.A. Clarke and Gordon V. Cormack. 2000. Shortest substring retrieval and ranking. *ACM Transactions on Information Systems (TOIS)*, 18(1):44–78.

Charles L.A. Clarke, G.V. Cormack, and F.J. Burkowski. 1995. An algebra for structured text search and a framework for its implementation. *The Computer Journal*, 38(1):43–56.

Egidio L. Terra and Charles L.A. Clarke. 2003. Frequency estimates for statistical word similarity measures. In *Proceedings of the Human Language Technology and North American Chapter of Association of Computational Linguistics Conference 2003 (HLT/NAACL 2003)*, pages 244–251.

Peter D. Turney. 2001. Mining the Web for synonyms: PMI-IR versus LSA on TOEFL. In *Proceedings of the Twelfth European Conference on Machine Learning (ECML-2001)*, pages 491–502.

Peter D. Turney. 2003. Coherent keyphrase extraction via Web mining. In *Proceedings of the Eighteenth International Joint Conference on Artificial Intelligence (IJCAI-03)*, pages 434–439.

Ian H. Witten and Eibe Frank. 1999. *Data Mining: Practical Machine Learning Tools and Techniques with Java Implementations*. Morgan Kaufmann, San Mateo, CA.

D. Yarowsky, S. Cucerzan, R. Florian, C. Schafer, and R. Wicentowski. 2001. The Johns Hopkins SENSEVAL2 system descriptions. In *Proceedings of SENSEVAL2*, pages 163–166.