

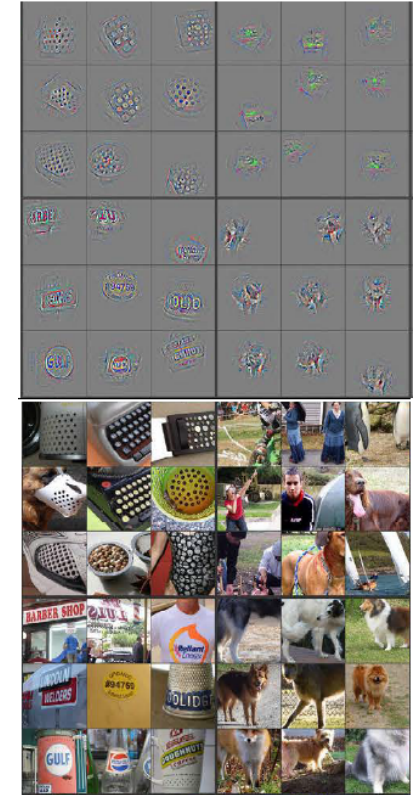

ECS289 VISUAL RECOGNITION

Intriguing properties of neural networks

Wei-Chih Chen(Michael)

2015/10/22

Introduction





- We know already:
 - (1) Math: loss function, backpropagation
 - (2) how to train: by mini batch, stochastic gradient descent
 - (3) Implement details: Max pooling, Relu, dropout, architecture
 - (4) Visualize features from each layers

Introduction

- We do NOT know:
Relation between each layers.
- This paper proposed:
 - Space rather than individual units contain semantic info.**
 - misclassify an image by applying imperceptible perturbation**

Space?

- Space rather than individual units contain semantic info
- Representation Φ as an function mapping an image x to feature space.



$m=784$
for MNIST

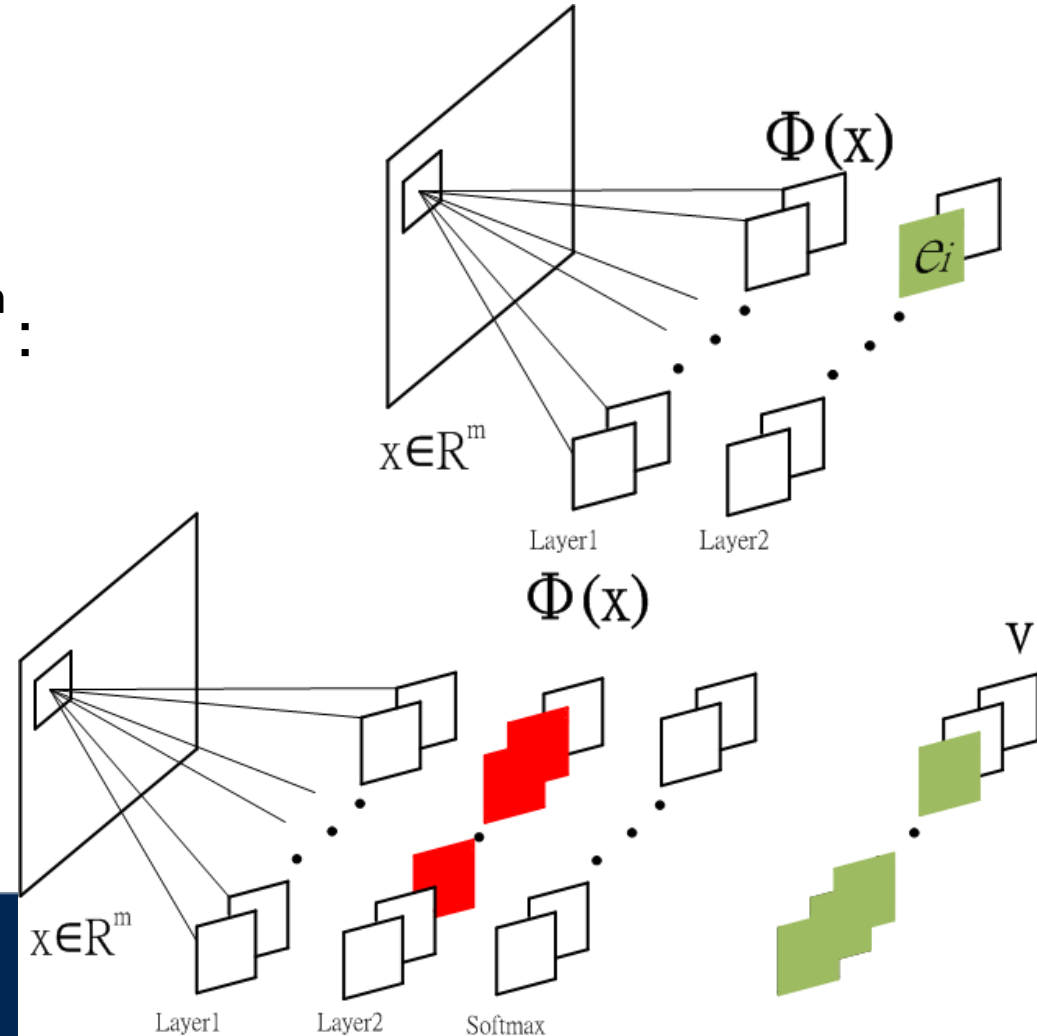
Space?


- Using the natural basis of the i-th hidden unit:

$$x' = \arg \max_{x \in \mathcal{I}} \langle \phi(x), e_i \rangle$$


- Feature vector direction $v \in \mathbb{R}^n$:

$$x' = \arg \max_{x \in \mathcal{I}} \langle \phi(x), v \rangle$$





- Natural basis direction



- Random direction

Result on MNIST

- Natural basis direction



(b) Unit sensitive to upper round stroke, or lower straight stroke.



(d) Unit sensitive to diagonal straight stroke.

- Random basis



(b) Direction sensitive to lower left loop.



(d) Direction sensitive to right, upper round stroke.

Result on ImageNet

- Natural basis direction



(a) Unit sensitive to white flowers.



(c) Unit sensitive to round, spiky flowers.

- Random basis




(a) Direction sensitive to white, spread flowers.



(c) Direction sensitive to spread shapes.



Conclude for first question

- The vector representations are well-defined up to rotation of the space, so the **individual units are unlikely to contain semantic information.**




Blind Spots

- Networks level contains semantic info.
- Output is highly nonlinear function of its input.
- Encoded a non-local generalization prior over input to put non-significant probability to some region without misclassified.



Local generalization

- Input \mathbf{x} , and generate \mathbf{x}_ϵ which satisfies $\|\mathbf{x} - \mathbf{x}_\epsilon\| < \epsilon$, ϵ is a small enough radius.
- \mathbf{x}_ϵ should NOT change underlying class.
- This paper want find *adversarial examples* that cause **MISLABEL**.




(a) All misLabeled to ostrich (b)

Function define



- Minimize $\|r\|_2$ subject to:

- $f(x + r) = l$
- $x + r \in [0, 1]^m$



- A minimizer by $D(x, l)$, if $x+r$ is close to x , then x is classified as l by f . Like, $D(x, f(x)) = f(x)$
- Using a box-constrained L-BFGS for optimization
- Our Goal: minimized r to get $D(x + r, f(x + r)) \neq l$

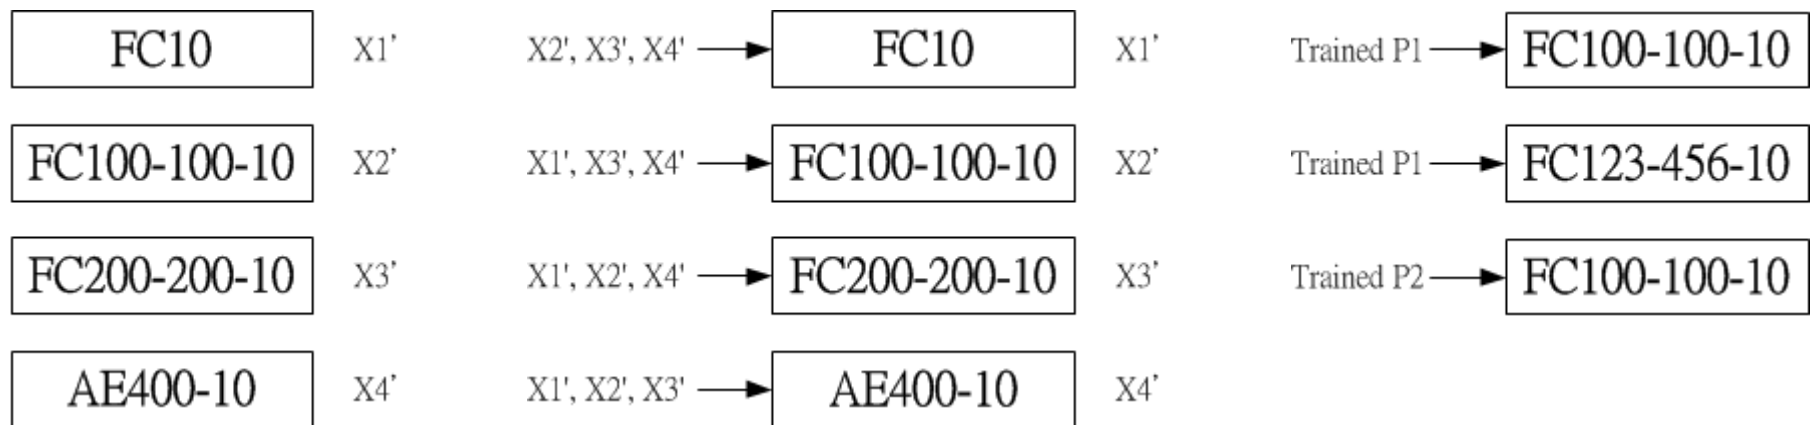
Goal






- Minimize $c|r| + \text{loss}_f(x + r, l)$ subject to $x + r \in [0, 1]^m$

Properties for distortion function D

- All networks(MNIST, AlexNet) can generate visually indistinguishable adversarial example.
- Cross model generalization: misclassified by other networks.
- Cross training-set generalization: misclassified by other disjoint training set.

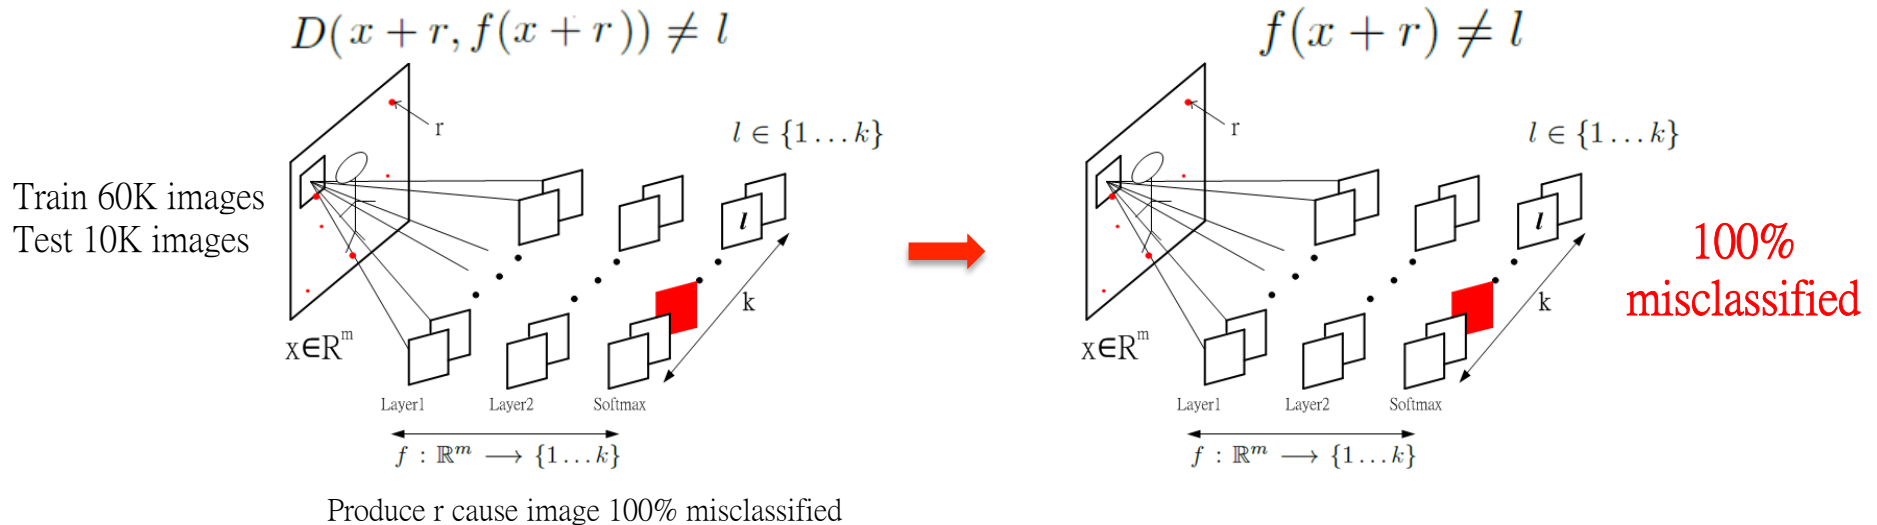


Tests to generate adversarial instance on MNIST(1/2)





Tests to generate adversarial instance on MNIST(2/2)

Model Name	Description	Training error	Test error	Av. min. distortion
FC10(10^{-4})	Softmax with $\lambda = 10^{-4}$	6.7%	7.4%	0.062
FC10(10^{-2})	Softmax with $\lambda = 10^{-2}$	10%	9.4%	0.1
FC10(1)	Softmax with $\lambda = 1$	21.2%	20%	0.14
FC100-100-10	Sigmoid network $\lambda = 10^{-5}, 10^{-5}, 10^{-6}$	0%	1.64%	0.058
FC200-200-10	Sigmoid network $\lambda = 10^{-5}, 10^{-5}, 10^{-6}$	0%	1.54%	0.065
AE400-10	Autoencoder with Softmax $\lambda = 10^{-6}$	0.57%	1.9%	0.086



Cross model generalization on MNIST(1/2)

	FC10(10^{-4})	FC10(10^{-2})	FC10(1)	FC100-100-10	FC200-200-10	AE400-10	Av. distortion
FC10(10^{-4})	100%	11.7%	22.7%	2%	3.9%	2.7%	0.062
FC10(10^{-2})	87.1%	100%	35.2%	35.9%	27.3%	9.8%	0.1
FC10(1)	71.9%	76.2%	100%	48.1%	47%	34.4%	0.14
FC100-100-10	28.9%	13.7%	21.1%	100%	6.6%	2%	0.058
FC200-200-10	38.2%	14%	23.8%	20.3%	100%	2.7%	0.065
AE400-10	23.4%	16%	24.8%	9.4%	6.6%	100%	0.086
Gaussian noise, stddev=0.1	5.0%	10.1%	18.3%	0%	0%	0.8%	0.1
Gaussian noise, stddev=0.3	15.6%	11.3%	22.7%	5%	4.3%	3.1%	0.3



Cross model generalization on MNIST(2/2)

0% Accuracy




(a) Even columns: adversarial examples for a linear (FC) classifier (stddev=0.06)

0% Accuracy



(b) Even columns: adversarial examples for a 200-200-10 sigmoid network (stddev=0.063)

51% Accuracy



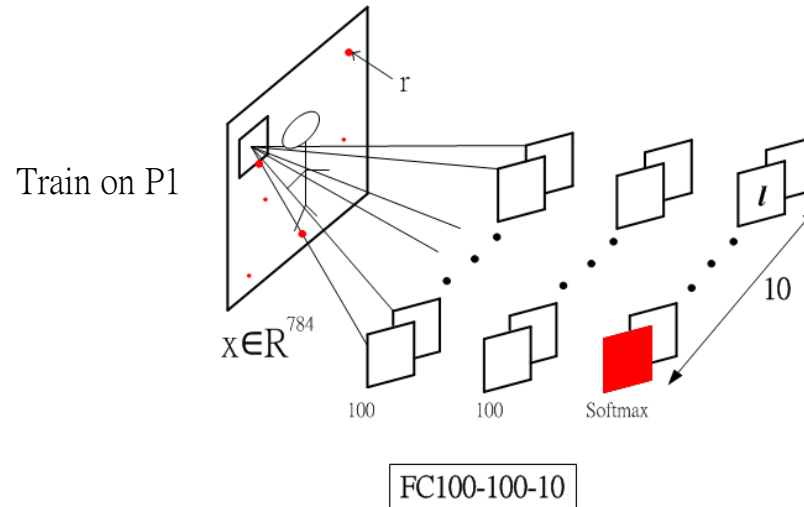

(c) Randomly distorted samples by Gaussian noise with stddev=1. Accuracy: 51%.

Cross training-set generalization on MNIST

- Cross training-set generalization – baseline

MNIST 60K images
Half P_1 and half P_2

Model	Error on P_1	Error on P_2	Error on Test	Min Av. Distortion
FC100-100-10: 100-100-10 trained on P_1	0%	2.4%	2%	0.062
FC123-456-10: 123-456-10 trained on P_1	0%	2.5%	2.1%	0.059
FC100-100-10' trained on P_2	2.3%	0%	2.1%	0.058




Cross training-set generalization on MNIST

- Cross training-set generalization error rate(magnify distortion)

	FC100-100-10	FC123-456-10	FC100-100-10'
Distorted for FC100-100-10 (av. stddev=0.062)	100%	26.2%	5.9%
Distorted for FC123-456-10 (av. stddev=0.059)	6.25%	100%	5.1%
Distorted for FC100-100-10' (av. stddev=0.058)	8.2%	8.2%	100%
Gaussian noise with stddev=0.06	2.2%	2.6%	2.4%
Distorted for FC100-100-10 amplified to stddev=0.1	100%	98%	43%
Distorted for FC123-456-10 amplified to stddev=0.1	96%	100%	22%
Distorted for FC100-100-10' amplified to stddev=0.1	27%	50%	100%
Gaussian noise with stddev=0.1	2.6%	2.8%	2.7%

$$x + 0.1 \frac{x' - x}{\|x' - x\|_2}$$



Conclusion

- Space rather than the individual units contain of the semantic information.
- Adversarial examples(imperceptible perturbation) misclassify by the network no matter type of network, cross model and cross training-set