
Incorporating Prior Knowledge into Boosting

Robert E. Schapire
Marie Rochery
Mazin Rahim
Narendra Gupta

SCHAPIRE@RESEARCH.ATT.COM
 MARIE.ROCHERY@EURECOM.FR
 MAZIN@RESEARCH.ATT.COM
 NGUPTA@RESEARCH.ATT.COM

AT&T Labs – Research, Shannon Laboratory, 180 Park Avenue, Florham Park, NJ 07932

Abstract

We describe a modification to the AdaBoost algorithm that permits the incorporation of prior human knowledge as a means of compensating for a shortage of training data. We give a convergence result for the algorithm. We describe experiments on four datasets showing that prior knowledge can substantially improve performance.

1. Introduction

Like many machine-learning methods, Freund and Schapire's [4] AdaBoost algorithm is entirely data-driven in the sense that the classifier it generates is derived exclusively from the evidence present in the training data itself. When data is abundant, this approach makes sense. However, in some applications, data may be severely limited, but there may be human knowledge that, in principle, might compensate for the lack of data.

In its standard form, boosting does not allow for the direct incorporation of such prior knowledge. In this paper, we describe a new modification of boosting that combines and balances human expertise with available training data. We aim for an approach that allows the human's rough judgments to be refined, reinforced and adjusted by the statistics of the training data, but in a manner that does not permit the data to entirely overwhelm human judgments.

The basic idea of our approach is to modify the loss function used by boosting so that the algorithm balances two terms, one measuring fit to the training data, and the other measuring fit to a human-built model. The actual algorithmic modification that this entails turns out to be very simple, only requiring the addition of weighted pseudo-examples to the training set. We allow prior knowledge that may be of any form that provides guesses, however rough, of the con-

ditional probability of class labels for each training example. We include one example of how such a model can be easily built for text categorization tasks from human-chosen keywords.

Our approach is based on the boosting-style algorithm for logistic regression described by Collins, Schapire and Singer [2], and we use their results to prove a simple convergence theorem for our algorithm.

The work in this paper arose in the development of spoken-dialogue systems at AT&T. In these systems, a computer must formulate an appropriate response to the utterances of a telephone caller. A key task is the extraction of the meaning of what the caller said to the extent that his or her utterance can be classified among a fixed set of categories. The construction of such a classifier is done using machine learning. However, in many cases, the system must be deployed before enough data has been collected; indeed, actual data cannot be easily collected until the system is actually deployed. The work in this paper permitted us to use human-crafted knowledge to compensate for this initial dearth of data until enough could be collected following deployment.

We describe experiments on datasets derived from these spoken-dialogue applications. Besides these proprietary datasets, we also conducted experiments on two benchmark datasets. In each case, we compared boosting with and without prior knowledge. The results show that prior knowledge can substantially improve performance, particularly when data is substantially limited.

2. Boosting and Logistic Regression

We begin with a review of logistic regression and the boosting-style algorithm for it described by Collins, Schapire and Singer [2]. Let \mathcal{X} and \mathcal{Y} be spaces of instances and labels, respectively. For now, we assume only two labels $\mathcal{Y} = \{-1, +1\}$. Let

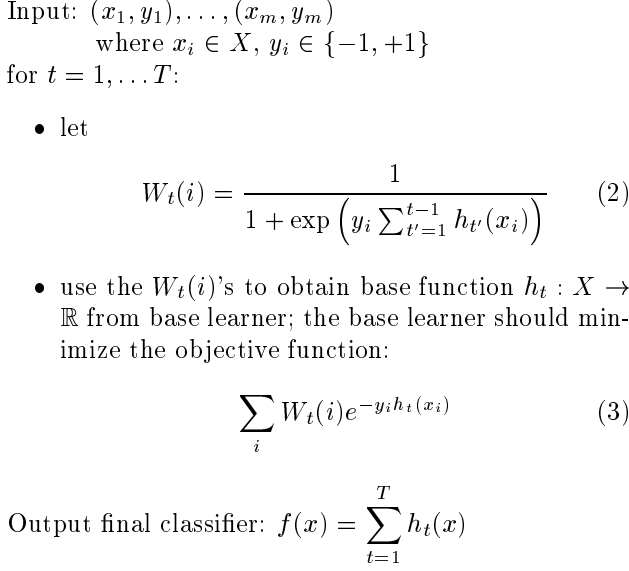


Figure 1. A binary boosting algorithm.

$(x_1, y_1), \dots, (x_m, y_m)$ be a given sequence of training examples from $\mathcal{X} \times \mathcal{Y}$. When discussing probabilities, we assume that all training and test examples are selected independently from some distribution \mathcal{D} on $\mathcal{X} \times \mathcal{Y}$.

Although our eventual goal is classification, we focus on estimating probabilities which can be converted into classifications in the obvious way by thresholding. Specifically, given training data, we wish to build a rule that estimates the conditional probability that $y = +1$ given x when test example (x, y) is chosen according to \mathcal{D} . In logistic regression, we do this by building a real-valued function $f : \mathcal{X} \rightarrow \mathbb{R}$ and estimating this probability by $\sigma(f(x))$ where

$$\sigma(z) = \frac{1}{1 + e^{-z}}.$$

Later, f will be of a particular form, namely, a linear combination of base functions. Once such a model has been postulated, we can attempt to find f by maximizing the conditional likelihood of the data, or equivalently, minimizing the negative log conditional likelihood which works out to be

$$\sum_i \ln(1 + \exp(-y_i f(x_i))). \quad (1)$$

Collins, Schapire and Singer [2] describe a variant of Freund and Schapire's [4] AdaBoost algorithm for minimizing Eq. (1) over functions f that are linear combinations of base functions. Pseudo-code for the algorithm, which we call AdaBoost.L, is shown in Figure 1.

Like AdaBoost, AdaBoost.L works in rounds. On each round, a set of weights $W_t(i)$ over the training set is computed as in Eq. (2) and used to find a base function $h_t : \mathcal{X} \rightarrow \mathbb{R}$. This base function should minimize Eq. (3) over some space of base functions; thus, we are using Schapire and Singer's [10] confidence-rated variant of AdaBoost. After T rounds, the sum of all the h_t 's is output as the final function f .

This procedure is in fact identical to confidence-rated AdaBoost if we instead compute $W_t(i)$ using the rule

$$W_t(i) = \exp\left(-y_i \sum_{t'=1}^{t-1} h_{t'}(x_i)\right).$$

2.1 Convergence

We can use the results and techniques of Collins, Schapire and Singer [2] to prove the convergence of this algorithm to the minimum of Eq. (1), provided the base functions have a particular form, namely, that the space \mathcal{H} of base functions is *semi-finite*, meaning that \mathcal{H} contains a finite set of functions \mathcal{G} and that

1. every function in \mathcal{H} can be written as a linear combination of the functions in \mathcal{G} , and
2. αg is in \mathcal{H} for every $\alpha \in \mathbb{R}$ and $g \in \mathcal{G}$.

Theorem 1 *Assume the base functions h_t in Fig. 1 minimize Eq. (3) over a semi-finite space \mathcal{H} . Then as $T \rightarrow \infty$, the loss in Eq. (1) for the final function f converges to the infimum of this loss over all linear combinations of functions in \mathcal{H} .*

Proof sketch: To prove the result, we only need to show that, on each round, AdaBoost.L makes at least as much progress as Collins, Schapire and Singer's [2]'s sequential-update algorithm applied to the finite set \mathcal{G} . In particular, we note that

$$\begin{aligned} \sum_i W_t(i) e^{-y_i h_t(x_i)} &= \min_{h \in \mathcal{H}} \sum_i W_t(i) e^{-y_i h(x_i)} \\ &\leq \min_{\alpha \in \mathbb{R}, g \in \mathcal{G}} \sum_i W_t(i) e^{-y_i \alpha g(x_i)} \\ &\leq \sum_i W_t(i) e^{-y_i \alpha_t g_t(x_i)} \end{aligned}$$

where α_t and g_t are the choices that would have been made by their algorithm. With these additional steps, their proof of convergence is easily modified. ■

The base learning algorithm that we use in our experiments for finding base functions is the same as in Schapire and Singer's BoosTexter system [11]. These

experiments all deal with text, and each base function tests for the presence or absence of a particular word, short phrase or other simple pattern, henceforth referred to simply as a *term*. If the term is present, then one value is output; otherwise, some other value is output. For instance, the base function might be: “If the word ‘yes’ occurs in the text, then output +1.731, else output −2.171.” Schapire and Singer [11] describe a base learning algorithm that efficiently finds the best base function of this form, i.e., the one minimizing Eq. (3). It can be seen that this space of base functions is semi-finite since there are only finitely many terms and since a rule of this form can be decomposed as $a_0 g_0 + a_1 g_1$ where $a_0, a_1 \in \mathbb{R}$ and g_1 (respectively, g_0) outputs 1 if the term is present (respectively, absent), and 0 otherwise.

3. Incorporating Prior Knowledge

We now describe our modification to boosting to incorporate prior knowledge. In our approach, a human expert must begin by constructing a rule π mapping each instance x to an estimated conditional probability distribution $\pi(y|x)$ over the possible label values $\{-1, +1\}$. We discuss below some methods for constructing such a rule.

Given this background or prior model and training data, we now have two possibly conflicting goals in constructing a predictor: (1) fit the data, and (2) fit the prior model. As before, we measure fit to the data using log conditional likelihood as in Eq. (1). To measure fit to the prior model, for each example x_i , we use relative entropy (also called Kullback-Leibler divergence) between the prior model distribution $\pi(\cdot|x_i)$ and the distribution over labels associated with our constructed logistic model $\sigma(f(x_i))$. More precisely, letting $\pi_+(x) = \pi(y = +1|x)$, we measure fit to the prior model by

$$\sum_i \text{RE}(\pi_+(x_i) \parallel \sigma(f(x_i))) \quad (4)$$

where $\text{RE}(p \parallel q) = p \ln(p/q) + (1-p) \ln((1-p)/(1-q))$ is binary relative entropy. The relative importance of the two terms is controlled by the parameter η .

Putting these together, we get the objective function

$$\sum_i [\ln(1 + \exp(-y_i f(x_i))) + \eta \text{RE}(\pi_+(x_i) \parallel \sigma(f(x_i)))]. \quad (5)$$

This can be rewritten as

$$C + \sum_i [\ln(1 + e^{-y_i f(x_i)})$$

$$+ \eta \pi_+(x_i) \ln(1 + e^{-f(x_i)}) + \eta (1 - \pi_+(x_i)) \ln(1 + e^{f(x_i)})] \quad (6)$$

where C is a term that is independent of f , and so can be disregarded. Note that this objective function has the same form as Eq. (1) over a larger set and with the addition of nonnegative weights on each term.

Thus, to minimize Eq. (6), we apply the AdaBoost.L procedure described in Section 2 to a larger weighted training set. This new set includes all of the original training examples (x_i, y_i) , each with unit weight. In addition, for each training example (x_i, y_i) , we create two new training examples $(x_i, +1)$ and $(x_i, -1)$ with weights $\eta \pi_+(x_i)$ and $\eta (1 - \pi_+(x_i))$, respectively. Thus, we triple the number of examples (although, by noticing that (x_i, y_i) occurs twice, we can actually get away with only doubling the training set). During training, these weights w_0 are now used in computing W_t so that

$$W_t(i) = \frac{w_0(i)}{1 + \exp\left(y_i \sum_{t'=0}^{t-1} h_{t'}(x_i)\right)}$$

(here, i ranges over all of the examples in the *new* training set). The modification of Theorem 1 for weighted training sets is straightforward.

One final modification that we make is to add a 0-th base function h_0 that is based on π_+ so as to incorporate π_+ right from the start. In particular, we take

$$h_0(x) = \sigma^{-1}(\pi_+(x)) = \ln\left(\frac{\pi_+(x)}{1 - \pi_+(x)}\right)$$

and include h_0 in computing the final classifier f .

3.1 Multiclass problems

Up until now, we have assumed a binary prediction problem with $\mathcal{Y} = \{-1, +1\}$. More generally, we follow Schapire and Singer’s [10, 11] approach to multiclass problems in which more than two classes are allowed and furthermore in which each example may belong to multiple classes. The intuitive idea is to reduce to binary questions which ask if each example is or is not in each of the classes.

In particular, suppose that there are k classes $\mathcal{Y} = \{1, 2, \dots, k\}$. Each label \mathbf{y}_i is now a vector in $\{-1, +1\}^k$ where the ℓ -th component indicates if the example is or is not in class ℓ . Our purpose now is to find a function $f : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, and $\sigma(f(x, \ell))$ is then the estimated probability that example x belongs to class ℓ . Treating each class separately, the objective function in Eq. (1) becomes

$$\sum_i \sum_{\ell} \ln\left(1 + e^{-y_{i\ell} f(x_i, \ell)}\right).$$

The boosting algorithm AdaBoost.L is modified straightforwardly: Maintaining weights on example-label pairs, Eq. (2) becomes

$$W_t(i, \ell) = \frac{1}{1 + \exp\left(y_{i\ell} \sum_{t'=1}^{t-1} h_{t'}(x_i, \ell)\right)},$$

and Eq. (3) becomes

$$\sum_i \sum_{\ell} W_t(i, \ell) e^{-y_{i\ell} h_t(x_i, \ell)}.$$

As was done by Schapire and Singer [11], our base learner finds rules that still test for the presence or absence of a term, but now outputs a whole vector of numbers (one for each class) depending on the result of this test.

Our prior knowledge now gives guessed estimates $\pi(\ell|x)$ of the conditional probability that example x belongs to class ℓ . We do not require that $\pi(\cdot|x)$ be a probability distribution. The objective function in Eqs. (5) and (6) becomes

$$\begin{aligned} & \sum_i \sum_{\ell} [\ln(1 + e^{-y_{i\ell} f(x_i, \ell)}) \\ & \quad + \eta \text{RE}(\pi(\ell|x_i) \parallel \sigma(f(x_i, \ell)))] \\ & = \sum_i \sum_{\ell} [\ln(1 + e^{-y_{i\ell} f(x_i, \ell)}) \\ & \quad + \eta \pi(\ell|x_i) \ln(1 + e^{-f(x_i, \ell)}) \\ & \quad + \eta (1 - \pi(\ell|x_i)) \ln(1 + e^{f(x_i, \ell)})] + C. \end{aligned}$$

So to handle this objective function, similar to the binary case, we create a new training set with weights over example-label pairs: The original examples (x_i, \mathbf{y}_i) occur with unit weight $w_0(i, \ell) = 1$. Each such example is replicated twice as $(x_i, +\mathbf{1})$ and $(x_i, -\mathbf{1})$ where $\mathbf{1}$ is the all ones vector. Letting $i + m$ and $i + 2m$ be the indices of the new replicated examples, their weights are, respectively,

$$\begin{aligned} w_0(i + m, \ell) &= \eta \pi(\ell|x_i) \\ \text{and } w_0(i + 2m, \ell) &= \eta (1 - \pi(\ell|x_i)). \end{aligned}$$

4. Experiments

In this section, we describe experiments comparing boosting with prior knowledge against boosting with no such knowledge, particularly when data is substantially limited. We did not compare to other text categorization methods, since this was not the purpose of the study; moreover, Schapire and Singer [11] carried out extensive experiments comparing boosting to several other methods on text categorization problems.

We used two publicly available text categorization datasets and two proprietary speech categorization datasets. The latter datasets come from the application that was the original motivation for this work. We chose the former datasets because they are large, and also because they naturally lent themselves to the easy construction of a human-crafted model. We could not use a substantially larger number of datasets because of the inherently intensive, subjective and non-automatic nature of building such models.

4.1 Benchmark datasets

In the first set of experiments, we used these two benchmark text-categorization datasets:

- *AP-Titles*: This is a corpus of Associated Press newswire headlines [8, 7]. The object is to classify the headlines by topic. We used the preparation of this dataset described by Schapire and Singer [11] consisting of 29,841 examples and 20 classes.
- *Newsgroups*: This dataset consists of Usenet articles collected by Lang [6] from 20 different newsgroups. The object is to predict which newsgroup a particular article was posted to. One thousand articles were collected for each newsgroup. However, after removing duplicates, the total number of articles dropped to 19,466.

Prior model. Our framework permits prior knowledge of any kind, so long as it provides estimates, however rough, of the probability of any example belonging to any class. Here we describe one possible technique for creating such a rough model.

For each dataset, one of the authors, with access to the list of categories but not to the data itself, thought up a handful of keywords for each class. These lists of keywords are shown in Tables 1 and 2. These keywords were produced through an entirely subjective process of free association with general knowledge of what the categories were about (and also the time period during which the data was connected), but no other information or access to the data. Although this step required direct human involvement, the rest of the process of generating a prior model was fully automatic.

We next used these keywords to build a very simple and naive model. We purposely used a model that is very far from perfect to see how the algorithm performs with prior knowledge that is as rough as can be expected in practice. We began by defining the conditional probability of a class ℓ given the presence or absence of a keyword w , denoted $\pi(\ell|w)$ or $\pi(\ell|\bar{w})$. We

<i>Class</i>	<i>Keywords</i>
japan	japan, tokyo, yen
bush	bush, george, president, election
israel	israel, jerusalem, peres, sharon, palestinian, israeli, arafat
britx	britain, british, england, english, london, thatcher
gulf	gulf, iraq, saudi, arab, iraqi, saddam, hussein, kuwait
german	german, germany, bonn, berlin, mark
weather	weather, rain, snow, cold, ice, sun, sunny, cloudy
dollargold	dollar, gold, price
hostages	hostages, ransom, holding, hostage
budget	budget, deficit, taxes
arts	art, painting, artist, music, entertainment, museum, theater
dukakis	dukakis, boston, taxes, governor
yugoslavia	yugoslavia
quayle	quayle, dan
ireland	ireland, ira, dublin
burma	burma
bonds	bond, bonds, yield, interest
nielsens	nielsens, rating, t v, tv
boxoffice	box office, movie
tickertalk	stock, bond, bonds, stocks, price, earnings

Table 1. The keywords used for each class on the *AP-Titles* dataset.

let

$$\pi(\ell|w) = \begin{cases} 0.9/n_w & \text{if } w \text{ is a keyword for } \ell \\ 0.1/(k - n_w) & \text{otherwise} \end{cases}$$

where n_w is the number of classes listing w as a keyword. In other words, if the keyword w is listed for a single class ℓ then seeing the word gives a 90% probability that the correct class is ℓ ; if w is listed for several classes, the 90% probability is divided equally among them. The remaining 10% probability is divided equally among all classes not listing w as a keyword.

If w is not present, we assign equal probability to all classes: $\pi(\ell|\overline{w}) = 1/k$. We also define the prior distribution of classes to be uniform: $\pi(\ell) = 1/k$.

Given these rules, we make the naive assumption that the keywords are conditionally independent of one another given the class. We can then use Bayes' rule to compute the probability (under π) of each class given the presence or absence of all the keywords. This becomes our estimate $\pi(\ell|x)$.

Experimental set-up. For each dataset and on each run, we first randomly permuted the data. We then trained boosting, with or without prior knowledge, on the first m examples, for $m =$

<i>Class</i>	<i>Keywords</i>
alt.atheism	god, atheism, christ, jesus, religion, atheist
comp.graphics	graphics, color, computer, computers, plot, screen
comp.os.ms-windows.misc	computer, computers, operating system, microsoft, windows, ms, dos
comp.sys.ibm.pc.hardware	computer, computers, ibm, pc, clone, hardware, cpu, disk
comp.sys.mac.hardware	computer, computers, mac, macintosh, hardware, cpu, disk
comp.win-dows.x	computer, computers, windows, x, unix
misc.forsale	for sale, asking, selling, price
rec.autos	car, drive, fast, jaguar, toyota, ford, honda, volkswagen, gm, chevrolet, tire, engine
rec.motorcycles	motorcycle, honda, harley, wheel, engine, throttle
rec.sport.baseball	baseball, hit, strike, ball, base, bases, homerun, runs, out, outs
rec.sport.hockey	hockey, stick, puck, goal, check
sci.crypt	cryptography, encrypt, cipher, decrypt, security, secret, key
sci.electronics	electronics, computer, computers, chip, electric
sci.med	medicine, doctor, science, heal, sick, cancer
sci.space	space, astronaut, nasa, rocket, space shuttle
soc.religion.christian	religion, christian, jesus, christ, god, catholic, protestant
talk.politics.guns	guns, gun, nra, brady, kill, shoot, shot
talk.politics.mideast	mideast, israel, jordan, arafat, palestinian, syria, lebanon, saudi, iraq, iran
talk.politics.misc	politics, clinton, president, congress, senate, congressman, senator
talk.religion.misc	religion, jewish, christian, catholic, protestant, god, believe

Table 2. The keywords used for each class on the *News-groups* dataset.

25, 50, 100, 200, ..., M , where M is 12800 for *AP-Titles* and 6400 for *Newsgroups*. The remaining examples (i.e., starting with example $M + 1$) are used for testing. We ran each experiment ten times and averaged the results. We fixed the number of rounds of boosting to 1000. We set the parameter η using the heuristic formula $2000m^{-1.66}$ (which was chosen to interpolate smoothly between guesses at appropriate values of η for a couple values of m). No experiments were conducted to determine if better performance could be achieved with a wiser choice of η .

Results. Figs. 2 and 3 show the results of these experiments. The figures show test error rate for boosting with and without prior knowledge measured as a




Figure 2. Comparison of test error rate using prior knowledge and data separately or together on the *AP-Titles* dataset, measured as a function of the number of training examples.

function of the number of training examples. The figures also show the error rate achieved using the prior model alone with no training examples at all. Here, error rate means fraction of test examples on which the top-scoring label was not one of the correct labels (recall that each example may belong to more than one class). The results are averaged over the ten runs.

For fairly small datasets, using prior knowledge gives dramatic improvements over straight boosting. On large training sets on the *Newsgroups* dataset, the imperfect nature of the prior knowledge eventually hurts performance, although this effect is not seen on *AP-Titles*.

4.2 Spoken-dialogue datasets

We next describe experiments on datasets from two AT&T spoken-dialogue applications. In both applications, the goal is to extract the meaning of utterances spoken by telephone callers. These utterances are then passed through an automatic speech recognizer. Our goal is to train a classifier that can categorize the resulting (very noisy) text. The classifier’s output would then be passed to the dialogue manager which carries on with the dialogue by formulated an appropriate response to the caller’s utterance.

The two applications are:

- *How May I Help You?* (*HMIHY*): Here, the goal is to identify a particular call type, such as collect call, a request for billing information, etc. There are 15 different classes. We did experiments with 50 to 1600 sentences in the training set and 2991




Figure 3. Comparison of test error rate using prior knowledge and data separately or together on the *Newsgroups* dataset, measured as a function of the number of training examples.

sentences in the test set. More information about this dataset is provided by Gorin, Riccardi and Wright [5].

- *HelpDesk*: This application provides information about AT&T’s Natural Voices text-to-speech engine. For instance, the caller can ask for a demo, price information, or a sales representative. There are 22 different classes. We trained models on 100 to 2675 sentences and tested on 2000 sentences.

The prior models were built in a similar fashion to those used in the preceding experiments, although we allowed the human more freedom in choosing probabilities, and more rules were used. See Rochery et al. [9] for further details. We performed similar experiments to those described in Section 4.1 measuring classification accuracy as a function of the number of examples used during training, and comparing models built only with some training examples and models built with both hand-crafted rules (prior knowledge) and training examples.

On the *HMIHY* dataset, we trained the models on 50, 100, 200, 300 rounds when the number of available training examples was respectively 50, 100, 200, 400 and up. The parameter η was selected empirically based on the number of available training examples. We set η to 1 when the number of training examples was less than or equal to 200, 0.1 when it was between 400 and 800, and 0.01 when it was greater. The dashed line in Figure 4 shows the classification accuracy for models built on hand-crafted rules and training examples whereas the solid lines show the classification




Figure 4. Comparison of performance using data and knowledge separately or together on the *HMIHY* task.

accuracy for models built either on training examples only or on hand-crafted rules only. An improvement in accuracy is observed when using hand-crafted rules and training examples together. This comes from the fact that some patterns of the hand-crafted rules are not in the data at all or are not in a sufficient number of sentences to have a statistical impact when training only on the data.

In this experiment, when fewer training examples were available (< 100 examples) exploiting human expertise provided classification accuracy levels that are equivalent to models trained on four times the amount of training data. When the number of training examples is larger (> 100), accuracy levels become equivalent to two times the amount of training data. When larger than 6000 sentences were available, both models were found to converge to similar classification accuracy.

Figure 5 shows a similar comparison for the *HelpDesk* task. We trained the models on 100, 200, 400, 600, 800, 1000 when the number of training examples was respectively 100, 200, 400, 800, 1600, and up. We set η to 0.1 when the number of training examples was less than or equal to 1600 and to 0.01 otherwise. Figure 5 shows an improvement in classification accuracy when hand-crafted rules are being used. This improvement is up to 9% absolute with 100 training examples and drops to 0.5% when more data becomes available.

In the figures, the knowledge-only curves are not perfectly flat. This comes from the fact that the models from the knowledge take into account the empirical distribution of classes of the available training examples rather than using a uniform distribution as was done in Section 4.1.




Figure 5. Comparison of performance using data and knowledge separately or together on the *HelpDesk* task.

A further experiment was performed to evaluate the accuracy of our classifier when new semantic classes are added following system training. This is the situation when new functionalities are needed following system deployment but with no data available. Fig. 6 shows the classification accuracy when four additional semantic classes are added to the *HMIHY* model after being trained on 11 classes. Although the system performance drops in general, the results demonstrate that incorporating human judgment helps to provide an initial boost in performance even when no data is present.

5. Variations and extensions

We have described the extension of one particular boosting algorithm to incorporate prior knowledge. However, the same basic technique can be applied to a great variety of boosting algorithms. For instance, we have used Schapire and Singer’s [10] confidence-rated boosting framework in which the base functions map to real numbers whose magnitude indicate a level of confidence. This choice is orthogonal to our basic method for incorporating prior knowledge. Although this approach can substantially speed up convergence when using a rather weak base learner, in some settings, one may wish to use a more standard base learner outputting “hard” predictions in $\{-1, +1\}$ and for which the goal is simply (weighted) error minimization; for this, a more basic version of AdaBoost can be used.

We also have chosen a particular method of extending binary AdaBoost to the multiclass case, an extension that Schapire and Singer [10] call AdaBoost.MH. We




Figure 6. Adding new semantic classes following model training.

could instead use one of the other multiclass extensions such as AdaBoost.MR [4, 10] as modified for logistic regression by Collins, Schapire and Singer [2].

In fact, our approach is not even limited to boosting algorithms. The basic idea of modifying the loss function used in logistic regression by adding pseudo-examples can be applied to any algorithm for logistic regression.

Note that our measure of fit to the prior model given in Eq. (4) is independent of the actual training labels y_i . This means that we need not limit this term to labeled data: if, as is often the case, we have access to abundant unlabeled data, we can use it instead for this term.

Another idea for future research is to follow the co-training approach studied by Blum and Mitchell [1] and Collins and Singer [3] in which we train two models, say f and g , which we force to give similar predictions on a large set of unlabeled data. In this case, the term in Eq. (4) might be replaced by something like

$$\sum_i \text{RE}(\sigma(g(x_i)) \parallel \sigma(f(x_i)))$$

where the sum is over the unlabeled dataset.

References

- [1] Avrim Blum and Tom Mitchell. Combining labeled and unlabeled data with co-training. In *Proceedings of the Eleventh Annual Conference on Computational Learning Theory*, pages 92–100, 1998.
- [2] Michael Collins, Robert E. Schapire, and Yoram Singer. Logistic regression, AdaBoost and Bregman distances. *Machine Learning*, 48(1/2/3), 2002.
- [3] Michael Collins and Yoram Singer. Unsupervised models for named entity classification. In *Empirical Methods in Natural Language Processing and Very Large Corpora*, 1999.
- [4] Yoav Freund and Robert E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1):119–139, August 1997.
- [5] A. L. Gorin, G. Riccardi, and J. H. Wright. How may I help you? *Speech Communication*, 23(1-2):113–127, October 1997.
- [6] K. Lang. Newsweeder: Learning to filter net-news. In *Proceedings of the Twelfth International Conference on Machine Learning*, pages 331–339, 1995.
- [7] David Lewis and Jason Catlett. Heterogeneous uncertainty sampling for supervised learning. In *Machine Learning: Proceedings of the Eleventh International Conference*, 1994.
- [8] David Lewis and William Gale. Training text classifiers by uncertainty sampling. In *Seventeenth Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, 1994.
- [9] M. Rochery, R. Schapire, M. Rahim, N. Gupta, G. Riccardi, S. Bangalore, H. Alshawi, and S. Douglas. Combining prior knowledge and boosting for call classification in spoken language dialogue. In *International Conference on Acoustics, Speech and Signal Processing*, 2002.
- [10] Robert E. Schapire and Yoram Singer. Improved boosting algorithms using confidence-rated predictions. *Machine Learning*, 37(3):297–336, December 1999.
- [11] Robert E. Schapire and Yoram Singer. Boos-Texter: A boosting-based system for text categorization. *Machine Learning*, 39(2/3):135–168, May/June 2000.