

Approximating the Semantics of Logic Programs by Recurrent Neural Networks *

STEFFEN HÖLLDOBLER AND HANS-PETER STÖRR

*Artificial Intelligence Institute
Computer Science Department
University of Technology Dresden
D-01062 Dresden, Germany*

{sh,haps}@inf.tu-dresden.de

YVONNE KALINKE**

*Neurocomputing Research Center
Queensland University of Technology
Brisbane, Australia, GPO Box 2434, QLD 4001*

yvonne@fit.qut.edu.au

;

Abstract. In [8] we have shown how to construct a 3-layer recurrent neural network that computes the iteration of the meaning function $T_{\mathcal{P}}$ of a given propositional logic program, what corresponds to the computation of the semantics of the program. In this article we define a notion of approximation for interpretations and prove that there exists a 3-layer feed forward neural network that approximates the calculation of $T_{\mathcal{P}}$ for a given (first order) recurrent logic program with an injective level mapping arbitrarily well. Extending the feed forward network by recurrent connections one obtains an recurrent neural network whose iteration approximates the fixed point of $T_{\mathcal{P}}$. This result is proved by taking advantage of the fact that for recurrent logic programs $T_{\mathcal{P}}$ is a contraction mapping on the complete metric space of the interpretations for the program. Mapping this metric space to the metric space \mathbb{R} a real valued function $f_{\mathcal{P}}$ can be found which corresponds to $T_{\mathcal{P}}$, is continuous as well a contraction, and — for this reason — can be approximated by an appropriately chosen class of feed forward neural networks.

Keywords: Recurrent Neural Networks, Semantics of Logic Programming, Approximation Capabilities of Neural Networks

*This is the revised version of a paper published at the 11th Australian Joint Conference on Artificial Intelligence.

**The author acknowledges support from the German Academic Exchange Service (DAAD) under grant no. D/97/29570.

1. Introduction

Many researchers are convinced that intelligent agents reason by generating models and deducing conclusions with respect to these models (see e.g. [10]). In real agents the reasoning process itself is performed by highly recurrent neural networks, whose precise structure and functionality is still not very well understood. Artificial neural or connectionist networks are just a rather crude model for such real neural networks. Nevertheless, they serve as a reasonable model and have been taken up by many researchers. Currently, one of the major open problems in this area is the question of how the full power of deductive processes can be implemented on connectionist networks (see e.g. [15]). In this article we will focus on this problem by establishing a strong link between model generation in the context of first order logic and recursive artificial neural networks. In particular, we will show that for a certain class of logic programs the least model of a given program can be approximated arbitrarily well by a recursive artificial neural network.

Model generation is a well established area within automated deduction (see e.g. [12, 6, 14]). In particular, the semantics of a (logic) program \mathcal{P} is often defined with the help of a so-called meaning function $T_{\mathcal{P}}$. In many but not all cases a logic program \mathcal{P} admits a least model, which can be computed as the least fixed point of $T_{\mathcal{P}}$, and research is focussed on identifying classes of programs for which such least models exist. For these classes $T_{\mathcal{P}}$ effectively specifies a model generation procedure. Examples are the class of definite programs (see e.g. [11]), where the correspondence between the least model of \mathcal{P} and the least fixed point of $T_{\mathcal{P}}$ can be shown by lattice-theoretic arguments [1], or the class of acceptable programs, where the just mentioned correspondence can be shown using metric methods [5].

It turned out that the computation of the least model for a program from one of the just mentioned classes can be performed by a recursive network of binary threshold units if the programs are propositional [8]: With each interpretation I a vector of units is identified such that the j th unit is active iff the j th propositional variable is mapped to true by I . Two such vectors serve as

input and respectively as output layer of a 3-layer feed forward network (or FNN for short), the hidden layer is constructed such that it contains a unit for each program clause. Such a unit becomes active as soon as the body of the clause is satisfied by the interpretation represented by the activation pattern of the input layer, and propagates its activation to the unit in the output layer representing the head of the clause. Fig. 1 depicts such a network for a small example. From its construction follows immediately that such a network computes the application of $T_{\mathcal{P}}$ to an interpretation I . Turning this network into a recursive one (RNN) by connecting corresponding units in the output and input layer with weight 1 allows to compute the least fixed point of $T_{\mathcal{P}}$ for (propositional) acceptable logic programs.

One should observe that the number of units and the number of connections in an RNN corresponding to a propositional program are bounded by $O(m + n)$ and $O(m \times n)$ respectively, where n is the number of clauses and m is the number of propositional variables occurring in the program. Furthermore, the application of $T_{\mathcal{P}}$ to a given interpretation is computed in 2 steps. As the sequential time to compute $T_{\mathcal{P}}$ is bound by $O(n \times m)$ (assuming that no literal occurs more than once in the conditions of a clause), the resulting parallel computational model is optimal.¹

The approach in [8] provides a new computational model for the computation of the least model of logic programs, which is massively parallel and optimal. Since FNNs are widely used and there are powerful learning algorithms like back-propagation, we also may use these techniques to adapt our program. We can, for instance, train a given FNN with knowledge not yet included in the program. If we then extract a new logic program from the trained network we get an extended program including new clauses (see [4]). Moreover, the approach establishes a strong relationship between propositional logic programming and recursive neural networks. Since, in contrast to RNNs, logic programs are well understood this may help to gain a better insight into RNNs, to formally analyze these networks and to give a declarative semantics for what these networks are doing.

But the question remains how this relationship may look like in the first order case? Since in this case an interpretation may be an infinite subset of




Fig. 1. A FNN of binary threshold units computing $T_{\mathcal{P}}$ for the program $\mathcal{P} = \{A \leftarrow \neg B; A \leftarrow C, D; D \leftarrow \neg E, F\}$. The numbers in the units denote thresholds whereas the numbers at the connections denote weights. Connections with weight 0 are not shown.

the set of ground atoms wrt the alphabet underlying a program, the construction developed in [8] may lead to infinite FNNs and RNNs if applied to first order logic programs. This is unacceptable in practice, where we can handle only finite networks. On the other hand, we may use real valued units with sigmoidal activation functions instead of the binary threshold units used in the propositional case.

FNNs with at least one hidden layer of real valued units with sigmoidal activation function are known to approximate continuous real valued functions as well as Borel measurable functions arbitrarily well [7, 9]. This gives rise to the following idea: Can we find a real valued function $f_{\mathcal{P}}$ corresponding to $T_{\mathcal{P}}$ such that these results can be used to approximate $f_{\mathcal{P}}$ and, thereby, $T_{\mathcal{P}}$ arbitrarily well? If such a function $f_{\mathcal{P}}$ exists and we find an FNN approximating $f_{\mathcal{P}}$ to a desired degree of accuracy, can we then turn the FNN into an RNN approximating the least fixed point of $T_{\mathcal{P}}$ arbitrarily well?

In this article we define a class of first order logic programs such that both questions are answered positively for this class. We also discuss what precisely is denoted by an approximation of

the least model of a first order logic program. For the first time this gives us a strong link between first order model generation and the computation performed by a recursive neural network.

The remainder of this article is organized as follows. After stating some preliminaries concerning logic programs and metric spaces in Section 2, we state in Section 3 that the meaning function $T_{\mathcal{P}}$ for a recurrent logic program has a unique fixed point that can be computed by iterating $T_{\mathcal{P}}$. In Section 4 we show how to encode the domain of the meaning function $T_{\mathcal{P}}$ into \mathbb{R} . In Section 5 we use this mapping to construct a real valued function $f_{\mathcal{P}}$ corresponding to the function $T_{\mathcal{P}}$. In Section 6 we show that the real valued function $f_{\mathcal{P}}$ is continuous for a certain class of programs and that an FNN with sigmoidal activation functions and at least one hidden layer can approximate the function $f_{\mathcal{P}}$. Thereafter we show how to extend the FNN to a RNN such that the iteration of $f_{\mathcal{P}}$ corresponds to the iteration of $T_{\mathcal{P}}$ and therefore computes the least fixed point of $T_{\mathcal{P}}$. Since in the case of recurrent programs such a fixed point does always exist we end up with a RNN computing the semantics of the logic program \mathcal{P} . Finally, we discuss our results and point out future work in Section 9.

The proofs of our results can be found in the appendix of the article.

2. Preliminaries

In the following two subsections we briefly recall basic notions and notations concerning logic programs according to [11] and metric spaces as in [5].

2.1. Logic Programs

A *logic program* \mathcal{P} is a collection of (universally closed) clauses of the form $A \leftarrow L_1, \dots, L_n$, where $n \geq 0$, A is a first order atom and L_i , $1 \leq i \leq n$, are first order literals; A is called *head* and L_1, \dots, L_n *body* of a clause. An atom, literal, clause or program is said to be *ground* if it does not contain an occurrence of a variable. A program \mathcal{P} is called *definite* if each clause occurring in \mathcal{P} contains only atoms. $B_{\mathcal{P}}$ denotes the *Herbrand base* wrt the alphabet underlying \mathcal{P} . A *level mapping* for a program \mathcal{P} is a function $|| : B_{\mathcal{P}} \rightarrow \mathbb{N}$,

where \mathbf{N} denotes the set of natural numbers. If $|A| = n$ we will say the *level* of the atom $A \in B_{\mathcal{P}}$ is n .²

An *interpretation* is a mapping from ground atoms to $\{\mathbf{1}, \mathbf{0}\}$.³ It is extended to literals, clauses and programs in the usual way. A *model* for \mathcal{P} is an interpretation which maps \mathcal{P} to $\mathbf{1}$. The *meaning function* $T_{\mathcal{P}} : 2^{B_{\mathcal{P}}} \rightarrow 2^{B_{\mathcal{P}}}$ is defined as follows: Let I be an interpretation and A a ground atom. $A \in T_{\mathcal{P}}(I)$ iff there exists a ground instance $A \leftarrow L_1, \dots, L_n$ of a clause in \mathcal{P} such that for all $1 \leq i \leq n$ we find $L_i \in I$.⁴

2.2. Metric Spaces

A *metric* or *distance function* on a space \mathcal{M} is a mapping $d : \mathcal{M} \times \mathcal{M} \rightarrow \mathbb{R}^{\geq 0}$ such that $d(x, y) = 0$ iff $x = y$, $d(x, y) = d(y, x)$, and $d(x, y) \leq d(x, z) + d(z, y)$, where $\mathbb{R}^{\geq 0}$ denotes the set of non-negative real numbers. Let (\mathcal{M}, d) be a metric space and $\mathcal{S} = s_1, s_2, \dots, s_i \in \mathcal{M}$, be a sequence on \mathcal{M} . \mathcal{S} *converges* if $\exists s \in \mathcal{M} : \forall \varepsilon > 0 : \exists N : \forall n \geq N : d(s_n, s) \leq \varepsilon$. \mathcal{S} is *Cauchy* if $\forall \varepsilon > 0 : \exists N : \forall n, m \geq N : d(s_n, s_m) \leq \varepsilon$. (\mathcal{M}, d) is *complete* if every Cauchy sequence converges. A mapping $f : \mathcal{M} \rightarrow \mathcal{M}$ is a *contraction* on (\mathcal{M}, d) if $\exists 0 < k < 1 : \forall x, y \in \mathcal{M} : d(f(x), f(y)) \leq k \cdot d(x, y)$.

Consider a logic program \mathcal{P} , a level mapping $||$ for \mathcal{P} and the set $2^{B_{\mathcal{P}}}$ of interpretations for \mathcal{P} . A distance function $d_{\mathcal{P}}$ on $2^{B_{\mathcal{P}}}$ associated with $||$ is defined as follows. Let I and J be two interpretations. If $I = J$ then $d_{\mathcal{P}}(I, J) = 0$. Otherwise, $d_{\mathcal{P}}(I, J) = 2^{-n}$, where I and J differ on some atom A of level n but agree on all atoms of lower level. As shown in [5] the distance function $d_{\mathcal{P}}$ associated with a level mapping $||$ for \mathcal{P} is a metric and the metric space $(2^{B_{\mathcal{P}}}, d_{\mathcal{P}})$ is complete.

For complete metric spaces and contraction mappings the following theorem is well-known:

Theorem 1. (Banach Contraction Theorem [17]) *A contraction mapping f on a complete metric space has a unique fixed point. The sequence $x, f(x), f(f(x)), \dots$ converges to this fixed point for any x .*

3. Recurrent Logic Programs

Recurrent logic programs provide a characteristic that, as we will show, allows to approximate their semantics by a RNN. The main reason for that is that, by using metric methods as proposed by Fitting for the class of acceptable logic programs ([5]), we can show that the meaning function $T_{\mathcal{P}}$ for a recurrent logic program has a unique fixed point that describes the semantics of the program.

Definition 1. A logic program is called *recurrent* wrt to a level mapping $||$, if for every ground instance $A \leftarrow L_1, \dots, L_n$ of a clause in \mathcal{P} :

$$\forall 1 \leq i \leq n : |A| > |L_i|.$$

For a recurrent logic program \mathcal{P} the meaning function is a contraction on the complete metric space $(2^{B_{\mathcal{P}}}, d_{\mathcal{P}})$ for the program.

Proposition 1. *Let \mathcal{P} be a recurrent logic program wrt a level mapping $||$ and $d_{\mathcal{P}}$ the distance function associated with the level mapping $||$. The meaning function $T_{\mathcal{P}}$ is a contraction on the complete metric space $(2^{B_{\mathcal{P}}}, d_{\mathcal{P}})$.*

Applying the Banach Contraction Theorem (1) it follows that for a recurrent logic program, the meaning function has a unique fixed point, and that fixed point is reached by iterating $T_{\mathcal{P}}$ starting from any interpretation $I \in B_{\mathcal{P}}$.

4. Mapping Interpretations to Real Numbers

Since we are interested in first order logic programs, the arguments of the meaning function $T_{\mathcal{P}}$ are interpretations which may consist of a countably infinite number of ground atoms. As already mentioned in the introduction the simple solution for the propositional case as presented in [8], where each ground atom is represented by a binary threshold unit in the input and output layer is no longer feasible. To extend the representational capability of our network we use real valued units with sigmoidal activation functions instead. Thus interpretations are to be represented by real numbers.

In this section we define an encoding R of the domain 2^{B_P} of the meaning function T_P into \mathbb{R} . Considering our aim of approximating T_P by the use of an FNN, we should make sure that the encoding of T_P in terms of a real valued function f_P is done in a way such that f_P can indeed be approximated. For example, if f_P would be continuous, then we could apply the result of [7] that continuous functions can be approximated arbitrarily well by an FNN.⁵ As we shall see, by restricting ourselves to a certain class of programs we can ensure that the function f_P encoding T_P is continuous.

We start from a level mapping $\|\cdot\| : B_P \rightarrow \mathbb{N}$ from ground atoms to natural numbers. We further require it to be injective, and express this by renaming it to $\|\cdot\|$. This restriction is discussed further in Section 7.

We use the level mapping $\|\cdot\|$ to define a mapping R from the set 2^{B_P} of interpretations for a logic program P to the real numbers:

Definition 2. The mapping $R : 2^{B_P} \rightarrow \mathbb{R}$ is defined as $R(I) = \sum_{A \in I} 4^{-\|A\|}$.

Thus an interpretation I is mapped to the real number whose n -th digit after the comma in the quaternary⁶ number system is 1 if there is an atom A in I with $\|A\| = n$,⁷ and 0 otherwise. The requirement of $\|\cdot\|$ to be injective ensures that R is injective and thus there is an inverse mapping R^{-1} . This inverse mapping is extended to a function $R^{-1} : \mathbb{R} \rightarrow 2^{B_P}$ by taking the value I of the nearest point $R(I)$ for $I \in 2^{B_P}$.⁸

There is a close relation between the metric d_P on interpretations and the distance of the corresponding real numbers:

Proposition 2. For two interpretations $I, J \in 2^{B_P}$ the following holds:

$$\frac{2}{3} * d_P(I, J)^2 \leq |R(I) - R(J)| \leq \frac{4}{3} * d_P(I, J)^2.$$

5. Mapping T_P to a Real Valued Function f_P

The encoding R maps the elements of the domain of T_P to real numbers. Hence, we immedi-




Fig. 2. The relation between T_P and f_P , where $\mathcal{D}_f = \{r \in \mathbb{R} \mid \exists I \in 2^{B_P} : r = R(I)\}$.

ately obtain a function f_P on \mathbb{R} such that the application of f_P to a real number $r = R(I)$ encoding the interpretation I is equivalent to applying T_P to $I = R^{-1}(r)$. Figure 2 depicts the relation between T_P and f_P . One should observe that the encoding R defined in Definition 2 is just an example for mapping 2^{B_P} to \mathbb{R} . We can use any mapping which transfers the contraction property of T_P to f_P , i.e. for which f_P is continuous on \mathbb{R} iff T_P is a contraction on 2^{B_P} .

Definition 3. Let P be a logic program, B_P the Herbrand base and T_P the meaning function associated with P . Let R be an encoding of 2^{B_P} in \mathbb{R} and the closed set $\mathcal{D}_f \subseteq \mathbb{R}$ the range of R . The real valued function \bar{f}_P corresponding to T_P is defined by

$$\bar{f}_P : \mathcal{D}_f \rightarrow \mathcal{D}_f : r \mapsto R(T_P(R^{-1}(r))).$$

The function \bar{f}_P is extended to a function $f_P : \mathbb{R} \rightarrow \mathbb{R}$ by linear interpolation:

$$f_P(r) = \begin{cases} \bar{f}_P(r) & \text{if } r \in \mathcal{D}_f \\ \bar{f}_P(\min(\mathcal{D}_f)) & \text{if } r < \min(\mathcal{D}_f) \\ \bar{f}_P(\max(\mathcal{D}_f)) & \text{if } r > \max(\mathcal{D}_f) \\ \frac{(M_r - r)}{M_r - m_r} \bar{f}_P(m_r) + \frac{(r - m_r)}{M_r - m_r} * \bar{f}_P(M_r) & \text{otherwise,} \end{cases}$$

where $m_r = \max(\mathcal{D}_f \cup (-\infty, r])$ and $M_r = \min(\mathcal{D}_f \cup [r, \infty))$ are the greatest point of \mathcal{D}_f below r and least point of \mathcal{D}_f above r respectively.⁹

To prove that f_P is continuous we first prove the following property of f_P :

Proposition 3. Let P be a recurrent logic program with an injective level mapping, T_P the meaning function associated with P and f_P the

real valued function corresponding to $T_{\mathcal{P}}$. Then $f_{\mathcal{P}}$ is a contraction on \mathbb{R} , i.e.

$$\forall r, r' \in \mathbb{R} : |f_{\mathcal{P}}(r') - f_{\mathcal{P}}(r)| \leq \frac{1}{2} |r' - r| .$$

As an immediate consequence of Proposition 3 we obtain:

Corollary 1. *Let \mathcal{P} be a recurrent logic program with an injective level mapping, $T_{\mathcal{P}}$ the meaning function associated with \mathcal{P} and $f_{\mathcal{P}}$ the real valued function corresponding to $T_{\mathcal{P}}$. Then the function $f_{\mathcal{P}}$ is continuous.*

6. Approximating the Meaning Function $T_{\mathcal{P}}$

Since $f_{\mathcal{P}}$ is a continuous real valued function we can apply the following theorem stating the approximation capability of a class of FNNs.

Theorem 2. (Funahashi [7]) *Let $\phi(x)$ be a non constant, bounded and monotone increasing continuous function. Let K be a compact subset (bounded closed subset) of \mathbb{R}^n and $f(x_1, \dots, x_n)$ be a continuous real valued function on K . Then for an arbitrary $\varepsilon > 0$, there exists an integer N and real constants c_i, θ_i ($i = 1, \dots, N$), w_{ij} ($i = 1, \dots, N, j = 1, \dots, n$) such that*

$$\tilde{f}(x_1, \dots, x_n) = \sum_{i=1}^N c_i \phi \left(\sum_{j=1}^n w_{ij} x_j - \theta_i \right)$$

satisfies

$$\max_{\vec{x} \in K} |f(x_1, \dots, x_n) - \tilde{f}(x_1, \dots, x_n)| < \varepsilon .$$

In other words, for an arbitrary $\varepsilon > 0$, there exists a 3-layer FNN whose output function for the hidden layer is $\phi(x)$, its output functions for input and output layers are linear and it has an input-output function $\tilde{f}(x_1, \dots, x_n)$ such that $\max_{\vec{x} \in K} |f(x_1, \dots, x_n) - \tilde{f}(x_1, \dots, x_n)| < \varepsilon$.

Applying Theorem 2 to Corollary 1 which states that $f_{\mathcal{P}}$ is continuous we obtain that the following theorem:

Theorem 3. *Let \mathcal{P} be a recurrent logic program with an injective level mapping, $T_{\mathcal{P}}$ the meaning function associated with \mathcal{P} and $f_{\mathcal{P}}$ the continuous real valued function corresponding to $T_{\mathcal{P}}$. Then for an arbitrary $\varepsilon > 0$, there exists an FNN with sigmoidal activation function for the hidden layer units and linear activation functions for the input and output layer units computing the function $f_{\mathcal{P}}$ which satisfies*

$$\max_{x \in [0,1]} |f_{\mathcal{P}}(x) - \tilde{f}_{\mathcal{P}}(x)| < \varepsilon .$$

The theorem states that given a certain accuracy we can construct an FNN that approximates $f_{\mathcal{P}}$ to this desired degree of accuracy. Using the mapping R^{-1} we thereby obtain an FNN capable of approximating the meaning function $T_{\mathcal{P}}$ for a recurrent logic program \mathcal{P} with injective level mapping. In the following section we exemplify this result and clarify the notion of an approximation of $T_{\mathcal{P}}$.

7. An Example

Consider the program

$$\begin{aligned} \mathcal{P} : & p(0). \\ & p(s(X)) \leftarrow p(X). \end{aligned}$$

and the injective level mapping $\|p(s^n(0))\| = n+1$ for $n \in \mathbb{N}$.¹⁰ The program is recurrent wrt the level mapping $\|\cdot\|$. In this case we find

$$\begin{aligned} f_{\mathcal{P}}(R(I)) &= 4^{-\|p(0)\|} + \sum_{p(X) \in I} 4^{-\|p(s(X))\|} \\ &= 4^{-\|p(0)\|} + \sum_{p(X) \in I} 4^{-(\|p(X)\|+1)} \\ &= \frac{1}{4} + \frac{1}{4} R(I) . \end{aligned}$$

The iteration of $T_{\mathcal{P}}$, which yields the semantics of the program \mathcal{P} , corresponds to the iteration of $\tilde{f}_{\mathcal{P}}$. What happens during this iteration? If we approximate $f_{\mathcal{P}}$ to an accuracy ε the evaluation of the approximation $\tilde{f}_{\mathcal{P}}$ yields a value $\tilde{f}_{\mathcal{P}}(x) \in [\frac{1+x}{4} - \varepsilon, \frac{1+x}{4} + \varepsilon]$. Thus, if x is in the interval $[a, b]$, the result $\tilde{f}_{\mathcal{P}}(x)$ is in the interval $[\frac{a-\frac{1-4\varepsilon}{3}}{4} + \frac{1-4\varepsilon}{3}, \frac{b-\frac{1-4\varepsilon}{3}}{4} + \frac{1+4\varepsilon}{3}]$. It is easy to see that in the limit the iteration of $\tilde{f}_{\mathcal{P}}$ yields a value within $[\frac{1-4\varepsilon}{3}, \frac{1+4\varepsilon}{3}]$ (though it does not necessarily converge to a fixed value within this interval.)

If we convert such a value $r \in [\frac{1-4\varepsilon}{3}, \frac{1+4\varepsilon}{3}]$ with R^{-1} back to $B_{\mathcal{P}}$ we see that $p(s^n(0)) \in R^{-1}(r)$ for $n < \log_4 \varepsilon - 1$, which coincides with the model $\{p(0), p(s(0)), p(s(s(0))), \dots\}$ of \mathcal{P} . The values for $p(s^n(0))$ wrt $R^{-1}(r)$ with $n \geq \log_4 \varepsilon - 1$ may differ from the intended model.

This clarifies our notion of approximation of an interpretation or a model: An interpretation I *approximates* an interpretation J to a degree N , if for all atoms $A \in B_{\mathcal{P}}$ with $\|A\| < N$, $A \in I$ iff $A \in J$. Equivalently, $d_{\mathcal{P}}(I, J) \leq 2^{-N}$.

One should observe that in our example we are not only able to approximate the calculation of $T_{\mathcal{P}}$ but are able to approximate the fixed point of $T_{\mathcal{P}}$ by iteration of the approximation.

Unfortunately, not all recurrent programs admit an *injective* level mapping. Consider for instance the logic program

$$\mathcal{Q} : q \leftarrow p(f(X)).$$

It is recurrent wrt the level mapping, which maps each ground instance of $p(f(X))$ to 1 and q to 2. But it does not admit an injective level mapping for which $T_{\mathcal{P}}$ is a contraction because the level of q has to be greater than the maximum of the level of all infinitely ground instances of $p(f(X))$. So the restriction to recurrent programs that admit an *injective* level mapping does not allow to treat programs that contain clauses with variables occurring in the body but not in the head of the clause.

Note, moreover, that the class of recurrent logic programs contains not only definite programs but negation is allowed and we can handle recurrent programs with negation as well.

8. Iteration of the Approximation $\tilde{f}_{\mathcal{P}}$

So far we have shown how to approximate the meaning function $T_{\mathcal{P}}$. However, we are mainly interested in the approximation of the least fixed point of $T_{\mathcal{P}}$. From Theorem 3 we have learned that there exists an FNN computing $\tilde{f}_{\mathcal{P}}$, which is the approximation of $f_{\mathcal{P}}$, which in turn represents $T_{\mathcal{P}}$. Given such an FNN we could turn this FNN into an RNN by adding recurrent connections with weight 1 between corresponding units in the output and input layer of the FNN. Does

such an RNN approximate the least fixed point of $T_{\mathcal{P}}$? Whereas Theorem 3 tells us the activation value $\tilde{f}_{\mathcal{P}}(r)$ is within ε of $f_{\mathcal{P}}(r)$ after the first iteration of the RNN, this small difference could lead to bigger difference each step. Our final theorem tells us this is not the case with our encoding R :

Theorem 4. *Let \mathcal{P} be a recurrent logic program with an injective level mapping, $T_{\mathcal{P}}$ the meaning function associated with \mathcal{P} and M the unique fixed point of $T_{\mathcal{P}}$. For an arbitrary $N \in \mathbb{N}$ there exists an RNN with sigmoidal activation function for the hidden layer units and linear activation functions for the input and output layer units computing a function $f_{\mathcal{P}}$ such that there exists an $n \in \mathbb{N}$ with*

$$d_{\mathcal{P}}(R^{-1}(\tilde{f}_{\mathcal{P}}^n(0)), M) \leq \frac{1}{2^N}.$$

This theorem tells us that $R^{-1}(\tilde{f}_{\mathcal{P}}^n(0))$ and M agree in atoms of a level less than N . In other words, we can approximate the fixed point of the meaning function $T_{\mathcal{P}}$ of a recurrent logic program with an injective level mapping arbitrarily well with a recurrent neural network.

9. Discussion and Future Work

Extending the result of [8] for the propositional case, we have proven in this article that it is possible to approximate the meaning function $T_{\mathcal{P}}$ arbitrarily well in constant time¹¹ by an 3-layer FNN (Theorem 3) and that it is possible to approximate the fixed point of $T_{\mathcal{P}}$ by forming the FNN into an RNN (Theorem 4).

However, Theorem 3 does not give any clue as to how to construct such a network and how to represent the arguments of the computations, i.e. sets of atoms. While our mapping R enables us to prove these theoretical results, it does not seem to be a practical solution for representing the interpretations in the process of the iteration of $T_{\mathcal{P}}$ because of its brittleness. It does not exploit the natural insensitivity of FNN to disturbances.

A very obvious and easy solution avoiding these problems is to construct a network with input and output nodes each representing an atom, i.e. an element of the Herbrand base of \mathcal{P} . But because we do not know in advance which element we will really need to represent the result of the compu-

tation, i.e. the special interpretations, we have to represent all the atoms of the Herbrand base up to the end of our memory and it may happen, that we will not need many of them for our computation.

A solution to this problem is to use a reduced description to represent the elements of our domain, such as LRAAM- or HRR-coded terms (Labeled Recursive Auto Associative Memory [16], Holographic Reduced Representation [13]). The problem with the LRAAM model will then be that we have to train it with many of the terms we will need during the computation. Since we do not know which terms we will need during the computation the same problem will arise as before. A better solution is to use HRR-coded terms. We do not need a training before the computation and we will be able to save the used terms at the time they appear in the

computation process. Then we can save them in a special memory which will allow us to reduce the decoding error and to save only the terms we really need during the computation.

Another problem in the practical application of the iteration of $T_{\mathcal{P}}$ with RNN is that the interpretations that must be represented can become infinite during one step of the computation of the function $T_{\mathcal{P}}$. This will fill our finite memory after one step of computing $T_{\mathcal{P}}$. A solution to this problem is to use finite representations for infinite interpretations as shown in [2].

The results presented in this article consider logic programs which are recursive and admit an injective level mapping. These conditions are sufficient, but it is an open question whether they are necessary as well. We expect that they can be weakened. This research problem will be addressed in the future.

Appendix

Proof of Proposition 1: We will show $d_{\mathcal{P}}(T_{\mathcal{P}}(I), T_{\mathcal{P}}(J)) \leq \frac{1}{2} d_{\mathcal{P}}(I, J)$. Assume $d_{\mathcal{P}}(I, J) = 2^{-n}$, so that I and J agree on all ground atoms A with $|A| < n$. To show $d_{\mathcal{P}}(T_{\mathcal{P}}(I), T_{\mathcal{P}}(J)) \leq 2^{-(n+1)}$ it is enough to show that $T_{\mathcal{P}}(I), T_{\mathcal{P}}(J)$ agree on all ground atoms A with $|A| < n + 1$.

Now suppose that I and J disagree on a ground atom A with $|A| < n + 1$. There are two cases that arise: $A \in T_{\mathcal{P}}(I)$ and $A \notin T_{\mathcal{P}}(J)$ and the other way around. The second case is symmetric to the first one, so we omit its proof.

Since $A \in T_{\mathcal{P}}(I)$ and $A \notin T_{\mathcal{P}}(J)$ there is a ground instance $A \leftarrow L_1, \dots, L_n$ of a clause in \mathcal{P} such that $I \models L_1 \wedge \dots \wedge L_n$ and $J \not\models L_1 \wedge \dots \wedge L_n$. Consequently, there is a literal L_k with $k \in [1, n]$ such that $I \models L_k$ and $J \not\models L_k$. But from the definition of recurrent logic programs we know that $|A| > |L_k|$. This is a contradiction to our assumption that I and J agree on all atoms with a level less than n , so the proposition is established. \square

Proof of Proposition 2: Let n be the lowest level of an atom for which the interpretations I and J disagree. Thus $d_{\mathcal{P}}(I, J) = 2^{-n}$ and hence $d_{\mathcal{P}}(I, J)^2 = 4^{-n}$.

Since all atoms A with level $\|A\| < n$ are mapped to the same truth value by I and J it follows from the definition of R that

$$4^{-n} - \sum_{i=n+1}^{\infty} 4^{-i} \leq |R(I) - R(J)| \leq 4^{-n} + \sum_{i=n+1}^{\infty} 4^{-i},$$

which yields $\frac{2}{3} * 4^{-n} \leq |R(I) - R(J)| \leq \frac{4}{3} * 4^{-n}$ and hence the proposition holds. \square

Proof of Proposition 3: We distinguish four cases wrt r and r' :

(i) $r, r' \in \mathcal{D}_f$:

Let $I = R^{-1}(r)$ and $I' = R^{-1}(r')$. From Proposition 2 and since $I, I' \in 2^{B_{\mathcal{P}}}$ it follows

$$2/3 * d_{\mathcal{P}}(I, I')^2 \leq |R(I) - R(I')| = |r - r'| \quad (\text{A1})$$

and

$$|f_{\mathcal{P}}(r) - f_{\mathcal{P}}(r')| = |R(T_{\mathcal{P}}(I)) - R(T_{\mathcal{P}}(I'))| \leq 4/3 * d_{\mathcal{P}}(T_{\mathcal{P}}(I), T_{\mathcal{P}}(I'))^2 . \quad (\text{A2})$$

Since by Proposition 1 $T_{\mathcal{P}}$ is a contraction, i.e. $d_{\mathcal{P}}(T_{\mathcal{P}}(I), T_{\mathcal{P}}(I')) < d_{\mathcal{P}}(I, I')$, and by the definition of $d_{\mathcal{P}}$, $d_{\mathcal{P}}(T_{\mathcal{P}}(I), T_{\mathcal{P}}(I')) \leq \frac{1}{2}d_{\mathcal{P}}(I, I')$ we conclude $d_{\mathcal{P}}(T_{\mathcal{P}}(I), T_{\mathcal{P}}(I'))^2 \leq \frac{1}{4}d(I, I')^2$. The proposition follows immediately using (A1) and (A2).

(ii) $r \in \mathcal{D}_f$, $r' \notin \mathcal{D}_f$:

According to Definition 3 we have to distinguish three cases:

(iia) $r' > \max(\mathcal{D}_f)$: We can apply case (i) to show

$$|f_{\mathcal{P}}(\max(\mathcal{D}_f)) - f_{\mathcal{P}}(r)| \leq \frac{1}{2} |\max(\mathcal{D}_f) - r| . \quad (\text{A3})$$

Because $r \leq \max(\mathcal{D}_f) < r'$ we have $f_{\mathcal{P}}(r') = f_{\mathcal{P}}(\max(\mathcal{D}_f))$ and $|\max(\mathcal{D}_f) - r| < |r' - r|$ and thus the proposition follows.

(iib) $r' > \min(\mathcal{D}_f)$: The proof of this case is similar to the previous case.

(iic) $\min(\mathcal{D}_f) < r' < \max(\mathcal{D}_f)$: According to Def. 3 we have

$$f_{\mathcal{P}}(r') = \frac{M_r - r'}{M_r - m_r} \bar{f}_{\mathcal{P}}(m_r) + \frac{r' - m_r}{M_r - m_r} \bar{f}_{\mathcal{P}}(M_r) \quad (\text{A4})$$

where $m_r = \max(\mathcal{D}_f \cup (-\infty, r])$ and $M_r = \min(\mathcal{D}_f \cup [r, \infty))$. Because of $m_r, M_r, r \in \mathcal{D}_f$ we can apply case (i) and find that

$$|f_{\mathcal{P}}(m_r) - f_{\mathcal{P}}(r)| \leq \frac{1}{2} |m_r - r| \quad \text{and} \quad |f_{\mathcal{P}}(M_r) - f_{\mathcal{P}}(r)| \leq \frac{1}{2} |M_r - r'| . \quad (\text{A5})$$

Thus, we have

$$R |f_{\mathcal{P}}(m_r) - f_{\mathcal{P}}(r)| + R' |f_{\mathcal{P}}(M_r) - f_{\mathcal{P}}(r)| \leq \frac{1}{2} (R |m_r - r| + R' |M_r - r'|) . \quad (\text{A6})$$

where $R = \frac{M_r - r'}{M_r - m_r}$ and $R' = \frac{r' - m_r}{M_r - m_r}$. Since $R + R' = 1$ as well as the triangle inequality hold, the left side of (A6) is greater or equal to

$$\left| \frac{M_r - r'}{M_r - m_r} \bar{f}_{\mathcal{P}}(m_r) + \frac{r' - m_r}{M_r - m_r} \bar{f}_{\mathcal{P}}(M_r) - f_{\mathcal{P}}(r) \right| = |f_{\mathcal{P}}(r') - f_{\mathcal{P}}(r)|$$

Furthermore, r' cannot lie in the interval (m_r, M_r) the right side of (A6) is equal to

$$\frac{1}{2} \left| \frac{M_r - r'}{M_r - m_r} m_r + \frac{r' - m_r}{M_r - m_r} M_r - r \right| = \frac{1}{2} |r' - r| ,$$

and thus the proposition follows.

(iii) $r \notin \mathcal{D}_f$, $r' \in \mathcal{D}_f$:

In analogy to case (ii).

(iv) $r \notin \mathcal{D}_f$, $r' \notin \mathcal{D}_f$:

The proof of this case is almost identic to the case (ii) except that we apply case(ii) instead of case(i) to establish (A3) and (A5). We have to split a fourth subcase off the subcase corresponding to (iic), though: If r is in the interval (m_r, M_r) ¹² we find

$$f_{\mathcal{P}}(r) = \frac{r - m_r}{M_r - m_r} f_{\mathcal{P}}(M_r) + \frac{M_r - r}{M_r - m_r} f_{\mathcal{P}}(m_r) \quad (\text{A7})$$

and

$$f_{\mathcal{P}}(r') = \frac{r' - m_r}{M_r - m_r} f_{\mathcal{P}}(M_r) + \frac{M_r - r'}{M_r - m_r} f_{\mathcal{P}}(m_r) . \quad (\text{A8})$$

Consequently, we have

$$|f_{\mathcal{P}}(r') - f_{\mathcal{P}}(r)| = \left| \frac{r' - r}{M_r - m_r} f_{\mathcal{P}}(M_r) - \frac{r' - r}{M_r - m_r} f_{\mathcal{P}}(m_r) \right| = \left| (r' - r) \frac{f_{\mathcal{P}}(M_r) - f_{\mathcal{P}}(m_r)}{M_r - m_r} \right|. \quad (\text{A9})$$

Since we can apply case (i) to establish $|f_{\mathcal{P}}(M_r) - f_{\mathcal{P}}(m_r)| \leq \frac{1}{2}|M_r - m_r|$, the proposition follows from (A9). \square

Proof of Theorem 4: Let $f_{\mathcal{P}}$ be the continuous real valued function corresponding to $T_{\mathcal{P}}$. According to Theorem 3 there exists an FNN computing the function $\tilde{f}_{\mathcal{P}}$ which satisfies

$$\max_{r \in [0,1]} |f_{\mathcal{P}}(r) - \tilde{f}_{\mathcal{P}}(r)| < \varepsilon \quad (\text{A10})$$

with $\varepsilon = \frac{1}{2^{2N+3}}$. If we connect the input and output unit using a connection with weight 1, and start with an initial activation 0, we obtain an RNN which computes $\tilde{f}_{\mathcal{P}}^n(0)$ in the n -th step.

Let M be the unique fixed point of $T_{\mathcal{P}}$. Since the value of $f_{\mathcal{P}}$ is within the interval $[0, 1]$ and $\tilde{f}_{\mathcal{P}}$ differs at most by ε , $|\tilde{f}_{\mathcal{P}}^0(0) - R(M)| < 1 + 2\varepsilon$ holds. Assume now $|\tilde{f}_{\mathcal{P}}^{n-1}(0) - R(M)| < \frac{1}{2^{n-1}} + 2\varepsilon$. From (A10) follows $|f_{\mathcal{P}}(\tilde{f}_{\mathcal{P}}^{n-1}(0)) - f_{\mathcal{P}}(R(M))| < \frac{1}{2^n} + \varepsilon$. Considering $f_{\mathcal{P}}(R(M)) = R(M)$ and (A10) we get $|\tilde{f}_{\mathcal{P}}(\tilde{f}_{\mathcal{P}}^{n-1}(0)) - R(M)| < \frac{1}{2^n} + 2\varepsilon$ and thus by induction

$$|\tilde{f}_{\mathcal{P}}^n(0) - R(M)| < \frac{1}{2^n} + 2\varepsilon. \quad (\text{A11})$$

If we take $n = 2N + 2$ and insert $\varepsilon = \frac{1}{2^{2N+3}}$ we get $|\tilde{f}_{\mathcal{P}}^n(0) - R(M)| < \frac{1}{3 \cdot 2^{2N}}$. Because of the definition of R^{-1} $|R(R^{-1}(\tilde{f}_{\mathcal{P}}^n(0))) - \tilde{f}_{\mathcal{P}}^n(0)| < \frac{1}{3 \cdot 2^{2N}}$ holds.¹³ Thus, $|R(R^{-1}(\tilde{f}_{\mathcal{P}}^n(0))) - R(M)| < \frac{2}{3 \cdot 2^{2N}}$ and by the application of Proposition 2 the theorem follows. \square

Notes

1. A parallel computational model requiring $p(n)$ processors and $t(n)$ time to solve a problem of size n is *optimal* if $p(n) \times t(n) = O(T(n))$, where $T(n)$ is the sequential time to solve this problem.
2. A level mapping is extended to literals by defining $|\neg A| = |A|$.
3. As usual we denote interpretations by the set $I \subseteq B_p$ of atoms mapped to 1.
4. $\neg A \in I$ iff $A \notin I$.
5. In this article we do not use the more general result of [9] that a FNN can approximate Borel-measurable arbitrarily well because the notion of approximation in [9] includes only *almost all* points of the domain of the function instead of *all* points, as in the case of [7].
6. The obvious choice of the binary number system does not work because of the ambiguity $0.1000\dots_2 = 0.01111\dots_2$. Additionally, the quaternary number system ensures f_p being a contraction (see Prop. 3).
7. There is at most one such atom because $\|\cdot\|$ is injective.
8. Such a point exists, because the set $\{R(I) | I \in 2^{B_p}\}$ is closed. If there are two such points, we take the lower one.
9. $m_r, M_r, \min(\mathcal{D}_f)$ and $\max(\mathcal{D}_f)$ exist since \mathcal{D}_f is closed.
10. $s^n(0)$ is an abbreviation for $\underbrace{s(s(\dots s(0)\dots))}_n$.
11. For an interpretation I , $T_p(I)$ is computed in two time steps: propagation from the input to the hidden layer and propagation from the hidden layer to the output layer.
12. $r' \in (m_r, M_r)$ is not possible in subcase (iic) because $r \in \mathcal{D}_f$.
13. The nearest point to $\tilde{f}_p^n(0)$ in $\{A(I) | I \in 2^{B_p}\}$ cannot be further away than $R(M)$.

References

1. K.R. Apt and M.H. Van Emden. *Contributions to the Theory of Logic Programming*. Journal of the ACM, **29**, pp. 841–862, 1982.
2. S.-E. Bornscheuer. Generating Rational Models. In M. Maher, editor, *Proceedings of the Joint International Conference and Symposium on Logic Programming (JICSLP)*, p. 547. MIT Press, 1996.
3. K. L. Clark. Negation as failure. In Gallaire and Nicolas, editors, *Workshop Logic and Databases*, CERT, Toulouse, France, 1977.
4. A.S. d'Avila Garcez, G. Zaverucha, and L.A.V. de Carvalho. Logic programming and inductive learning in artificial neural networks. In Ch. Herrmann, F. Reine, and A. Strohmaier, editors, *Knowledge Representation in Neural Networks*, pp. 33–46, Berlin, Logos Verlag, 1997.
5. M. Fitting. Metric methods – three examples and a theorem. *Journal of Logic Programming*, **21**(3), pp. 113–127, 1994.
6. M. Fujita and R. Hasegawa and M. Koshimura and H. Fujita. *Model Generation Theorem Provers on a Parallel Inference Machine*. Proceedings of the International Conference on Generation Computer Systems, 1992.
7. K.-I. Funahashi. On the approximate realization of continuous mappings by neural networks. *Neural Networks*, **2**, pp. 183–192, 1989.
8. S. Hölldobler and Y. Kalinke. Towards a massively parallel computational model for logic programming. In *Proceedings of the ECAI94 Workshop on Combining Symbolic and Connectionist Processing*, pp. 68–77, ECCAI, 1994.
9. K. Hornik and M. Stinchcombe and H. White. *Multi-layer feedforward networks are universal approximators*. Neuronal Networks, **2**, pp. 359–366, 1989.
10. P.N. Johnson-Laird and R.M.J. Byrne. *Deduction*. LEA, Hove and London, 1991.
11. J. W. Lloyd. *Foundations of Logic Programming*. Springer, 1987.
12. R. Manthey and F. Bry. SATCHMO: A Theorem Prover Implemented in Prolog. In: E. Lusk and R. Overbeek, editors, LLNCS 310, Springer, pp. 415–434, 1988.
13. T. A. Plate. *Distributed Representations and Nested Compositional Structure*. PhD thesis, Department of Computer Science, University of Toronto, 1994.
14. J. Slaney. Scott: A model-guided theorem prover. In *Proceedings of the International Joint Conference on Artificial Intelligence*, pp. 109–114, 1993.
15. P. Smolensky. On the Proper Treatment of Connectionism. *Behavioral and Brain Sciences*, **11**, pp. 1–74, 1988.
16. A. Sperduti. Labeling RAAM. Technical Report TR-93-029, International Computer Science Institute, Berkeley, Ca, 1992.
17. S. Willard. *General Topology*. Addison-Wesley, 1970.