Learning to Resolve Natural Language Ambiguities: A Unified Approach

Dan Roth
Department of Computer Science
University of Illinois at Urbana-Champaign
Urbana 61801
danr@cs.uiuc.edu

Abstract

We analyze a few of the commonly used statistics based and machine learning algorithms for natural language disambiguation tasks and observe that they can be recast as learning linear separators in the feature space. Each of the methods makes a priori assumptions, which it employs, given the data, when searching for its hypothesis. Nevertheless, as we show, it searches a space that is as rich as the space of *all* linear separators. We use this to build an argument for a data driven approach which merely searches for a *good* linear separator in the feature space, without further assumptions on the domain or a specific problem.

We present such an approach - a sparse network of linear separators, utilizing the Winnow learning algorithm - and show how to use it in a variety of ambiguity resolution problems. The learning approach presented is attribute-efficient and, therefore, appropriate for domains having very large number of attributes.

In particular, we present an extensive experimental comparison of our approach with other methods on several well studied lexical disambiguation tasks such as context-sensitive spelling correction, prepositional phrase attachment and part of speech tagging. In all cases we show that our approach either outperforms other methods tried for these tasks or performs comparably to the best.

Introduction

Many important natural language inferences can be viewed as problems of resolving ambiguity, either semantic or syntactic, based on properties of the surrounding context. Examples include part-of speech tagging, word-sense disambiguation, accent restoration, word choice selection in machine translation, context-sensitive spelling correction, word selection in speech recognition and identifying discourse markers. In each of these problems it is necessary to disambiguate two or more [semantically, syntactically or structurally]-distinct forms which have been fused together into the same representation in some medium. In a prototypical instance of this problem, word sense disambiguation,

---

distinct semantic concepts such as `interest rate` and has `interest in Math` are conflated in ordinary text. The surrounding context - word associations and syntactic patterns in this case - are sufficient to identify the correct form.

Many of these are important stand-alone problems but even more important is their role in many applications including speech recognition, machine translation, information extraction and intelligent human-machine interaction. Most of the ambiguity resolution problems are at the lower level of the natural language inferences chain; a wide range and a large number of ambiguities are to be resolved simultaneously in performing any higher level natural language inference.

Developing learning techniques for language disambiguation has been an active field in recent years and a number of statistics based and machine learning techniques have been proposed. A partial list consists of Bayesian classifiers (Gale, Church, & Yarowsky 1993), decision lists (Yarowsky 1994), Bayesian hybrids (Golding 1995), HMMs (Charniak 1993), inductive logic methods (Zelle & Mooney 1996), memory-based methods (Zavrel, Daelemans, & Veenstra 1997) and transformation-based learning (Brill 1995). Most of these have been developed in the context of a specific task although claims have been made as to their applicativity to others.

In this paper we cast the disambiguation problem as a learning problem and use tools from computational learning theory to gain some understanding of the assumptions and restrictions made by different learning methods in shaping their search space.

The learning theory setting helps in making a few interesting observations. We observe that many algorithms, including naive Bayes, Brill's transformation based method, Decision Lists and the Back-off estimation method can be re-cast as learning linear separators in their feature space. As learning techniques for linear separators these techniques are limited in that, in general, they cannot learn *all* linearly separable functions. Nevertheless, we find, they still search a space that is as complex, in terms of its VC dimension, as the space of all linear separators. This has implications to the generalization ability of their hypotheses. Together with

the fact that different methods seem to use different a priori assumptions in guiding their search for the linear separator, it raises the question of whether there is an alternative - search for the best linear separator in the feature space, without resorting to assumptions about the domain or any specific problem.

Partly motivated by these insights, we present a new algorithm, and show how to use it in a variety of disambiguation tasks. The architecture proposed, *SNOW*, is a Sparse Network Of linear separators which utilizes the Winnow learning algorithm. A target node in the network corresponds to a candidate in the disambiguation task; all subnetworks learn autonomously from the same data, in an on line fashion, and at run time, they compete for assigning the correct meaning. The architecture is data-driven (in that its nodes are allocated as part of the learning process and depend on the observed data) and supports efficient on-line learning. Moreover, The learning approach presented is attribute-efficient and, therefore, appropriate for domains having very large number of attributes. All together, We believe that this approach has the potential to support, within a single architecture, a large number of simultaneously occurring and interacting language related tasks.

To start validating these claims we present experimental results on three disambiguation tasks. *Prepositional phrase attachment (PPA)* is the task of deciding whether the Prepositional Phrase (PP) attaches to the noun phrase (NP), as in Buy the car with the steering wheel or to the verb phrase (VP), as in Buy the car with his money. *Context-sensitive Spelling correction (Spell)* is the task of fixing spelling errors that result in valid words, such as It's not to late, where too was mistakenly typed as to. *Part of speech tagging (POS)* is the task of assigning each word in a given sentence the part of speech it assumes in this sentence. For example, assign N or V to talk in the following pair of sentences: Have you listened to his (him) talk ?. In all cases we show that our approach either outperforms other methods tried for these tasks or performs comparably to the best.

This paper focuses on analyzing the learning problem and on motivating and developing the learning approach; therefore we can only present the bottom line of the experimental studies and the details are deferred to companion reports.

The Learning Problem

Disambiguation tasks can be viewed as general classification problems. Given an input sentence we would like to assign it a single property out of a set of potential properties. Formally, given a sentence s and a predicate p defined on the sentence, we let $C = \{c_1, c_2, \ldots c_m\}$ be the collection of possible values this predicate can assume in s. It is assumed that one of the elements in C is the *correct* assignment, $c(s, p)$. c_i can take values from $\{site, cite, sight\}$ if the predicate p is the correct spelling of any occurrence of a word from this set in the sentence; it can take values from

$\{v, n\}$ if the predicate p is the attachment of the PP to the preceding VP (v) or the preceding NP (n), or it can take values from $\{industrial, living\ organism\}$ if the predicate is the meaning of the word plant in the sentence. In some cases, such as part of speech tagging, we may apply a collection P of different predicates to the same sentence, when tagging the first, second, kth word in the sentence, respectively. Thus, we may perform a classification operation on the sentence multiple times. However, in the following definitions it would suffice to assume that there is a single pre-defined predicate operating on the sentence s; moreover, since the predicate studied will be clear from the context we omit it and denote the correct classification simply by $c(s)$.

A *classifier h* is a function that maps the set S of all sentences[1], given the task defined by the predicate p, to a single value in C, $h : S \to C$.

In the setting considered here the classifier h is selected by a training procedure. That is, we assume[2] a class of functions $\mathcal{H}$, and use the training data to select a member of this class. Specifically, given a *training corpus S_{tr}* consisting of *labeled example* $(s, c(s))$, a learning algorithm selects a hypothesis $h \in \mathcal{H}$, the classifier.

The performance of the classifier is measured empirically, as the fraction of correct classifications it performs on a set S_{ts} of test examples. Formally,

$$Perf(f) = |\{s \in S_{ts}|h(s) = c(s)\}|/|\{s \in S_{ts}\}|. \quad (1)$$

A sentence s is represented as a collection of *features*, and various kinds of feature representation can be used. For example, typical features used in correcting context-sensitive spelling are *context words* – which test for the presence of a particular word within $\pm k$ words of the target word, and *collocations* – which test for a pattern of up to ℓ contiguous words and/or part-of-speech tags around the target word.

It is useful to consider features as sequences of *tokens* (e.g., words in the sentence, or pos tags of the words). In many applications (e.g., n-gram language models), there is a clear ordering on the features. We define here a natural partial order $\prec$ as follows: for features f, g define $f \prec g \equiv f \subseteq g$, where on the right end side features are viewed simply as sets of tokens[3]. A feature f is of *order k* if it consists of k tokens.

A definition of a disambiguation problem consists of the task predicate p, the set C of possible classifications and the set $\mathcal{F}$ of features. $\mathcal{F}^{(k)}$ denotes the features of order k. Let $|\mathcal{F}| = n$, and x_i be the ith feature. x_i can either be present (*active*) in a sentence s (we then say that $x_i = 1$), or absent from it ($x_i = 0$). Given that,

---

[1]The basic unit studied can be a paragraph or any other unit, but for simplicity we will always call it a sentence.

[2]This is usually not made explicit in statistical learning procedures, but is assumed there too.

[3]There are many ways to define features and order relations among them (e.g., restricting the number of tokens in a feature, enforcing sequential order among them, etc.). The following discussion does not depend on the details; one option is presented to make the discussion more concrete.

a sentence s can be represented as the set of all *active* features in it $s = (x_{i_1}, x_{i_2}, \ldots x_{i_m})$.

From the stand point of the general framework the exact mapping of a sentence to a feature set will not matter, although it is crucially important in the specific applications studied later in the paper. At this point it is sufficient to notice that the a sentence can be mapped into a binary feature vector. Moreover, w.l.o.g we assume that $|C| = 2$; moving to the general case is straight forward. From now on we will therefore treat classifiers as Boolean functions, $h : \{0,1\}^n \to \{0,1\}$.

Approaches to Disambiguation

Learning approaches are usually categorized as statistical (or probabilistic) methods and symbolic methods. However, all learning methods are statistical in the sense that they attempt to make inductive generalization from observed data and use it to make inferences with respect to previously unseen data; as such, the statistical based theories of learning (Vapnik 1995) apply equally to both. The difference may be that symbolic methods do not explicitly use probabilities in the hypothesis. To stress the equivalence of the approaches further in the following discussion we will analyze two "statistical" and two "symbolic" approaches.

In this section we present four widely used disambiguation methods. Each method is first presented as known and is then re-cast as a problem of learning a linear separator. That is, we show that, there is a linear condition $\sum_{x_i \in \mathcal{F}} w_i x_i > \theta$ such that, given a sentence $s = (x_{i_1}, x_{i_2}, \ldots x_{i_m})$, the method predicts $c = 1$ if the condition holds for it, and $c = 0$ otherwise.

Given an example $s = (x_1, x_2 \ldots x_m)$ a probabilistic classifier h works by choosing the element of C that is most probable, that is $h(s) = argmax_{c_i \in C} Pr(c_i | x_1, x_2, \ldots x_m)$[4], where the probability is the empirical probability estimated from the labeled training data. In general, it is unlikely that one can estimate the probability of the event of interest $(c_i | x_1, x_2, \ldots x_m)$ directly from the training data. There is a need to make some probabilistic assumptions in order to evaluate the probability of this event indirectly, as a function of "more frequent" events whose probabilities can be estimated more robustly. Different probabilistic assumptions give rise to difference learning methods and we describe two popular methods below.

The naive Bayes estimation (NB) The naive Bayes estimation (e.g., (Duda & Hart 1973)) assumes that given the class value $c \in C$ the features values are statistically independent. With this assumption and using Bayes rule the Bayes optimal prediction is given by: $h(s) = argmax_{c_i \in C} \Pi_{i=1}^m Pr(x_j | c_i) P(c_i)$.

The prior probabilities $p(c_i)$ (i.e., the fraction of training examples labeled with c_i) and the conditional probabilities $Pr(x_j | c_i)$ (the fraction of the training ex-

amples labeled c_i in which the jth feature has value x_j) can be estimated from the training data fairly robustly[5], giving rise to the naive Bayes predictor. According to it, the optimal decision is $c = 1$ when

$$P(c=1)\Pi_i P(x_i | c=1)/P(c=0)\Pi_i P(x_i | c=0) > 1.$$

Denoting $p_i \equiv P(x_i = 1 | c = 1), q_i \equiv P(x_i = 1 | c = 0)$, $P(c = r) \equiv P(r)$, we can write this condition as

$$\frac{P(1)\Pi_i p_i^{x_i}(1-p_i)^{1-x_i}}{P(0)\Pi_i q_i^{x_i}(1-q_i)^{1-x_i}} = \frac{P(1)\Pi_i(1-p_i)(\frac{p_i}{1-p_i})^{x_i}}{P(0)\Pi_i(1-q_i)(\frac{q_i}{1-q_i})^{x_i}} > 1,$$

and by taking log we get that using naive Bayes estimation we predict $c = 1$ if and only if

$$\log \frac{P(1)}{P(0)} + \sum_i \log \frac{1-p_i}{1-q_i} + \sum_i (\log \frac{p_i}{1-p_i} \frac{1-q_i}{q_i}) x_i > 0.$$

We conclude that the decision surface of the naive Bayes algorithm is given by a linear function in the feature space. Points which reside on one side of the hyperplane are more likely to be labeled 1 and points on the other side are more likely to be labeled 0.

This representation immediately implies that this predictor is optimal also in situations in which the conditional independence assumption does no hold. However, a more important consequence to our discussion here is the fact that not all linearly separable functions can be represented using this predictor (Roth 1998).

The back-off estimation (BO) Back-off estimation is another method for estimating the conditional probabilities $Pr(c_i | s)$. It has been used in many disambiguation tasks and in learning models for speech recognition (Katz 1987; Chen & Goodman 1996; Collins & Brooks 1995). The back-off method suggests to estimate $Pr(c_i | x_1, x_2, \ldots, x_m)$ by interpolating the more robust estimates that can be attained for the conditional probabilities of more general events. Many variation of the method exist; we describe a fairly general one and then present the version used in (Collins & Brooks 1995), which we compare with experimentally.

When applied to a disambiguation task, BO assumes that the sentence itself (the basic unit processed) is a feature[6] of maximal order $f = f^{(k)} \in \mathcal{F}$. We estimate

$$Pr(c_i | s) = Pr(c_i | f^{(k)}) = \sum_{\{f \in \mathcal{F} | f \prec f^{(k)}\}} \lambda_f Pr(c_i | f).$$

---

[4] As usual, we use the notation $Pr(c_i | x_1, x_2, \ldots x_m)$ as a shortcut for $Pr(c = c_i | x_1 = a_1, x_2 = a_2, \ldots x_m = a_m)$.

[5] Problems of sparse data may arise, though, when a specific value of x_i observed in testing has occurred infrequently in the training, in conjunction with c_j. Various *smoothing* techniques can be employed to get more robust estimations but these considerations will not affect our discussion and we disregard them.

[6] The assumption that the maximal order feature is the classified sentence is made, for example, in (Collins & Brooks 1995). In general, the method deals with multiple features of the maximal order by assuming their conditional independence, and superimposing the NB approach.

The sum is over all features f which are more general (and thus occur more frequently) than $f^{(k)}$. The conditional probabilities on the right are empirical estimates measured on the training data, and the coefficients λ_f are also estimated given the training data. (Usually, these are maximum likelihood estimates evaluated using iterative methods, e.g. (Samuelsson 1996)).

Thus, given an example $s = (x_1, x_2 \ldots x_m)$ the BO method predicts $c = 1$ if and only if

$$\sum_{i=1}^{|\mathcal{F}|} \lambda_i (Pr(c = 1|x_i) - Pr(c = 0|x_i))x_i > 0,$$

a linear function over the feature space.

For computational reasons, various simplifying assumptions are made in order to estimate the coefficients λ_f; we describe here the method used in (Collins & Brooks 1995)[7]. We denote by $\mathcal{N}(f^{(j)})$ the number of occurrences of the jth order feature $f^{(j)}$ in the training data. Then BO estimates $P = Pr(c_i|f^{(k)})$ as follows:

If $\mathcal{N}(f^{(k)}) > 0$, $\qquad P = Pr(c_i|f^{(k)})$
Else if $\sum_{f \in \mathcal{F}^{(k-1)}} \mathcal{N}(f) > 0$, $P = \frac{1}{|\mathcal{F}^{(k-1)}|} \sum Pr(c_i|f^{(k-1)})$
Else if ...
Else if $\sum_{f \in \mathcal{F}^{(1)}} \mathcal{N}(f) > 0$, $\quad P = \frac{1}{|\mathcal{F}^{(1)}|} \sum Pr(c_i|f^{(1)})$

In this case, it is easy to write down the linear separator defining the estimate in an explicit way. Notice that with this estimation, given a sentence s, only the highest order features active in it are considered. Therefore, one can define the weights of the jth order feature in an inductive way, making sure that it is larger than the sum of the weights of the smaller order features. Leaving out details, it is clear that we get a simple representation of a linear separator over the feature space, that coincides with the BO algorithm.

It is important to notice that the assumptions made in the BO estimation method result in a linear decision surface that is, in general, *different* from the one derived in the NB method.

Transformation Based Learning (TBL) Transformation based learning (Brill 1995) is a machine learning approach for rule learning. It has been applied to a number of natural language disambiguation tasks, often achieving state-of-the-art accuracy.

The learning procedure is a mistake-driven algorithm that produces a set of rules. Irrespective of the learning procedure used to derive the TBL representation, we focus here on the final hypothesis used by TBL and how it is evaluated, given an input sentence, to produce a prediction. We assume, w.l.o.g, $|C| = 2$.

The hypothesis of TBL is an ordered list of transformations. A *transformation* is a rule with an antecedent

t and a consequent[8] $c \in C$. The antecedent t is a condition on the input sentence. For example, in Spell, a condition might be word W occurs within $\pm k$ of the target word. That is, applying the condition to a sentence s defines a feature $t(s) \in \mathcal{F}$. Phrased differently, the application of the condition to a given sentence s, checks whether the corresponding feature is active in this sentence. The condition holds if and only if the feature is active in the sentence.

An ordered list of transformations (the TBL hypothesis), is evaluated as follows: given a sentence s, an initial label $c \in C$ is assigned to it. Then, each rule is applied, in order, to the sentence. If the feature defined by the condition of the rule applies, the current label is replaced by the label in the consequent. This process goes on until the last rule in the list is evaluated. The last label is the output of the hypothesis.

In its most general setting, the TBL hypothesis is not a classifier (Brill 1995). The reason is that the truth value of the condition of the ith rule may change while evaluating one of the preceding rules. However, in many applications and, in particular, in Spell (Mangu & Brill 1997) and PPA (Brill & Resnik 1994) which we discuss later, this is not the case. There, the conditions do not depend on the labels, and therefore the output hypothesis of the TBL method can be viewed as a classifier. The following analysis applies only for this case.

Using the terminology introduced above, let $(x_{i_1}, c_{i_1}), (x_{i_2}, c_{i_2}), \ldots (x_{i_k}, c_{i_k})$ be the ordered sequence of rules defining the output hypothesis of TBL. (Notice that it is quite possible, and happens often in practice, for a feature to appear more than once in this sequence, even with different consequents). While the above description calls for evaluating the hypothesis by sequentially evaluating the conditions, it is easy to see that the following simpler procedure is sufficient:

Search the ordered sequence in a reversed order. Let x_{i_j} be the first active feature in the list (i.e., the largest j). Then the hypothesis predicts c_{i_j}.

Alternatively, the TBL hypothesis can be represented as a (positive) 1-Decision-List (p1-DL) (Rivest 1987), over the set $\mathcal{F}$ of features[9]. Given the p1-DL represen-

If $\qquad$ x_{i_k} is active then predict c_k.
Else $\qquad$ If $x_{i_{k-1}}$ is active then predict c_{k-1}.
Else ...
Else $\qquad$ If x_1 is active then predict c_1.
Else $\qquad$ Predict the initial value

Figure 1: TBL as a p1-Decision List

---

[8]The consequent is sometimes described as a *transformation* $c_i \rightarrow c_j$, with the semantics – if the current label is c_i, relabel it c_j. When $|C| = 2$ it is equivalent to simply using c_j as the consequent.

[9]Notice, the order of the features is reversed. Also, multiple occurrences of features can be discarded, leaving only the last rule in which this feature occurs. By "positive" we mean that we never condition on the absence of a feature, only on its presence.

---

[7]There, the empirical ratios are smoothed; experimentally, however, this yield only a slight improvement, going from 83.7% to 84.1% so we present it here in the pure form.

tation (Fig 1), we can now represent the hypothesis as a linear separator over the set $\mathcal{F}$ of features. For simplicity, we now name the class labels $\{-1, +1\}$ rather than $\{0, 1\}$. Then, the hypothesis predicts $c = 1$ if and only if $\sum_{j=1}^{k} 2^j \cdot c_{i_j} \cdot x_{i_j} > 0$. Clearly, with this representation the active feature with the highest index dominates the prediction, and the representations are equivalent[10].

Decision Lists (p1-DL) It is easy to see (details omitted), that the above analysis applies to p1-DL, a method used, for example, in (Yarowsky 1995). The BO and p1-DL differ only in that they keep the rules in reversed order, due to different evaluation methods.

The Linear Separator Representation

To summarize, we have shown:
claim: *All the methods discussed – NB, BO, TBL and p1-DL search for a decision surface which is a linear function in the feature space.*

This is not to say that these methods assume that the data is linearly separable. Rather, all the methods assume that the feature space is divided by a linear condition (i.e., a function of the form $\sum_{x_i \in \mathcal{F}} w_i x_i > \theta$) into two regions, with the property that in one of the defined regions the more likely prediction is 0 and in the other, the more likely prediction is 1.

As pointed out, it is also instructive to see that these methods yield *different* decision surfaces and that they cannot represent every linearly separable function.

Theoretical Support for the Linear Separator Framework

In this section we discuss the implications these observations have from the learning theory point of view.

In order to do that we need to resort to some of the basic ideas that justify inductive learning. Why do we hope that a classifier learned from the training corpus will perform well (on the test data) ? Informally, the basic theorem of learning theory (Valiant 1984; Vapnik 1995) guarantees that, if the training data and the test data are sampled from the same distribution[11], good performance on the training corpus guarantees good performance on the test corpus.

If one knows something about the model that generates the data, then estimating this model may yield good performance on future examples. However, in the problems considered here, no reasonable model is known, or is likely to exist. (The fact that the assumptions discussed above disagree with each other, in general, may be viewed as a support for this claim.)

In the absence of this knowledge a learning method merely attempts to make correct predictions. Under these conditions, it can be shown that the error of a classifier selected from class $\mathcal{H}$ on (previously unseen) test data, is bounded by the sum of its training error and a function that depends linearly on the complexity of $\mathcal{H}$. This complexity is measured in terms of a combinatorial parameter - the VC-dimension of the class $\mathcal{H}$ (Vapnik 1982) - which measures the richness of the function class. (See (Vapnik 1995; Kearns & Vazirani 1992)) for details).

We have shown that all the methods considered here look for a linear decision surface. However, they do make further assumptions which seem to restrict the function space they search in. To quantify this line of argument we ask whether the assumptions made by the different algorithms significantly reduce the complexity of the hypothesis space. The following claims show that this is not the case; the VC dimension of the function classes considered by all methods are as large as that of the full class of linear separators.

Fact 1: *The VC dimension of the class of linear separators over n variables is $n + 1$.*

Fact 2: *The VC dimension of the class of p1-DL over n variables[12] is $n + 1$.*

Fact 3: *The VC dimension of the class of linear separators derived by either NB or BO over n variables is bounded below by n.*

Fact 1 is well known; 2 and 3 can be derived directly from the definition (Roth 1998).

The implication is that a method that merely searches for the optimal linear decision surface given the training data may, in general, outperform all these methods also on the test data. This argument can be made formal by appealing to a result of (Kearns & Schapire 1994), which shows that even when there is no perfect classifier, the optimal linear separator on a polynomial size set of training examples is optimal (in a precise sense) also on the test data.

The optimality criterion we seek is described in Eq. 1. A linear classifier that minimizes the number of disagreements (the sum of the false positives and false negatives classifications). This task, however, is known to be NP-hard (Höffgen & Simon 1992), so we need to resort to heuristics. In searching for good heuristics we are guided by computational issues that are relevant to the natural language domain. An essential property of an algorithm is being feature-efficient. Consequently, the approach describe in the next section makes use of the *Winnow* algorithm which is known to produce good results when a linear separator exists, as well as under certain more relaxed assumptions (Littlestone 1991).

---

[10] In practice, there is no need to use this representation, given the efficient way suggested above to evaluate the classifier. In addition, very few of the features in $\mathcal{F}$ are active in every example, yielding more efficient evaluation techniques (e.g., (Valiant 1998))

[11] This is hard to define in the context of natural language; typically, this is understood as texts of similar nature; see a discussion of this issue in (Golding & Roth 1996).

[12] In practice, when using p1-DL as the hypothesis class (i.e., in TBL) an effort is made to discard many of the features and by that reduce the complexity of the space; however, this process, which is data driven and does not a-priori restrict the function class can be employed by other methods as well (e.g., (Blum 1995)) and is therefore orthogonal to these arguments.

The *SNOW* Approach

The *SNOW* architecture is a network of threshold gates. Nodes in the first layer of the network are allocated to input features in a data-driven way, given the input sentences. Target nodes (i.e., the element $c \in C$) are represented by nodes in the second layer. Links from the first to the second layer have weights; each target node is thus defined as a (linear) function of the lower level nodes. (A similar architecture which consists of an additional layer is described in (Golding & Roth 1996). Here we do not use the "cloud" level described there.)

For example, in Spell, target nodes represent members of the confusion sets; in POS, target nodes correspond to different *pos tags*. Each target node can be thought of as an autonomous network, although they all feed from the same input. The network is *sparse* in that a target node need not be connected to all nodes in the input layer. For example, it is not connected to input nodes (features) that were never active with it in the same sentence, or it may decide, during training to disconnect itself from some of the irrelevant inputs.

Learning in *SNOW* proceeds in an on-line fashion[13]. Every example is treated autonomously by each target subnetworks. Every labeled example is treated as positive for the target node corresponding to its label, and as negative to all others. Thus, every example is used once by all the nodes to refine their definition in terms of the others and is then discarded. At prediction time, given an input sentence which activates a subset of the input nodes, the information propagates through all the subnetworks; the one which produces the highest activity gets to determine the prediction.

A local learning algorithm, Winnow (Littlestone 1988), is used at each target node to learn its dependence on other nodes. Winnow is a mistake driven on-line algorithm, which updates its weights in a multiplicative fashion. Its key feature is that the number of examples it requires to learn the target function grows linearly with the number of *relevant* attributes and only logarithmically with the total number of attributes. Winnow was shown to learn efficiently any linear threshold function and to be robust in the presence of various kinds of noise, and in cases where no linear-threshold function can make perfect classifications and still maintain its abovementioned dependence on the number of total and relevant attributes (Littlestone 1991; Kivinen & Warmuth 1995).

Notice that even when there are only two target nodes and the cloud size (Golding & Roth 1996) is 1 *SNOW* behaves differently than pure Winnow. While each of the target nodes is learned using a positive Winnow algorithm, a winner-take-all policy is used to determine the prediction. Thus, we do not use the learning algorithm here simply as a discriminator. One reason is that the *SNOW* architecture, influenced by the Neuroidal system (Valiant 1994), is being used in a system

---

[13]Although for the purpose of the experimental study we do not update the network while testing.

Table 1: **Spell System comparison.** The second column gives the number of test cases. All algorithms were trained on 80% of Brown and tested on the other 20%; Baseline simply identifies the most common member of the confusion set during training, and guesses it every time during testing.

Sets	Cases	Baseline	NB	TBL	*SNOW*
14	1503	71.1	89.9	88.5	93.5
21	4336	74.8	93.8		96.4

developed for the purpose of learning knowledge representations for natural language understanding tasks, and is being evaluated on a variety of tasks for which the node allocation process is of importance.

Experimental Evidence

In this section we present experimental results for three of the most well studied disambiguation problems, Spell, PPA and POS. We present here only the bottom-line results of an extensive study that appears in companion reports (Golding & Roth 1998; Krymolovsky & Roth 1998; Roth & Zelenko 1998).

Context Sensitive Spelling Correction Context-sensitive spelling correction is the task of fixing spelling errors that result in valid words, such as *It's not to late*, where *too* was mistakenly typed as *to*.

We model the ambiguity among words by *confusion sets*. A confusion set $C = \{c_1, \ldots, c_n\}$ means that each word c_i in the set is ambiguous with each other word. All the results reported here use the same pre-defined set of confusion sets (Golding & Roth 1996).

We compare *SNOW* against TBL (Mangu & Brill 1997) and a naive-Bayes based system (NB). The latter system presents a few augmentations over the simple naive Bayes (but still shares the same basic assumptions) and is among the most successful methods tried for the problem (Golding 1995). An indication that a Winnow-based algorithm performs well on this problem was presented in (Golding & Roth 1996). However, the system presented there was more involved than *SNOW* and allows more expressive output representation than we allow here. The output representation of all the approaches compared is a linear separator.

The results presented in Table 1 for NB and *SNOW* are the (weighted) average results of 21 confusion sets, 19 of them are of size 2, and two of size 3. The results presented for the TBL[14] method are taken from (Mangu & Brill 1997) and represent an average on a subset of 14 of these, all of size 2.

Prepositional Phrase Attachment The problem is to decide whether the Prepositional Phrase (PP) attaches to the noun phrase, as in Buy the car with

---

[14]Systems are compared on the same feature set. TBL was also used with an enhanced feature set (Mangu & Brill 1997) with improved results of 93.3% but we have not run the other systems with this set of features.

Table 2: **PPA System comparison.** All algorithms were trained on 20801 training examples from the WSJ corpus tested 3097 previously unseen examples from this corpus; all the system use the same feature set.

Test cases	Baseline	NB	TBL	BO	*SNOW*
3097	59.0	83.0	81.9	84.1	83.9

the steering wheel or the verb phrase, as in Buy the car with his money. Earlier works on this problem (Ratnaparkhi, Reynar, & Roukos 1994; Brill & Resnik 1994; Collins & Brooks 1995) consider as input the four head words involved in the attachment - the VP head, the first NP head, the preposition and the second NP head (in this case, buy, car, with and steering wheel, respectively). These four-tuples, along with the attachment decision constitute the labeled input sentence and are used to generate the feature set. The features recorded are all sub-sequences of the 4-tuple, total of 15 for every input sentence. The data set used by all the systems in this in this comparison was extracted from the Penn Treebank WSJ corpus by (Ratnaparkhi, Reynar, & Roukos 1994). It consists of 20801 training examples and 3097 separate test examples. In a companion paper we describe an extensive set of experiments with this and other data sets, under various conditions. Here we present only the bottom line results that provide direct comparison with those available in the literature[15]. The results presented in Table 2 for NB and *SNOW* are the results of our system on the 3097 test examples. The results presented for the TBL and BO are on the same data set, taken from (Collins & Brooks 1995).

Part of Speech Tagging A part of speech tagger assigns each word in a sentence the part of speech that it assumes in that sentence. See (Brill 1995) for a survey of much of the work that has been done on POS in the past few years. Typically, in English there will be between 30 and 150 different parts of speech depending on the tagging scheme. In the study presented here, following (Brill 1995) and many other studies there are 47 different tags. Part-of-speech tagging suggests a special challenge to our approach, as the problem is a multi-class prediction problem (Roth & Zelenko 1998). In the *SNOW* architecture, we devote one linear separator to each pos tag and each sub network learns to separate its corresponding pos tag from all others. At run time, all class nodes process the given sentence, applying many classifiers simultaneously. The classifiers then compete for deciding the pos of this word, and the node that records the highest activity for a given word in a sentence determines its pos. The methods compared use

---

[15]*SNOW* was evaluated with an enhanced feature set (Krymolovsky & Roth 1998) with improved results of 84.8%. (Collins & Brooks 1995) reports results of 84.4% on a different enhanced set of features, but other systems were not evaluated on these sets.

Table 3: **POS System comparison.** The first column gives the number of test cases. All algorithms were trained on 550,000 words of the tagged WSJ corpus. Baseline simply predicts according to the most common pos tag for the word in the training corpus.

Test cases	Baseline	TBL	*SNOW*
250,000	94.4	96.9	96.8

context and collocation features as in (Brill 1995).

Given a sentence, each word in the sentence is assigned an initial tag, based on the most common part of speech in the training corpus. Then, for each word in the sentence, the network processes the sentence, and makes a suggestion for the pos of this word. Thus, the input for the predictor is noisy, since the initial assignment is not accurate for many of the words. This process can repeat a few times, where after predicting the pos of a word in the sentence we re-compute the new feature-based representation of the sentence and predict again. Each time the input to the predictors is expected to be slightly less noisy. In the results presented here, however, we present the performance without the recycling process, so that we maintain the linear function expressivity (see (Roth & Zelenko 1998) for details).

The results presented in Table 3 are based on experiments using 800,000 words of the Penn Treebank Tagged WSJ corpus. About 550,000 words were used for training and 250,000 for testing. *SNOW* and TBL were trained and tested on the same data.

Conclusion

We presented an analysis of a few of the commonly used statistics based and machine learning algorithms for ambiguity resolution tasks. We showed that all the algorithms investigated can be re-cast as learning linear separators in the feature space. We analyzed the complexity of the function space in which each of these method searches, and show that they all search a space that is as complex as the space of all linear separators. We used these to argue motivate our approach of learning a sparse network of linear separators (*SNOW*), which learns a network of linear separator by utilizing the Winnow learning algorithm. We then presented an extensive experimental study comparing the *SNOW* based algorithms to other methods studied in the literature on several well studied disambiguation tasks. We present experimental results on Spell, PPA and POS. In all cases we show that our approach either outperformed other methods tried for these tasks or performs comparably to the best. We view this as a strong evidence to that this approach provides a unified framework for the study of natural language disambiguation tasks.

The importance of providing a unified framework stems from the fact the essentially all ambiguity resolution problems that are addressed here are at the lower level of the natural language inferences chain. A large

number of different kinds of ambiguities are to be resolved simultaneously in performing any higher level natural language inference (Cardie 1996). Naturally, these processes, acting on the same input and using the same "memory", will interact. A unified view of ambiguity resolution within a single architecture, is valuable if one wants understand how to put together a large number of these inferences, study interactions among them and make progress towards using these in performing higher level inferences.

References

Blum, A. 1995. Empirical support for Winnow and weighted-majority based algorithms: results on a calendar scheduling domain. In *Proc. 12th International Conference on Machine Learning*, 64–72. Morgan Kaufmann.

Brill, E., and Resnik, P. 1994. A rule-based approach to prepositional phrase attachment disambiguation. In *Proc. of COLING*.

Brill, E. 1995. Transformation-based error-driven learning and natural language processing: A case study in part of speech tagging. *Computational Linguistics* 21(4):543–565.

Cardie, C. 1996. *Embedded Machine Learning Systems for natural language processing: A general framework.* Springer. 315–328.

Charniak, E. 1993. *Statistical Language Learning.* MIT Press.

Chen, S., and Goodman, J. 1996. An empirical study of smoothing techniques for language modeling. In *Proc. of the Annual Meeting of the ACL.*

Collins, M., and Brooks, J. 1995. Prepositional phrase attachment through a backed-off model. In *Proceedings of Third the Workshop on Very Large Corpora.*

Duda, R. O., and Hart, P. E. 1973. *Pattern Classification and Scene Analysis.* Wiley.

Gale, W. A.; Church, K. W.; and Yarowsky, D. 1993. A method for disambiguating word senses in a large corpus. *Computers and the Humanities* 26:415–439.

Golding, A. R., and Roth, D. 1996. Applying winnow to context-sensitive spelling correction. In *Proc. of the International Conference on Machine Learning*, 182–190.

Golding, A. R., and Roth, D. 1998. A winnow-based approach to word correction. *Machine Learning*. Special issue on Machine Learning and Natural Language; to appear.

Golding, A. R. 1995. A bayesian hybrid method for context-sensitive spelling correction. In *Proceedings of the 3rd workshop on very large corpora, ACL-95.*

Höffgen, K., and Simon, H. 1992. Robust trainability of single neurons. In *Proc. 5th Annu. Workshop on Comput. Learning Theory*, 428–439. New York, New York: ACM Press.

Katz, S. M. 1987. Estimation of probabilities from sparse data for the language model component of a speech recognizer. *IEEE Transactions on Acoustics, speech, and Signal Processing* 35(3):400–401.

Kearns, M., and Schapire, R. 1994. Efficient distribution-free learning of probabilistic concepts. *Journal of Computer and System Sciences* 48:464–497.

Kearns, M., and Vazirani, U. 1992. *Introduction to computational Learning Theory.* MIT Press.

Kivinen, J., and Warmuth, M. K. 1995. Exponentiated gradient versus gradient descent for linear predictors. In *Proceedings of the Annual ACM Symp. on the Theory of Computing.*

Krymolovsky, Y., and Roth, D. 1998. Prepositional phrase attachment. In Preparation.

Littlestone, N. 1988. Learning quickly when irrelevant attributes abound: A new linear-threshold algorithm. *Machine Learning* 2:285–318.

Littlestone, N. 1991. Redundant noisy attributes, attribute errors, and linear threshold learning using Winnow. In *Proc. 4th Annu. Workshop on Comput. Learning Theory*, 147–156. San Mateo, CA: Morgan Kaufmann.

Mangu, L., and Brill, E. 1997. Automatic rule acquisition for spelling correction. In *Proc. of the International Conference on Machine Learning*, 734–741.

Ratnaparkhi, A.; Reynar, J.; and Roukos, S. 1994. A maximum entropy model for prepositional phrase attachment. In *ARPA.*

Rivest, R. L. 1987. Learning decision lists. *Machine Learning* 2(3):229–246.

Roth, D., and Zelenko, D. 1998. Part of speech tagging using a network of linear separators. Submitted.

Roth, D. 1998. Learning to resolve natural language ambiguities:a unified approach. Long Version, in Preparation.

Samuelsson, C. 1996. Handling sparse data by successive abstraction. In *Proc. of COLING.*

Valiant, L. G. 1984. A theory of the learnable. *Communications of the ACM* 27(11):1134–1142.

Valiant, L. G. 1994. *Circuits of the Mind.* Oxford University Press.

Valiant, L. G. 1998. Projection learning. In *Proc. of the Annual ACM Workshop on Computational Learning Theory*, xxx–xxx.

Vapnik, V. N. 1982. *Estimation of Dependences Based on Empirical Data.* New York: Springer-Verlag.

Vapnik, V. N. 1995. *The Nature of Statistical Learning Theory.* New York: Springer-Verlag.

Yarowsky, D. 1994. Decision lists for lexical ambiguity resolution: application to accent restoration in Spanish and French. In *Proc. of the Annual Meeting of the ACL*, 88–95.

Yarowsky, D. 1995. Unsupervised word sense disambiguation rivaling supervised methods. In *Proceedings of ACL-95.*

Zavrel, J.; Daelemans, W.; and Veenstra, J. 1997. Resolving pp attachment ambiguities with memory based learning.

Zelle, J. M., and Mooney, R. J. 1996. Learning to parse database queries using inductive logic proramming. In *Proc. National Conference on Artificial Intelligence*, 1050–1055.