END-TO-END ATTENTION-BASED LARGE VOCABULARY SPEECH RECOGNITION

Dzmitry Bahdanau[*], *Jan Chorowski*[†], *Dmitriy Serdyuk*[‡], *Philémon Brakel*[‡] *and Yoshua Bengio*[‡1]

[*]Jacobs University Bremen
[†]University of Wrocław
[‡] Université de Montréal
[1] CIFAR Fellow

ABSTRACT

Many of the current state-of-the-art Large Vocabulary Continuous Speech Recognition Systems (LVCSR) are hybrids of neural networks and Hidden Markov Models (HMMs). Most of these systems contain separate components that deal with the acoustic modelling, language modelling and sequence decoding. We investigate a more direct approach in which the HMM is replaced with a Recurrent Neural Network (RNN) that performs sequence prediction directly at the character level. Alignment between the input features and the desired character sequence is learned automatically by an attention mechanism built into the RNN. For each predicted character, the attention mechanism scans the input sequence and chooses relevant frames. We propose two methods to speed up this operation: limiting the scan to a subset of most promising frames and pooling over time the information contained in neighboring frames, thereby reducing source sequence length. Integrating an n-gram language model into the decoding process yields recognition accuracies similar to other HMM-free RNN-based approaches.

Index Terms— neural networks, LVCSR, attention, speech recognition, ASR

1. INTRODUCTION

Deep neural networks have become popular acoustic models for state-of-the-art large vocabulary speech recognition systems (Hinton et al., 2012a). However, in these systems most of the other components, such as Hidden Markov Models (HMMs), Gaussian Mixture Models (GMMs) and n-gram language models, are the same as in their predecessors. These combinations of neural networks and statistical models are often referred to as *hybrid systems*. In a typical hybrid system, a deep neural network is trained to replace the Gaussian Mixture Model (GMM) emission distribution of an HMM by predicting for each input frame the most likely HMM state. These state labels are obtained from a trained GMM-HMM system that has been used to perform forced alignment. In other words, a two-stage training process is required, in which the older GMM approach is still used as a starting point. An

obvious downside of this hybrid approach is that the acoustic model is not directly trained to minimize the final objective of interest. Our aim was to investigate neural LVCSR models that can be trained with a more direct approach by replacing the HMMs with a Attention-based Recurrent Sequence Generators (ARSG) such that they can be trained end-to-end for sequence prediction.

Recently, some work on end-to-end neural network LVCSR systems has shown promising results. A neural network model trained with Connectionist Temporal Classification (CTC) (Graves et al., 2006) achieved promising results on the Wall Street Journal corpus (Graves and Jaitly, 2014; Hannun et al., 2014b). A similar setup was used to obtain state-of-the-art results on the Switchboard task as well (Hannun et al., 2014a). Both of these models were trained to predict sequences of characters and were later combined with a word level language model. Furthermore, when the language model was implemented as a CTC-specific Weighted Finite State Transducer, decoding accuracies competitive with DNN-HMM hybrids were obtained (Miao et al., 2015).

At the same time, a new direction of neural network research has emerged that deals with models that learn to focus their "attention" to specific parts of their input. Systems of this type have shown very promising results on a variety of tasks including machine translation (Bahdanau et al., 2015), caption generation (Xu et al., 2015), handwriting synthesis (Graves, 2013), visual object classification (Mnih et al., 2014) and phoneme recognition (Chorowski et al., 2014, 2015).

In this work, we investigate the application of an Attention-based Recurrent Sequence Generator (ARSG) as a part of an end-to-end LVCSR system. We start from the system proposed in (Chorowski et al., 2015) and make the following contributions:

1. We show how training on long sequences can be made feasible by limiting the area explored by the attention to a range of most promising locations. This reduces the total training complexity from quadratic to linear, largely solving the scalability issue of the approach. This has already been proposed (Chorowski et al., 2015) under the name "windowing", but was used only

at the decoding stage in that work.

2. In the spirit of he Clockwork RNN (Koutnik et al., 2014) and hierarchical gating RNN (Chung et al., 2015), we introduce a recurrent architecture that successively reduces source sequence length by pooling frames neighboring in time. [1]

3. We show how a character-level ARSG and $n-$gram word-level language model can be combined into a complete system using the Weighted Finite State Transducers (WFST) framework.

2. ATTENTION-BASED RECURRENT SEQUENCE GENERATORS FOR SPEECH

The system we propose is a neural network that can map sequences of speech frames to sequences of characters. While the whole system is differentiable and can be trained directly to perform the task at hand, it can still be divided into different functional parts that work together to learn how to *encode* the speech signal into a suitable feature representation and to *decode* this representation into a sequence of characters. We used RNNs for both the encoder and decoder[2] parts of the system. The decoder combines an RNN and an attention mechanism into an Attention-based Recurrent Sequence Generator that is able to learn the alignment between its input and its output. Therefore, we will first discuss RNNs, and subsequently, how they can be combined with attention mechanisms to perform sequence alignment.

2.1. Recurrent Neural Networks

There has been quite some research into Recurrent Neural Networks (RNNs) for speech recognition (Robinson et al., 1996; Lippmann, 1989) and this can probably be explained to a large extent by the elegant way in which they can deal with sequences of variable length.

Given a sequence of feature vectors $(\mathbf{x}_1, \cdots, \mathbf{x}_T)$, a standard RNN computes a corresponding sequence of hidden state vectors $(\mathbf{h}_1, \cdots, \mathbf{h}_T)$ using

$$\mathbf{h}_t = g(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h), \qquad (1)$$

where $\mathbf{W}_{xh}$ and $\mathbf{W}_{hh}$ are matrices of trainable parameters that represent the connection weights of the network and $\mathbf{b}_h$ is a vector of trainable bias parameters. The function $g(\cdot)$ is often a non-linear squashing function like the hyperbolic tangent and applied element-wise to its input. The hidden states can be used as features that serve as inputs to a layer that performs a task like classification or regression. Given that this

---

[1] This mechanism has been recently independently proposed in (Chan et al., 2015).

[2] The word "decoder" refers to a network in this context, not to the final recognition algorithm.

output layer and the objective to optimize are differentiable, the gradient of this objective with respect to the parameters of the network can be computed with backpropagation through time. Like feed-forward networks, RNNs can process discrete input data by representing it as 1-hot-coding feature vectors.

An RNN can be used as a statistical model over sequences of labels. For that, it is trained it to predict the probability of the next label conditioned on the part of the sequence it has already processed. If $(y_1, \cdots, y_T)$ is a sequence of labels, an RNN can be trained to provide the conditional distribution the next label using

$$p(y_t|y_1, \cdots, y_{t-1}) = p(y_t|\mathbf{h}_t)$$
$$= \text{softmax}(\mathbf{W}_{hl}\mathbf{h}_t + \mathbf{b}_l),$$

where $\mathbf{W}_{hl}$ is a matrix of trainable connection weights, $\mathbf{b}_l$ is a vector of bias parameters and $\text{softmax}_i(\mathbf{a}) = \frac{\exp(a_i)}{\sum_j \exp(a_j)}$. The likelihood of the complete sequence is now given by $p(y_1) \prod_{t=2}^{T} p(y_t|y_1, \cdots, y_{t-1})$. This distribution can be used to *generate* sequences by either sampling from the distribution $p(y_t|y_1, \cdots, y_{t-1})$ or choosing the most likely labels iteratively.

Equation 1 defines the simplest RNN, however in practice usually more advanced equations define the dependency of $\mathbf{h}_t$ on $\mathbf{h}_{t-1}$. Famous examples of these so-called recurrent transitions are Long Short Term Memory (Hochreiter and Schmidhuber, 1997) and Gated Recurrent Units (GRU) (Cho et al., 2014), which are both designed to better handle long-term dependencies. In this work we use GRU for it has a simpler architecture and is easier to implement efficiently. The hidden states $\mathbf{h}_t$ are computed using the following equations:

$$\mathbf{z}_t = \sigma(\mathbf{W}_{xz}\mathbf{x}_t + \mathbf{U}_{hz}\mathbf{h}_{t-1}),$$
$$\mathbf{r}_t = \sigma(\mathbf{W}_{xr}\mathbf{x}_t + \mathbf{U}_{hr}\mathbf{h}_{t-1}),$$
$$\tilde{\mathbf{h}}_t = \tanh(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{U}_{rh}(\mathbf{r}_t \otimes \mathbf{h}_{t-1})),$$
$$\mathbf{h}_t = (1 - \mathbf{z}_t)\mathbf{h}_{t-1} + \mathbf{z}_t\tilde{\mathbf{h}}_t,$$

where $\tilde{\mathbf{h}}_t$ are candidate activations, $\mathbf{z}_t$ and $\mathbf{r}_t$ are update and reset gates respectively. The symbol $\otimes$ signifies element-wise multiplication.

To obtain a model that uses information from both future frames and past frames, one can pass the input data through two recurrent neural networks that run in opposite directions and concatenate their hidden state vectors. Recurrent neural network of this type are often referred to as *bidirectional* RNNs.

Finally, it has been shown that better results for speech recognition tasks can be obtained by stacking multiple layers of recurrent neural networks on top of each other (Graves et al., 2013). This can simply be done by treating the sequence of state vectors $(\mathbf{h}_1, \cdots, \mathbf{h}_T)$ as the input sequence for the next RNN in the pile. Figure 1 shows an example of two bidirectional RNNs that have been stacked on top of each other to construct a deep architecture.

Fig. 1. Two Bidirectional Recurrent Neural Networks stacked on top of each other.


Fig. 2. A pooling over time BiRNN: the upper layer runs twice slower then the lower one. It can average, or subsample (as shown in the figure) the hidden states of the layer below it.

2.2. Encoder-Decoder Architecture

Many challenging tasks involve inputs and outputs which may have variable length. Examples are machine translation and speech recognition, where both input and output have variable length; and image caption generation, where the captions may have variable lengths.

Encoder-decoder networks are often used to deal with variable length input and output sequences (Cho et al., 2014; Sutskever et al., 2014). The encoder is a network that transforms the input into an intermediate representation. The decoder is typically an RNN that uses this representation in order to generate the outputs sequences as described in 2.1.

In this work, we use a deep BiRNN as an encoder. Thus, the representation is a sequence of BiRNN state vectors $(\mathbf{h_1}, \ldots, \mathbf{h_L})$. For a standard deep BiRNN, the sequence $(\mathbf{h_1}, \ldots, \mathbf{h_L})$ is as long as the input of the bottom-most layer, which in the context of speech recognition means one $\mathbf{h}_i$ for every 10ms of the recordings. We found that for our decoder

(see 2.3) such representation is overly precise and contains much redundant information. This led us to add pooling between BiRNN layers as illustrated by Figure 2.

2.3. Attention-equipped RNNs

The decoder network in our system is an Attention-based Recurrent Sequence Generator (ARSG). This subsection introduces ARSGs and explains the motivation behind our choice of an ARSG for this study.

While RNNs can process and generate sequential data, the length of the sequence of hidden state vectors is always equal to the length of the input sequence. One can aim to learn the alignment between these two sequences to model a distribution $p(y_1, \cdots, y_T | \mathbf{h}_1, \cdots, \mathbf{h}_L)$ for which there is no clear functional dependency between T and L.

An ARSG produces an output sequence $(y_1, \cdots, y_T)$ one element at a time, simultaneously aligning each generated element to the encoded input sequence $(\mathbf{h}_1, \cdots, \mathbf{h}_L)$. It is composed of an RNN and an additional subnetwork called 'attention mechanism'. The attention selects the temporal locations over the input sequence that should be used to update the hidden state of the RNN and to make a prediction for the next output value. Typically, the selection of elements from the input sequence is a weighted sum $\mathbf{c}_t = \sum_l \alpha_{tl} \mathbf{h}_l$, where α_{tl} are called the attention weights and we require that $\alpha_{tl} \geq 0$ and that $\sum_l \alpha_{tl} = 1$. See Figure 3 for a schematic representation of an ARSG.

The attention mechanism used in this work is an improved version of the hybrid attention with convolutional features from (Chorowski et al., 2015), which is described by the following equations:

$$\mathbf{F} = \mathbf{Q} * \boldsymbol{\alpha}_{t-1} \qquad (2)$$

$$e_{tl} = \mathbf{w}^\top \tanh(\mathbf{W}\mathbf{s}_{t-1} + \mathbf{V}\mathbf{h}_l + \mathbf{U}\mathbf{f}_l + \mathbf{b}) \qquad (3)$$

$$\alpha_{tl} = \exp(e_{tl}) \left/ \sum_{l=1}^{L} \exp(e_{tl}) \right. . \qquad (4)$$

where $\mathbf{W}$, $\mathbf{V}$, $\mathbf{U}$, $\mathbf{Q}$ are parameter matrices, $\mathbf{w}$ and $\mathbf{b}$ are parameter vectors, $*$ denotes convolution, $\mathbf{s}_{t-1}$ stands for the previous state of the RNN component of the ARSG. We explain how it works starting from the end: (4) shows how the weights α_{tl} are obtained by normalizing the scores e_{tl}. As illustrated by (3), the score depends on the previous state $\mathbf{s}_{t-1}$, the content in the respective location $\mathbf{h}_l$ and the vector of so-called convolutional features $\mathbf{f}_l$. The name "convolutional" comes from the convolution along the time axis used in (2) to compute the matrix $\mathbf{F}$ that comprises all feature vectors $\mathbf{f}_l$.

Simply put, the attention mechanism described above combines information from three sources to decide where to focus at the step t: the decoding history contained in $\mathbf{s}_{t-1}$, the content in the candidate location $\mathbf{h}_l$ and the focus from the previous step described by attention weights $\boldsymbol{\alpha}_{t-1}$. It is shown in (Chorowski et al., 2015) that making the attention

location-aware, that is using $\boldsymbol{\alpha}_{t-1}$ in the equations defining $\boldsymbol{\alpha}_t$, is crucial for reliable behaviour on long input sequences.

A disadvantage of the approach from (Chorowski et al., 2015) is the complexity of the training procedure, which is $O(LT)$ since weights α_{tl} have to be computed for all pairs of input and output positions. The same paper showcases a windowing approach that reduces the complexity of decoding to $O(L + T)$. In this work we apply the windowing at the training stage as well. Namely, we constrain the attention mechanism to only consider positions from the range $(m_{t-1} - w_l, \ldots, m_{t-1} + w_r)$, where m_{t-1} is the median of $\boldsymbol{\alpha}_{t-1}$, interpreted in this context as a distribution. The values w_l and w_r define how much the window expands to the left and to the right respectively. This modification makes training significantly faster.

Apart from the speedup it brings, windowing can be also very helpful for starting the training procedure. From our experience, it becomes increasingly harder to train an ARSG completely from scratch on longer input sequences. We found that providing a very rough estimate of the desired alignment at the early training stage is an effective way to quickly bring network parameters in a good range. Specifically, we forced the network to only choose from positions in the range $R_t = (s_{min} + tv_{min}, \ldots, s_{max} + tv_{max})$. The numbers s_{min}, s_{max}, v_{min}, v_{max} were roughly estimated from the training set so that the number of leading silent frames for training utterances were between s_{min} and s_{max} and so that the speaker speed, i.e. the ratio between the transcript and the encoded input lengths, were between v_{min} and v_{max}. We aimed to make the windows R_t as narrow as possible, while keeping the invariant that the character y_t was pronounced within the window R_t. The resulting sequence of windows is quickly expanding, but still it was sufficient to quickly move the network out of the random initial mode, in which it had often aligned all characters to a single location in the audio data. We note, that the median-centered windowing could not be used for this purpose, since it relies on the quality of the previous alignment to define the window for the new one.

3. INTEGRATION WITH A LANGUAGE MODEL

Although an ARSG by construction implicitly learns how an output symbol depends on the previous ones, the transcriptions of the training utterances are typically insufficient to learn a high-quality language model. For this reason, we investigate how an ARSG can be integrated with a language model. The main challenge is that in speech recognition word-based language models are used, whereas our ARSG models a distribution over character sequences.

We use the Weighted Finite State Transducer (WFST) framework (Mohri et al., 2002; Allauzen et al., 2007) to build a character-level language model from a word-level one. A WFST is a finite automaton, whose transitions have weight and input and output labels. It defines a cost of transducing


Fig. 3. Schematic representation of the Attention-based Recurrent Sequence Generator. At each time step t, an MLP combines the hidden state $\mathbf{s}_{t-1}$ with all the input vectors $\mathbf{h}_l$ to compute the attention weights α_{tl}. Subsequently, the new hidden state $\mathbf{s}_t$ and prediction for output label y_t can be computed.

an input sequence into an output sequence by considering all pathes with corresponding sequences of input and output labels. The composition operation can be used to combine FSTs that define different levels of representation, such as characters and words in our case.

We compose the language model Finite State Transducer (FST) G with a lexicon FST L that simply spells out the letters of each word. More specifically, we build an FST $T = \min(\det(L \circ G))$ to define the log-probability for character sequences. We push the weights of this FST towards the starting state to help hypothesis pruning during decoding.

For decoding we look for a transcript y that minimizes the cost L which combines the encoder-decoder (ED) and the language model (LM) outputs as follows:

$$L = -\log p_{ED}(y|x) - \beta \log p_{LM}(y) - \gamma T \qquad (5)$$

where β and γ are tunable parameters. The last term γT is important, because without it the LM component dominates and the cost L is minimized by too short sequences. We note that the same criterion for decoding was proposed in (Hannun et al., 2014b) for a CTC network.

Integrating an FST and an ARSG in a beam-search decoding is easy because they share the property that the current state depends only on the previous one and the input symbol. Therefore one can use a simple left-to-right beam search algorithm similar to the one described in (Sutskever et al., 2014) to approximate the value of y that minimizes L.

The determinization of the FST becomes impractical for moderately large FSTs, such as the trigram model shipped with the Wall Street Journal corpus (see Subsection 5.1). To handle non-deterministic FSTs we assume that its weights

are in the logarithmic semiring and compute the total log-probability of all FST paths corresponding to a character prefix from the beam. This probability can be quickly recomputed when a new character is added to the prefix.

4. RELATED WORK

A popular method to train networks to perform sequence prediction is Connectionist Temporal Classification (Graves et al., 2006). It has been used with great success for both phoneme recognition (Graves et al., 2013) and character-based LVCSR (Graves and Jaitly, 2014; Hannun et al., 2014b,a; Miao et al., 2015). CTC allows recurrent neural networks to predict sequences that are shorter than the input sequence by summing over all possible alignments between the output sequence and the input of the CTC module. This summation is done using dynamic programming in a way that is similar to the forward and backward passes that are used to do inference in an HMM. In the CTC approach, output labels are conditionally independent given the alignment and the output sequences. In the context of speech recognition, this means that a CTC network lacks a language model, which greatly boosts the system performance when added to a trained CTC network (Hannun et al., 2014b; Miao et al., 2015).

An extension of CTC is the RNN Transducer which combines two RNNs into a sequence transduction system (Graves, 2012; Boulanger-Lewandowski et al., 2013). One network is similar to a CTC network and runs at the same time-scale as the input sequence, while a separate RNN models the probability of the next label output label conditioned on the previous ones. Like in CTC, inference is done with a dynamic programming method similar to the backward-forward algorithm for HMMs, but taking into account the constraints defined by both of the RNNs. Unlike CTC, RNN transduction systems can also generate output sequences that are longer than the input. RNN Transducers have led to state-of-the-art results in phoneme recognition on the TIMIT dataset (Graves et al., 2013) which were recently matched by an ASRG network (Chorowski et al., 2015).

The RNN Transducer and ARSG approaches are roughly equivalent in their capabilities. In both approaches an implicit language model is learnt jointly with the rest of the network. The main difference between the approaches is that in ARSG the alignment is explicitly computed by the network, as opposed to dealing with a distribution of alignments in the RNN Transducer. We hypothesize that this difference might have a major impact on the further development of these methods.

Finally, we must mention two very recently published works that partially overlap with the content of this paper. In (Chan et al., 2015) Encoder-Decoder for character-based recognition, with the model being quite similar to ours. In particular, in this work pooling between the BiRNN layers is also proposed. Also, in (Miao et al., 2015) using FSTs

to build a character-level model from an n-gram model is advocated. We note, that the research described in this paper was carried independently and without communication with the authors of both aforementioned works.

5. EXPERIMENTS

5.1. Data

We trained and evaluated our models [3] on the Wall Street Journal (WSJ) corpus (available at the Linguistic Data Consortium as LDC93S6B and LDC94S13B). Training was done on the 81 hour long SI-284 set of about 37K sentences. As input features, we used 40 mel-scale filterbank coefficients together with the energy. These 41 dimensional features were extended with their first and second order temporal derivatives to obtain a total of 123 feature values per frame. Evaluation was done on the "eval92" evaluation set. Hyperparameter selection was performed on the "dev93" set. For language model integration we have used the 20K closed vocabulary setup and the bigram and trigram language model that were provided with the data set. We use the same text preprocessing as in (Hannun et al., 2014b), leaving only 32 distinct labels: 26 characters, apostrophe, period, dash, space, noise and end-of-sequence tokens.

5.2. Training

Our model used 4 layers of 250 forward + 250 backward GRU units in the encoder, with the top two layers reading every second of hidden states of the network below it (see Figure 2). Therefore, the encoder reduced the utterance length by the factor of 4. A centered convolution filter of width 200 was used in the attention mechanism to extract a single feature from the previous step alignment as described in (4).

The AdaDelta algorithm (Zeiler, 2012) with gradient clipping was used for optimization. We initialized all the weights randomly from an isotropic Gaussian distribution with variance 0.1.

We used a rough estimate of the proper alignment for the first training epoch as described in Section 2.3. After that the training was restarted with the windowing described in the same section. The window parameters were $w_L = w_R = 100$, which corresponds to considering a large 8 second long span of audio data at each step, taking into account the pooling done between layers. Training with the AdaDelta hyperparameters $\rho = 0.95$, $\epsilon = 10^{-8}$ was continued until log-likelihood stopped improving. Finally, we annealed the best model in terms of CER by restarting the training with $\epsilon = 10^{-10}$.

We found regularization necessary for the best performance. The column norm constraint 1 was imposed on all weight matrices (Hinton et al., 2012b). This corresponds to

---

[3]Our code is available at https://github.com/rizar/attention-lvcsr

constraining the norm of the weights of all the connections incoming to a unit.

5.3. Decoding and Evaluation

As explained in Section 3, we used beam search to minimize the combined cost L defined by (5). We finished when k terminated sequences cheaper than any non-terminated sequence in the beam were found. A sequence was considered terminated when it ended with the special end-of-sequence token, which the network was trained to generate in the end of each transcript.

To measure the best performance we used the beam size 200, however this brought us only $\approx 10\%$ relative improvement over beam size 10. We used parameter settings $\alpha = 0.5$ and $\gamma = 1$ with a language model and $\gamma = 0.1$ without one. It was necessary to use an asymmetric window for the attention when decoding with large γ. More specifically, we reduced w_L to 10. Without this trick, the cost L could be infinitely minimized by looping across the input utterance, for the penalty for jumping back in time included in $\log p(y|x)$ was not high enough.

5.4. Results

Results of our experiments are gathered in Table 5.4. Our model shows performance superior to that of CTC systems when no external language model is used. The improvement from adding an external language model is however much larger for CTC-based systems. The final peformance of our model is better than the one reported in (Hannun et al., 2014b) (13.0% vs 14.1%), but worse than the the one from (Miao et al., 2015) (11.3% vs 9.0%) when the same language models are used.

6. DISCUSSION

A major difference between the CTC and ARSG approaches is that a language model is implicitly learnt in the latter. Indeed, one can see that an RNN sequence model as explained in 2.1 is literally contained in an ARSG as a subnetwork. We believe that this is the reason for the greater performance of the ARSG-based system when no external LM is used. However, this implicit language model was trained on a relatively small corpus of WSJ transcripts containing less than 4 million characters. It has been reported that RNNs overfit on corpora of such size (Graves, 2013) and in our experiments we had to combat overfitting as well. Using the weight clipping brought a consistent performance gain but did not change the big picture. For these reasons, we hypothesize that overfitting of the internal RNN language model was one of the main reasons why our model did not reach the performance level reported in (Miao et al., 2015), where a CTC network is used.

Table 1. Character Error Rate (CER) and Word Error Rate (WER) scores for our setup on the Wall Street Journal Corpus in comparison with other results from the literature. Note that our results are not directly comparable with those of networks predicting phonemes instead of characters, since phonemes are easier targets.

Model	CER	WER
Encoder-Decoder	6.4	18.6
Encoder-Decoder + bigram LM	5.3	11.7
Encoder-Decoder + trigram LM	4.8	10.8
Encoder-Decoder + extended trigram LM	3.9	9.3
Graves and Jaitly (2014)		
CTC	9.2	30.1
CTC, expected transcription loss	8.4	27.3
Hannun et al. (2014)		
CTC	10.0	35.8
CTC + bigram LM	5.7	14.1
Miao et al. (2015),		
CTC for phonemes + lexicon	-	26.9
CTC for phonemes + trigram LM	-	7.3
CTC + trigram LM	-	9.0

That being said, we treat it as an advantage of the ARSG that it supports joint training of a language model with the rest of the network. For one, WSJ contains approximately only 80 hours of training data, and overfitting might be less of an issue for corpora containing hundreds or even thousands hours of annotated speech. For two, an RNN language model trained on a large text corpus could be integrated in an ARSG from the beginning of training by using the states of this language model as an additional input of the ARSG. We suppose that this would block the incentive of memorizing the training utterances, and thereby reduce the overfitting. In addition, no extra n-gram model would be required. We note that a similar idea has been already proposed in (Gulcehre et al., 2015) for machine translation.

Finally, trainable integration with an n-gram language model could also be investigated.

6.1. Conclusion

In this work we showed how an Encoder-Decoder network with an attention mechanism can be used to build a LVCSR system. The resulting approach is significantly simpler than the dominating HMM-DNN one, with fewer training stages, much less auxiliary data and less domain expertise involved. Combined with a trigram language model our system shows decent, although not yet state-of-the-art performance.

We present two methods to improve the computational complexity of the investigated model. First, we propose pooling over time between BiRNN layers to reduce the length of the encoded input sequence. Second, we propose to use windowing during training to ensure that the decoder network

performs a constant number of operations for each output character. Together these two methods facilitate application of attention-based models to large-scale speech recognition.

Unlike CTC networks, our model has an intrinsic language-modeling capability. Furthermore, it has a potential to be trained jointly with an external language model. Investigations in this direction are likely to be a part of our future work.

Acknowledgments

References

Allauzen, C., Riley, M., Schalkwyk, J., Skut, W., and Mohri, M. (2007). OpenFst: A general and efficient weighted finite-state transducer library. In Holub, J. and Žďárek, J., editors, *Implementation and Application of Automata*, number 4783 in Lecture Notes in Computer Science, pages 11–23. Springer Berlin Heidelberg.

Bahdanau, D., Cho, K., and Bengio, Y. (2015). Neural machine translation by jointly learning to align and translate. In *International Conference on Learning Representations*.

Bastien, F., Lamblin, P., Pascanu, R., Bergstra, J., Goodfellow, I. J., Bergeron, A., Bouchard, N., and Bengio, Y. (2012). Theano: new features and speed improvements. Deep Learning and Unsupervised Feature Learning NIPS Workshop.

Bergstra, J., Breuleux, O., Bastien, F., Lamblin, P., Pascanu, R., Desjardins, G., Turian, J., Warde-Farley, D., and Bengio, Y. (2010). Theano: a CPU and GPU math expression compiler. In *Proceedings of the Python for Scientific Computing Conference (SciPy)*.

Boulanger-Lewandowski, N., Bengio, Y., and Vincent, P. (2013). High-dimensional sequence transduction. In *ICASSP*, pages 3178–3182. IEEE.

Chan, W., Jaitly, N., Le, Q. V., and Vinyals, O. (2015). Listen, attend and spell. *arXiv:1508.01211 [cs, stat]*.

Cho, K., van Merrienboer, B., Gulcehre, C., Bougares, F., Schwenk, H., and Bengio, Y. (2014). Learning phrase representations using RNN encoder-decoder for statistical machine translation. In *Empirical Methods of Natural Language Processing*.

Chorowski, J., Bahdanau, D., Cho, K., and Bengio, Y. (2014). End-to-end continuous speech recognition using attention-based recurrent NN: First results. *arXiv:1412.1602 [cs, stat]*.

Chorowski, J., Bahdanau, D., Serdyuk, D., Cho, K., and Bengio, Y. (2015). Attention-based models for speech recognition. In *NIPS*. to appear.

Chung, J., Gulcehre, C., Cho, K., and Bengio, Y. (2015). Gated feedback recurrent neural networks. In *ICML-15*.

Graves, A. (2012). Sequence transduction with recurrent neural networks. *arXiv preprint arXiv:1211.3711*.

Graves, A. (2013). Generating sequences with recurrent neural networks. *arXiv:1308.0850*.

Graves, A., Fernández, S., Gomez, F., and Schmidhuber, J. (2006). Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks. In *ICML-06*.

Graves, A. and Jaitly, N. (2014). Towards end-to-end speech recognition with recurrent neural networks. In *ICML-14*.

Graves, A., Mohamed, A.-r., and Hinton, G. (2013). Speech recognition with deep recurrent neural networks. In *ICASSP*, pages 6645–6649. IEEE.

Gulcehre, C., Firat, O., Xu, K., Cho, K., Barrault, L., Lin, H.-C., Bougares, F., Schwenk, H., and Bengio, Y. (2015). On using monolingual corpora in neural machine translation. *arXiv preprint arXiv:1503.03535*.

Hannun, A., Case, C., Casper, J., Catanzaro, B., Diamos, G., Elsen, E., Prenger, R., Satheesh, S., Sengupta, S., Coates, A., et al. (2014a). Deepspeech: Scaling up end-to-end speech recognition. *arXiv preprint arXiv:1412.5567*.

Hannun, A. Y., Maas, A. L., Jurafsky, D., and Ng, A. Y. (2014b). First-pass large vocabulary continuous speech recognition using bi-directional recurrent dnns. *arXiv preprint arXiv:1408.2873*.

Hinton, G., Deng, L., Yu, D., Dahl, G. E., Mohamed, A.-r., Jaitly, N., Senior, A., Vanhoucke, V., Nguyen, P., Sainath, T. N., and Kingsbury, B. (2012a). Deep neural networks for acoustic modeling in speech recognition: The shared views of four research groups. *Signal Processing Magazine, IEEE*, 29(6):82–97.

Hinton, G. E., Srivastava, N., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. R. (2012b). Improving neural networks by preventing co-adaptation of feature detectors. *arXiv preprint arXiv:1207.0580*.

Hochreiter, S. and Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9(8):1735–1780.

Koutnik, J., Greff, K., Gomez, F., and Schmidhuber, J. (2014). A clockwork RNN. In *ICML-14*.

Lippmann, R. P. (1989). Review of neural networks for speech recognition. *Neural computation*, 1(1):1–38.

Miao, Y., Gowayyed, M., and Metze, F. (2015). EESEN: End-to-end speech recognition using deep RNN models and WFST-based decoding. *arXiv:1507.08240 [cs]*.

Mnih, V., Heess, N., Graves, A., et al. (2014). Recurrent models of visual attention. In *NIPS*, pages 2204–2212.

Mohri, M., Pereira, F., and Riley, M. (2002). Weighted finite-state transducers in speech recognition. *Computer Speech & Language*, 16(1):69–88.

Robinson, T., Hochberg, M., and Renals, S. (1996). The use of recurrent neural networks in continuous speech recognition. In *Automatic speech and speaker recognition*, pages 233–258. Springer.

Sutskever, I., Vinyals, O., and Le, Q. V. (2014). Sequence to sequence learning with neural networks. In *NIPS*.

van Merriënboer, B., Bahdanau, D., Dumoulin, V., Serdyuk, D., Warde-Farley, D., Chorowski, J., and Bengio, Y. (2015). Blocks and fuel: Frameworks for deep learning. *arXiv:1506.00619 [cs, stat]*.

Xu, K., Ba, J., Kiros, R., Cho, K., Courville, A., Salakhutdinov, R., Zemel, R., and Bengio, Y. (2015). Show, attend and tell: Neural image caption generation with visual attention. In *ICML-15*.

Zeiler, M. D. (2012). Adadelta: An adaptive learning rate method. *arXiv preprint arXiv:1212.5701*.