

A Scalable Algorithm for Minimal Unsatisfiable Core Extraction*

Nachum Dershowitz
School of Computer Science
Tel Aviv University, Ramat Aviv, Israel
nachumd@post.tau.ac.il

Ziyad Hanna
Design Technology Solutions Group
Intel Corporation, Haifa, Israel
ziyad.hanna@intel.com

Alexander Nadel
School of Computer Science
Tel Aviv University, Ramat Aviv, Israel
ale1@post.tau.ac.il
Design Technology Solutions Group
Intel Corporation, Haifa, Israel
alexander.nadel@intel.com

July 4, 2018

Abstract

The task of extracting an unsatisfiable core for a given Boolean formula has been finding more and more applications in recent years. The only existing approach that scales well for large real-world formulas exploits the ability of modern SAT solvers to produce resolution refutations. However, the resulting unsatisfiable cores are suboptimal. We propose a new algorithm for minimal unsatisfiable core extraction, based on a deeper exploration of resolution-refutation properties. We provide experimental results on formal verification benchmarks confirming that our algorithm finds smaller cores than suboptimal algorithms; and that it runs faster than those algorithms that guarantee minimality of the core.

1 Introduction

Many real-world problems, arising in formal verification of hardware and software, planning and other areas, can be formulated as constraint satisfaction problems, which can be translated into Boolean formulas in conjunctive normal form (CNF). Modern Boolean satisfiability (SAT) solvers, such as Chaff [1, 2] and MiniSAT [3], which implement enhanced versions of the Davis-Putnam-Longeman-Loveland (DPLL) backtrack-search algorithm, are often able to determine whether a large formula is satisfiable or unsatisfiable. When a formula is unsatisfiable, it is often required to find an *unsatisfiable core*—that is, a small unsatisfiable subset of the formula’s clauses. Example applications include functional verification of hardware [5], FPGA routing [6], and abstraction refinement [7]. For example, in FPGA routing, an unsatisfiable instance implies that the channel is unroutable. Localizing a small unsatisfiable core is necessary to determine the underlying reasons for the failure. An unsatisfiable core is a *minimal unsatisfiable core (MUC)*, if it becomes satisfiable whenever any one of its clauses is removed. It is always desirable to find a minimal unsatisfiable core, but this problem is very hard (it is D^P -complete; see [4]).

*This research was supported by the Israel Science Foundation (grant no. 250/05). The work of Alexander Nadel was carried out in partial fulfillment of the requirements for a Ph.D.

In this paper, we propose an algorithm that is able to find a minimal unsatisfiable core for large “real-world” formulas. Benchmark families, arising in formal verification of hardware (such as [24]), are of particular interest for us. The only approach for unsatisfiable core extraction that scales well for formal verification benchmarks was independently proposed in [17] and in [18]. We refer to this method as the *EC (Empty-clause Cone)* algorithm. EC exploits the ability of modern SAT solvers to produce a resolution refutation, given an unsatisfiable formula. Most state-of-the-art SAT solvers, beginning with GRASP [19] and including Chaff [1, 2] and MiniSAT [3], implement a DPLL backtrack search enhanced by a failure-driven assertion loop [19]. These solvers explore the assignment tree and create new *conflict clauses* at the leaves of the tree, using resolution on the initial clauses and previously created conflict clauses. This process stops when either a satisfying assignment is found or when the empty clause (\square) is derived. In the latter case, SAT solvers are able to produce a *resolution refutation*—a directed acyclic graph (dag), whose vertices are associated with clauses, and whose edges describe resolution relations between clauses. The sources of the refutation are the initial clauses and the empty clause \square is a sink. EC traverses a reversed refutation, starting with \square and taking initial clauses, connected to \square , as the unsatisfiable core. Invoking EC until a fixed point is reached [17] allows one to reduce the unsatisfiable core even more. We refer to this algorithm as *EC-fp*. However, the resulting cores can be reduced further.

The basic flow of the algorithm for minimal unsatisfiable core extraction proposed in this paper is composed of the following steps:

1. Produce a resolution refutation Π of a given formula using a SAT solver.
2. Drop from Π all clauses not connected to \square . At this point, all the initial clauses remaining in Π comprise an unsatisfiable core.
3. For every initial clause C remaining in Π , check whether it belongs to a minimal unsatisfiable core (MUC) in the following manner:

Remove C from Π , along with all conflict clauses for which C was required to derive them.

Pass all the remaining clauses (including conflict clauses) to a SAT solver.

- If they are satisfiable, then C belongs to a minimal unsatisfiable core, so continue with another initial clause.
- If the clauses are unsatisfiable, then C does not belong to a MUC, so replace Π by a new valid resolution refutation not containing C .

4. Terminate when all the initial clauses remaining in Π comprise a MUC.

Related work is discussed in the next section. Section 3 is dedicated to refutation-related definitions. Our basic *Complete Resolution Refutation (CRR)* algorithm is described in Sect. 4, and a pruning technique, enhancing CRR and called *Resolution Refutation-based Pruning (RRP)*, is described in Sect. 5. Experimental results are analyzed in Sect. 6. This is followed up by a brief conclusion.

2 Related Work

Algorithms for unsatisfiable core based on the ability of modern SAT solvers to produce resolution refutations [17, 18] are the most relevant for our purposes for two reasons. First, this approach allows one to deal with real-world examples arising in formal verification. Second, it serves as the basis of our algorithm. We have already described the EC and EC-fp algorithms in the introduction. Here we briefly consider other approaches.

Theoretical work (e.g., [8]) has concentrated on developing efficient algorithms for formulas with small *deficiency* (the number of clauses minus the number of variables). However, real-world formulas have arbitrary (and usually large) deficiency. A number of works considered the harder problem of finding the smallest minimal unsatisfiable core [12, 14], or even finding all minimally unsatisfiable formulas [13]. As one can imagine, these algorithms are not scalable for even moderately large real-world formulas.

In [9, 10], an “adaptive core search” is applied for finding a small unsatisfiable core. The algorithm starts with a very small satisfiable subformula, consisting of *hard* clauses. The unsatisfiable core is built by an iterative process that expands or contracts the current core by a fixed percentage of clauses. The procedure succeeds in finding small, though not necessarily minimal, unsatisfiable cores for the problem instances it was tested on, but these are very small and artificially generated.

Another approach that allows for finding small, but not necessarily minimal, unsatisfiable cores is called AMUSE [11]. In this approach, selector variables are added to each clause and the unsatisfiable core is found by a branch-and-bound algorithm on the updated formula. Selector variables allow the program to implicitly search for unsatisfiable cores using an enhanced version of DPLL on the updated formula. The authors note their methods ability to locate different unsatisfiable cores, as well as its inability to cope with large formulas.

The above described algorithms do not guarantee minimality of the cores extracted. One folk algorithm for minimal unsatisfiable core extraction, which we dub *Naïve*, works as follows: For every clause C in an unsatisfiable formula F , Naïve checks if it belongs to the minimal unsatisfiable core, by invoking a SAT solver on $F \setminus C$. Clause C does not belong to MUC if and only if the solver finds that $F \setminus C$ is unsatisfiable, in which case C is removed from F . In the end, F contains a minimal unsatisfiable core.

The only non-trivial algorithm existing in the current literature that guarantees minimality is MUP [15]. MUP is mainly a prover of minimal unsatisfiability, as opposed to an unsatisfiable core extractor. It decides the minimal unsatisfiability of a CNF formula through BDD manipulation. When MUP is used as a core extractor, it removes one clause at a time until the remaining core is minimal. MUP is able to prove minimal unsatisfiability of some particularly hard classical problems quickly, whereas even just proving unsatisfiability is a challenge for modern SAT solvers. However, the formulas described in [15] are small and arise in areas other than formal verification. We will see in Section 6 that MUP is significantly outperformed by Naïve on formal verification benchmarks.

3 Resolution Refutations

We begin with a resolution refutation of a given unsatisfiable formula, defined as follows:

Definition 1 (Resolution refutation) *Let F be an unsatisfiable CNF formula (set of clauses) and let $\Pi(V, E)$ be a dag whose vertices are clauses.¹ Suppose $V = V^i \cup V^c$, where V^i are all the sources of Π , referred to as initial clauses, and $V^c = C_1^c, \dots, C_m^c$ is an ordered set of non-source vertices, referred to as conflict clauses. Then, the dag $\Pi(V, E)$ is a resolution refutation of F if:*

1. $V^i = F$;
2. *for every conflict clause C_i^c , there exists a resolution derivation $\{D_1, D_2, \dots, D_k, C_i^c\}$, such that:*
 - (a) *for every $j = 1, \dots, k$, D_j is either an initial clause or a prior conflict clause C_f^c , $f < i$, and*
 - (b) *there are edges $D_1 \rightarrow C_i^c, \dots, D_k \rightarrow C_i^c \in E$ (these are the only edges in E);*
3. *the sink vertex C_m^c is the only empty clause in V .*

For the subsequent discussion, it will be helpful to capture the notion of vertices that are “reachable”, or “backward reachable”, from a given clause in a given dag.

Definition 2 (Reachable vertices) *Let Π be a dag. A vertex D is reachable from C if there is a path (of 0 or more edges) from C to D . The set of all vertices reachable from C in Π is denoted $Re(\Pi, C)$. The set of all vertices unreachable from C in Π is denoted by $\overline{Re}(\Pi, C)$*

¹From now on, we shall use the terms “vertex” and “clause” interchangeably in the context of resolution refutation.

Definition 3 (Backward reachable vertices) Let Π be a dag. A vertex D is backward reachable from C if there is a path (of 0 or more edges) from D to C . The set of all vertices backward reachable from C in Π is denoted by $BRe(\Pi, C)$. The set of all vertices not backward reachable from C in Π is denoted $\overline{BRe}(\Pi, C)$.

For example, consider the resolution refutation of Fig. 1. We have $Re(\Pi, C_5^i) = \{C_5^i, C_2^c, C_3^c, C_4^c, C_5^c\}$ and $BRe(\Pi, C_4^c) = \{C_4^c, C_5^i, C_6^i\}$.

Resolution refutations trace all resolution derivations of conflict clauses, including the empty clause. Generally, not all clauses of a refutation are required to derive \square , but only such that are backward reachable from \square . It is not hard to see that even if all other clauses and related edges are omitted, the remaining graph is still a refutation. We refer to such refutations as *non-redundant* (see Definition 4). The refutation in Fig. 1 is non-redundant.

To retrieve a non-redundant subgraph of a refutation, it is sufficient to take $BRe(\Pi, \square)$ as the vertex set and to restrict the edge set E to edges having both ends in $BRe(\Pi, \square)$. We denote a non-redundant subgraph of a refutation Π by $\Pi \upharpoonright_{BRe(\Pi, \square)}$. Observe that $\Pi \upharpoonright_{BRe(\Pi, \square)}$ is a valid non-redundant refutation.

Definition 4 (Non-redundant resolution refutation) A resolution refutation Π is non-redundant if there is a path in Π from every clause to \square .

Lastly, we define a relative hardness of a resolution refutation.

Definition 5 (Relative hardness) The relative hardness of a resolution refutation is the ratio between the total number of clauses and the number of initial clauses.

4 The Complete Resolution Refutation (CRR) Algorithm

Our goal is to find a minimal unsatisfiable core of a given unsatisfiable formula F . The proposed *CRR* method is displayed as Algorithm 1.

Algorithm 1 (CRR). Returns a MUC, given an unsatisfiable formula F .

```

1: Build a non-redundant refutation  $\Pi(V^i \cup V^c, E)$ 
2: while unmarked clauses exist in  $V^i$  do
3:    $C \leftarrow \text{PickUnmarkedClause}(V^i)$ 
4:   Invoke a SAT solver on  $\overline{Re}(\Pi, C)$ 
5:   if  $\overline{Re}(\Pi, C)$  is satisfiable then
6:     Mark  $C$  as a MUC member
7:   else
8:     Let  $G = \overline{Re}(\Pi, C)$ 
9:     Build resolution refutation  $\Pi'(V_G^i \cup V_G^c, E_G)$ 
10:     $V^i \leftarrow V^i \setminus \{C\}$ 
11:     $V^c \leftarrow (V^c \setminus Re(\Pi, C)) \cup V_G^c$ 
12:     $E \leftarrow (E \setminus Re^E(\Pi, C)) \cup E_G$ 
13:     $\Pi(V^i \cup V^c, E) \leftarrow \Pi(V^i \cup V^c, E) \upharpoonright_{BRe(\Pi, \square)}$ 
14: return  $V^i$ 

```

First, CRR builds a non-redundant resolution refutation. Invoking a SAT solver for constructing a (possibly redundant) resolution refutation $\Pi(V, E)$ and restricting it to $\Pi \upharpoonright_{BRe(\Pi, \square)}$ is sufficient for this purpose.

Suppose $\Pi(V^i \cup V^c, E)$ is a non-redundant refutation. CRR checks, for every unmarked clause C left in V^i , whether C belongs to the minimal unsatisfiable core. Initially, all clauses are unmarked. At each stage of the algorithm, CRR maintains a valid refutation of F .

Recall from Definition 2 that $\overline{Re}(\Pi, C)$ is the set of all vertices in Π unreachable from C . By construction of Π , the $\overline{Re}(\Pi, C)$ clauses were derived independently of C . To check whether C belongs to the minimal

unsatisfiable core, we provide the SAT solver with $\overline{Re}(\Pi, C)$, including the conflict clauses. We are trying to *complete the resolution refutation*, not using C as one of the sources. Observe that \square is always reachable from C , since Π is a non-redundant refutation; thus \square is never passed as an input to the SAT solver. We let the SAT solver try to derive \square , using $\overline{Re}(\Pi, C)$ as the input formula, or else prove that $\overline{Re}(\Pi, C)$ is satisfiable.

In the latter case, we conclude that C must belong to the minimal unsatisfiable core, since we found a model for an unsatisfiable subset of initial clauses minus C . Hence, if the SAT solver returns *satisfiable*, the algorithm marks C (line 6) and moves to the next initial clause. However, if the SAT solver returns *unsatisfiable*, we cannot simply remove C from F and move to the next clause, since we need to keep a valid resolution refutation for our algorithm to work properly. We describe the construction of a valid refutation (lines 8–13) next.

Let $G = \overline{Re}(\Pi, C)$. The SAT solver produces a new resolution refutation $\Pi'(V_G^i \cup V_G^c, E_G)$ for G , whose sources are the clauses $\overline{Re}(\Pi, C)$. We cannot use Π' as the refutation for the subsequent iterations, since the sources of the refutation may only be initial clauses of F . However, the “superfluous” sources of Π' are conflict clauses of Π , unreachable from C , and thus are derivable from $V^i \setminus C$ using resolution relations, corresponding to edges of Π . Hence, it is sufficient to augment Π' with such edges of Π that connect $V^i \setminus C$ and $\overline{Re}(\Pi, C)$ to obtain a valid refutation whose initial clauses belong to F . Algorithm CRR constructs a new refutation, whose sources are $V^i \setminus C$; the conflict clauses are $\overline{Re}(\Pi, C) \cup V_G^c$ and the edges are $(E \setminus (V_1, V_2) | (V_1 \in Re(\Pi, C) \text{ or } V_2 \in Re(\Pi, C))) \cup E_G$. This new refutation might be redundant, since $\Pi'(V_G^i \cup V_G^c, E_G)$ is not guaranteed to be non-redundant. Therefore, prior to checking the next clause, we reduce the new refutation to a non-redundant one. Observe that in the process of reduction to a non-redundant subgraph, some of the initial clauses of F , may be omitted; hence, each time a clause C is found not to belong to the minimal unsatisfiable core, we potentially drop not only C , but also other clauses.

We demonstrate the process of completing a refutation on the example from Fig. 1. Suppose we are checking whether C_1^i belongs to the minimal unsatisfiable core. In this case, $G = \overline{Re}(\Pi, C_1^i) = \{C_2^i, C_3^i, C_4^i, C_5^i, C_6^i, C_7^i, C_2^c, C_4^c\}$. The SAT solver receives G as the input formula. It is not hard to check that G is unsatisfiable. One refutation of G is $\Pi'(V_G^i \cup V_G^c, E_G)$, where $V_G^i = \{C_2^i, C_2^c, C_7^i, C_4^c\}$, $V_G^c = (D_1 = \square, D_2 = a \vee b)$, and $E_G = \{(C_2^i, D_2), (C_2^c, D_2), (D_2, D_1), (C_7^i, D_1), (C_4^c, D_1)\}$. Therefore, C_1^i , C_3^i , C_5^i and related edges are excluded from the refutation of F , whereas D_2 , D_1 and related edges are added to the refutation of F . In this case, the resulting refutation is non-redundant.

We did not define how the function *PickUnmarkedClause* should pick clauses (line 3). Our current implementation picks clauses in the order in which clauses appear in the given formula. Development of sophisticated heuristics is left for future research.

Another direction that may lead to a speed-up of CRR is adjusting the SAT solver for the purposes of CRR algorithm, considering that the SAT solver is invoked thousands of times on rather easy instances. Integrating the data structures of CRR and the SAT solver, tuning SAT solver’s heuristics for CRR, and holding the refutation in-memory, rather than on disk (as suggested in [17] for EC), can be helpful.

5 Resolution-Refutation-Based Pruning

In this section, we propose an enhancement of Algorithm CRR by developing resolution refutation-based pruning techniques for when a SAT solver invoked on $\overline{Re}(\Pi, C)$ to check whether it is possible to complete a refutation without C . We refer to the pruning technique, proposed in this section, *Resolution Refutation-based Pruning (RRP)*. We presume that the reader is familiar with the functionality of a modern SAT solver. (An overview is given in [20].)

An assignment σ *falsifies* a clause C , if every literal of C is *false* under σ . An assignment σ *falsifies* a set of clauses P if every clause $C \in P$ is falsified by σ . We claim that a model for $\overline{Re}(\Pi, C)$ can only be found under such a partial assignment that falsifies every clause in some path from C to the empty clause in $Re(\Pi, C)$. The intuitive reason is that every other partial assignment satisfies C and must falsify $\overline{Re}(\Pi, C)$, since F is unsatisfiable. A formal statement and proof is provided in Theorem 1 below.

Consider the example of Fig. 1. Suppose the currently visited clause is C_5^i . There exist two paths from




Figure 1: Resolution refutation example

C_5^i to the empty clause C_5^c , namely $\{C_5^i, C_4^c, C_5^c\}$ and $\{C_5^i, C_2^c, C_3^c, C_5^c\}$. A model for $\overline{Re}(\Pi, C_5^i)$ can only be found in a subspace under the partial assignment $\{a = 1, c = 0\}$, falsifying all the clauses of the first path. The clauses of the second path cannot be falsified, since a must be *true* to falsify clause C_5^i and *false* to falsify clause C_3^c .

Denote a subtree connecting C and \square by $\Pi|_C$. The proposed pruning technique, RRP, is integrated into the decision engine of the SAT solver. The solver receives $\Pi|_C$, together with the input formula $\overline{Re}(\Pi, C)$. The decision engine of the SAT solver explores $\Pi|_C$ in a depth-first manner, picking unassigned variables in the currently explored path as decision variables and assigning them *false*. As usual, Boolean Constraint Propagation (BCP) follows each assignment. Backtracking in $\Pi|_C$ is tightly related to backtracking in the assignment space. Both happen when a satisfied clause in $\Pi|_C$ is found or when a new conflict clause is discovered during BCP. After a particular path in $\Pi|_C$ has been falsified, a general-purpose decision heuristic is used until the SAT solver either finds a satisfying assignment or proves that no such assignment can be found under the currently explored path. This process continues until either a model is found or the decision engine has completed exploring $\Pi|_C$. In the latter case, one can be sure that no model for $\overline{Re}(\Pi, C)$ exists. However, the SAT solver should continue its work to produce a refutation.

We need to describe in greater detail the changes in the decision and conflict analysis engines of the SAT solver required to implement RRP. The decision engine first invokes function RRP_Decide , depicted in Fig. 2, as a state transition relation. Each transition edge has a label consisting of a condition under which the state transition occurs and an operation, executed upon transition. The state can be one of the following:

- (**Norm**) normal;
- (**Sat**) the currently explored clause is satisfied;
- (**False**) the currently explored clause is falsified;
- (**EoT**) subgraph $\Pi|_C$ has been explored;
- (**EoF**) all clauses in the currently explored path are falsified.

The states are managed globally, that is, if RRP_Decide moves to state S , it will start in state S when invoked next. A pointer D to the currently visited clause of $\Pi|_C$ is also managed globally. The state

D is not satisfied nor falsified / Return an unassigned literal

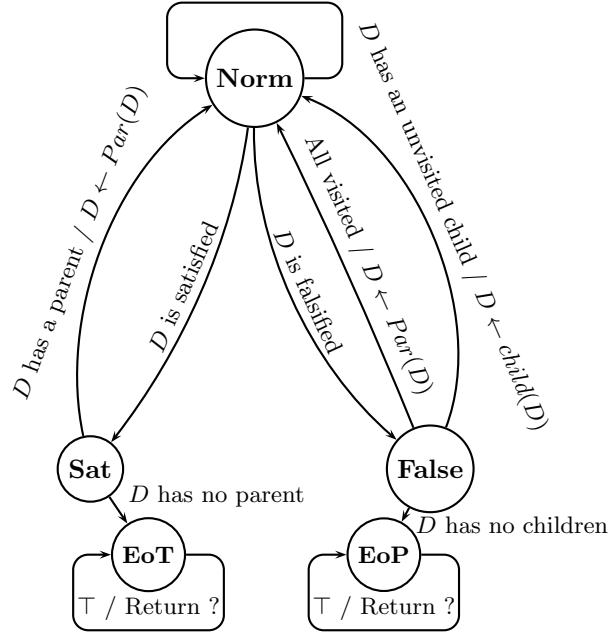


Figure 2: Function RRP_Decide , represented as a transition relation. This function is invoked by the decision engine of a SAT solver, implementing the RRP pruning technique.

transition relation is initialized prior to the first invocation of the decision engine. Pointer D is initialized to C and the initial state is **Norm**.

State **Norm** corresponds to a situation when the algorithm does not know what the status of D is. If D is neither satisfied nor falsified, RRP_Decide returns a negation of some literal of D , which will serve as the next decision variable. If D is satisfied, the algorithm moves to **Sat**. Observe that a clause may become satisfied only as a result of BCP. Encountering a satisfied clause means that the currently explored path cannot be falsified, and we can backtrack. Suppose we are in **Sat**, meaning that D is satisfied. If D has a parent, the algorithm backtracks by moving D to point to its parent, and goes back to **Norm**; otherwise, the tree is explored and the algorithm moves to **EoT**. In this case, RRP_Decide returns an unknown value and a general-purpose heuristic must be used. Consider now the case when the state is **Norm** and D is falsified. The algorithm moves to **False**. Here, one of the three conditions hold:

- (a) D has an unvisited child S . In this case D is updated to point to S and RRP_Decide moves back to **Norm**.
- (b) All children of D are visited. In this case, we backtrack by moving D back to its parent and go back to **Norm**.
- (c) D has no children. In this case, all the clauses in the currently explored path are falsified. The algorithm moves to **EoF**; RRP_Decide returns an unknown value; and a general-purpose heuristic must be used.

To complete the picture, we describe the changes to the conflict analysis engine required to implement RRP. One of the main tasks of conflict analysis in modern SAT solvers is to decide to which level in the decision tree the algorithm should backtrack. Let this decision level be bl . When invoked in RRP mode, the conflict analysis engine must also find whether it is required to backtrack in $\Pi|_C$, and to which clause. The goal is to backtrack to the highest clause B in the currently explored path in $\Pi|_C$, such that B has unassigned literals. Recall that D is a pointer to the currently visited clause of $\Pi|_C$. Denote by $mdl(D)$ the

Table 1: Comparing algorithms for unsatisfiable core extraction. Columns **Instance**, **Var** and **Cls** contain instance name, number of variables, and clauses, respectively. The next seven columns contain execution times (in seconds) and core sizes (in number of clauses) for each algorithm. The cut-off time was 24 hours (86,400 sec.). Column **Rel. Hard.** contains the relative hardness of the final resolution refutation, produced by CRR+RRP. Bold times are the best among algorithms guaranteeing minimality.

Instance	Var	Cls	Subopt.		CRR		Naïve		MUP	Rel. Hard.
			EC	EC-fp	RRP	plain	EC-fp	AMUSE		
4pipe	4237		9	171	3527	4933	24111	time-out	time-out	1.4
4pipe_1_ooo	4647	80213	23305	17724	17184	17180	17182			
			10	332	4414	10944	25074	time-out	mem-out	1.7
4pipe_2_ooo	4941	74554	24703	14932	12553	12515	12374			
			13	347	5190	12284	49609	time-out	mem-out	1.7
4pipe_3_ooo	5233	82207	25741	17976	14259	14192	14017			
			14	336	6159	15867	41199	time-out	mem-out	1.6
4pipe_4_ooo	5525	89473	30375	20034	16494	16432	16419			
			16	341	6369	16317	47394	time-out	mem-out	1.6
		96480	31321	21263	17712	17468	17830			
3pipe_k	2391		2	20	411	493	2147	12544	mem-out	1.5
4pipe_k	5095	27405	10037	6953	6788	6786	6784			
			8	121	3112	3651	15112	time-out	time-out	1.5
5pipe_k	9330	79489	24501	17149	17052	17078	17077			
			16	169	13836	17910	83402	time-out	mem-out	1.4
		189109	47066	36571	36270	36296	36370			
barrel5	1407		2	19	93	86	406	326	mem-out	1.8
barrel6	2306	5383	3389	3014	2653	2653	2653			
			35	322	351	423	4099	4173	mem-out	1.8
barrel7	3523	8931	6151	5033	4437	4437	4437			
			124	1154	970	1155	6213	24875	mem-out	1.9
barrel8	5106	13765	9252	7135	6879	6877	6877			
			384	9660	2509	2859	time-out	time-out	mem-out	1.8
		20083	14416	11249	10076	10075				
longmult4	1966		0	0	8	7	109	152	13	2.6
longmult5	2397	6069	1247	1246	972	972	972	976	972	
			0	1	74	31	196	463	35	3.6
longmult6	2848	7431	1847	1713	1518	1518	1518	1528	1518	
			2	13	288	311	749	2911	5084	5.6
longmult7	3319	8853	2639	2579	2187	2187	2187	2191	2187	
			17	91	6217	3076	6154	32791	68016	14.2
		10335	3723	3429	2979	2979	2979	2993	2979	

maximal decision level of D 's literals. If $bl \geq mdl(D)$, the algorithm does nothing; otherwise, it finds the first predecessor of D in $\Pi \upharpoonright_C$, such that $bl < mdl(B)$ and sets $D \leftarrow B$.

We found experimentally that the optimal performance for RRP is achieved when it explores $\Pi \upharpoonright_C$ starting from \square toward C (and not vice-versa). In other words, prior to the search, the SAT solver reverses all the edges of $\Pi \upharpoonright_C$ and sets the pointer D to \square , rather than to C . (By default, the current version of RRP explores the graph only until a predefined depth of 50.) The next literal from the currently visited clause is chosen by preferring an unassigned literal with the maximal number of appearances in recent conflict clause derivations (similar to Berkmin's [21] heuristic for SAT). The next visited child is chosen arbitrary. Further tuning of the algorithm is left for future research.

Theorem 1 *Let $\Pi(V^i, V^c)$ be a non-redundant resolution refutation. Let $C \in V^i$ be an initial clause and σ be an assignment. Then, if $\sigma \models \overline{Re}(\Pi, C)$, there is a path $P = \{C, \dots, C_m^c\}$ in $Re(\Pi, C)$, connecting C to the empty clause², such that σ falsifies every clause in P .*

Proof. Suppose, on the contrary, that no such path exists. Then, there exists a satisfiable vertex cut U in Π . But the empty clause is derivable from U , since it is a vertex cut; thus U unsatisfiable, a contradiction. \square

²The empty clause always belongs to $Re(\Pi, C)$, since $\Pi(V^i, V^c)$ is non-redundant.

6 Experimental Results

We have implemented CRR and RRP in the framework of the VE solver. VE, a variant of the industrial solver Eureka, is a modern SAT solver, which implements the following state-of-the-art algorithms and techniques for SAT: it uses UIP conflict clause recording [1], enhanced by conflict clause minimization [22], frequent search restarts [21, 2], an aggressive clause deletion strategy [21, 2], and decision stack shrinking [20, 2]. VE uses Berkmin’s decision heuristic [21] until 4000 conflicts are detected, and then switches to the CBH heuristic, described in [23]. We used benchmarks from four well-known unsatisfiable families, taken from bounded model checking (*barrel*, *longmult*) [25] and microprocessor verification (*fvp-unsat.2.0*, *pipe-unsat.1.0*) [24]. All the instances we used appear in the first column of Table 1. The experiments on families *barrel* and *fvp-unsat.2.0* were carried out on a machine with 4Gb of memory and two Intel Xeon CPU 3.06 processors. A machine with the same amount of memory and two Intel Xeon CPU 3.20 processors was used for experiments with the families *longmult* and *pipe-unsat.1.0*.

Table 1 summarizes the results of a comparison of the performance of two algorithms for suboptimal unsatisfiable core extraction and five algorithms for minimal unsatisfiable core extraction in terms of execution time and core sizes.

First, we compare algorithms for minimal unsatisfiable core extraction, namely, Naïve, MUP, plain CRR, and CRR enhanced by RRP. In preliminary experiments, we found that Naïve demonstrates its best performance on formulas that are first trimmed down by a suboptimal algorithm for unsatisfiable core extraction. We tried Naïve in combination with EC, EC-fp and AMUSE and found that EC-fp is the best front-end for Naïve. In our main experiments, we used Naïve, combined with EC-fp, and Naïve combined with AMUSE. We have also found that MUP demonstrates its best performance when combined with EC-fp, while CRR performs the best when the first refutation is constructed by EC, rather than EC-fp. Consequently, we provide results for MUP combined with EC-fp and CRR combined with EC. MUP requires a so-called “decomposition tree”, in addition to the CNF formula. We used the c2d package [?] for decomposition tree construction.

The sizes of the cores do not vary much between MUC algorithms, so we concentrate on a performance comparison. One can see that the combination of EC-fp and Naïve outperforms the combination of AMUSE and Naïve, as well as MUP. Plain CRR outperforms Naïve on every benchmark, whereas CRR+RRP outperforms Naïve on 15 out of 16 benchmarks (the exception being the hardest instance of *longmult*). This demonstrates that our algorithms are justified practically. Usually, the speed-up of these algorithms over Naïve varies between 4 and 10x, but it can be as large as 34x (for the hardest instance of *barrel* family) and as small as 2x (for the hardest instance of *longmult*). RRP improves performance on most instances. The most significant speed-up of RRP is about 2.5x, achieved on hard instances of Family *fvp-unsat.2.0*. The only family for which RRP is usually unhelpful is *longmult*.

A natural question is why the complex instances of family *longmult* are hard for CRR, and even harder for RRP. The key difference between *longmult* and other families is the hardness of the resolution proof. The relative hardness of a resolution refutation produced by CRR+RRP varies between 1.4 to 2 for every instance of every family, except *longmult*, where it reaches 14.2 for the *longmult7* instance. When the refutation is too complex, the exploration of $\overline{Re}(\Pi, C)$ executed by RRP is too complicated; thus, plain CRR is advantageous over CRR+RRP. Also, when the refutation is too complex, it is costly to perform traversal operations, as required by CRR. This explains why the advantage of CRR over Naïve is as small as 2X.

Comparing CRR+RRP on one side and EC and EC-fp on the other, we find that CRR+RRP always produce smaller cores than both EC and EC-fp. The average gain on all instances of cores produced by CRR+RRP over cores produced by EC and EC-fp is 53% and 11%, respectively. The biggest average gain of CRR+RRP over EC-fp is achieved on Families *fvp-unsat.2.0* and *longmult* (18% and 17%, respectively). Unsurprisingly, both EC and EC-fp are usually much faster than CRR+RRP. However, on the three hardest instances of the barrel family, CRR+RRP outperforms EC-fp in terms of execution time.

7 Conclusions

We have proposed an algorithm for minimal unsatisfiable core extraction. It builds a resolution refutation using a SAT solver and finds a first approximation of a minimal unsatisfiable core. Then it checks, for every remaining initial clause C , if it belongs to the minimal unsatisfiable core. The algorithm reuses conflict clauses and resolution relations throughout its execution. We have demonstrated that our algorithm is faster than currently existing algorithms by a factor of 6 or more on large problems with non-overly hard resolution proofs, and that it can find minimal unsatisfiable cores for real-world formal verification benchmarks.

Acknowledgments

We thank Jinbo Huang for his help in obtaining and using MUP and Zaher Andraus for his help in receiving AMUSE.

References

- [1] Moskewicz, M.W., Madigan, C.F., Zhao, Y., Zhang, L., Malik, S.: Chaff: engineering an efficient SAT solver. In: Proceedings of the 38th Design Automation Conference (DAC'01). (2001)
- [2] Fu, Z., Mahajan, Y., Malik, S.: ZChaff2004: an efficient SAT solver. In: Proceedings of Seventh International Conference on Theory and Applications of Satisfiability Testing (SAT'04). (2004)
- [3] Eén, N., Sörensson, N.: An extensible SAT-solver. In: Proceedings of Sixth International Conference on Theory and Applications of Satisfiability Testing (SAT'03). (2003)
- [4] Papadimitriou, C.H., Yannakakis, M.: The complexity of facets (and some facets of complexity). In: Proceedings of the Fourteenth Annual ACM Symposium on Theory of Computing (STOC'82). (1982)
- [5] McMillan, K.L.: Symbolic model checking: an approach to the state explosion problem. PhD thesis, Pittsburgh, PA, USA (1992)
- [6] Nam, G.J., Aloul, F., Sakallah, K., Rutenbar, R.: A comparative study of two Boolean formulations of FPGA detailed routing constraints. In: Proceedings of the 2001 International Symposium on Physical Design (ISPD'01). (2001)
- [7] McMillan, K.L., Amla, N.: Automatic abstraction without counterexamples. In: Proceedings of the Ninth International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'03). (2003)
- [8] Szeider, S.: Minimal unsatisfiable formulas with bounded clause-variable difference are fixed-parameter tractable. *Journal of Computer and System Sciences* **69**(4) (2004) 656–674
- [9] Bruni, R., Sassano, A.: Restoring Satisfiability or Maintaining Unsatisfiability by Finding Small Unsatisfiable Subformulae. In: Proceedings of the Workshop on Theory and Application of Satisfiability Testing (SAT'01). (2001)
- [10] Bruni, R.: Approximating minimal unsatisfiable subformulae by means of adaptive core search. *Discrete Applied Mathematics* **130**(2) (2003) 85–100
- [11] Oh, Y., Mneimneh, M.N., Andraus, Z.S., Sakallah, K.A., Markov, I.L.: AMUSE: a minimally-unsatisfiable subformula extractor. In: Proceedings of the 41th Design Automation Conference (DAC'04). (2004)
- [12] Lynce, I., Marques-Silva, J.P.: On computing minimum unsatisfiable cores. In: Proceedings of the Seventh International Conference on Theory and Applications of Satisfiability Testing (SAT'04). (2004)

- [13] Liffton, M.H., Sakallah, K.A.: On finding all minimally unsatisfiable subformulas. In: Proceedings of the Eighth International Conference on Theory and Applications of Satisfiability Testing (SAT'05). (2005)
- [14] Mneimneh, M.N., Lynce, I., Andraus, Z.S., Marques-Silva, J.P., Sakallah, K.A.: A branch and bound algorithm for extracting smallest minimal unsatisfiable formulas. In: Proceedings of the Eighth International Conference on Theory and Applications of Satisfiability Testing (SAT'05). (2005)
- [15] Huang, J.: MUP: A minimal unsatisfiability prover. In: Proceedings of the Tenth Asia and South Pacific Design Automation Conference (ASP-DAC'05). (2005)
- [16] Berre, D.L., Simon, L.: (Results of the SAT 2005 competition) <http://www.lri.fr/~simon/contest/results/sizes-industrial.png>.
- [17] Zhang, L., Malik, S.: Extracting small unsatisfiable cores from unsatisfiable Boolean formula. In: Proceedings of Sixth International Conference on Theory and Applications of Satisfiability Testing (SAT'03). (2003)
- [18] Goldberg, E., Novikov, Y.: Verification of proofs of unsatisfiability for CNF formulas. In: Proceedings of Design, Automation and Test in Europe Conference and Exhibition (DATE'03). (2003)
- [19] Marques-Silva, J.P., Sakallah, K.A.: GRASP: A search algorithm for propositional satisfiability. *IEEE Transactions on Computers* **48**(5) (1999) 506–521
- [20] Nadel, A.: Backtrack search algorithms for propositional logic satisfiability: Review and innovations. Master's thesis, Hebrew University of Jerusalem, Jerusalem, Israel (2002)
- [21] Goldberg, E., Novikov, Y.: BerkMin: A fast and robust SAT-solver. In: Design, Automation, and Test in Europe (DATE '02). (2002)
- [22] Eén, N., Sörensson, N.: MiniSat v1.13 - a SAT solver with conflict-clause minimization. (http://www.cs.chalmers.se/Cs/Research/FormalMethods/MiniSat/cgi/MiniSat_v1.13_short.ps.gz)
- [23] Dershowitz, N., Hanna, Z., Nadel, A.: A clause-based heuristic for SAT solvers. In: Proceedings of the Eighth International Conference on Theory and Applications of Satisfiability Testing (SAT'05). (2005)
- [24] Velez, M., Bryant, R.: Effective use of Boolean satisfiability procedures in the formal verification of superscalar and VLIW microprocessors. In: Proceedings of the 38th Design Automation Conference (DAC'01). (2001)
- [25] Biere, A., Cimatti, A., Clarke, E.M., Zhu, Y.: Symbolic Model Checking without BDDs. In: Proceedings of the Fifth International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS'99). (1999)
- [26] Berre, D.L., Simon, L.: Fifty-five solvers in Vancouver: The SAT 2004 competition. In: Proceedings of the Seventh International Symposium on Theory and Applications of Satisfiability Testing (SAT'04). (2004)