

Improving Document Ranking with Dual Word Embeddings

Eric Nalisnick
University of California
Irvine, USA
enalsni@uci.edu

Bhaskar Mitra
Microsoft
Cambridge, UK
bmitra@microsoft.com

Nick Craswell, Rich Caruana
Microsoft
Redmond, USA
nickcr, rcaruana@microsoft.com

ABSTRACT

This paper investigates the popular neural word embedding method *Word2vec* as a source of evidence in document ranking. In contrast to NLP applications of *word2vec*, which tend to use only the input embeddings, we retain both the input and the output embeddings, allowing us to calculate a different word similarity that may be more suitable for document ranking. We map the query words into the *input* space and the document words into the *output* space, and compute a relevance score by aggregating the cosine similarities across all the query-document word pairs. We postulate that the proposed *Dual Embedding Space Model* (DESM) provides evidence that a document is *about* a query term, in addition to and complementing the traditional term frequency based approach.

Categories and Subject Descriptors

H.3 [Information Storage and Retrieval]: H.3.3 Information Search and Retrieval

Keywords: Document ranking; Word embeddings; Word2vec

1. INTRODUCTION

A key challenge for information retrieval (IR) is to distinguish whether a document merely references a term or is *about* that entity. The traditional approach is to count repetitions of query terms [6] in the document. However, as previously noted [5], the probabilistic model of IR can also consider additional terms that correlate with relevance. The two passages in Fig. 1 are indistinguishable for the query term *Albuquerque* under term counting but the presence of related terms like "population" and "metropolitan" points to the passage about the city. We propose to identify such related terms using *word2vec*¹.

Word2vec [2, 3] learns word embeddings via maximizing the log conditional probability of the word given the context word(s) occurring within a fixed-sized window. Therefore the learnt embeddings contain useful knowledge about word co-occurrence. A crucial detail often overlooked is that two different sets of vectors are learnt by the model corresponding to the input and the output words, henceforth referred to as the *IN* and *OUT* embeddings. By default, *word2vec* discards the *OUT* vectors at the end of training. However, for certain IR tasks we postulate that we should use both the *IN* and the *OUT* embeddings jointly. Table 1 shows that the nearest neighbours of the word "yale" using IN-OUT vector co-

Albuquerque is the most populous city in the U.S. state of New Mexico. The high-altitude city serves as the county seat of Bernalillo County, and it is situated in the central part of the state, straddling the Rio Grande. The city population is 557,169 as of the July 1, 2014, population estimate from the United States Census Bureau, and ranks as the 32nd-largest city in the U.S. The Metropolitan Statistical Area (or MSA) has a population of 902,797 according to the United States Census Bureau's most recently available estimate for July 1, 2013.

(a)

Allen suggested that they could program a BASIC interpreter for the device; after a call from Gates claiming to have a working interpreter, MITS requested a demonstration. Since they didn't actually have one, Allen worked on a simulator for the Altair while Gates developed the interpreter. Although they developed the interpreter on a simulator and not the actual device, the interpreter worked flawlessly when they demonstrated the interpreter to MITS in Albuquerque, New Mexico in March 1975; MITS agreed to distribute it, marketing it as Altair BASIC.

(b)

Figure 1: Two different passages from Wikipedia that mentions *Albuquerque* (highlighted in orange) exactly once. Highlighted in green are all the words that have an IN-OUT similarity score with the word "Albuquerque" above a fixed threshold and provides supporting evidence that (a) is *about* Albuquerque, whereas (b) happens to only *mention* the city.

sine similarity produces words that often co-occur with "yale" (e.g., "faculty" and "alumni") as opposed to the IN-IN similarity which gives functionally similar words (e.g., "harvard" and "nyu"). We use this property of the IN-OUT embeddings to propose a novel *Dual Embedding Space Model* (DESM) for document ranking.

2. DUAL EMBEDDING SPACE MODEL

Given q_i and d_j as the embedding vectors for the i^{th} and the j^{th} term of the query and the document, respectively, we define the *Dual Embedding Space Model* as:

$$DESM(Q, D) = \frac{1}{|Q|} \sum_{q_i \in Q} \frac{q_i^T \bar{D}}{\|q_i\| \|\bar{D}\|}, \quad (1)$$

where,

$$\bar{D} = \frac{1}{|D|} \sum_{d_j \in D} \frac{d_j}{\|d_j\|} \quad (2)$$

\bar{D} is the centroid of all the normalized document word vectors serving as a single embedding for the whole document. Note that taking the centroid of the document word vectors is equivalent to computing the similarity between all query-document word pairs.

¹<https://code.google.com/archive/p/word2vec>

Table 1: The nearest neighbours for the words "yale", "seahawks" and "eminem" based on the IN-IN and the IN-OUT vector cosine similarities. The IN-IN cosine similarities are high for words that are similar by function or type (*typical*), and the IN-OUT similarities are high between words that co-occur in the same query or document frequently (*topical*).

| yale | | seahawks | | eminem | |
|---------|-------------|----------|------------|----------|-----------|
| IN-IN | IN-OUT | IN-IN | IN-OUT | IN-IN | IN-OUT |
| yale | yale | seahawks | seahawks | eminem | eminem |
| harvard | faculty | 49ers | highlights | rihanna | rap |
| nyu | alumni | broncos | jerseys | ludacris | featuring |
| cornell | orientation | packers | tshirts | kanye | tracklist |
| tulane | haven | nfl | seattle | beyonce | diss |
| tufts | graduate | steelers | hats | 2pac | performs |

Also, the document embeddings can be pre-computed which is important for runtime efficiency. We only need to sum the score contributions across the query terms at the time of ranking.

As previously mentioned, the word2vec model contains two separate embedding spaces (IN and OUT) which gives us at least two variants of the DESM, corresponding to retrieval in the IN-OUT space or the IN-IN space².

$$DESM_{IN-OUT}(Q, D) = \frac{1}{|Q|} \sum_{q_i \in Q} \frac{q_{IN,i}^T \overline{D_{OUT}}}{\|q_{IN,i}\| \|D_{OUT}\|} \quad (3)$$

$$DESM_{IN-IN}(Q, D) = \frac{1}{|Q|} \sum_{q_i \in Q} \frac{q_{IN,i}^T \overline{D_{IN}}}{\|q_{IN,i}\| \|D_{IN}\|} \quad (4)$$

We expect the $DESM_{IN-OUT}$ to behave differently than the $DESM_{IN-IN}$ because of the difference in their notions of word relatedness as shown in Table 1.

One of the challenges of the embedding models is that they can only be applied to a fixed size vocabulary. We leave the exploration of possible strategies to deal with out-of-vocab (OOV) words for future investigation. In this paper, all the OOV words are ignored for computing the DESM score, but not for computing the TF-IDF feature, a potential advantage for the latter.

3. EXPERIMENTS

We train a *Continuous Bag-of-Words* (CBOW) model on a query corpus consisting of 618,644,170 queries and a vocabulary size of 2,748,230 words. The queries are sampled from Bing's large scale search logs from the period of August 19, 2014 to August 25, 2014. We repeat all our experiments using another CBOW model trained on a corpus of document body text with 341,787,174 distinct sentences sampled from the Bing search index and a corresponding vocabulary size of 5,108,278 words. Empirical results for both the models are presented in Table 2. Although our results are based exclusively on the CBOW model, the proposed methodology should be applicable to vectors produced by the *Skip-Gram* model, as both produce qualitatively and quantitatively similar embeddings.

We compare the retrieval performance of DESM against BM25 [6], a traditional count-based method, and Latent Semantic Analysis (LSA) [1], a traditional vector-based method. For the BM25 baseline we use the values of 1.7 for the k_1 and 0.95 for the b parameters based on a parameter sweep on a validation set. The LSA model is trained on the body text of 366,470 randomly sampled documents from Bing's index with a vocabulary size of 480,608 words. The evaluation set consists of 7,741 queries randomly sampled from

²It is also possible to define $DESM_{OUT-OUT}$ and $DESM_{OUT-IN}$, but we expect them to behave similar to $DESM_{IN-IN}$ and $DESM_{IN-OUT}$, respectively.

Table 2: The $DESM_{IN-OUT}$ performs significantly better than both the BM25 and the LSA baselines, as well as the $DESM_{IN-IN}$ on NDCG computed at positions three and ten. Also, the DESMs using embeddings trained on the query corpus performs better than if trained on document body text. The highest NDCG values for every column is highlighted in bold and all the statistically significant ($p < 0.05$) differences over the BM25 baseline are marked with the asterisk (*).

| | NDCG@3 | NDCG@10 |
|-------------------------------------|---------------|---------------|
| BM25 | 29.14 | 44.77 |
| LSA | 28.25* | 44.24* |
| DESM (IN-IN, trained on body text) | 29.59 | 45.51* |
| DESM (IN-IN, trained on queries) | 29.72 | 46.36* |
| DESM (IN-OUT, trained on body text) | 30.32* | 46.57* |
| DESM (IN-OUT, trained on queries) | 31.14* | 47.89* |

Bing's query logs from the period of October, 2014 to December, 2014. For each sampled query, a set of candidate documents is constructed by retrieving the top results from Bing over multiple scrapes during a period of a few months. In total the final evaluation set contains 171,302 unique documents across all queries which are then judged by human evaluators on a five point relevance scale.

We report the normalized discounted cumulative gain (NDCG) at different rank positions as a measure of performance for the different models. The results show that the $DESM_{IN-OUT}$ outperforms both the BM25 and the LSA baselines, as well as the $DESM_{IN-IN}$ at all rank positions. The embeddings trained on the query corpus also achieves better results than the embeddings trained on body text. We provide additional analysis and experiment results in [4].

4. DISCUSSION AND CONCLUSION

We formulated a *Dual Embedding Space Model* (DESM) that leverages the often discarded output embeddings learned by the word2vec model. Our model exploits both the input and the output embeddings to capture topic-based semantic relationships. The examples in Table 1 show that different nearest neighbours can be found by using proximity in the IN-OUT vs the IN-IN spaces. In our experiments ranking via proximity in the IN-OUT space performs better for retrieval than the IN-IN based ranking. This finding emphasizes that the performance of the word2vec model is application dependent and that quantifying semantic relatedness via cosine similarity in the IN space should not be a default practice.

References

- [1] S. C. Deerwester, S. T. Dumais, T. K. Landauer, G. W. Furnas, and R. A. Harshman. Indexing by latent semantic analysis. *JASIS*, 41(6):391–407, 1990.
- [2] T. Mikolov, K. Chen, G. Corrado, and J. Dean. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*, 2013.
- [3] T. Mikolov, I. Sutskever, K. Chen, G. S. Corrado, and J. Dean. Distributed representations of words and phrases and their compositionality. In *Proc. NIPS*, pages 3111–3119, 2013.
- [4] B. Mitra, E. Nalisnick, N. Craswell, and R. Caruana. A dual embedding space model for document ranking. *arXiv:1602.01137*, 2016.
- [5] S. Robertson. Understanding inverse document frequency: on theoretical arguments for idf. *Journal of documentation*, 60(5): 503–520, 2004.
- [6] S. Robertson and H. Zaragoza. *The probabilistic relevance framework: BM25 and beyond*. Now Publishers Inc, 2009.