Minimizing Manual Annotation Cost In Supervised Training From Corpora

Sean P. Engelson and Ido Dagan

Department of Mathematics and Computer Science Bar-Ilan University 52900 Ramat Gan, Israel {engelson,dagan}@bimacs.cs.biu.ac.il

Abstract

Corpus-based methods for natural language processing often use supervised training, requiring expensive manual annotation of training corpora. This paper investigates methods for reducing annotation cost by sample selection. In this approach, during training the learning program examines many unlabeled examples and selects for labeling (annotation) only those that are most informative at each stage. This avoids redundantly annotating examples that contribute little new information. This paper extends our previous work on committee-based sample selection for probabilistic classifiers. We describe a family of methods for committee-based sample selection, and report experimental results for the task of stochastic part-ofspeech tagging. We find that all variants achieve a significant reduction in annotation cost, though their computational efficiency differs. In particular, the simplest method, which has no parameters to tune, gives excellent results. We also show that sample selection yields a significant reduction in the size of the model used by the tagger.

1 Introduction

Many corpus-based methods for natural language processing (NLP) are based on supervised training—acquiring information from a manually annotated corpus. Therefore, reducing annotation cost is an important research goal for statistical NLP. The ultimate reduction in annotation cost is achieved by unsupervised training methods, which do not require an annotated corpus at all (Kupiec, 1992; Merialdo,

1994; Elworthy, 1994). It has been shown, however, that some supervised training prior to the unsupervised phase is often beneficial. Indeed, fully unsupervised training may not be feasible for certain tasks. This paper investigates an approach for optimizing the supervised training (learning) phase, which reduces the annotation effort required to achieve a desired level of accuracy of the trained model.

In this paper, we investigate and extend the committee-based sample selection approach to minimizing training cost (Dagan and Engelson, 1995). When using sample selection, a learning program examines many unlabeled (not annotated) examples, selecting for labeling only those that are most informative for the learner at each stage of training (Seung, Opper, and Sompolinsky, 1992; Freund et al., 1993; Lewis and Gale, 1994; Cohn, Atlas, and Ladner, 1994). This avoids redundantly annotating many examples that contribute roughly the same information to the learner.

Our work focuses on sample selection for training probabilistic classifiers. In statistical NLP, probabilistic classifiers are often used to select a preferred analysis of the linguistic structure of a text (for example, its syntactic structure (Black et al., 1993), word categories (Church, 1988), or word senses (Gale, Church, and Yarowsky, 1993)). As a representative task for probabilistic classification in NLP, we experiment in this paper with sample selection for the popular and well-understood method of stochastic part-of-speech tagging using Hidden Markov Models.

We first review the basic approach of committee-based sample selection and its application to part-of-speech tagging. This basic approach gives rise to a family of algorithms (including the original algorithm described in (Dagan and Engelson, 1995)) which we then describe. First, we describe the 'simplest' committee-based selection algorithm, which

has no parameters to tune. We then generalize the selection scheme, allowing more options to adapt and tune the approach for specific tasks. The paper compares the performance of several instantiations of the general scheme, including a batch selection method similar to that of Lewis and Gale (1994). In particular, we found that the simplest version of the method achieves a significant reduction in annotation cost, comparable to that of other versions. We also evaluate the computational efficiency of the different variants, and the number of unlabeled examples they consume. Finally, we study the effect of sample selection on the size of the model acquired by the learner.

2 Probabilistic Classification

This section presents the framework and terminology assumed for probabilistic classification, as well as its instantiation for stochastic bigram part-of-speech tagging.

A probabilistic classifier classifies input examples e by classes $c \in C$, where C is a known set of possible classes. Classification is based on a score function, $F_M(c,e)$, which assigns a score to each possible class of an example. The classifier then assigns the example to the class with the highest score. F_M is determined by a probabilistic model M. In many applications, F_M is the conditional probability function, $P_M(c|e)$, specifying the probability of each class given the example, but other score functions that correlate with the likelihood of the class are often used.

In stochastic part-of-speech tagging, the model assumed is a Hidden Markov Model (HMM), and input examples are sentences. The class c, to which a sentence is assigned is a sequence of the parts of speech (tags) for the words in the sentence. The score function is typically the joint (or conditional) probability of the sentence and the tag sequence¹. The tagger then assigns the sentence to the tag sequence which is most probable according to the HMM.

The probabilistic model M, and thus the score function F_M , are defined by a set of parameters, $\{\alpha_i\}$. During training, the values of the parameters are estimated from a set of statistics, S, extracted from a training set of annotated examples. We denote a particular model by $M = \{a_i\}$, where each a_i is a specific value for the corresponding α_i .

In bigram part-of-speech tagging the HMM model M contains three types of parameters: transition probabilities $P(t_i \rightarrow t_j)$ giving the probability of tag

 t_j occurring after tag t_i , lexical probabilities P(t|w) giving the probability of tag t labeling word w, and tag probabilities P(t) giving the marginal probability of a tag occurring.² The values of these parameters are estimated from a tagged corpus which provides a training set of labeled examples (see Section 4.1).

3 Evaluating Example Uncertainty

A sample selection method needs to evaluate the expected usefulness, or information gain, of learning from a given example. The methods we investigate approach this evaluation implicitly, measuring an example's informativeness as the uncertainty in its classification given the current training data (Seung, Opper, and Sompolinsky, 1992; Lewis and Gale, 1994; MacKay, 1992). The reasoning is that if an example's classification is uncertain given current training data then the example is likely to contain unknown information useful for classifying similar examples in the future.

We investigate the *committee-based* method, where the learning algorithm evaluates an example by giving it to a *committee* containing several variant models, all 'consistent' with the training data seen so far. The more the committee members agree on the classification of the example, the greater our certainty in its classification. This is because when the training data entails a specific classification with high certainty, most (in a probabilistic sense) classifiers consistent with the data will produce that classification.

The committee-based approach was first proposed in a theoretical context for learning binary nonprobabilistic classifiers (Seung, Opper, and Sompolinsky, 1992; Freund et al., 1993). In this paper, we extend our previous work (Dagan and Engelson, 1995) where we applied the basic idea of the committee-based approach to probabilistic classification. Taking a Bayesian perspective, the posterior probability of a model, P(M|S), is determined given statistics S from the training set (and some prior distribution for the models). Committee members are then generated by drawing models randomly from P(M|S). An example is selected for labeling if the committee members largely disagree on its classification. This procedure assumes that one can sample from the models' posterior distribution, at least approximately.

To illustrate the generation of committeemembers, consider a model containing a single binomial parameter α (the probability of a success),

 $^{^{1}\}mathrm{This}$ gives the Viterbi model (Merialdo, 1994), which we use here.

This version of the method uses Bayes' theorem $(P(w_i|t_i) \propto \frac{P(t_i|w_i)}{P(t_i)})$ (Church, 1988).

with estimated value a. The statistics S for such a model are given by N, the number of trials, and x, the number of successes in those trials.

Given N and x, the 'best' parameter value may be estimated by one of several estimation methods. For example, the maximum likelihood estimate for α is $a=\frac{x}{N}$, giving the model $M=\{a\}=\{\frac{x}{N}\}$. When generating a committee of models, however, we are not interested in the 'best' model, but rather in sampling the distribution of models given the statistics. For our example, we need to sample the posterior density of estimates for α , namely $P(\alpha=a|S)$. Sampling this distribution yields a set of estimates scattered around $\frac{x}{N}$ (assuming a uniform prior), whose variance decreases as N increases. In other words, the more statistics there are for estimating the parameter, the more similar are the parameter values used by different committee members.

For models with multiple parameters, parameter estimates for different committee members differ more when they are based on low training counts, and they agree more when based on high counts. Selecting examples on which the committee members disagree contributes statistics to currently uncertain parameters whose uncertainty *also* affects classification.

It may sometimes be difficult to sample P(M|S) due to parameter interdependence. Fortunately, models used in natural language processing often assume independence between most model parameters. In such cases it is possible to generate committee members by sampling the posterior distribution for each independent group of parameters separately.

4 Bigram Part-Of-Speech Tagging

4.1 Sampling model parameters

In order to generate committee members for bigram tagging, we sample the posterior distributions for transition probabilities, $P(t_i \rightarrow t_j)$, and for lexical probabilities, P(t|w) (as described in Section 2).

Both types of the parameters we sample have the form of multinomial distributions. Each multinomial random variable corresponds to a conditioning event and its values are given by the corresponding set of conditioned events. For example, a transition probability parameter $P(t_i \rightarrow t_j)$ has conditioning event t_i and conditioned event t_j .

Let $\{u_i\}$ denote the set of possible values of a given multinomial variable, and let $S = \{n_i\}$ denote a set of statistics extracted from the training set for that variable, where n_i is the number of times that the value u_i appears in the training set for the variable, defining $N = \sum_i n_i$. The parameters

whose posterior distributions we wish to estimate are $\alpha_i = P(u_i)$.

The maximum likelihood estimate for each of the multinomial's distribution parameters, α_i , is $\hat{\alpha}_i = \frac{n_i}{N}$. In practice, this estimator is usually smoothed in some way to compensate for data sparseness. Such smoothing typically reduces slightly the estimates for values with positive counts and gives small positive estimates for values with a zero count. For simplicity, we describe here the approximation of $P(\alpha_i = a_i|S)$ for the unsmoothed estimator³.

We approximate the posterior $P(\alpha_i = a_i|S)$ by first assuming that the multinomial is a collection of independent binomials, each of which corresponds to a single value u_i of the multinomial; we then normalize the values so that they sum to 1. For each such binomial, we approximate $P(\alpha_i = a_i|S)$ as a truncated normal distribution (restricted to [0,1]), with estimated mean $\mu = \frac{n_i}{N}$ and variance $\sigma^2 = \frac{\mu(1-\mu)}{N}$. To generate a particular multinomial distribution,

To generate a particular multinomial distribution, we randomly choose values for the binomial parameters α_i from their approximated posterior distributions (using the simple sampling method given in (Press et al., 1988, p. 214)), and renormalize them so that they sum to 1. Finally, to generate a random HMM given statistics S, we choose values independently for the parameters of each multinomial, since all the different multinomials in an HMM are independent.

4.2 Examples in bigram training

Typically, concept learning problems are formulated such that there is a set of training examples that are independent of each other. When training a bigram model (indeed, any HMM), this is not true, as each word is dependent on that before it. This problem is solved by considering each sentence as an individual example. More generally, it is possible to break the text at any point where tagging is unambiguous. We thus use unambiguous words (those with only one possible part of speech) as example boundaries in bigram tagging. This allows us to train on smaller

³In the implementation we smooth the MLE by interpolation with a uniform probability distribution, following Merialdo (1994). Approximate adaptation of $P(\alpha_i = a_i|S)$ to the smoothed version of the estimator is simple.

⁴The normal approximation, while easy to implement, can be avoided. The posterior probability $P(\alpha_i = a_i|S)$ for the multinomial is given exactly by the Dirichlet distribution (Johnson, 1972) (which reduces to the Beta distribution in the binomial case). In this work we assumed a uniform prior distribution for each model parameter; we have not addressed the question of how to best choose a prior for this problem.

examples, focusing training more on the truly informative parts of the corpus.

5 Selection Algorithms

Within the committee-based paradigm there exist different methods for selecting informative examples. Previous research in sample selection has used either sequential selection (Seung, Opper, and Sompolinsky, 1992; Freund et al., 1993; Dagan and Engelson, 1995), or batch selection (Lewis and Catlett, 1994; Lewis and Gale, 1994). We describe here general algorithms for both sequential and batch selection.

Sequential selection examines unlabeled examples as they are supplied, one by one, and measures the disagreement in their classification by the committee. Those examples determined to be sufficiently informative are selected for training. Most simply, we can use a committee of size two and select an example when the two models disagree on its classification. This gives the following, parameter-free, two member sequential selection algorithm, executed for each unlabeled input example e:

- 1. Draw 2 models randomly from P(M|S), where S are statistics acquired from previously labeled examples;
- 2. Classify e by each model, giving classifications c_1 and c_2 ;
- 3. If $c_1 \neq c_2$, select e for annotation;
- 4. If e is selected, get its correct label and update S accordingly.

This basic algorithm needs no parameters. If desired, it is possible to tune the frequency of selection, by changing the variance of P(M|S) (or the variance of $P(\alpha_i = a_i|S)$ for each parameter), where larger variances increase the rate of disagreement among the committee members. We implemented this effect by employing a temperature parameter t, used as a multiplier of the variance of the posterior parameter distribution.

A more general algorithm results from allowing (i) a larger number of committee members, k, in order to sample P(M|S) more precisely, and (ii) more refined example selection criteria. This gives the following **general sequential selection algorithm**, executed for each unlabeled input example e:

- 1. Draw k models $\{M_i\}$ randomly from P(M|S) (possibly using a temperature t);
- 2. Classify e by each model M_i giving classifications $\{c_i\}$;
- 3. Measure the disagreement D over $\{c_i\}$;

- Decide whether to select e for annotation, based on the value of D;
- 5. If e is selected, get its correct label and update S accordingly.

It is easy to see that two member sequential selection is a special case of general sequential selection, where any disagreement is considered sufficient for selection. In order to instantiate the general algorithm for larger committees, we need to define (i) a measure for disagreement (Step 3), and (ii) a selection criterion (Step 4).

Our approach to measuring disagreement is to use the *vote entropy*, the entropy of the distribution of classifications assigned to an example ('voted for') by the committee members. Denoting the number of committee members assigning c to e by V(c,e), the vote entropy is:

$$D = -\frac{1}{\log k} \sum_{c} \frac{V(c, e)}{k} \log \frac{V(c, e)}{k}$$

(Dividing by $\log k$ normalizes the scale for the number of committee members.) Vote entropy is maximized when all committee members disagree, and is zero when they all agree.

In bigram tagging, each example consists of a sequence of several words. In our system, we measure D separately for each word, and use the average entropy over the word sequence as a measurement of disagreement for the example. We use the average entropy rather than the entropy over the entire sequence, because the number of committee members is small with respect to the total number of possible tag sequences. Note that we do not look at the entropy of the distribution given by each single model to the possible tags (classes), since we are only interested in the uncertainty of the final classification (see the discussion in Section 7).

We consider two alternative selection criteria (for Step 4). The simplest is thresholded selection, in which an example is selected for annotation if its vote entropy exceeds some threshold θ . The other alternative is randomized selection, in which an example is selected for annotation based on the flip of a coin biased according to the vote entropy—a higher vote entropy entailing a higher probability of selection. We define the selection probability as a linear function of vote entropy: p = gD, where g is an entropy gain parameter. The selection method we used in our earlier work (Dagan and Engelson, 1995) is randomized sequential selection using this linear selection probability model, with parameters k, t and g.

An alternative to sequential selection is batch selection. Rather than evaluating examples individually for their informativeness a large batch of examples is examined, and the m best are selected for annotation. The **batch selection algorithm**, executed for each batch B of N examples, is as follows:


- 1. For each example e in B:
 - (a) Draw k models randomly from P(M|S);
 - (b) Classify e by each model, giving classifications $\{c_i\}$;
 - (c) Measure the disagreement D_e for e over $\{c_i\}$;
- 2. Select for annotation the m examples from B with the highest D_e ;
- 3. Update S by the statistics of the selected examples.

This procedure is repeated sequentially for successive batches of N examples, returning to the start of the corpus at the end. If N is equal to the size of the corpus, batch selection selects the m globally best examples in the corpus at each stage (as in (Lewis and Catlett, 1994)). On the other hand, as N decreases, batch selection becomes closer to sequential selection.

6 Experimental Results

This section presents results of applying committee-based sample selection to bigram part-of-speech tagging, as compared with complete training on all examples in the corpus. Evaluation was performed using the University of Pennsylvania tagged corpus from the ACL/DCI CD-ROM I. For ease of implementation, we used a complete (closed) lexicon which contains all the words in the corpus.

The committee-based sampling algorithm was initialized using the first 1,000 words from the corpus, and then sequentially examined the following examples in the corpus for possible labeling. The training set consisted of the first million words in the corpus, with sentence ordering randomized to compensate for inhomogeneity in corpus composition. The test set was a separate portion of the corpus, consisting of 20,000 words. We compare the amount of training required by different selection methods to achieve a given tagging accuracy on the test set, where both the amount of training and tagging accuracy are measured over ambiguous words.⁵






Figure 1: Training versus accuracy. In batch, random, and thresholded runs, k=5 and t=50. (a) Number of ambiguous words selected for labeling versus classification accuracy achieved. (b) Accuracy versus number of words examined from the corpus (both labeled and unlabeled).

The effectiveness of randomized committee-based selection for part-of-speech tagging, with 5 and 10 committee members, was demonstrated in (Dagan and Engelson, 1995). Here we present and compare results for batch, randomized, thresholded, and two member committee-based selection.

Figure 1 presents the results of comparing the several selection methods against each other. The plots shown are for the best parameter settings that we found through manual tuning for each method. Figure 1(a) shows the advantage that sample selection gives with regard to annotation cost. For example,

⁵Note that most other work on tagging has measured accuracy over all words, not just ambiguous ones. Complete training of our system on 1,000,000 words gave us an accuracy of 93.5% over ambiguous words, which cor-

responds to an accuracy of 95.9% over all words in the test set, comparable to other published results on bigram tagging.




Figure 2: Evaluating batch selection, for m = 5. (a) Accuracy achieved versus batch size at different numbers of selected training words. (b) Accuracy versus number of words examined from the corpus for different batch sizes.

Figure 3: The size of the trained model, measured by the number of frequency counts > 0, plotted (y-axis) versus classification accuracy achieved (x-axis). (a) Lexical counts (freq(t, w)) (b) Bigram counts (freq $(t_1 \rightarrow t_2)$).

0.92

0.93

0.93

complete training requires annotated examples containing 98,000 ambiguous words to achieve a 92.6% accuracy (beyond the scale of the graph), while the selective methods require only 18,000–25,000 ambiguous words to achieve this accuracy. We also find that, to a first approximation, all selection methods considered give similar results. Thus, it seems that a refined choice of the selection method is not crucial for achieving large reductions in annotation cost.

This equivalence of the different methods also largely holds with respect to computational efficiency. Figure 1(b) plots classification accuracy versus number of words examined, instead of those selected. We see that while all selective methods are less efficient in terms of examples examined than complete training, they are comparable to each other. Two member selection seems to have a clear, though small, advantage.

In Figure 2 we investigate further the properties of batch selection. Figure 2(a) shows that accuracy increases with batch size only up to a point, and then starts to decrease. This result is in line with theoretical difficulties with batch selection (Freund et al., 1993) in that batch selection does not account for the distribution of input examples. Hence, once batch size increases past a point, the input distribution has too little influence on which examples are selected, and hence classification accuracy decreases. Furthermore, as batch size increases, computational efficiency, in terms of the number of examples examined to attain a given accuracy, decreases tremendously (Figure 2(b)).

The ability of committee-based selection to fo-

cus on the more informative parts of the training corpus is analyzed in Figure 3. Here we examined the number of lexical and bigram counts that were stored (i.e, were non-zero) during training, using the two member selection algorithm and complete training. As the graphs show, the sample selection method achieves the same accuracy as complete training with fewer lexical and bigram counts. This means that many counts in the data are less useful for correct tagging, as replacing them with smoothed estimates works just as well.⁶ Committee-based selection ignores such counts, focusing on parameters which improve the model. This behavior has the practical advantage of reducing the size of the model significantly (by a factor of three here). Also, the average count is lower in a model constructed by selective training than in a fully trained model, suggesting that the selection method avoids using examples which increase the counts for already known parameters.

7 Discussion

Why does committee-based sample selection work? Consider the properties of those examples that are selected for training. In general, a selected training example will contribute data to several statistics, which in turn will improve the estimates of several parameter values. An *informative* example is therefore one whose contribution to the statistics leads to a significantly useful improvement of model parameter estimates. Model parameters for which acquiring additional statistics is most beneficial can be characterized by the following three properties:

- 1. The current estimate of the parameter is uncertain due to insufficient statistics in the training set. Additional statistics would bring the estimate closer to the true value.
- Classification of examples is sensitive to changes in the current estimate of the parameter. Otherwise, even if the current value of the parameter is very uncertain, acquiring additional statistics will not change the resulting classifications.
- 3. The parameter affects classification for a large proportion of examples in the input. Parameters that affect only few examples have low overall utility.

The committee-based selection algorithms work because they tend to select examples that affect parameters with the above three properties. Property 1 is addressed by randomly drawing the parameter values for committee members from the posterior distribution given the current statistics. When the statistics for a parameter are insufficient, the variance of the posterior distribution of the estimates is large, and hence there will be large differences in the values of the parameter chosen for different committee members. Note that property 1 is not addressed when uncertainty in classification is only judged relative to a *single* model⁷ (as in, eg, (Lewis and Gale, 1994)).

Property 2 is addressed by selecting examples for which committee members highly disagree in *classification* (rather than measuring disagreement in parameter estimates). Committee-based selection thus addresses properties 1 and 2 simultaneously: it acquires statistics just when uncertainty in current parameter estimates entails uncertainty regarding the *appropriate* classification of the example. Our results show that this effect is achieved even when using only two committee members to sample the space of likely classifications. By *appropriate* classification we mean the classification given by a perfectly-trained model, that is, one with accurate parameter values.

Note that this type of uncertainty regarding the identity of the appropriate classification, is different than uncertainty regarding the correctness of the classification itself. For example, sufficient statistics may yield an accurate 0.51 probability estimate for a class c in a given example, making it certain that c is the appropriate classification. However, the certainty that c is the correct classification is low, since there is a 0.49 chance that c is the wrong class for the example. A single model can be used to estimate only the second type of uncertainty, which does not correlate directly with the utility of additional training.

Finally, property 3 is addressed by independently examining input examples which are drawn from the input distribution. In this way, we implicitly model the distribution of model parameters used for classifying input examples. Such modeling is absent in batch selection, and we hypothesize that this is the reason for its lower effectiveness.

⁶As noted above, we smooth the MLE estimates by interpolation with a uniform probability distribution (Merialdo, 1994).

⁷The use of a single model is also criticized in (Cohn, Atlas, and Ladner, 1994).

8 Conclusions

Annotating large textual corpora for training natural language models is a costly process. We propose reducing this cost significantly using committee-based sample selection, which reduces redundant annotation of examples that contribute little new information. The method can be applied in a semi-interactive process, in which the system selects several new examples for annotation at a time and updates its statistics after receiving their labels from the user. The implicit modeling of uncertainty makes the selection system generally applicable and quite simple to implement.

Our experimental study of variants of the selection method suggests several practical conclusions. First, it was found that the simplest version of the committee-based method, using a two-member committee, yields reduction in annotation cost comparable to that of the multi-member committee. The two-member version is simpler to implement, has no parameters to tune and is computationally more efficient. Second, we generalized the selection scheme giving several alternatives for optimizing the method for a specific task. For bigram tagging, comparative evaluation of the different variants of the method showed similar large reductions in annotation cost, suggesting the robustness of the committee-based approach. Third, sequential selection, which implicitly models the expected utility of an example relative to the example distribution, worked in general better than batch selection. The latter was found to work well only for small batch sizes, where the method mimics sequential selection. Increasing batch size (approaching 'pure' batch selection) reduces both accuracy and efficiency. Finally, we studied the effect of sample selection on the size of the trained model, showing a significant reduction in model size.

8.1 Further research

Our results suggest applying committee-based sample selection to other statistical NLP tasks which rely on estimating probabilistic parameters from an annotated corpus. Statistical methods for these tasks typically assign a probability estimate, or some other statistical score, to each alternative analysis (a word sense, a category label, a parse tree, etc.), and then select the analysis with the highest score. The score is usually computed as a function of the estimates of several 'atomic' parameters, often binomials or multinomials, such as:

• In word sense disambiguation (Hearst, 1991; Gale, Church, and Yarowsky, 1993): P(s|f),

where s is a specific sense of the ambiguous word in question w, and f is a feature of occurrences of w. Common features are words in the context of w or morphological attributes of it.

- In prepositional-phrase (PP) attachment (Hindle and Rooth, 1993): P(a|f), where a is a possible attachment, such as an attachment to a head verb or noun, and f is a feature, or a combination of features, of the attachment. Common features are the words involved in the attachment, such as the head verb or noun, the preposition, and the head word of the PP.
- In statistical parsing (Black et al., 1993): P(r|h), the probability of applying the rule r at a certain stage of the top down derivation of the parse tree given the history h of the derivation process.
- In text categorization (Lewis and Gale, 1994;
 Iwayama and Tokunaga, 1994): P(t|C), where
 t is a term in the document to be categorized,
 and C is a candidate category label.

Applying committee-based selection to supervised training for such tasks can be done analogously to its application in the current paper⁸. Furthermore, committee-based selection may be attempted also for training non-probabilistic classifiers, where explicit modeling of information gain is typically impossible. In such contexts, committee members might be generated by randomly varying some of the decisions made in the learning algorithm.

Another important area for future work is in developing sample selection methods which are independent of the eventual learning method to be applied. This would be of considerable advantage in developing selectively annotated corpora for general research use. Recent work on heterogeneous uncertainty sampling (Lewis and Catlett, 1994) supports this idea, using one type of model for example selection and a different type for classification.

Acknowledgments. We thank Yoav Freund and Yishay Mansour for helpful discussions. The first author gratefully acknowledges the support of the Fulbright Foundation.

⁸Measuring disagreement in full syntactic parsing is complicated. It may be approached by similar methods to those used for parsing evaluation, which measure the disagreement between the parser's output and the correct parse.

References

- [Black et al.1993] Black, Ezra, Fred Jelinek, John Lafferty, David Magerman, Robert Mercer, and Salim Roukos. 1993. Towards history-based grammars: using richer models for probabilistic parsing. In *Proc. of the Annual Meeting of the ACL*, pages 31–37.
- [Church1988] Church, Kenneth W. 1988. A stochastic parts program and noun phrase parser for unrestricted text. In *Proc. of ACL Conference on Applied Natural Language Processing*.
- [Cohn, Atlas, and Ladner1994] Cohn, David, Les Atlas, and Richard Ladner. 1994. Improving generalization with active learning. Machine Learning, 15:201–221.
- [Dagan and Engelson1995] Dagan, Ido and Sean Engelson. 1995. Committee-based sampling for training probabilistic classifiers. In Proc. Int'l Conference on Machine Learning, July.
- [Elworthy1994] Elworthy, David. 1994. Does Baum-Welch re-estimation improve taggers? In Proc. of ACL Conference on Applied Natural Language Processing, pages 53–58.
- [Freund et al.1993] Freund, Y., H. S. Seung, E. Shamir, and N. Tishby. 1993. Information, prediction, and query by committee. In Advances in Neural Information Processing, volume 5. Morgan Kaufmann.
- [Gale, Church, and Yarowsky1993] Gale, William, Kenneth Church, and David Yarowsky. 1993. A method for disambiguating word senses in a large corpus. *Computers and the Humanities*, 26:415–439.
- [Hearst1991] Hearst, Marti. 1991. Noun homograph disambiguation using local context in large text corpora. In Proc. of the Annual Conference of the UW Center for the New OED and Text Research, pages 1–22.
- [Hindle and Rooth1993] Hindle, Donald and Mats Rooth. 1993. Structural ambiguity and lexical relations. *Computational Linguistics*, 19(1):103– 120.
- [Iwayama and Tokunaga1994] Iwayama, M. and T. Tokunaga. 1994. A probabilistic model for text categorization based on a single random variable with multiple values. In Proceedings of the 4th Conference on Applied Natural Language Processing.
- [Johnson1972] Johnson, Norman L. 1972. Continuous Multivariate Distributions. John Wiley & Sons, New York.

- [Kupiec1992] Kupiec, Julian. 1992. Robust partof-speech tagging using a hidden makov model. Computer Speech and Language, 6:225–242.
- [Lewis and Catlett1994] Lewis, David D. and Jason Catlett. 1994. Heterogeneous uncertainty sampling for supervised learning. In *Proc. Int'l Conference on Machine Learning*.
- [Lewis and Gale1994] Lewis, David D. and William A. Gale. 1994. A sequential algorithm for training text classifiers. In *Proc. of the ACM SIGIR Conference*.
- [MacKay1992] MacKay, David J. C. 1992. Information-based objective functions for active data selection. Neural Computation, 4.
- [Merialdo1994] Merialdo, Bernard. 1994. Tagging text with a probabilistic model. Computational Linguistics, 20(2):155–172.
- [Press et al.1988] Press, William H., Brian P. Flannery, Saul A. Teukolsky, and William T. Vetterling. 1988. Numerical Recipes in C. Cambridge University Press.
- [Seung, Opper, and Sompolinsky1992] Seung, H. S., M. Opper, and H. Sompolinsky. 1992. Query by committee. In Proc. ACM Workshop on Computational Learning Theory.