
Conditional Random Fields: Probabilistic Models for Segmenting and Labeling Sequence Data

Paper by John Lafferty, Andrew McCallum, and Fernando Pereira
ICML 2001

Presentation by Joe Drish
May 9, 2002

Main Goals

- Present a new framework for labeling sequence data: Conditional Random Fields (CRFs)
- Describe the label bias problem
- Motivate the use of CRFs to solve the label bias problem
- Define the structure and properties of CRFs
- [Describe two training procedures for learning CRF parameters]
- [Sketch a proof of the convergence of the two training procedures]
- Compare experimentally to hidden Markov models (HMMs) and maximum entropy Markov models (MEMMs)

Labeling Sequence Data

- Given observed data sequences $X = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$
- A corresponding label sequence \mathbf{y}_k for each data sequence \mathbf{x}_k , and $Y = \{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\}$

Prediction Task

Given a sequence \mathbf{x} and model θ predict \mathbf{y}

Learning Task

Given training sets X and Y , learn the best model θ

Notation

set of sequences: X , uppercase

sequence of observations: \mathbf{x} , lowercase boldface, also called a data sequence

sequence of labels: \mathbf{y} , lowercase boldface, also called a sequence of states

single observation: x , lowercase, also called a data value

Example label and observation sequence

label y	<head>	X-NNTP-Poster: NewsHound v1.33	observation x
	<head>		
	<head>	Archive-name: acorn/faq/part2	
	<head>	Frequency: monthly	
	<head>		
	<question>	2.6) What configuration of serial cable should	
	<question>	I use?	
	<answer>	Here follows a diagram of the necessary	
	<answer>	connections programs to work properly. They	
	<answer>	are as far as know agreed upon by commercial	
	<answer>	comms software developers fo	
	<answer>		
	<answer>	Pins 1, 4, and 8 must be connected together	
	<answer>	is to avoid the well known serial port chip bugs.	

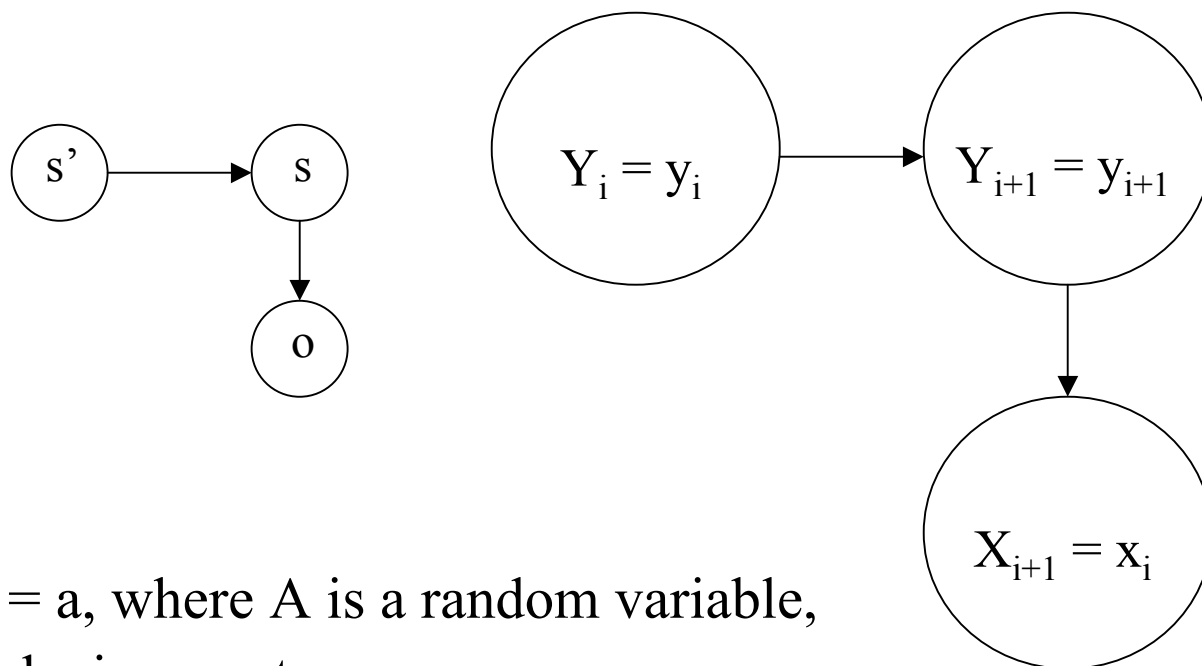
label sequence y

observation sequence x

Generative Modeling (HMMs)

Given training set X with label sequences Y :

- Train a model θ that maximizes $P(X, Y \mid \theta)$
- For a new data sequence \mathbf{x} , the predicted label \mathbf{y} maximizes $P(\mathbf{y} \mid \mathbf{x}) = P(\mathbf{y} \mid \mathbf{x}, \theta)P(\mathbf{x} \mid \theta)$

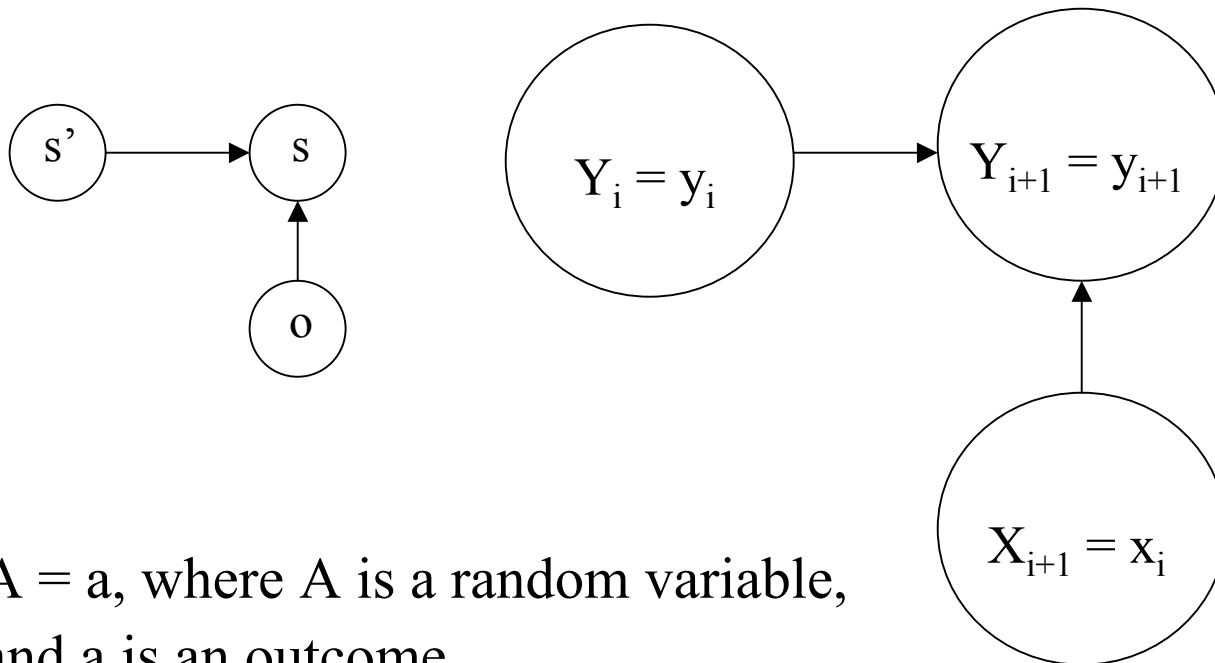


$A = a$, where A is a random variable,
and a is an outcome

Discriminative Modeling (MEMMs)

Given training set X with label sequences Y :

- Train a model θ that maximizes $P(Y \mid X, \theta)$
- For a new data sequence \mathbf{x} , the predicted label \mathbf{y} maximizes $P(\mathbf{y} \mid \mathbf{x}, \theta)$




$A = a$, where A is a random variable,
and a is an outcome


rib/rob models

Training data: $\{ \langle \text{rib}, 123 \rangle, \langle \text{rob}, 453 \rangle \}$


HMM



MEMM




CRF



Label Bias Problem

Consider this MEMM:



The label sequence 1,2 should score higher when ri is observed compared to ro.
Or, we expect $P(1 \text{ and } 2 \mid ri)$ to be greater than $P(1 \text{ and } 2 \mid ro)$.

Mathematically,

$$P(1 \text{ and } 2 \mid ro) = P(2 \mid 1 \text{ and } ro)P(1 \mid ro) = P(2 \mid 1 \text{ and } o)P(1 \mid r)$$

$$P(1 \text{ and } 2 \mid ri) = P(2 \mid 1 \text{ and } ri)P(1 \mid ri) = P(2 \mid 1 \text{ and } i)P(1 \mid r)$$


Since $P(2 \mid 1 \text{ and } x) = 1$ for all x , $P(1 \text{ and } 2 \mid ro) = P(1 \text{ and } 2 \mid ri)$

In the training data, label value 2 is the only label value observed after label value 1

Therefore $P(2 \mid 1) = 1$, so $P(2 \mid 1 \text{ and } x) = 1$ for all x

Changing the Set of States

Example:



$$P(14 \text{ and } 2 \mid ri) = P(2 \mid 14 \text{ and } ri)P(14 \mid ri) = P(2 \mid 14 \text{ and } i)P(14 \mid r) = (1)(1) = 1$$

$$P(14 \text{ and } 2 \mid ro) = P(2 \mid 14 \text{ and } ro)P(14 \mid ro) = P(2 \mid 14 \text{ and } o)P(14 \mid r) = (0)(1) = 0$$

This is a solution to the label bias problem.

But, changing the set of states would be impractical.

Conditional Random Fields (CRFs)

Disadvantages of MEMMs

$P(\mathbf{y} \mid \mathbf{x})$ = product of factors, one for each label

Each factor can depend only on previous label, and not future labels

So, let

$$P(\mathbf{y} \mid \mathbf{x}) = \exp\left(\sum_k f_k(\mathbf{y}, \mathbf{x})\right)$$


where each f_k is a property of part of \mathbf{y} and \mathbf{x}

Example: $f_k(\mathbf{x}, \mathbf{y}) = 1$ if X_i is uppercase and label Y_i is a proper noun.

Random Field Example

Let $G = (Y, E)$ be a graph where each vertex Y_v is a random variable
Suppose $P(Y_v \mid \text{all other } Y) = P(Y_v \mid \text{neighbors}(Y_v))$ then Y is a random field

Example:




- $P(Y_5 \mid \text{all other } Y) = P(Y_5 \mid Y_4, Y_6)$

Conditional Random Field Example

Suppose $P(Y_v | X, \text{all other } Y) = P(Y_v | X, \text{neighbors}(Y_v))$

then X with Y is a **conditional** random field



- $P(Y_3 | X, \text{all other } Y) = P(Y_3 | X, Y_2, Y_4)$
- Think of X as observations and Y as labels

Conditional Distribution

If Y is a tree, the distribution over the label sequence $Y = y$, given $X = x$, is:

$$p_{\Theta}(\mathbf{y} \mid \mathbf{x}) \propto \exp \left(\sum_{e \in E, k} \lambda_k f_k(e, \mathbf{y} \mid_e, \mathbf{x}) + \sum_{v \in V, k} \mu_k g_k(v, \mathbf{y} \mid_v, \mathbf{x}) \right) \quad (1)$$

- \mathbf{x} is a data sequence outcome
- \mathbf{y} is a label sequence outcome
- v is a vertex from vertex set V = set of label random variables
- e is an edge from edge set E over V
- f_k and g_k are given and fixed features; each g_k is a property of x and a vertex v for the label random variable associated with v .
- k is the number of features
- $\Theta = (\lambda_1, \lambda_2, \dots; \mu_1, \mu_2, \dots)$; λ_k and μ_k are parameters to be estimated
- $y|_e$ is the components of y defined by edge e
- $y|_v$ is one component of y defined by vertex v

Matrix Random Variable


- Add special start and stop states $\mathbf{y}_0 = \text{start}$ and $\mathbf{y}_{n+1} = \text{stop}$
- e_i is the edge with labels $(\mathbf{Y}_{i-1}, \mathbf{Y}_i)$
- v_i is the vertex with label \mathbf{Y}_i
- curly \mathbf{Y} = set of possible label values
- The matrix random variable has the range $|\text{curly } \mathbf{Y}| \times |\text{curly } \mathbf{Y}|$
- If $\mathbf{Y}_i = y_i$ then $y_i \in \text{curly } \mathbf{Y}$

Suppose that $p_{\Theta}(\mathbf{Y} \mid \mathbf{X})$ is a CRF given by (1). Assume \mathbf{Y} is a chain.

For each position i in the observation sequence \mathbf{x} , we define a matrix random variable $M_i(\mathbf{x}) = [M_i(y', y \mid \mathbf{x})]$ as:

$$M_i(y', y \mid \mathbf{x}) = \exp(\Lambda_i(y', y \mid \mathbf{x}))$$

$$\begin{aligned} \Lambda_i(y', y \mid \mathbf{x}) = & \sum_k \lambda_k f_k(e_i, \mathbf{Y} \mid_{e_i}, \mathbf{Y} \mid_{e_i} = (y', y), \mathbf{x}) + \\ & \sum_k \mu_k g_k(v_i, \mathbf{Y} \mid_{v_i} = y, \mathbf{x}) \end{aligned}$$



Conditional Probability for CRFs

The conditional probability of a label sequence \mathbf{y} is

$$p_{\Theta}(\mathbf{y} \mid \mathbf{x}) = \frac{\prod_{i=1}^{n+1} M_i(\mathbf{y}_{i-1}, \mathbf{y}_i \mid \mathbf{x})}{\left(\prod_{i=1}^{n+1} M_i(\mathbf{x}) \right)_{\text{start, stop}}}$$

- n is length of sequence $\mathbf{y} = y_1 \dots y_n$
- $y_0 = \text{start}$ and $y_{n+1} = \text{stop}$

Parameter Estimation

Input: training data $D = \{(\mathbf{x}^{(i)}, y^{(i)})\}$, where $i = 1 \dots N$ with empirical distribution $\tilde{p}(\mathbf{x}, y)$

Output: parameters $\Theta = (\lambda_1, \lambda_2, \dots; \mu_1, \mu_2, \dots)$

Maximize: the log-likelihood objective function:


$$\begin{aligned} O(\Theta) &= \sum_{i=1}^N \log p_{\Theta}(\mathbf{y}^{(i)} \mid \mathbf{x}^{(i)}) \\ &\propto \sum_{\mathbf{x}, y} \tilde{p}(\mathbf{x}, y) \log p_{\Theta}(\mathbf{y} \mid \mathbf{x}) \end{aligned}$$

- This is an expectation.


rib/rob models

Training data $\{<\text{rib}, 123>, <\text{rob}, 453>\}$


HMM



MEMM



CRF



Modeling the label bias problem

- Each state emits its designated symbol with probability $29/32$ and the other symbols with probability $1/32$ (on picture)
- They train MEMM and CRF with the same topologies
- A run consists of 2,000 training examples and 500 test examples, trained to convergence
- CRF error is 4.6%, and MEMM error is 42%
- MEMM fails to discriminate between the two branches

Mixed-order HMM

State transition probabilities are given by

$$p_{\alpha}(\mathbf{y}_i \mid \mathbf{y}_{i-1}, \mathbf{y}_{i-2}) = \alpha p_2(\mathbf{y}_i \mid \mathbf{y}_{i-1}, \mathbf{y}_{i-2}) + (1 - \alpha) p_1(\mathbf{y}_i \mid \mathbf{y}_{i-1})$$

Emission probabilities are given by

$$p_{\alpha}(\mathbf{x}_i \mid \mathbf{y}_i, \mathbf{x}_{i-1}) = \alpha p_2(\mathbf{x}_i \mid \mathbf{y}_i, \mathbf{x}_{i-1}) + (1 - \alpha) p_1(\mathbf{x}_i \mid \mathbf{y}_i)$$

$\alpha = 0$ is a standard first-order HMM.

A first-order HMM has transition probabilities given by

$$p_{\alpha}(y_i \mid y_{i-1}, y_{i-2}) = (1 - \alpha) p_1(y_i \mid y_{i-1})$$

And emission probabilities given by

$$p_{\alpha}(\mathbf{x}_i \mid \mathbf{y}_i, \mathbf{x}_{i-1}) = (1 - \alpha) p_1(\mathbf{x}_i \mid \mathbf{y}_i)$$

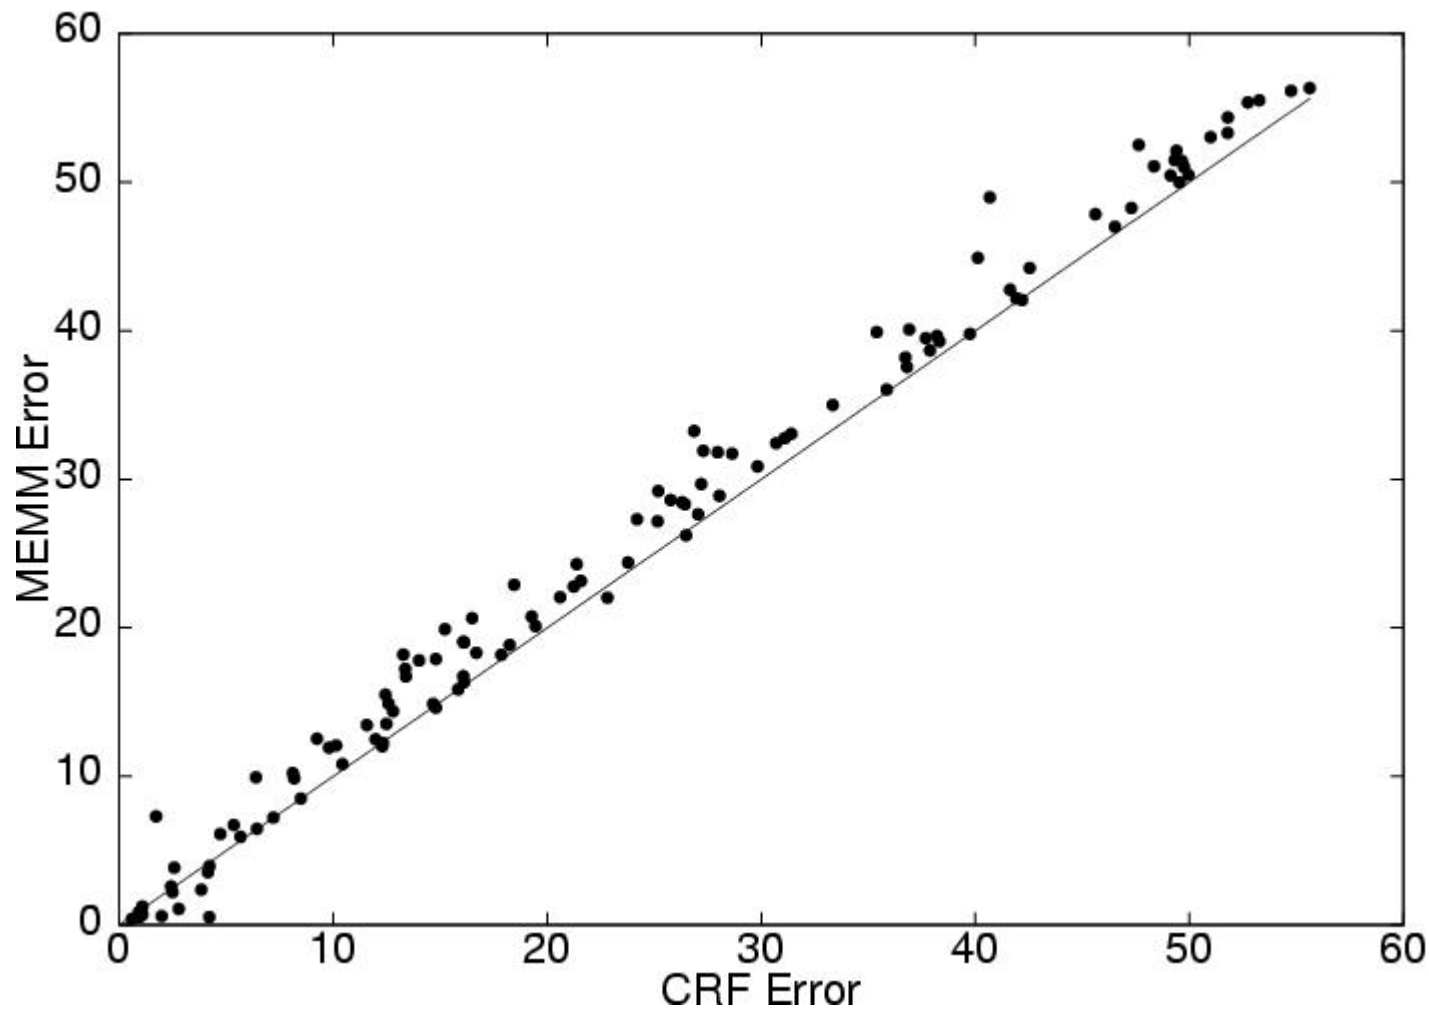
Modeling mixed-order sources

Experimental Setup

- For each randomly generated model, a sample of 1,000 sequences of length 25 is generated for training and testing
- On each randomly generated training set, a CRF is trained using Algorithm S, which is described in the paper
- On the same data an MEMM is trained using iterative scaling
- The Viterbi algorithm is used to label a test set
- *the advantages of the additional representational power of CRFs and MEMMs relative to HMMs.*


MEMM versus CRF

- CRF usually outperforms the MEMM




MEMM versus HMM

- The HMM outperforms the MEMM



CRF versus HMM

Each open square represents a data set with $\alpha < 1/2$, and a solid circle indicates a data set with $\alpha \geq 1/2$; When the data is mostly second order ($\alpha \geq 1/2$), the discriminatively trained CRF usually outperforms the HMM



UPenn Tagging Task

- 45 tags (syntactic), 1M words training

DT	NN	NN		NN	VBZ	RB		JJ
The	asbestos	fiber	;	crocidolite	;	is	unusually	resilient

IN	PRP	VBZ	DT	NNS	IN	RB	JJ	NNS		
once	it	enters	the	lungs	;	with	;	even	brief	exposures

TO	PRP	VBG		NNS	WDT	VBP	RP	NNS		JJ
to	it	causing		symptoms	that	show	up	decades		later ;

	NNS		VBD
researchers		said	

POS tagging experiment 1

- Compared HMMs, MEMMs, and CRFs on Penn treebank POS tagging
- Trained first-order HMM, MEMM, and CRF models as in the synthetic data experiments
- Introduced parameters $\mu_{y,x}$ for each tag-word pair and $\lambda_{y'y}$ for each tag-tag pair in the training set
- oov = out-of-vocabulary, and are not observed in the training set

model	error	oov error
HMM	5.69%	45.99%
MEMM	6.37%	54.61%
CRF	5.55%	48.05%

POS tagging experiment 2

- Compared HMMs, MEMMs, and CRFs on Penn treebank POS tagging
- Each word in a given input sentence must be labeled with one of 45 syntactic tags
- Add a small set of orthographic features: whether a spelling begins with a number or upper case letter, whether it contains a hyphen, and if it contains one of the following suffixes: -ing, -ogy, -ed, -s, -ly, -ion, -tion, -ity, -ies

			using spelling features			
model	error	oov error	error	delta	oov error	delta
HMM	5.69%	45.99%				
MEMM	6.37%	54.61%	4.81%	-25%	26.99%	-50%
CRF	5.55%	48.05%	4.27%	-24%	23.76%	-50%

Presentation Message

- Discriminative models are prone to the label bias problem
- CRFs provide the benefits of discriminative models
- CRFs solve the label bias problem well, and demonstrate good performance