1-2013

Estimation of Monthly Volatility: An Empirical Comparison of Realized Volatility, GARCH and ACD-ICV Methods

Shouwei Liu

Yiu Kuen TSE
Singapore Management University, yktse@smu.edu.sg

Follow this and additional works at: http://ink.library.smu.edu.sg/soe_research

Part of the Finance Commons

Estimation of Monthly Volatility: An Empirical Comparison of Realized Volatility, GARCH and ACD-ICV Methods

Shouwei Liu
School of Economics, Singapore Management University

Yiu-Kuen Tse
School of Economics, Singapore Management University

February 2012

Abstract: We apply the ACD-ICV method proposed by Tse and Yang (2011) for the estimation of intraday volatility to estimate monthly volatility, and empirically compare this method against the realized volatility (RV) and generalized autoregressive conditional heteroskedasticity (GARCH) methods. Our Monte Carlo results show that the ACD-ICV method performs well against the other two methods. Evidence on the Chicago Board Options Exchange volatility index (VIX) shows that it predicts the ACD-ICV volatility estimates better than it predicts the RV estimates. While the RV method is popular for the estimation of monthly volatility, its performance is inferior to the GARCH method.

Corresponding Author: Yiu-Kuen Tse, School of Economics, Singapore Management University, 90 Stamford Road, Singapore 178903, email: `yktse@smu.edu.sg`.

1 Introduction

Many studies in the empirical finance literature on the risk-return relationship involve the estimation of the monthly return volatility of stocks. Higher frequency (such as daily) data are usually not used because returns are typically rather noisy, rendering the risk-return relationship in high frequency difficult to establish. Furthermore, many studies examine the effects of macroeconomic variables on asset pricing, and these variables are only available monthly or quarterly. For example, Schwert (1989) constructed vector autoregression models involving monthly data of short-term interest rates, long-term yields of high-quality and medium-quality bonds, inflation rates and industrial production to analyze the dynamic structure of stock volatility. For some recent studies requiring estimates of monthly stock volatility, see Goyal and Santa-Clara (2003), Bali, Cakici, Yan and Zhang (2005), Guo and Savickas (2008), Ludvigson and Ng (2007), Jiang and Tian (2010) and Zhang (2010).

Since the seminal work of French, Schwert and Stambaugh (1987) and Schwert (1989), researchers often use the sum of the squared *daily* stock returns over a month (or some modifications of it) as an estimate of the monthly volatility of the stock. Later, Andersen, Bollerslev, Diebold and Ebens (2001) and Andersen, Bollerslev, Diebold and Labys (2001) proposed to use the sum of the squared returns of tick data to estimate intraday volatility, and called this estimate the realized volatility (RV). Since then the literature on RV has expanded very quickly.

The estimation of intraday volatility using RV methods typically requires sampling over 5-min intervals or shorter, with sampling over 1- or 2-min intervals not uncommon. Given the asymptotic theories established in the litertaure and the availability of large numbers of return observations over short durations in a trading day, the RV methods have a firm theoretical underpinning as a tool for estimating intraday volatility. In contrast, in applying the RV methods to estimate monthly volatility using daily data there are only approximately 21 return observations to compute each monthly estimate. Thus, estimation errors may be a concern and may weaken the validity of the statistical inference. As high-frequency data have become increasingly available we extend the use of monthly RV estimation to these data.

Another line of research applies the autoregressive conditional heteroskedasticity (ARCH) model of Engle (1982) and the generalized ARCH (GARCH) model of Bollerslev (1986) to estimate monthly volatility. French, Schwert and Stambaugh (1987) estimated monthly volatility using a GARCH-in-mean model, while Fu (2009) estimated monthly idiosyncratic risk using the exponential GARCH (EGARCH) model of Nelson (1991). Monthly volatility can be estimated using GARCH type of models on monthly data, or using these models on daily data from which aggregates of daily conditional variances form a monthly estimate.

Recently, Tse and Yang (2011) proposed a method to estimate high-frequency volatility using the autoregressive conditional duration (ACD) model of Engle and Russell (1998), called the ACD-ICV method. They estimate high-frequency volatility (over a day or shorter intervals) by integrating the instantaneous conditional return variance per unit time obtained from the ACD models. Unlike the RV methods, which sample data over regular intervals, the ACD-ICV method samples price events based on high-frequency transaction price changes exceeding a threshold. ACD models for the durations between sequential price events are estimated using the quasi maximum likelihood method, and the conditional variance over a given intraday interval is computed by integrating the instantaneous conditional variance within the interval. The Monte Carlo results of Tse and Yang (2011) show that the ACD-ICV method gives lower root mean-squared error than the RV methods in estimating intraday volatility. While Tse and Yang (2011) focused on the estimation of intraday volatility, in this paper we apply the ACD-ICV method to estimate monthly volatility.

The literature so far has little to say about the choice of the estimation method for monthly volatility. While the RV approach seems to dominate the literature, the use of the GARCH type of models is not uncommon. In addition, the ACD-ICV method may be a useful alternative, as it has been shown to perform well for the estimation of intraday volatility. In this paper we compare the performance of the RV, GARCH and ACD-ICV methods using Monte Carlo (MC) experiments and empirical data from the New York Stock Exchange (NYSE). Our MC results show that the ACD-ICV method outperforms the RV method in giving lower root mean-squared error. Indeed, it turns out that the GARCH method performs better than the RV method in the MC experiments. We also examine the use of the Chicago

Board Options Exchange (CBOE) volatility index (VIX) as a predictor of the volatility for the next 30 days estimated by the ACD-ICV, RV and GARCH methods using the S&P500 index. Our results show that VIX predicts the 30-day ACD-ICV volatility estimates better than it predicts the RV estimates.

The rest of the paper proceeds as follows. Section 2 summarizes the monthly volatility estimates considered in this paper. In Section 3 we report some MC results on the comparison of the RV, GARCH and ACD-ICV methods. The MC study suggests that the best results for the ACD-ICV method appear to be obtained when the range of the return for defining the price event is about 0.15% to 0.35%. It also shows that the ACD-ICV method performs very well against the RV method. Section 4 reports our results for the estimation of monthly volatility using some empirical data from the NYSE. In Section 5 we examine the use of VIX as a predictor of the market volatility over the next 30 days, with market volatility estimated using the ACD-ICV, RV and GARCH methods. Finally, Section 6 concludes.

2   Methods of Estimating Monthly Volatility

Volatility estimation over monthly or quarterly intervals dated back to the 1970s. Researchers in earlier work adopted the 12-month rolling standard-deviation estimate as the volatility estimate of the centered month, as in Officer (1973), Fama (1976), and Merton (1980). Schwert (1989) employed a two-step rolling regression to construct monthly volatility, which allows the conditional mean return to vary over time and allows different weights for the lagged absolute unexpected returns. Since the work of French, Schwert and Stambaugh (1987), Schwert (1989), Schwert (1990a), Schwert (1990b) and Schwert and Seguin (1990), the use of the sum of the squared daily returns over a month, called the RV method, has emerged as the most popular method for the estimation of monthly stock volatility. On the other hand, monthly volatility can also be estimated using GARCH models estimated with monthly data, or by aggregating daily conditional variances over a month estimated from GARCH models with daily data. In this section we summarize the methods of estimating monthly volatility examined in this study.

2.1   RV Method

Let r_i denote the return on day i of the month, $\bar{r}$ denote the average daily return of the month and N denote the number of trading days in the month. The basic RV estimate of the variance of the month,

denoted by V_D, is defined as

$$V_D = \sum_{i=1}^{N} (r_i - \bar{r})^2, \tag{1}$$

which has been widely adopted in the literature.[1] However, as this estimator uses only about 21 observations its accuracy may be a concern. As high-frequency data have become easily available, we consider the use of transaction data to estimate monthly RV. For the purpose of using as much data as possible, shorter sampling intervals are preferred. However, returns over short sampling intervals may be contaminated by market microstructure noise. To balance between these two conflicting goals, we use 5-min price data to calculate the RV and denote it by V_R.[2]

2.2   GARCH Method

Another popular method for constructing monthly volatility is to use the GARCH model and its extensions. In this paper, we adopt the EGARCH(1, 1) model defined by

$$\log \sigma_t^2 = \omega + \alpha \frac{\varepsilon_{t-1}}{\sigma_{t-1}} + \beta \left| \frac{\varepsilon_{t-1}}{\sigma_{t-1}} \right| + \gamma \log \sigma_{t-1}^2, \tag{2}$$

where ε_t is the residual of the return and σ_t^2 is the conditional variance. We assume that the standardized residual $z_t = \varepsilon_t / \sigma_t$ follows the generalized error distribution (GED) with density function

$$f_Z(z) = \frac{\nu \exp\left[-\frac{1}{2} \left| \frac{z}{\lambda} \right|^\nu \right]}{\lambda \, 2^{1+\frac{1}{\nu}} \, \Gamma\left(\frac{1}{\nu} \right)}, \quad -\infty < z < \infty, \, 0 < \nu \leq \infty, \tag{3}$$

where $\nu > 0$, $\Gamma(\cdot)$ denotes the gamma function and $\lambda \equiv [2^{(-2/\nu)} \Gamma(1/\nu)/\Gamma(3/\nu)]^{1/2}$. We estimate the EGARCH(1, 1) model using daily data and compute the daily conditional variance $\hat{\sigma}_t^2$, for $t = 1, \cdots, N$, over N days of the month. We then aggregate the estimated daily conditional variances and denote it by V_G, so that[3]

$$V_G = \sum_{t=1}^{N} \hat{\sigma}_t^2. \tag{4}$$

---

[1]French, Schwert and Stambaugh (1987) proposed a method to correct for the serial correlation in the return series. However, we find this method to be inferior to V_D and we will not report its results in this paper.

[2]This is in contrast to Jiang and Tian (2005) and Becker, Clements and White (2007), who used 30-min data to compute the RV. To extend the use of equation (1) to intraday returns, we have to consider the treatment of overnight price jumps. We tried using the first price and the 5th-min price of the day to compute the overnight return. As the results are similar we report only the case of using the first price. In V_R only the return (not deviation from mean return) is used.

[3]In the MC study we also consider estimating the GARCH model using monthly data, with the monthly volatility estimates directly computed without aggregation. The results, however, are poor and will not be reported.

2.3 ACD-ICV Method

The ACD model was first proposed by Engle and Russell (1998) to analyze the durations of transaction data. Tse and Yang (2011) proposed to estimate intraday volatility by integrating the instantaneous conditional variance per unit time estimated from the ACD model, resulting in the ACD-ICV method. In this section, we first review the ACD-ICV method proposed by Tse and Yang (2011), followed by an outline of the modification of this method for the estimation of monthly volatility.

Let $t_0, t_1, \cdots, t_N$ denote a sequence of times for which t_i is the time of occurrence of the ith price event, which is said to have occurred if the cumulative change in the logarithmic transaction price since the last price event is at least of a preset amount δ (whether upwards or downwards), called the price range. Thus, $x_i = t_i - t_{i-1}$, for $i = 1, 2, \cdots, N$, are the intervals between consecutive price events, called the price durations. Let Φ_i be the information set upon the transaction at time t_i, and denote $\psi_i = \mathrm{E}(x_i \mid \Phi_{i-1})$, which is the conditional expectation of the transaction duration. If the standardized durations $\epsilon_i = x_i/\psi_i$ are independently and identically distributed as exponential variables, the integrated conditional variance in the interval (t_0, t_N), denoted by ICV, is given by

$$\mathrm{ICV} = \delta^2 \sum_{i=0}^{N-1} \frac{t_{i+1} - t_i}{\psi_{i+1}}. \tag{5}$$

The implementation of the ACD-ICV method to estimate monthly volatility depends on the price data available. In this paper we use higher-frequency transaction data. We ignore the overnight close of the market and treat the first trade of each day as continuously away from the last trade of the previous trading day. Given the return range δ we compile $t_1, \cdots, t_{N-1}$ based on the continuous transaction data as the price-event times over a month.

The use of equation (5) requires estimates of the conditional expected duration ψ_{i+1}. Following Tse and Yang (2011) we adopt the augmented ACD (AACD) model (see Fernandes and Grammig (2006)) defined by

$$\psi_i^\lambda = \omega + \alpha \psi_{i-1}^\lambda [|\epsilon_{i-1} - b| + c(\epsilon_{i-1} - b)]^\upsilon + \beta \psi_{i-1}^\lambda. \tag{6}$$

We estimate the AACD model using the quasi maximum likelihood estimation (QMLE) method, with the standardized duration assumed to be exponentially distributed. The ACD-ICV estimate of monthly

volatility, denoted by V_A is computed as

$$V_A = \delta^2 \sum_{i=0}^{N-1} \frac{t_{i+1} - t_i}{\hat{\psi}_{i+1}}, \tag{7}$$

where $\hat{\psi}_{i+1}$ is the QMLE of ψ_{i+1}. The choice of δ affects the fit of the ACD$(1,1)$ model for price duration, and hence the performance of V_A as an estimate of the monthly ICV. In our empirical application we vary δ from 0.15% through 0.35% in steps of 0.05%, and denote the resulting estimates by V_{Aj} for $j = 1, \cdots, 5$, respectively, to correspond to δ being 0.15%, 0.20%, 0.25%, 0.30% and 0.35%.

3   Monte Carlo Study

We adopt the usual assumption in the empirical literature that the logarithmic stock price follows a Brownian semimartingale (BSM). Denoting the stock price at time t by $p(t)$ and defining $\tilde{p}(t) = \log p(t)$, we assume

$$\tilde{p}(t) = \int_0^t \mu(t) \, dt + \int_0^t \sigma(t) \, dW(t), \tag{8}$$

where $\mu(t)$ is the instantaneous drift rate, $\sigma^2(t)$ is the instantaneous variance and $W(t)$ is a standard Brownian process. For each MC sample, we generate data over 5 years, with a total of 60 months of observations.

For the drift term $\mu(t)$ we consider two different artificial processes, which are plotted in Figure 1. For the variance process $\sigma^2(t)$ we consider two methods: deterministic volatility and stochastic volatility models, which will be described in the next two subsections. Given the drift term $\mu(t)$ and the variance term $\sigma^2(t)$, we generate the logarithmic price series $\tilde{p}(t)$ by the equation

$$\tilde{p}(t + \Delta t) = \tilde{p}(t) + \mu(t) \, \Delta t + \sigma(t) \sqrt{\Delta t} \, \epsilon, \tag{9}$$

where $\epsilon \sim N(0,1)$. We take Δt to be one second and the starting price to be \$100. We further add to the generated series a jump component, which is assumed to follow a Poisson process with a mean of 0.4 per five minutes. When a jump occurs, it takes value of –\$0.05, –\$0.03, \$0.03 and \$0.05 with probabilities of 0.25 each. Finally, we also consider a price process consisting of a BSM with a white noise. Defining the noise-to-signal (NSR) ratio as NSR $= [\text{Var}\{\varepsilon(t)\}/\text{Var}\{\sigma(t)\}]^{\frac{1}{2}}$ we set NSR $= 0.6$.[4]

---

[4]The BSM is assumed to generate the efficient price, while the transaction price consists of the white noise.

From the generated price series p_t we round the price to the nearest cent and sample the rounded price by 1 cent to obtain the transaction price. Depending on the estimation method considered we also sub-sample the transaction price at 5-min and 1-day intervals.

3.1 Deterministic Volatility Model

For the deterministic instantaneous variance term $\sigma^2(t)$, we assume two artificial processes: a sinusoidal function and an empirical function, which are plotted in Figure 2. We call the model in the first panel DV Model 1, which is a sinusoidal function. The empirical function described in the second panel is obtained by spline-smoothing the empirical volatility function estimated using the GE data with a EGARCH(1, 1) model.

3.2 Stochastic Volatility Model

For the stochastic volatility model we adopt the set-up due to Heston (1993) as follows

$$d\,\tilde{p}(t) = \left(\mu(t) - \frac{\sigma^2(t)}{2} \right) dt + \sigma(t)\,dW_1(t), \tag{10}$$

and

$$d\,\sigma^2(t) = \kappa(\alpha - \sigma^2(t))\,dt + \gamma\sigma(t)\,dW_2(t), \tag{11}$$

where $W_1(t)$ and $W_2(t)$ are standard Brownian processes with a correlation coefficient of ρ. Two different sets of parameters are adopted for the Heston model. First, we set $\kappa = 4$, $\alpha = 0.04$, $\gamma = 0.4$ and $\rho = -0.5$, which will be called SV Model 1. Second, we vary SV Model 1 by setting the reversion-rate parameter κ to 5 and the volatility-rate parameter γ to 0.5, which is called SV Model 2, and was used by Aït-Sahalia and Mancini (2008). In addition we further increase the reversion-rate parameter κ to 6 with the volatility-rate parameter γ being 0.6, which is called SV Model 3.

3.3 Overnight Price Jump

While BSM may approximate price movements when the market is open, the process is disrupted when the market is closed. To this effect, it is important to examine how overnight price jumps affect the performance of the monthly volatility estimates. While the estimation of intraday volatility can be studied without taking account of overnight price jumps, this issue cannot be overlooked when the

objective is to estimate monthly volatility. Our data of the ten stocks show that the maximum absolute price jump of eight stocks are larger than 10%.

We consider two models for the overnight returns: the generalized normal-distribution model and the t-distribution model. Denoting the overnight return by Y, the generalized normal-distribution model assumes that $Y \sim GN(\mu, \alpha, \beta)$, with density function

$$f_Y(y) = \frac{\beta}{2\alpha \Gamma\left(\frac{1}{\beta}\right)} e^{-\frac{|y-\mu|}{\alpha}}. \tag{12}$$

On the other hand, the t-distribution model assumes that the density function of Y is

$$f_Y(y) = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\sigma\sqrt{\nu\pi}\,\Gamma\left(\frac{\nu}{2}\right)} \left[\frac{\nu + \left(\frac{y-\mu}{\sigma}\right)^2}{\nu}\right]^{-\frac{\nu+1}{2}}, \tag{13}$$

where μ and σ are the location and scale parameters, respectively, and ν is the degrees of freedom.

In our MC experiment, we consider $\mu = 0$, $\alpha = 0.0026$ and $\beta = 0.69$ for the generalized normal distribution model, and $\sigma = 0.004$, $\mu = 0$ and $\nu = 2.5$ for the t-distribution model.[5]

3.4 Monte Carlo Results

Tables 1 through 4 summarize the MC results on the comparison of the performance of different monthly volatility estimates, presenting their mean error (ME) and root mean-squared error (RMSE). All results are estimated using 1000 MC replications. To save space, only results for V_{A1}, V_{A3} and V_{A5} are reported.

Tables 1 shows the results for the case when there are no overnight price jumps. It can be seen that V_R performs better than the ACD-ICV method, which in turn is better than V_D and V_G. It should be noted, however, that V_A uses tick data, while V_D and V_G use only daily data. Surprisingly, V_G outperforms V_D in all reported cases, although the latter is more widely used in the literature.

The results for the cases when there are overnight price jumps are summarized in Tables 2, 3 and 4. The V_A estimates perform the best for both deterministic and stochastic volatility models, providing lower RMSE versus V_R, which uses 5-min data. Rather surprisingly, for the deterministic volatility

---

[5]These parameters are chosen to resemble the estimates obtained from the ten stocks. We also consider four other cases of distributions with thinner tails for robustness check.

models, V_G based on daily data performs well against V_R based on 5-min data. The results are, however, different for the stochastic volatility models, for which V_R clearly outperforms V_G. As in the case with no overnight price jumps, if only daily data are available for estimation, GARCH estimates outperform RV estimates.[6] Figure 3 presents examples of monthly volatility plots generated from the three stochastic volatility models with their estimates over a sample of 60 months. It can be seen that all estimates trace the true volatility quite closely, although V_R appears to be more volatile.

4 Empirical Results for NYSE Data

We apply the RV, GARCH and ACD-ICV methods to estimate monthly volatility using empirical data from the Trade and Quotation (TAQ) database provided through the Wharton Research Data Services (WRDS). We select ten actively traded stocks listed on the NYSE without company merger and acquisition from 2003 through 2007, with 60 months of data. The price changes due to stock splits are adjusted according to the capitalization of the company. The selected stocks and their codes are: Bank of America (BAC), General Electric (GE), Merck (MRK), Johnson & Johnson (JNJ), JP Morgan (JPM), Wal Mart (WMT), IBM (IBM), Pfizer (PFE), AT & T (T) and Chevron (CVX).

We estimate the monthly volatilities for the ten stocks in our sample. Figure 4 plots the estimates over the sample period. To avoid jamming the figures, only the estimates V_{A3} (for $\delta = 0.25\%$), V_R and V_G are presented. It can be seen that all estimates track each other quite closely, and there does not appear to be any systematic bias among the different methods. The RV method, however, exhibits a few extreme values of high volatility estimates and generally have the largest fluctuations among the three methods. It is interesting to observe that the volatility paths of the different stocks show significant co-movements.

Table 5 summarizes the pairwise correlation coefficients of the volatility estimates V_{A3}, V_R and V_G for the ten stocks. It can be seen that the correlations are highest for V_{A3}, followed by V_G and then V_R. Of the 45 pairs of correlations 93.3% are maximized when V_{A3} is used as the volatility estimate, the remaining 6.7% are maximized when V_G is used as the volatility estimate, with none for

---

[6]Results for alternative model parameters for robustness check are similar. These results are not presented here, but can be obtained from the authors on request.

V_R. Many studies in the literature examine the effects of macroeconomic variables on stock volatility, and generally points to the co-movements of volatility across stocks. The ACD-ICV estimates support a higher volatility co-movement versus estimates based on the RV method. It will be interesting to further investigate volatility co-movements using the ACD-ICV estimates, in particular in relation to macroeconomic variables such as inflation, exchange rate, GDP growth and interest-rate movements.

5   Volatility of S&P500

Implied volatility computed from option prices has often been used as a predictor for future historical volatility. The S&P500 Index volatility has been a case of particular research interest in the literature due to the popular reference to the CBOE volatility index VIX. Whaley (2009) provided a description of VIX and discussed some of its properties. In this section we examine the use of VIX as a predictor for future historical volatility when RV, GARCH and ACD-ICV estimates are used as proxies for historical volatility.

VIX is calculated and disseminated in real time by CBOE. It is a forward-looking index of the expected return volatility of the S&P500 Index over the next 30 calendar days and is implied from the prices of S&P500 Index options. VIX is quoted in percentage points as the annualized standard deviation of the return of the S&P500 Index over the next 30 days. It is based upon a model-free formula using a wide range of selected near- and near-term put and call options. Studies in the literature on the forecasting performance of implied volatility often use RV as the proxy for historical volatility. Jiang and Tian (2005) and Becker, Clements and White (2007) used RV computed over 30-day intervals as proxy for 30-day historical volatility in their studies on the information content of VIX on the volatility of the S&P500. Recently, Chung, Tsai, Wang and Weng (2011) considered both VIX and VIX options as predictors for the RV of the S&P500, although they did not specify the RV method used. We shall investigate the forecasting performance of VIX for the volatility of S&P500 when historical volatility is estimated by GARCH and ACD-ICV, and compare the results against using RV.

We downloaded daily closing values of VIX from the website of the CBOE. S&P500 tick data were obtained from The Institute for Financial Markets (IFM) Data Center. 5-min S&P500 data were

extracted from 8:30 to 15:00 (Chicago time) each day. The sample period is from 1998 through 2007, with 2516 daily observations.

We select N time points in the sample period that are at least 30 calendar days apart and denote them by t_i, for $i = 1, \cdots, N$. Altogether there are 117 nonoverlapping 30-day intervals in our sample ($N = 117$). Let VIX_i be the closing value of VIX on day t_i. We denote Y_i as an estimate of the historical volatility over the 30-day period starting from time t_i, and consider the following regressions of historical volatility estimates on volatility forecasts using VIX

$$Y_i = \alpha + \beta \, \text{VIX}_i + \xi_i, \tag{14}$$

and

$$Y_i^2 = \alpha + \beta \, \text{VIX}_i^2 + \xi_i, \tag{15}$$

for $i = 1, \cdots, N$. The results are summarized in Table 6. It can be seen that the highest R^2 is for the regressions with the ACD-ICV measures as the dependent variables. Rather remarkably, the regressions with V_R as the dependent variable produce the lowest R^2. The results show that VIX is a more successful predictor of future volatility if volatility is estimated by the ACD-ICV method, but not the RV method. Figure 5 plots VIX and some historical volatility estimates. There are some periods for which VIX over-predicts volatility as estimated by V_R. This over-prediction, however, is not evident if historical volatility is estimated by the ACD-ICV method. Overall, our results show that VIX has higher prediction value if its performance is measured against historical estimates using the ACD-ICV method.

6 Conclusion

In this paper we extend the ACD-ICV method proposed by Tse and Yang (2011) to estimate stock volatility over longer intervals such as a month. Estimation of low-frequency volatility is important for studies involving macroeconomic data that are available only monthly or quarterly. As returns over longer intervals are less susceptible to the contamination of noise over short intervals they may be preferred in studies on asset pricing. Our MC study suggests that price events defined by return

thresholds of about 0.15% to 0.35% are appropriate for the ACD-ICV method. Based on the transaction data, the ACD-ICV method outperforms the RV method in our MC experiments. On the other hand, if daily data are used, the GARCH method based on aggregating the daily estimates of the conditional variance is superior to the RV method, which is widely used in the literature.

Our empirical results using ten NYSE stocks show that the ACD-ICV, RV and GARCH estimates track each other quite closely. The RV estimates, however, have larger fluctuations and exhibit occasionally extreme volatility estimates. Co-movements of volatility across different stocks are highest according to the ACD-ICV estimates. Our empirical study on VIX and the S&P500 index shows that VIX is a more successful predictor of future volatility if volatility is estimated by the ACD-ICV method than by the RV method. Overall we have shown that using the ACD-ICV method on high-frequency data (tick transaction data) provides superior estimates of low-frequency volatility (over monthly intervals) to the RV method.

References

[1] Aït-Sahalia, Y., and L. Mancini, 2008, Out of sample forecasts of quadratic variation, Journal of Econometrics, 147, 17-33.

[2] Andersen, T.G., T. Bollerslev, F.X. Diebold, and H. Ebens, 2001, The distribution of Realized volatility, Journal of Financial Economics, 61, 43-76.

[3] Andersen, T.G., T. Bollerslev, F.X. Diebold, and P. Labys, 2001, The distribution of exchange rate volatility, Journal of American Statistical Assoction, 96, 42-55.

[4] Bali, T.G., N. Cakici, X.S. Yan, and Z. Zhang, 2005, Does idiosyncratic risk really matter? Journal of Finance, 60, 905-929.

[5] Becker, R., A.E. Clements, and S.I. White, 2007, Does implied volatility provide any information beyond that captured in model-based volatility forecasts? Journal of Banking and Finance, 31, 2535-2549.

[6] Bollerslev, T., 1986, Generalized autoregressive conditional heteroskedasticity, Journal of Econometrics, 31, 307-327.

[7] Chung, S.L., W.C. Tsai, Y.H. Wang, and P.S. Weng, 2011, The information content of the S&P500 index and VIX options on the dynamics of the S&P500 index, working paper, 21st Asia-Pacific Futures Research Symposium.

[8] Engle, R.F., 1982, Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation, Econometrica, 50, 987-1007.

[9] Engle, R.F., and J.R. Russell, 1998, Autoregressive conditional duration: A new model for irregularly spaced transaction data, Econometrica, 66, 1127-1162.

[10] Fama, E.F., 1976, Inflation uncertainty and expected returns on Treasury bills, Journal of Political Economy, 84, 427-448.

[11] Fernandes, M., and J. Grammig, 2006, A family of autoregressive conditional duration models, Journal of Econometrics, 130, 1-23.

[12] French, K.R., G.W. Schwert, and R.F. Stambaugh, 1987, Expected stock returns and volatility, Journal of Financial Economics, 19, 3-29.

[13] Fu, F., 2009, Idiosyncratic risk and the cross-section of expected stock returns, Journal of Financial Economics, 91, 24-37.

[14] Goyal, A., and P. Santa-Clara, 2003, Idiosyncratic risk matters! Journal of Fiance, 58, 975-1007.

[15] Guo, H., and R. Savickas, 2008, Average idiosyncratic volatility in G7 countries, Reviews of Financial Studies, 21, 1259-1296.

[16] Heston, S.L., 1993, A closed-form solution for options with stochastic volatility with application to bond and currency options, Review of Financial Studies, 6, 327-343.

[17] Jiang, G.J., and Y.S. Tian, 2005, The model-free implied volatility and its information content, Review of Financial Studies, 18, 1305-1342.

[18] Jiang, G.J., and Y.S. Tian, 2010, Forecasting volatility using long memory and comovements: An application to option valuation under SFAF 123R, Journal of Financial and Quantitative Analysis, 45, 502-533.

[19] Ludvigson, S.C., and S. Ng, 2007, The empirical risk-return relation: A factor analysis approach, Journal of Financial Economics, 83, 171-222.

[20] Merton, R.C., 1980, On estimating the expected return on the market: An exploratory investigation, Journal of Financial Economics, 8, 323-361.

[21] Nelson, D.B., 1991, Conditional heteroskedasticity in assert returns: a new approach, Econometrica, 59, 347-370.

[22] Officer, R., 1973, The variability of the market factor of the New York Stock Exchange, Journal of Business, 46, 434-454.

[23] Schwert, G.M., 1989, Why does stock market volatility change over time? Journal of Finance, 44, 1115-1153.

[24] Schwert, G.M., 1990a, Stock market volatility, Financial Analysts Journal, 46, 23-34.

[25] Schwert, G.M., 1990b, Stock Volatility and the Crash of '87, Review of Financial studies, 3, 77-102.

[26] Schwert, G.M., and P.J. Seguin, 1990, Heteroskedasticity in stock returns, Journal of Finance, 45, 1129-1155.

[27] Tse, Y.K., and T. Yang, 2011, Estimation of high-frequency volatility: An autoregressive conditional duration approach, working paper, Singapore Management University.

[28] Whaley, R.E., 2009, Understanding the VIX, Journal of Portfolio Management, 35, 98-105.

[29] Zhang, C., 2010, A Reexamination of the Causes of Time-Varying Stock Return Volatilities, Journal of Financial and Quantitative Analysis, 45, 663-684.

Table 1: Monte Carlo results without overnight jumps.

	MV1		MV2		SV1		SV2		SV3	
	\multicolumn{10}{c}{Volatility Model}									
method	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE
Panel A: Drift Model 1										
V_{A1}	0.538	0.654	0.578	0.681	0.894	1.115	0.839	1.061	0.816	1.041
V_{A3}	0.350	0.603	0.377	0.596	0.616	0.994	0.599	0.977	0.624	0.990
V_{A5}	0.272	0.669	0.298	0.639	0.507	1.060	0.540	1.089	0.587	1.131
V_R	0.043	0.272	0.034	0.325	0.025	0.519	0.028	0.486	0.031	0.466
V_D	-0.514	2.395	-0.624	2.892	-1.037	4.675	-0.957	4.358	-0.906	4.165
V_G	0.133	1.745	0.191	1.916	0.446	3.769	0.535	3.797	0.654	3.825
Panel B: Drift Model 2										
V_{A1}	0.554	0.666	0.596	0.699	0.949	1.172	0.882	1.106	0.856	1.080
V_{A3}	0.356	0.607	0.387	0.604	0.654	1.030	0.631	1.007	0.655	1.018
V_{A5}	0.286	0.673	0.306	0.643	0.534	1.091	0.562	1.109	0.606	1.149
V_R	0.044	0.272	0.035	0.325	0.026	0.520	0.029	0.486	0.032	0.466
V_D	-0.514	2.395	-0.624	2.892	-1.037	4.675	-0.957	4.358	-0.906	4.165
V_G	0.133	1.747	0.203	1.915	0.446	3.768	0.536	3.797	0.656	3.825

Notes: ME = mean error, RMSE = root mean-squared error. The results are based on 1000 MC replications of 5-year monthly volatility. All figures are annualized standard deviation in percentage. V_{A1}, V_{A3} and V_{A5} are the ACD-ICV volatility estimates with $\delta = 0.15\%$, 0.25% and 0.35%, respectively. V_R and V_D are the realized volatility estimates computed using 5-min and daily returns, respectively. V_G is the GARCH estimate based on daily data.

Table 2: Monte Carlo results with overnight jumps following the Generalized Normal Distribution

	MV1		MV2		SV1		SV2		SV3	
	\multicolumn{10}{c}{Volatility Model}									
method	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE
Panel A: Drift Model 1										
V_{A1}	0.249	1.890	0.502	1.321	0.827	1.659	0.606	1.850	0.424	2.061
V_{A3}	0.006	1.892	0.227	1.280	0.406	1.513	0.218	1.762	0.077	2.046
V_{A5}	-0.077	1.849	0.111	1.288	0.237	1.526	0.106	1.820	0.024	2.085
V_R	-0.182	2.860	-0.124	2.563	-0.066	2.026	-0.087	2.164	-0.104	2.273
V_D	-0.885	4.042	-0.919	4.211	-1.224	5.403	-1.170	5.170	-1.138	5.039
V_G	0.436	2.328	0.390	2.301	0.330	3.808	0.349	3.864	0.385	3.927
Panel B: Drift Model 2										
V_{A1}	0.274	1.895	0.530	1.327	0.904	1.735	0.672	1.901	0.485	2.103
V_{A3}	0.020	1.889	0.243	1.279	0.459	1.562	0.263	1.795	0.116	2.070
V_{A5}	-0.069	1.847	0.121	1.285	0.275	1.558	0.139	1.845	0.050	2.104
V_R	-0.181	2.860	-0.123	2.563	-0.064	2.026	-0.086	2.164	-0.103	2.272
V_D	-0.884	4.042	-0.919	4.211	-1.224	5.403	-1.169	5.170	-1.138	5.039
V_G	0.439	2.327	0.413	2.329	0.327	3.804	0.350	3.863	0.385	3.926

Notes: ME = mean error, RMSE = root mean-squared error. The results are based on 1000 MC replications of 5-year monthly volatility. All figures are annualized standard deviation in percentage. V_{A1}, V_{A3} and V_{A5} are the ACD-ICV volatility estimates with $\delta = 0.15\%$, 0.25% and 0.35%, respectively. V_R and V_D are the realized volatility estimates computed using 5-min and daily returns, respectively. V_G is the GARCH estimate based on daily data.

Table 3: Monte Carlo results with overnight jumps following the *t* Distribution

					Volatility Model					
	MV1		MV2		SV1		SV2		SV3	
method	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE
Panel A: Drift Model 1										
V_{A1}	-1.165	2.106	-0.767	1.402	-0.065	1.496	-0.341	1.779	-0.562	2.061
V_{A3}	-1.411	2.256	-1.041	1.581	-0.484	1.558	-0.732	1.872	-0.913	2.179
V_{A5}	-1.493	2.274	-1.154	1.670	-0.659	1.648	-0.846	1.962	-0.974	2.230
V_R	-1.705	3.537	-1.456	3.145	-1.038	2.462	-1.143	2.647	-1.225	2.791
V_D	-2.369	4.685	-2.215	4.717	-2.167	5.716	-2.197	5.534	-2.230	5.443
V_G	-1.092	2.460	-0.905	2.405	-0.595	3.759	-0.654	3.801	-0.680	3.847
Panel B: Drift Model 2										
V_{A1}	-1.143	2.094	-0.741	1.380	0.011	1.552	-0.276	1.818	-0.502	2.082
V_{A3}	-1.398	2.246	-1.027	1.570	-0.432	1.591	-0.687	1.891	-0.874	2.188
V_{A5}	-1.484	2.269	-1.146	1.661	-0.617	1.669	-0.813	1.976	-0.945	2.238
V_R	-1.704	3.537	-1.455	3.144	-1.037	2.462	-1.142	2.646	-1.224	2.790
V_D	-2.369	4.685	-2.215	4.717	-2.167	5.716	-2.196	5.534	-2.231	5.443
V_G	-1.116	2.469	-0.922	2.398	-0.595	3.758	-0.657	3.799	-0.680	3.845

Notes: ME = mean error, RMSE = root mean-squared error. The results are based on 1000 MC replications of 5-year monthly volatility. All figures are annualized standard deviation in percentage. V_{A1}, V_{A3} and V_{A5} are the ACD-ICV volatility estimates with $\delta = 0.15\%$, 0.25% and 0.35%, respectively. V_R and V_D are the realized volatility estimates computed using 5-min and daily returns, respectively. V_G is the GARCH estimate based on daily data.

Table 4: Monte Carlo results with overnight jumps randomly drawn from the empirical jumps

					Volatility Model					
	MV1		MV2		SV1		SV2		SV3	
method	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE	ME	RMSE
Panel A: Drift model 1										
V_{A1}	0.240	1.939	0.487	1.377	0.729	1.608	0.524	1.828	0.356	2.060
V_{A3}	0.012	1.930	0.226	1.330	0.349	1.506	0.176	1.773	0.039	2.054
V_{A5}	-0.065	1.880	0.120	1.333	0.195	1.534	0.081	1.838	0.001	2.097
V_R	-0.286	3.459	-0.200	3.142	-0.099	2.509	-0.127	2.667	-0.150	2.788
V_D	-0.986	4.525	-0.994	4.633	-1.256	5.663	-1.209	5.464	-1.184	5.358
V_G	0.369	2.359	0.386	2.341	0.342	3.862	0.357	3.923	0.387	3.993
Panel B: Drift model 2										
V_{A1}	0.263	1.944	0.513	1.375	0.791	1.665	0.577	1.870	0.403	2.092
V_{A3}	0.022	1.932	0.238	1.331	0.393	1.546	0.208	1.801	0.070	2.076
V_{A5}	-0.055	1.878	0.129	1.332	0.229	1.558	0.106	1.855	0.021	2.112
V_{R1}	-0.285	3.459	-0.200	3.142	-0.098	2.509	-0.126	2.667	-0.149	2.788
V_D	-0.986	4.525	-0.994	4.632	-1.256	5.663	-1.209	5.465	-1.185	5.358
V_G	0.327	2.332	0.407	2.369	0.341	3.861	0.357	3.922	0.387	3.992

Notes: ME = mean error, RMSE = root mean-squared error. The results are based on 1000 MC replications of 5-year monthly volatility. All figures are annualized standard deviation in percentage. V_{A1}, V_{A3} and V_{A5} are the ACD-ICV volatility estimates with $\delta = 0.15\%$, 0.25% and 0.35%, respectively. V_R and V_D are the realized volatility estimates computed using 5-min and daily returns, respectively. V_G is the GARCH estimate based on daily data.

Table 5: Correlations of volatility estimates of different stocks

	BAC	GE	MRK	JNJ	JPM	WMT	IBM	PFE	T	CVX
				Volatility estimate: V_{A3}						
BAC	1.000	**0.714**	**0.734**	0.441	**0.898**	**0.819**	**0.867**	**0.550**	**0.642**	**0.658**
GE		1.000	**0.603**	**0.843**	**0.903**	0.713	**0.828**	**0.772**	**0.806**	0.316
MRK			1.000	**0.517**	**0.720**	**0.648**	**0.704**	**0.691**	**0.632**	**0.744**
JNJ				1.000	**0.687**	0.455	**0.649**	**0.802**	**0.734**	0.218
JPM					1.000	**0.843**	**0.925**	**0.690**	**0.759**	**0.541**
WMT						1.000	**0.870**	**0.512**	0.547	**0.599**
IBM							1.000	**0.678**	**0.701**	**0.616**
PFE								1.000	**0.704**	**0.375**
T									1.000	**0.385**
CVX										1.000

	BAC	GE	MRK	JNJ	JPM	WMT	IBM	PFE	T	CVX
				Volatility estimate: V_R						
BAC	1.000	0.600	0.106	0.253	0.848	0.672	0.667	0.036	0.251	0.564
GE		1.000	0.042	0.776	0.859	0.685	0.716	0.109	0.558	0.207
MRK			1.000	0.062	0.058	0.005	-0.043	0.171	0.065	0.023
JNJ				1.000	0.577	0.384	0.530	0.187	0.535	-0.007
JPM					1.000	0.794	0.808	0.057	0.471	0.480
WMT						1.000	0.669	-0.048	0.327	0.514
IBM							1.000	0.031	0.435	0.451
PFE								1.000	0.226	0.037
T									1.000	-0.052
CVX										1.000

	BAC	GE	MRK	JNJ	JPM	WMT	IBM	PFE	T	CVX
				Volatility estimate: V_G						
BAC	1.000	0.704	0.057	0.247	0.817	0.559	0.655	0.147	0.418	0.120
GE		1.000	0.175	0.686	0.859	**0.812**	0.697	0.349	0.708	-0.049
MRK			1.000	0.332	0.125	0.005	-0.008	0.371	0.163	-0.082
JNJ				1.000	0.517	**0.536**	0.445	0.556	0.641	-0.158
JPM					1.000	0.788	0.781	0.268	0.679	0.174
WMT						1.000	0.600	0.179	**0.604**	-0.069
IBM							1.000	0.182	0.554	0.299
PFE								1.000	0.231	0.055
T									1.000	-0.114
CVX										1.000

Note: Maximized correlation coefficients are in bold face.

Table 6: Regression results of 30-day volatility estimates on VIX

Y:	V_{A1}	V_{A2}	V_{A3}	V_{A4}	V_{A5}	V_R	V_D	V_G

Panel A: $Y_i = \alpha + \beta\,\mathrm{VIX}_i + \xi_i$

	V_{A1}	V_{A2}	V_{A3}	V_{A4}	V_{A5}	V_R	V_D	V_G
α	0.0407	0.0267	0.0223	0.0134	0.0066	-0.0580	0.0145	0.0299
	(4.0632)	(2.5170)	(2.0564)	(1.2418)	(0.5818)	(-2.3715)	(1.1098)	(3.1344)
β	0.7896	0.8449	0.8627	0.8988	0.9213	1.1078	0.6896	0.6330
	(17.5326)	(17.6902)	(17.6637)	(18.5012)	(18.0094)	(10.0801)	(11.7403)	(14.7747)
R^2	0.7254	0.7289	0.7283	0.7463	0.7360	0.4645	0.5412	0.6520

Panel B: $Y_i^2 = \alpha + \beta\,\mathrm{VIX}_i^2 + \xi_i$

	V_{A1}	V_{A2}	V_{A3}	V_{A4}	V_{A5}	V_R	V_D	V_G
α	0.0125	0.0094	0.0075	0.0058	0.0037	-0.0184	0.0059	0.0070
	(4.1117)	(2.8253)	(2.1425)	(1.6487)	(0.9559)	(-1.4203)	(1.7788)	(3.1855)
β	0.7032	0.7584	0.7972	0.8259	0.8600	1.2749	0.4891	0.4602
	(14.2688)	(14.0031)	(14.0789)	(14.3752)	(13.6903)	(6.0709)	(9.0759)	(12.8213)
R^2	0.6359	0.6271	0.6296	0.6394	0.6164	0.2361	0.4123	0.5848

Notes: Y_i is a historical estimate of the volatility of the ith 30-day interval. Numbers in parentheses are t-statistics.
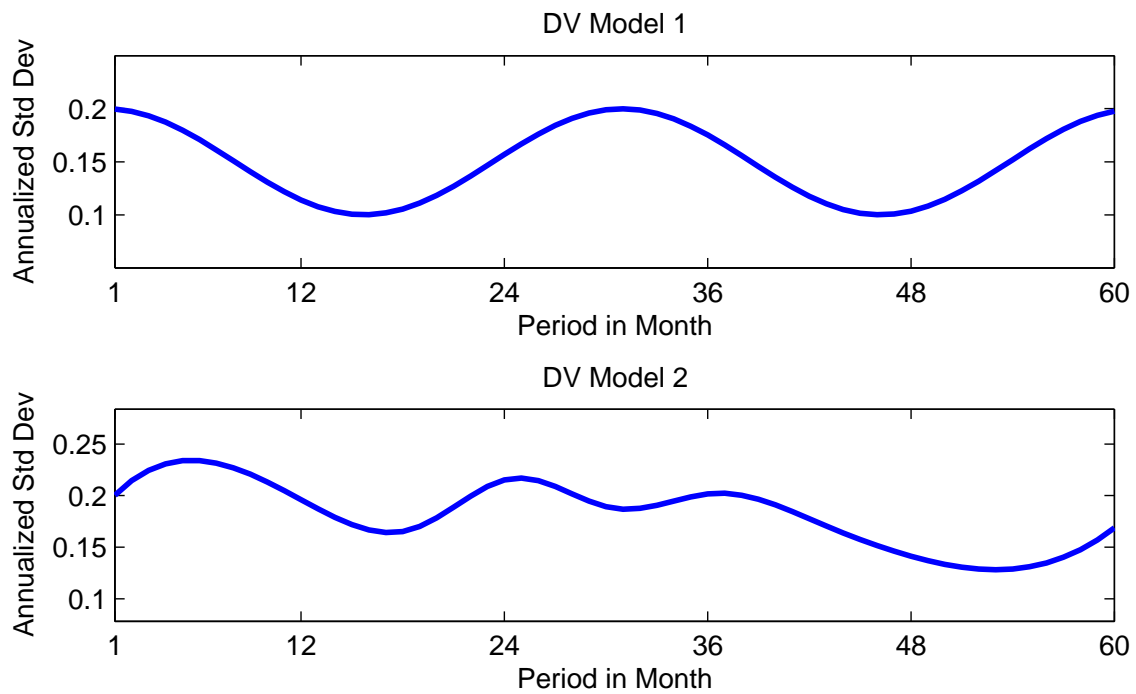
Figure 1: The drift term

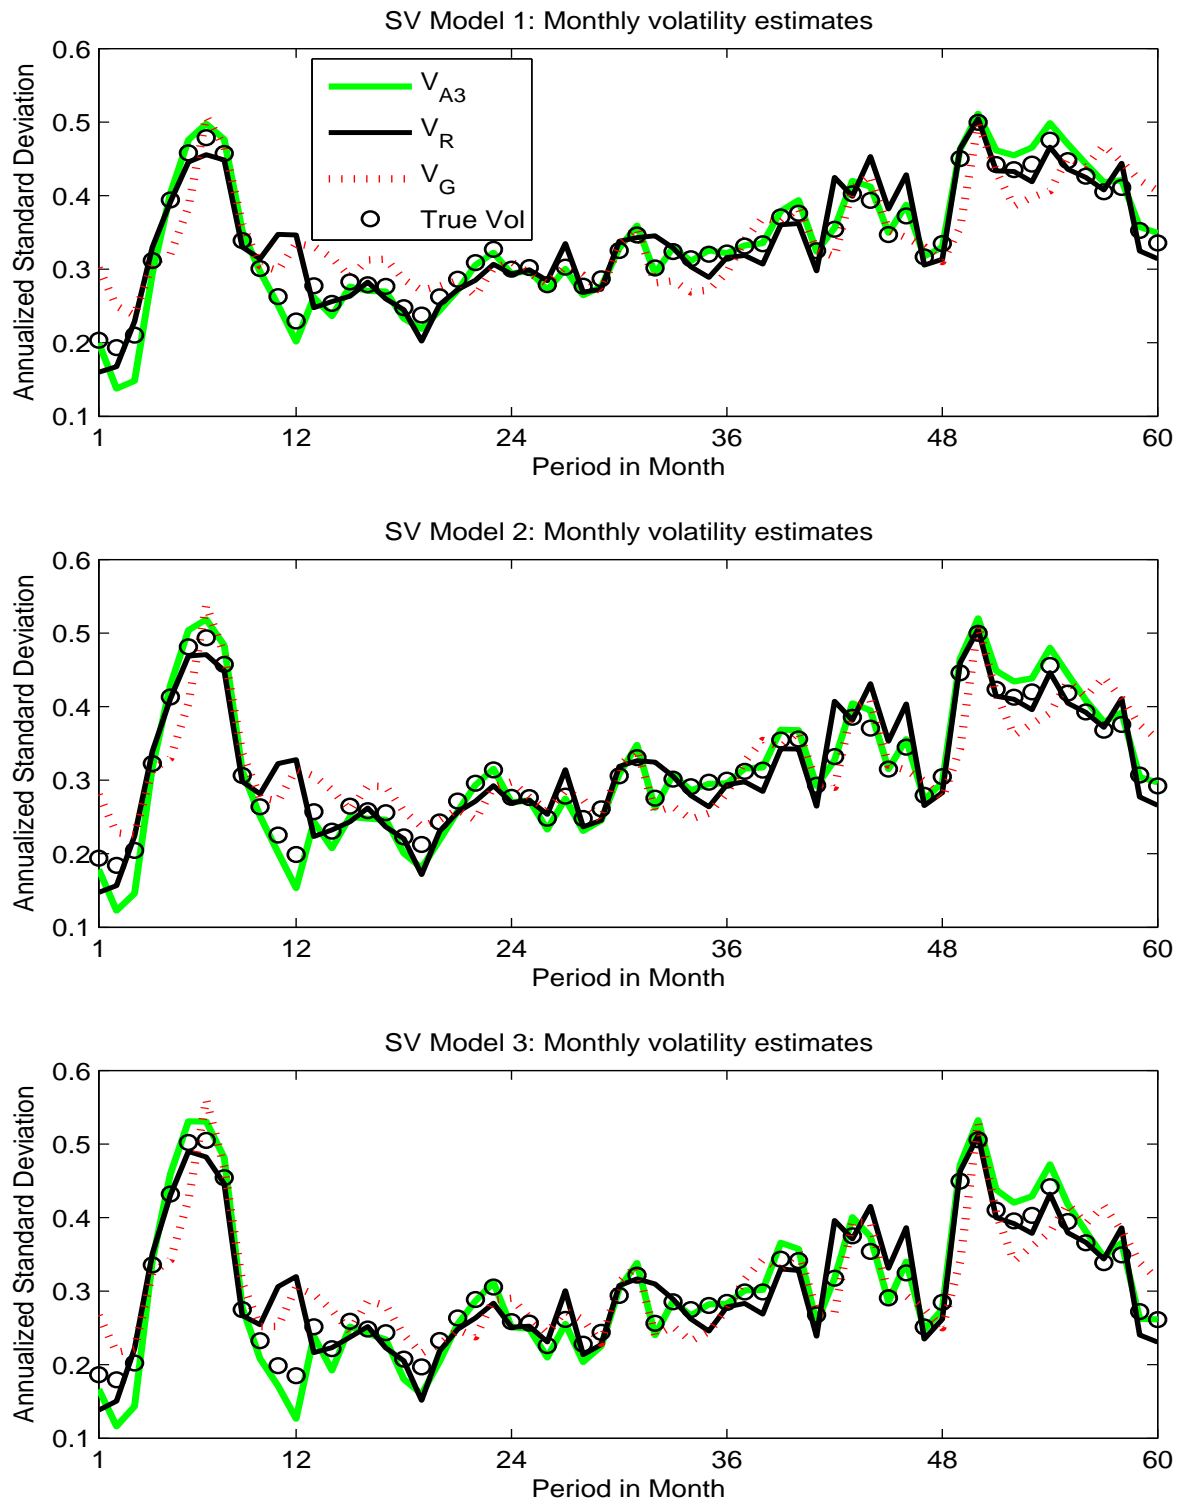**Figure 2:** Deterministic volatility models

Figure 3: Estimation of stochastic volatility with empirical overnight price jumps

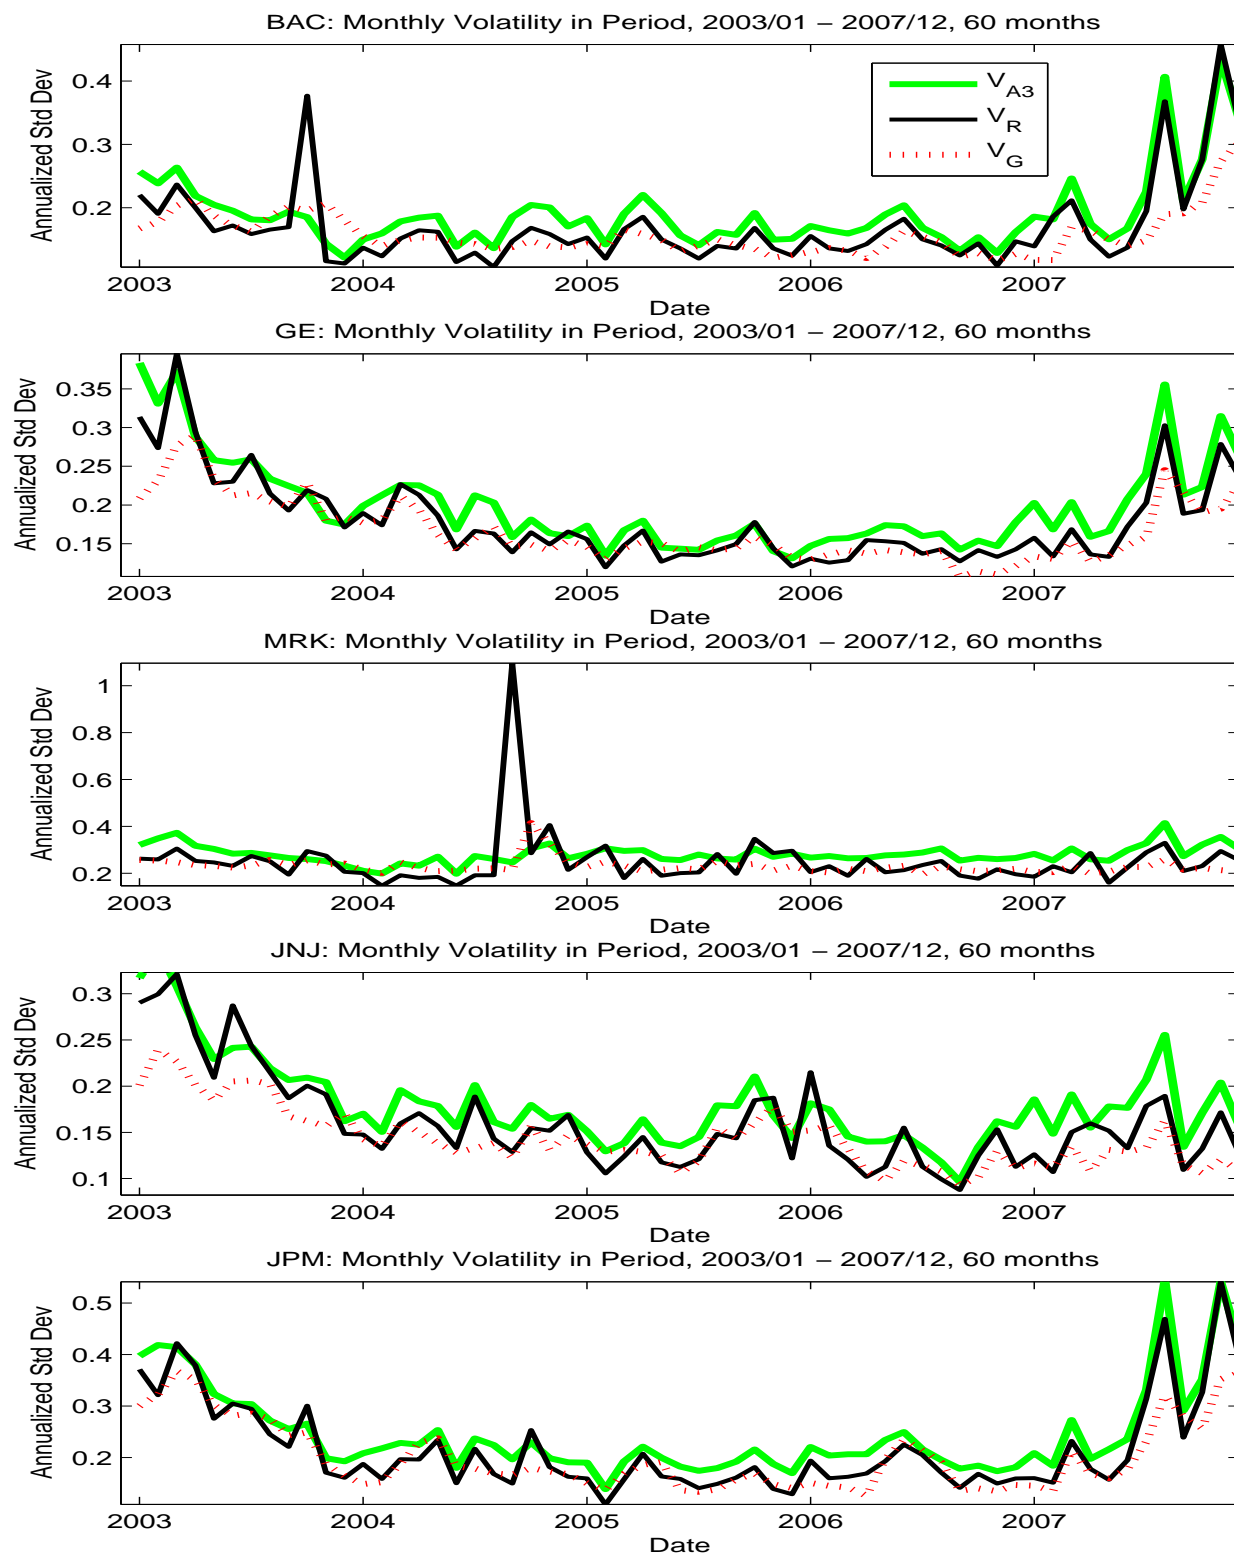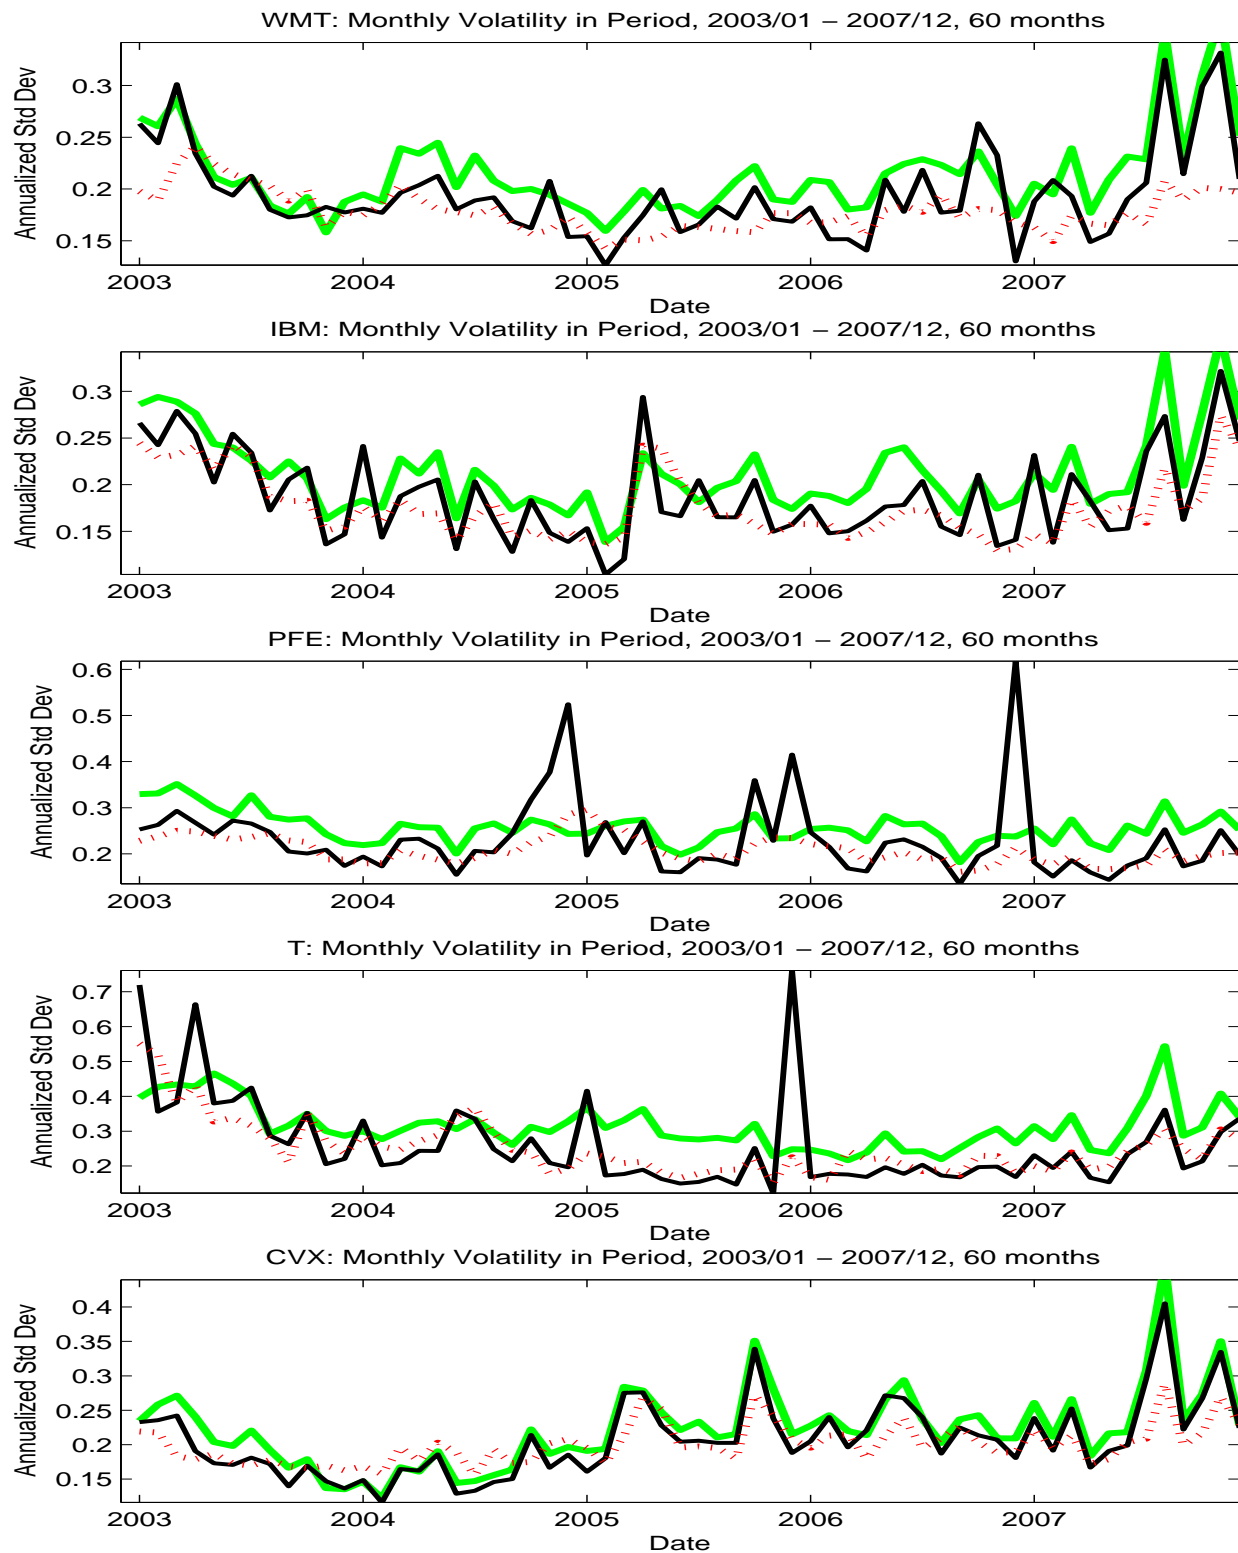**Fig 4a:** Empirical estimates of monthly volatility

Fig 4b: Empirical estimates of monthly volatility

Figure 5: VIX and S&P500 30-day volatility estimates